

Machine  
Learning  
& Compilers

**Optimisation heuristics  
are too course.**

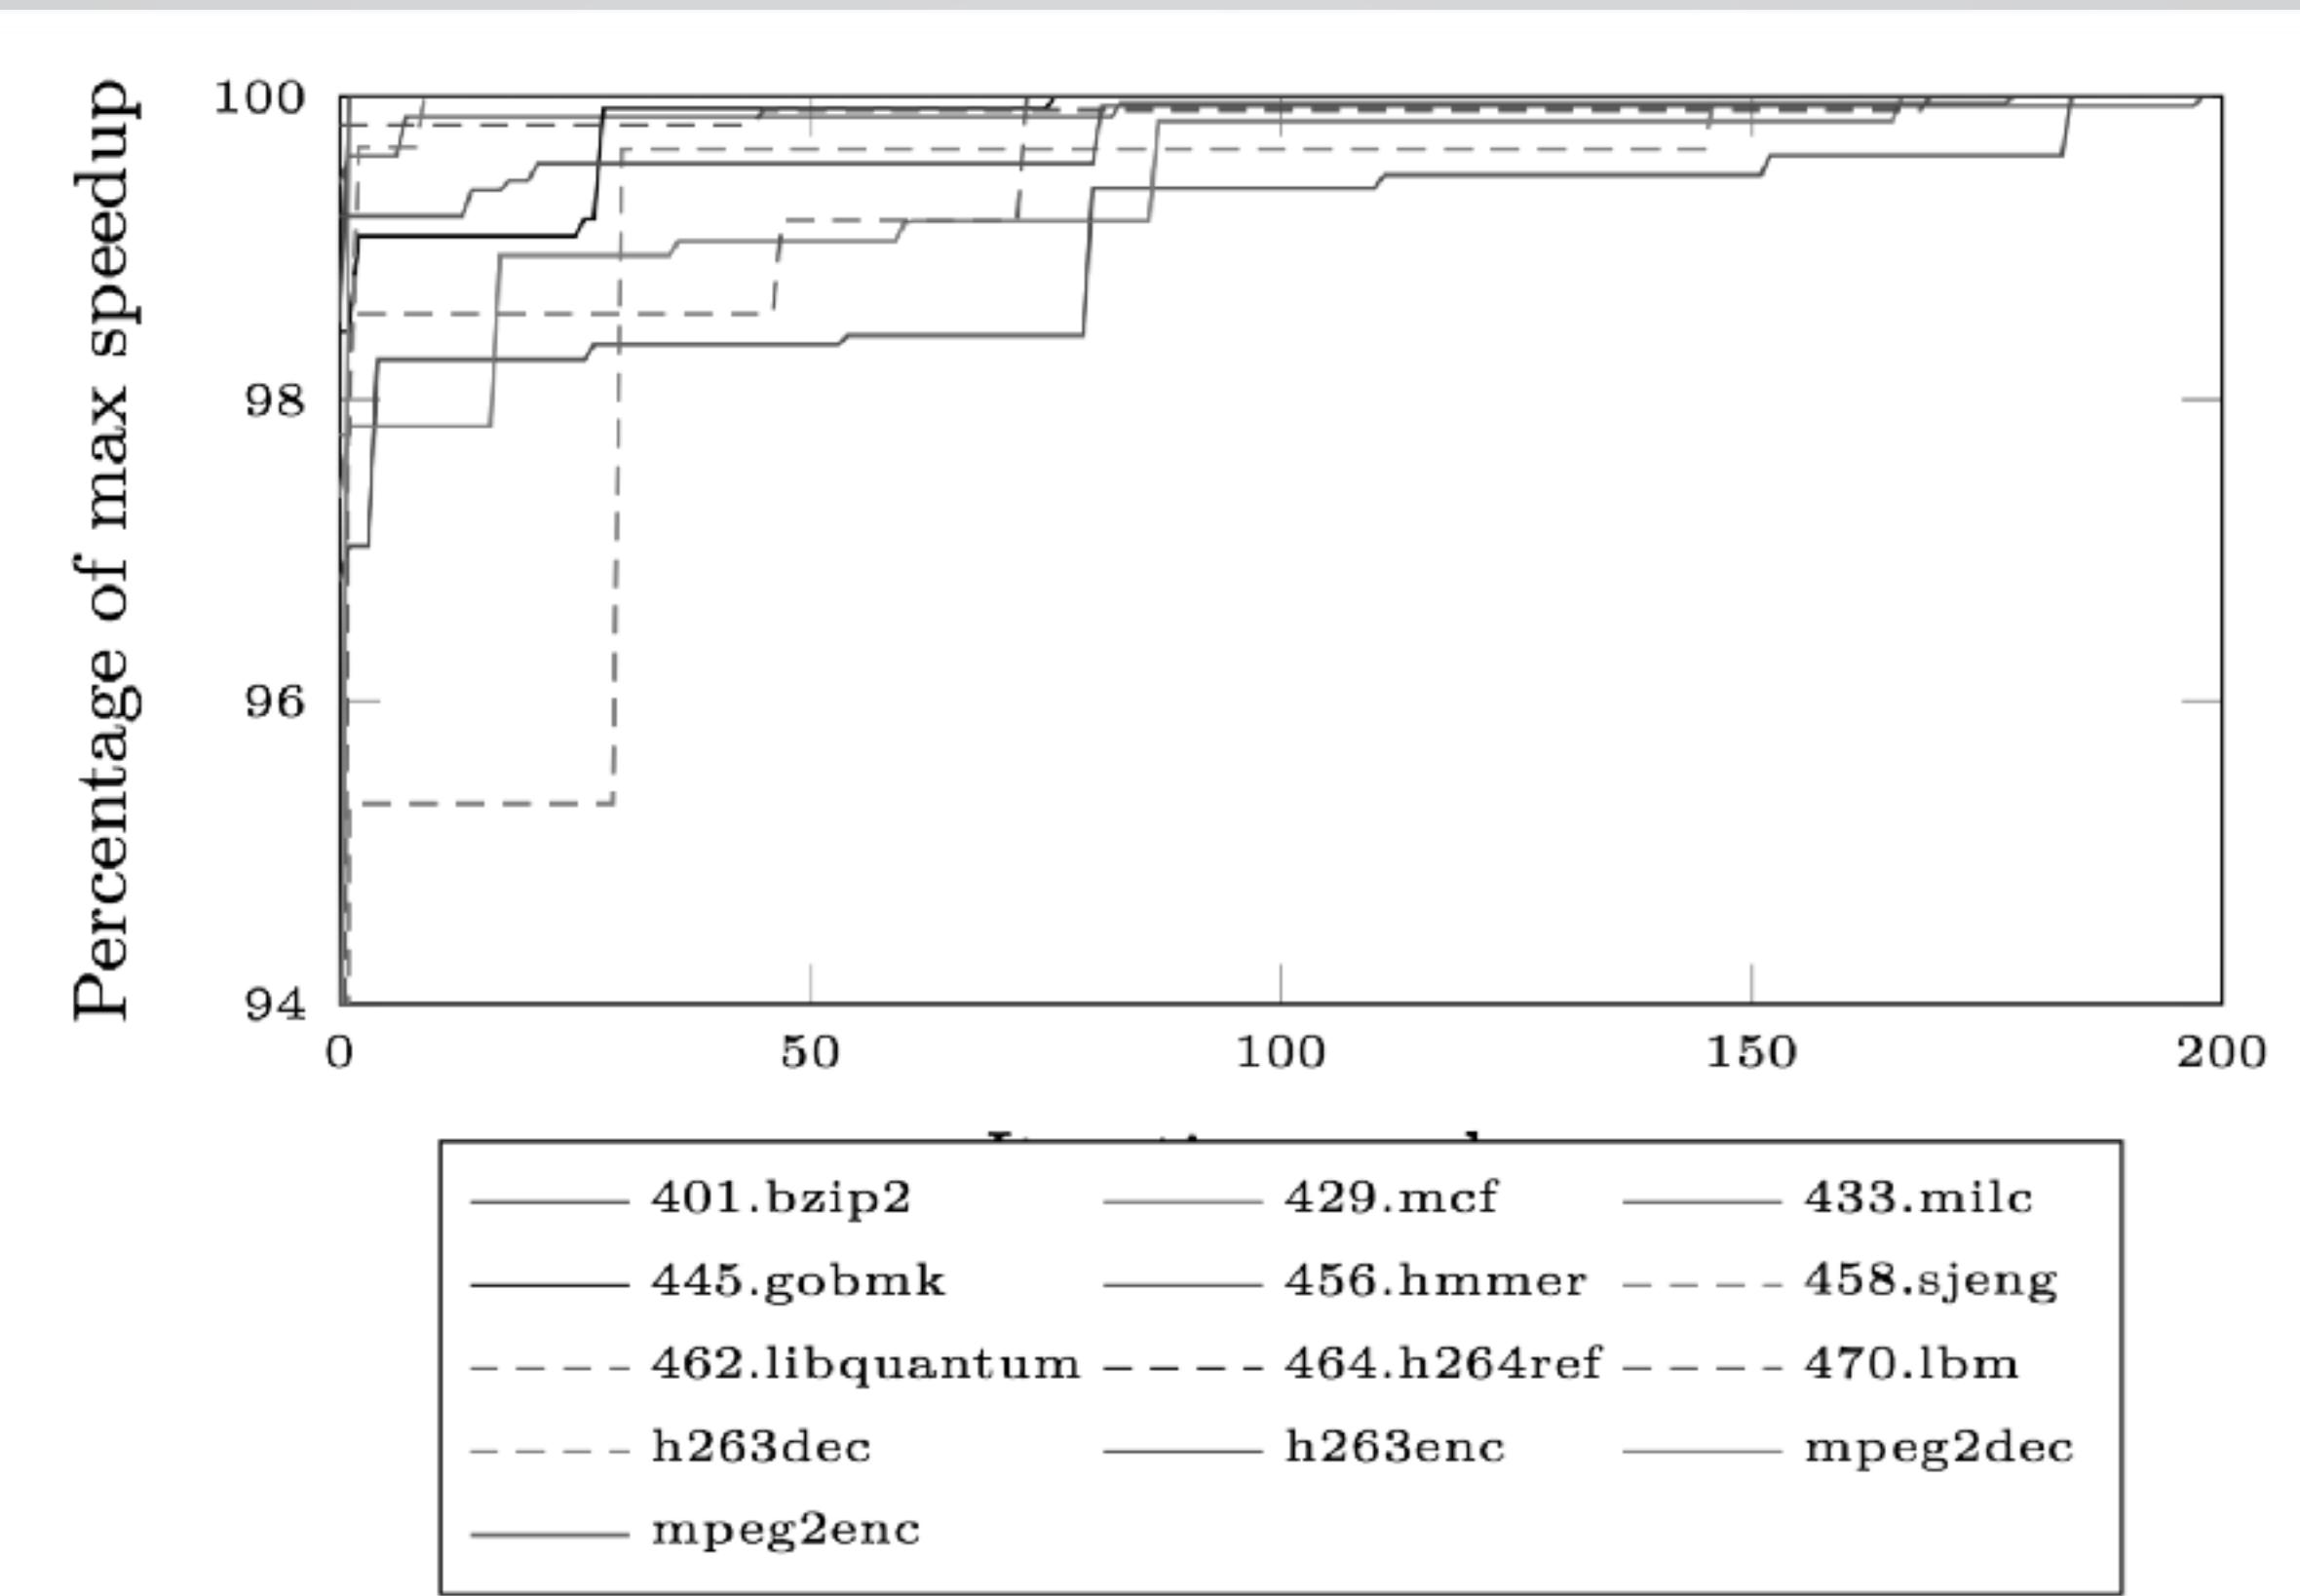
(aka. it doesn't take much to do better than -O3)

# Better than -O3

```
#!/bin/sh

while true; do
    sort --random-sort < "cflags.txt" | head -n 20 | xargs gcc -O1 app.c
    time ./a.out
done
```

after 200 attempts, ~5% improvement



**Figure 6:** The convergence towards best speedup found for each benchmark.

\*actual homework

**200 attempts sounds like a lot...**

( $200 \times 10 \times 30 \text{ s} \approx 16 \text{ hrs}$ )

# Drop in the ocean

250 GCC flags  $\geq 2^{250}$  options  $\approx 10^{75}$

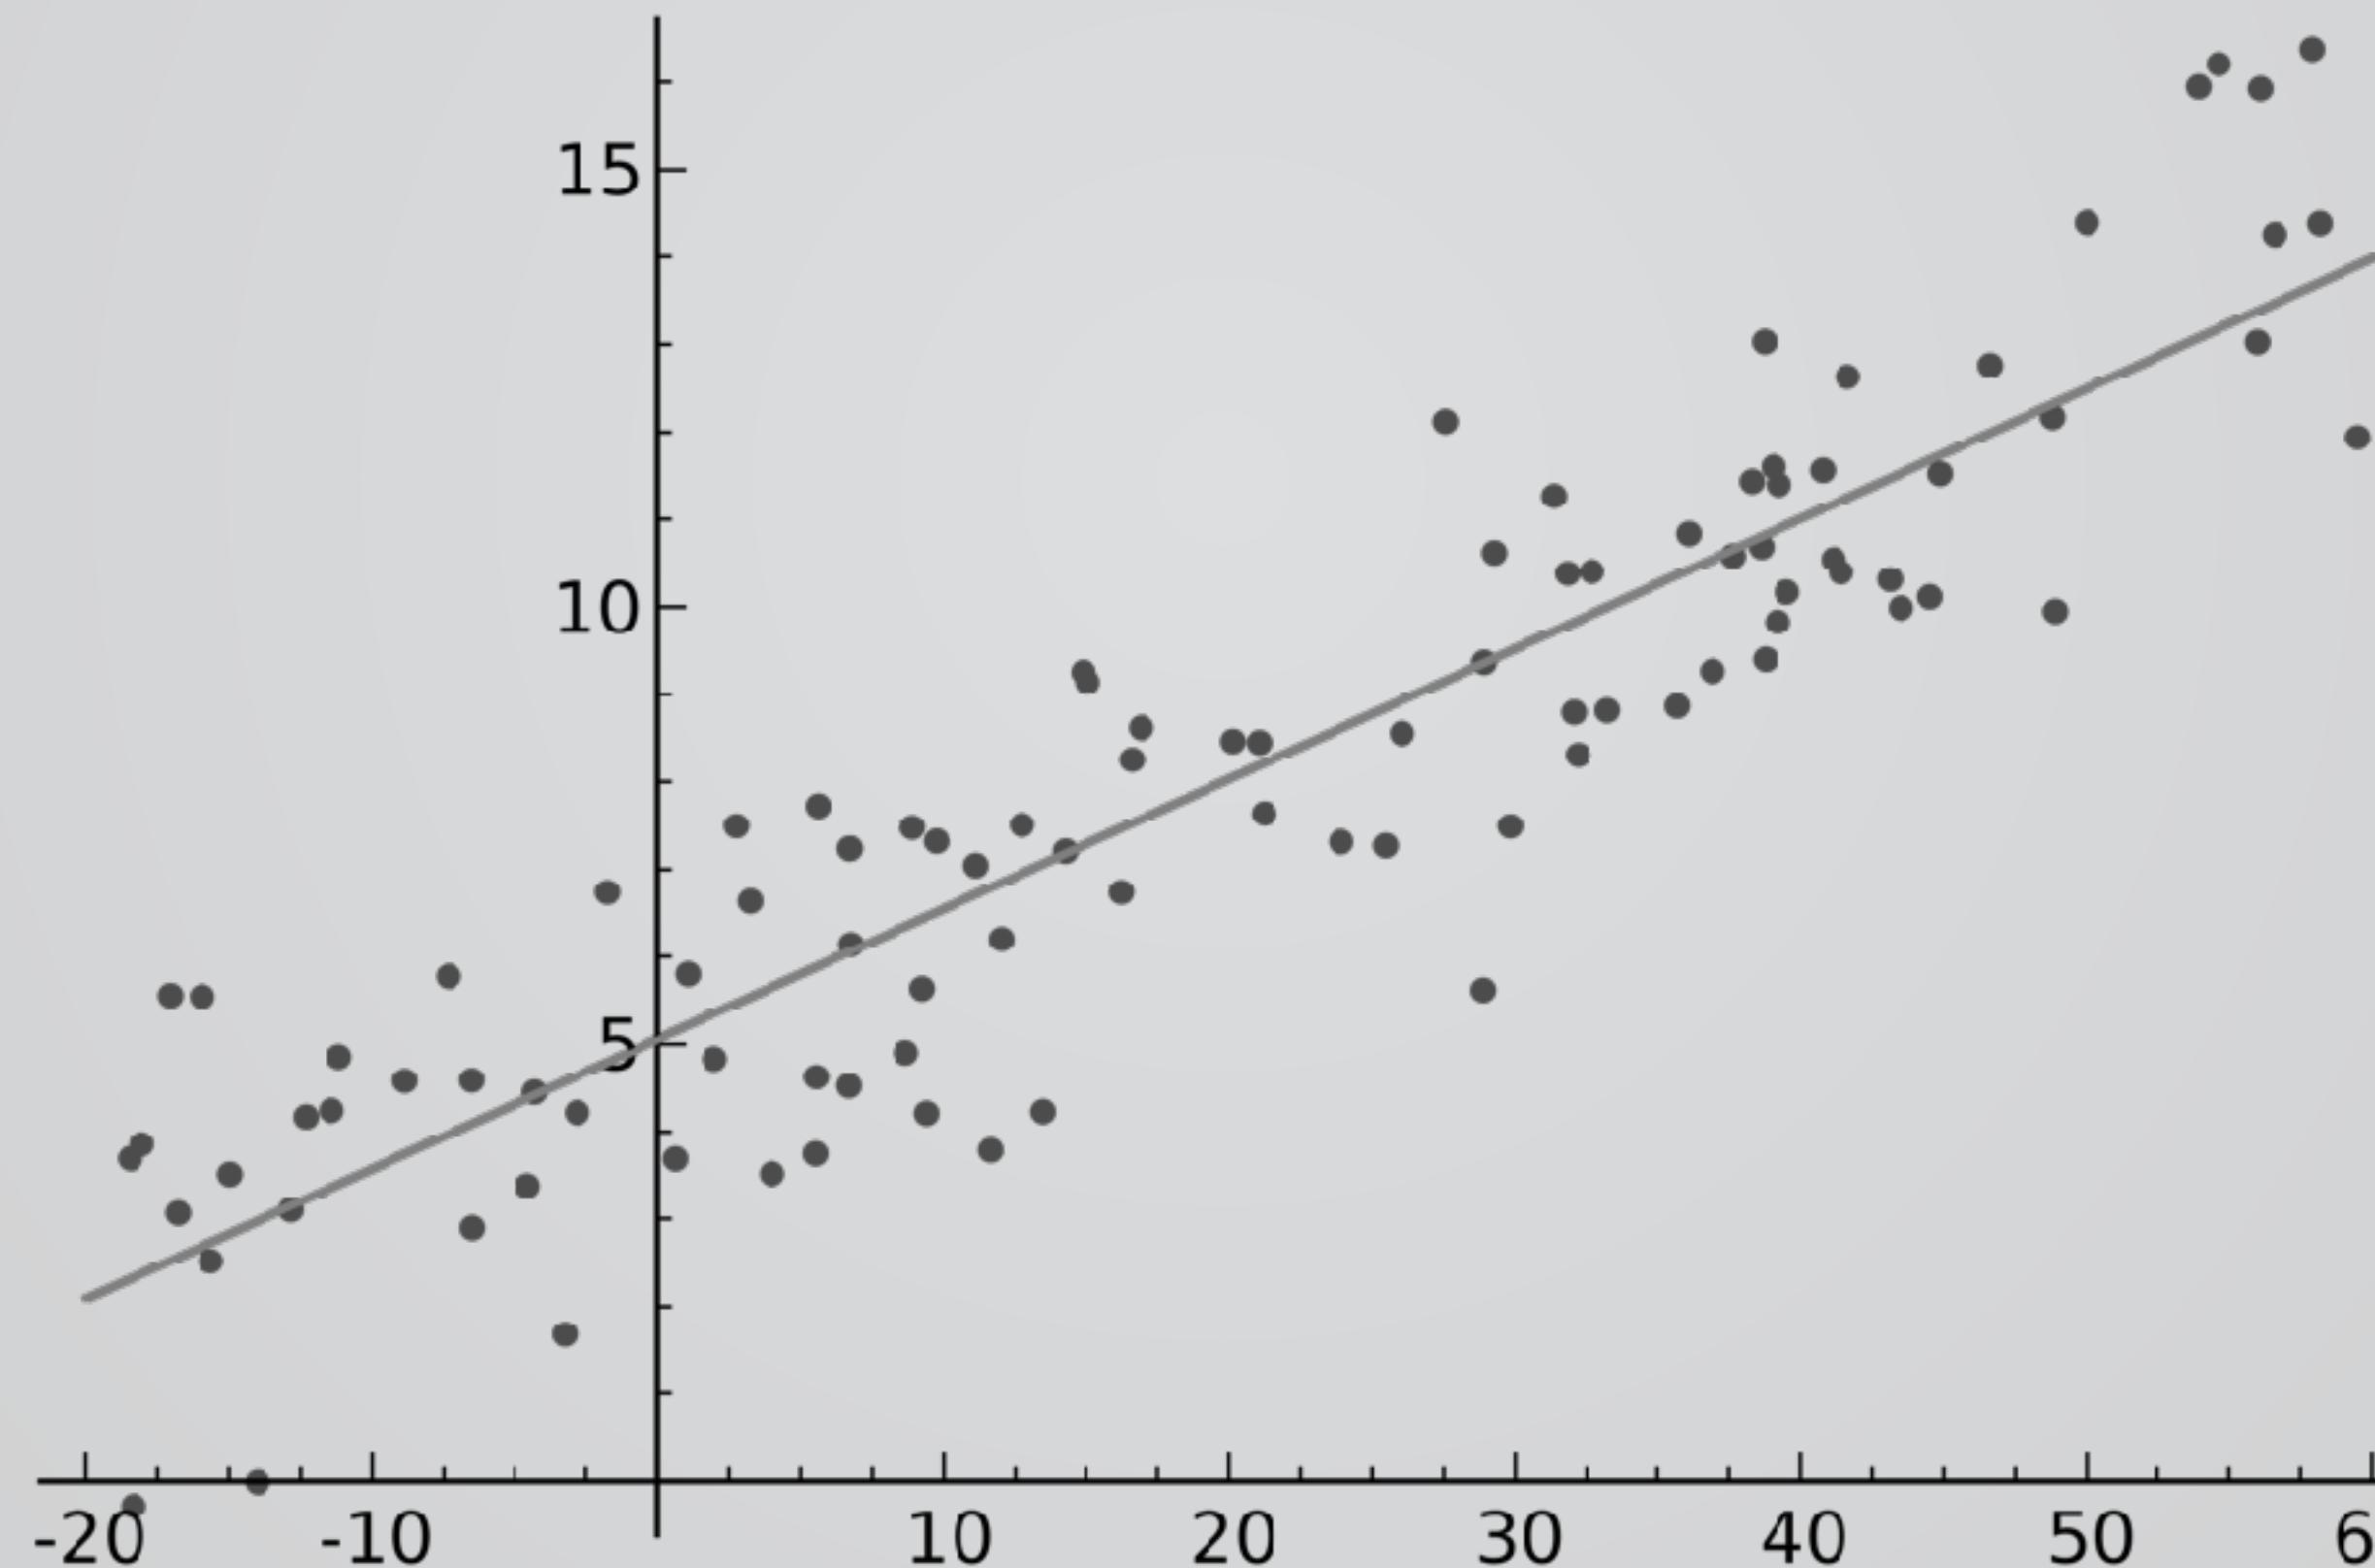
Atoms in the universe  $\approx 10^{80}$

**Exhaustive search is  
not really\* practical.**

\*in this universe

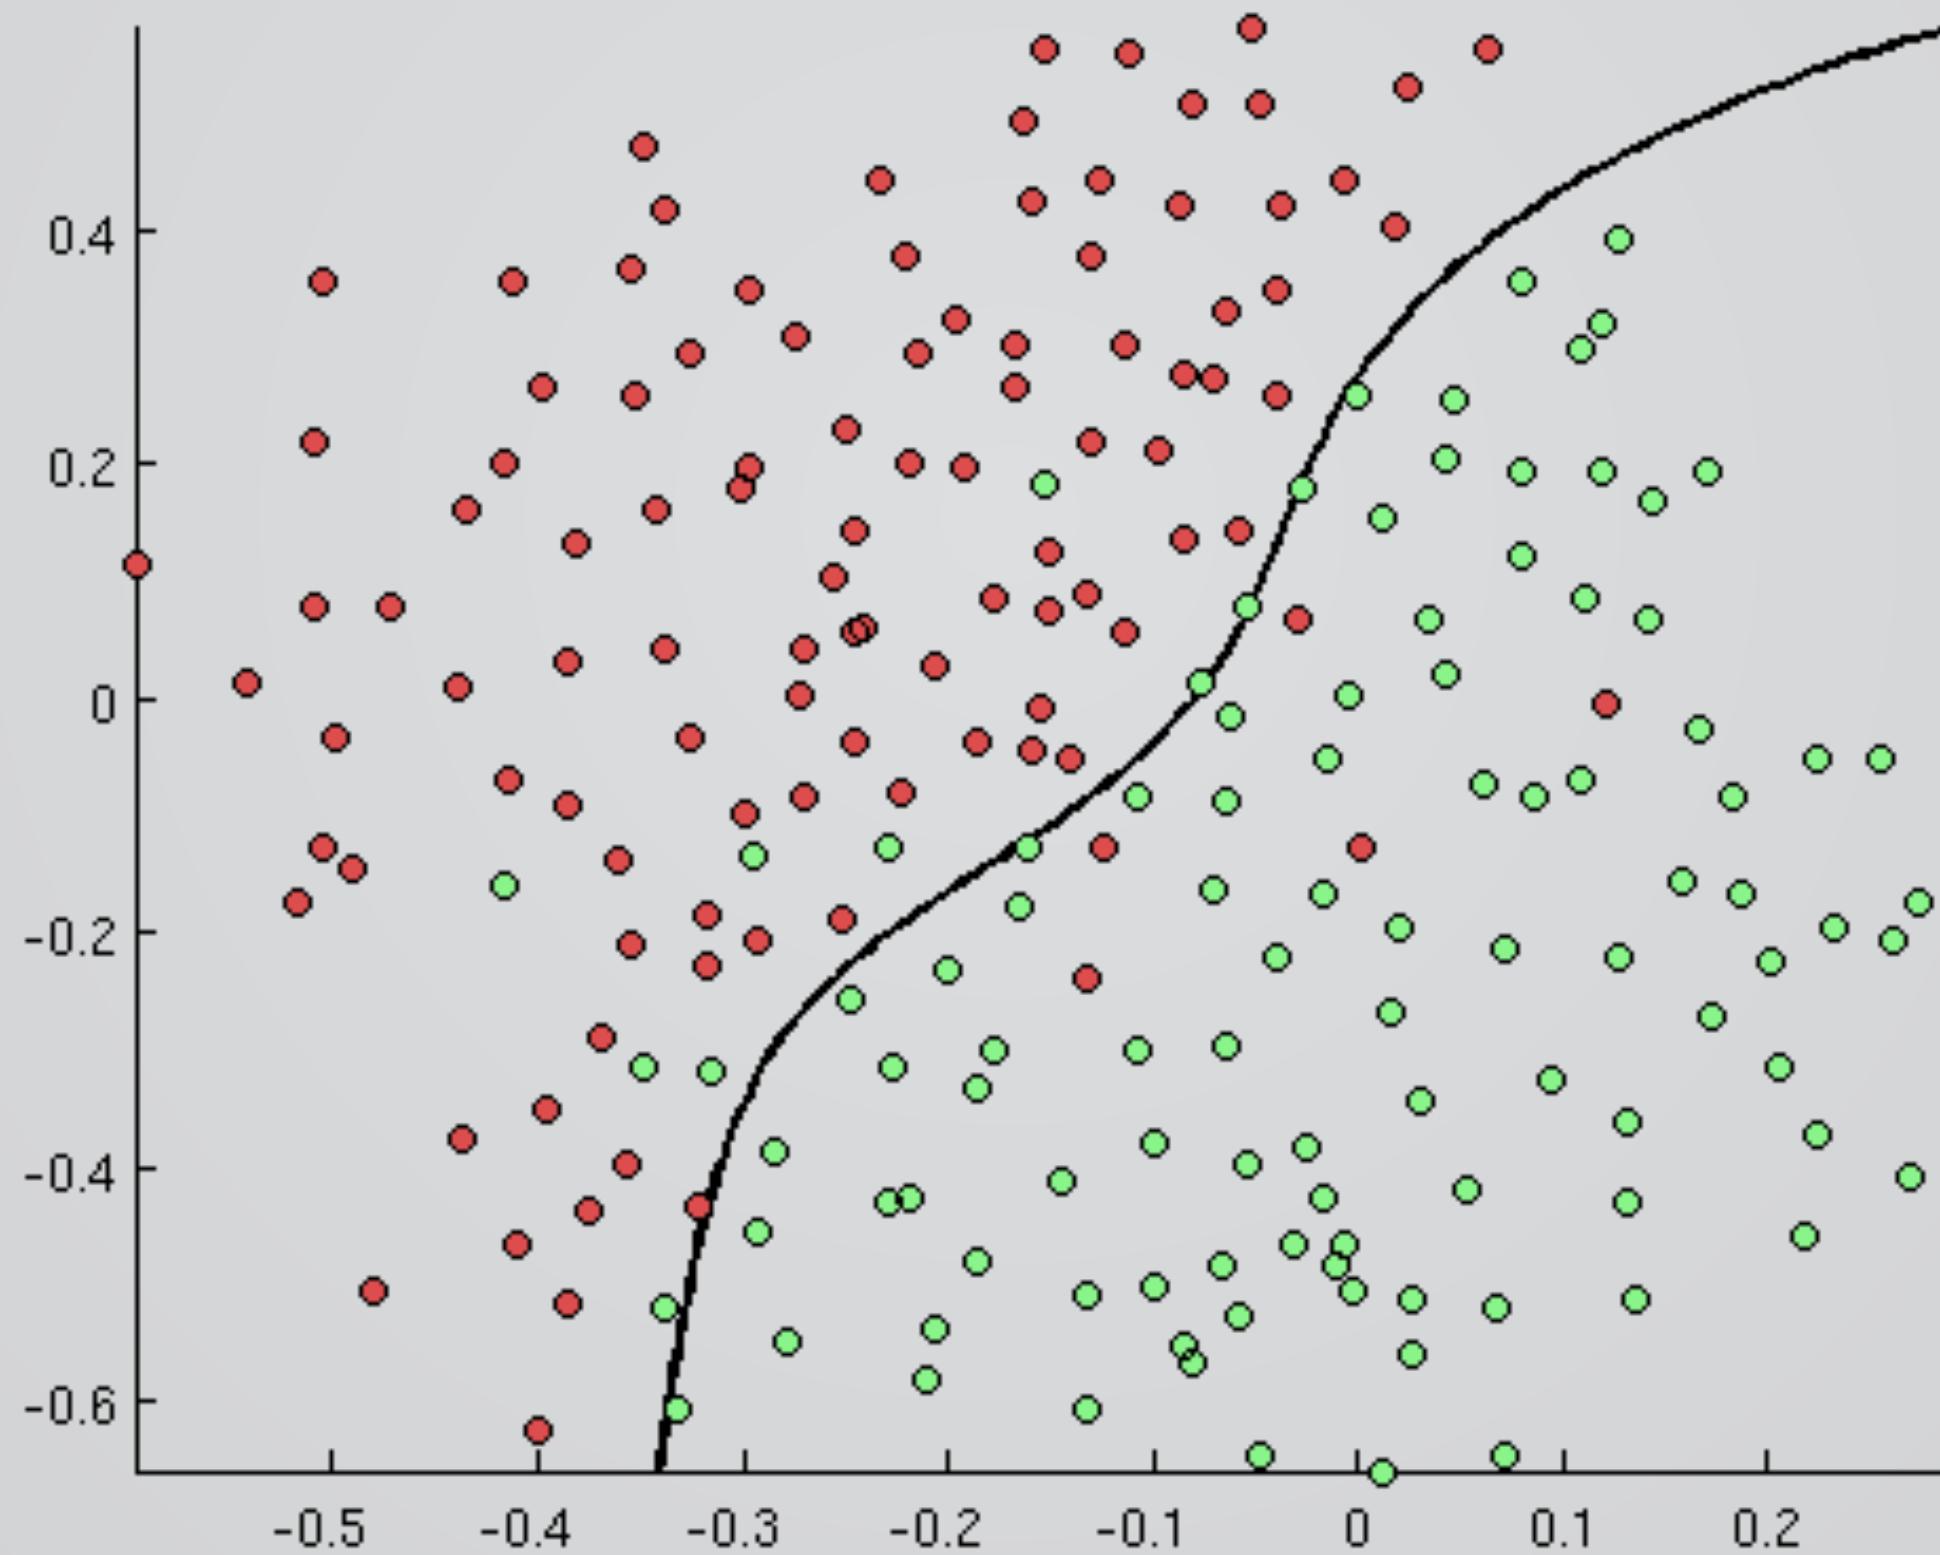
# Machine Learning

Estimate  $y = f(x)$

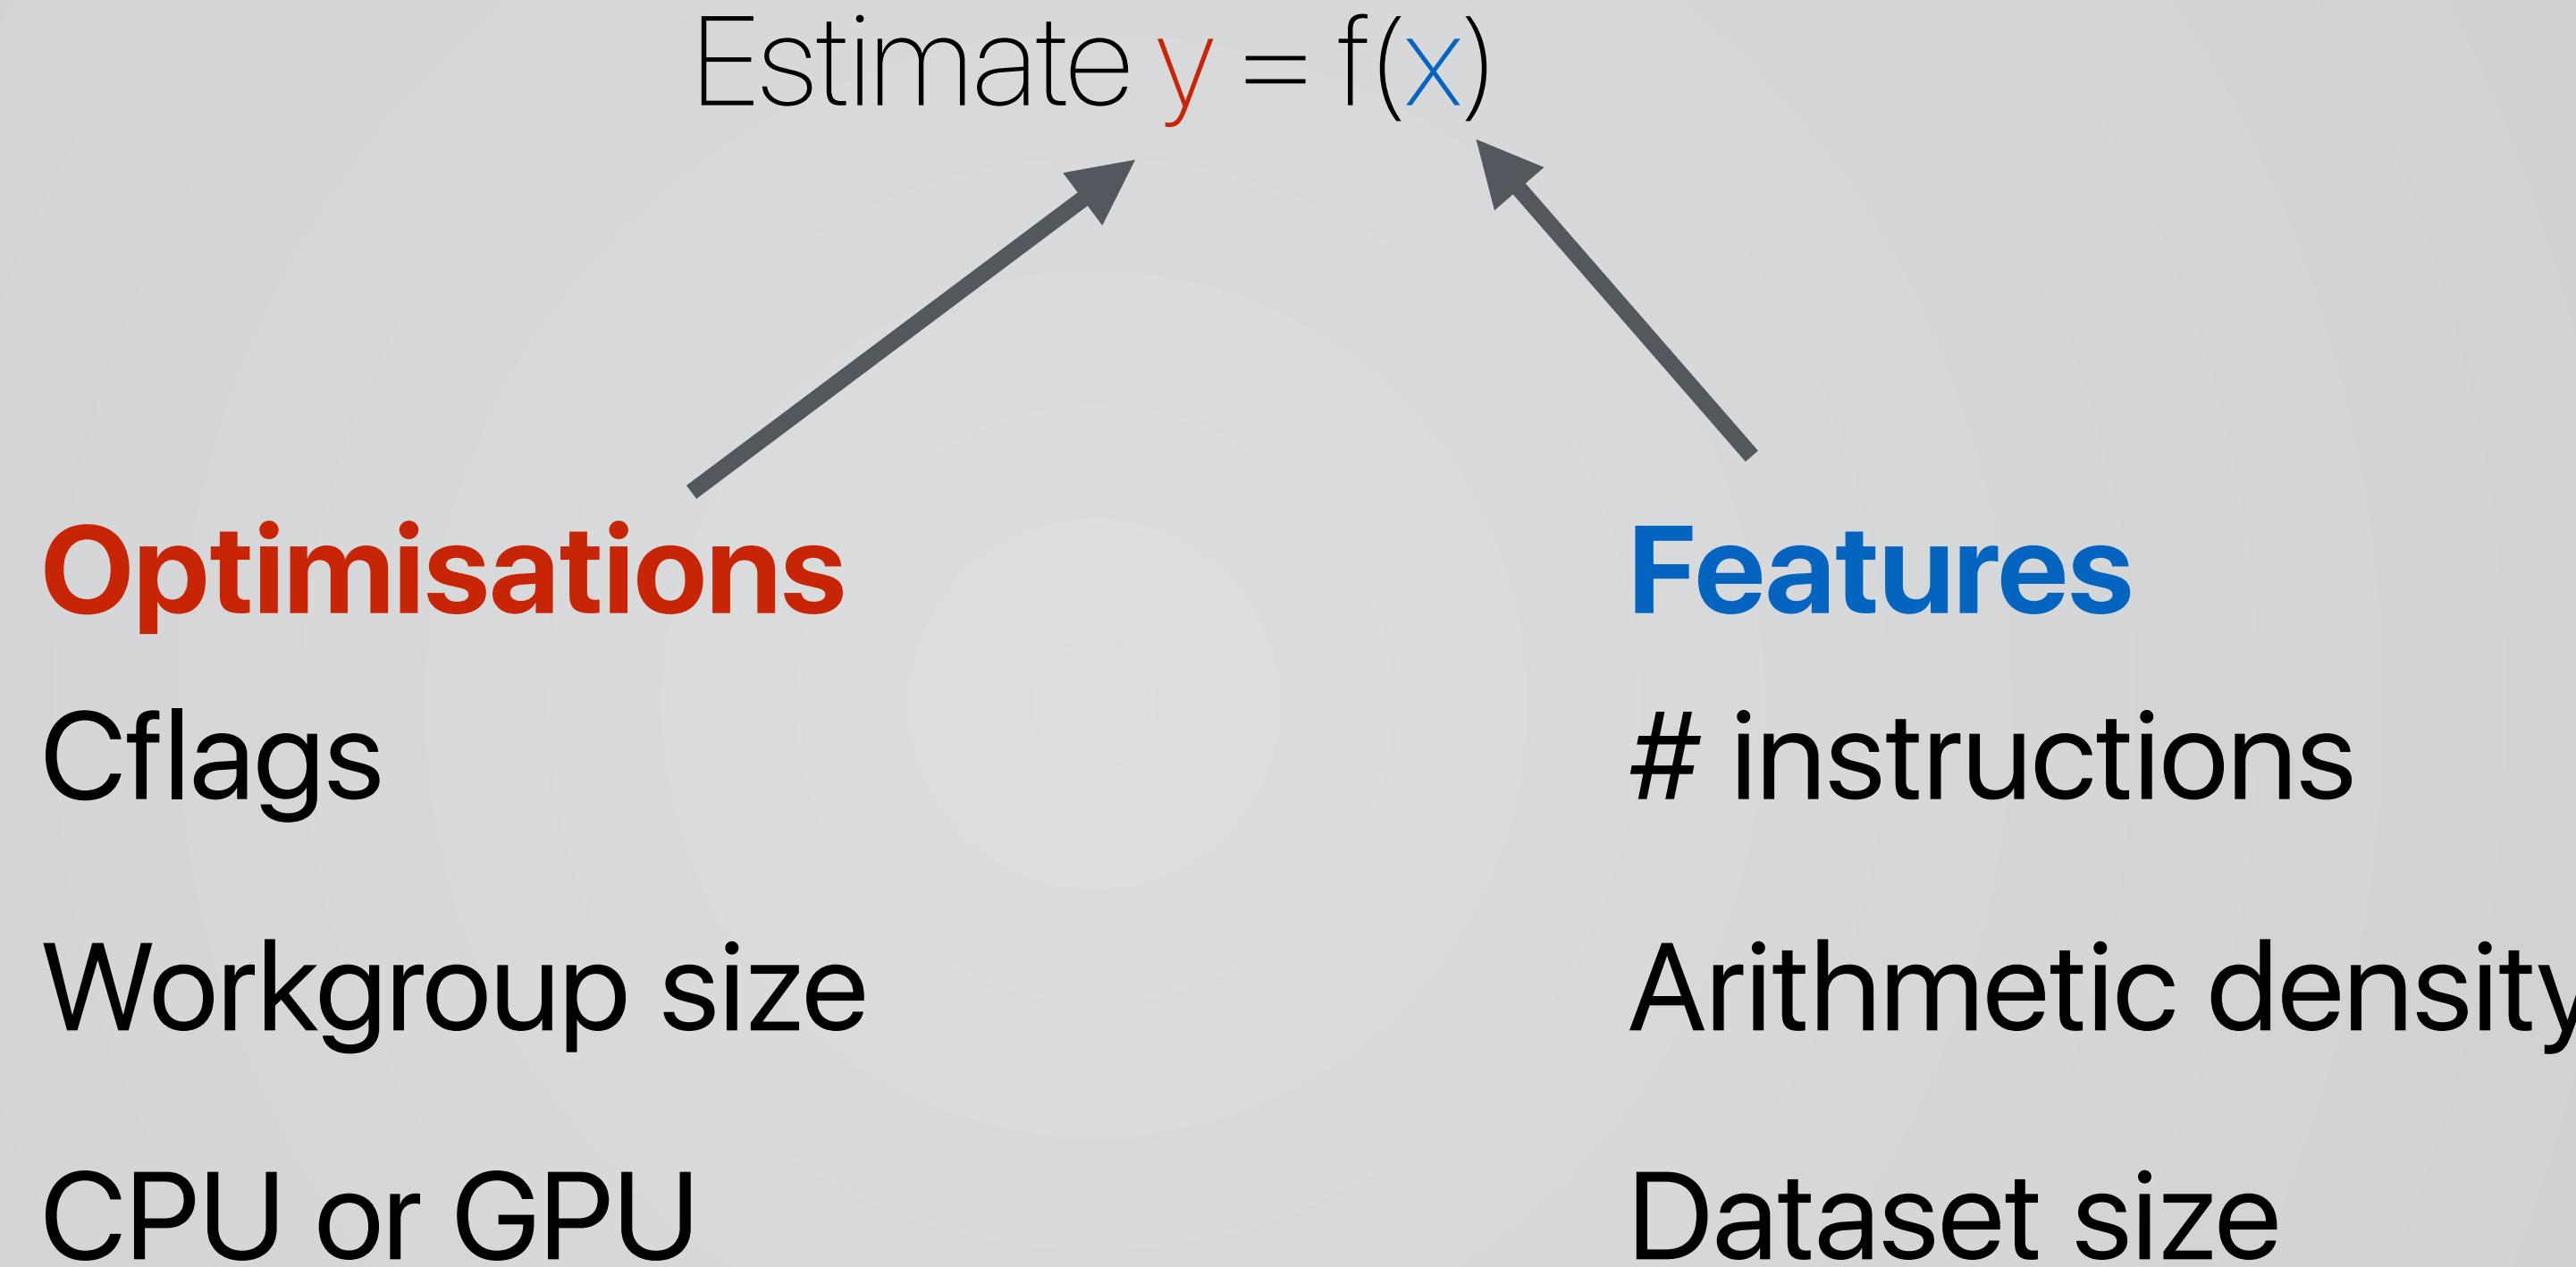


# Machine Learning

Estimate  $y = f(x)$

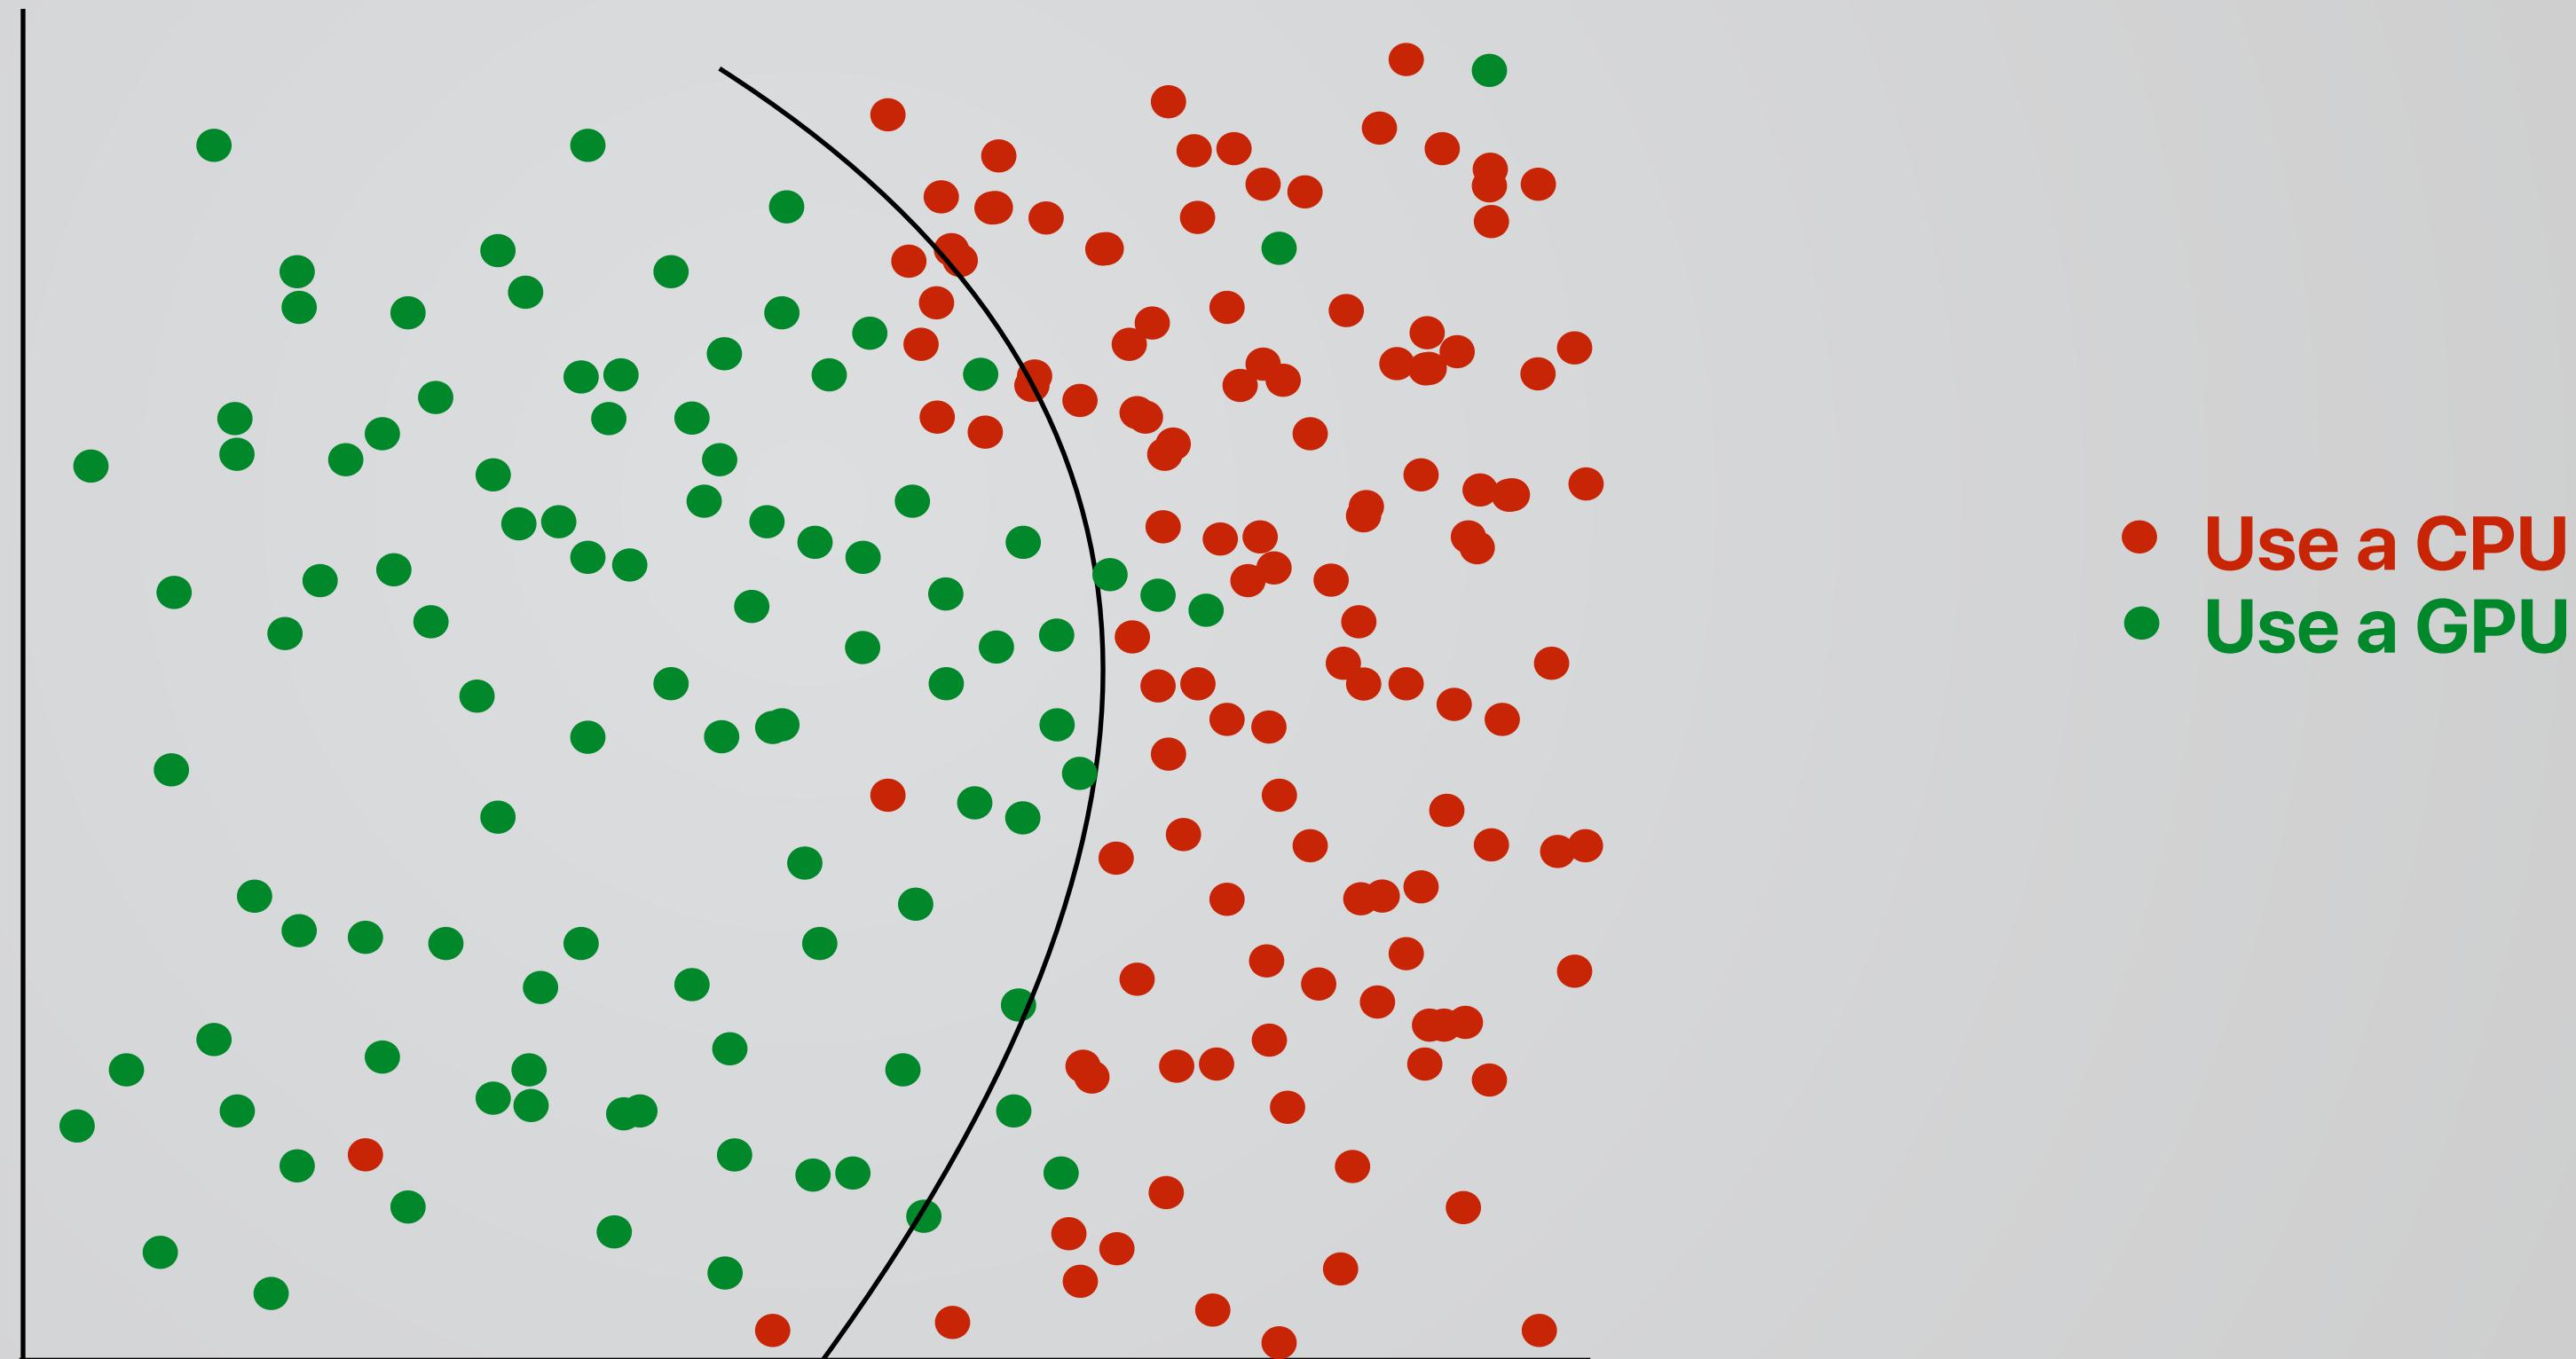


# Machine Learning



# Machine Learning

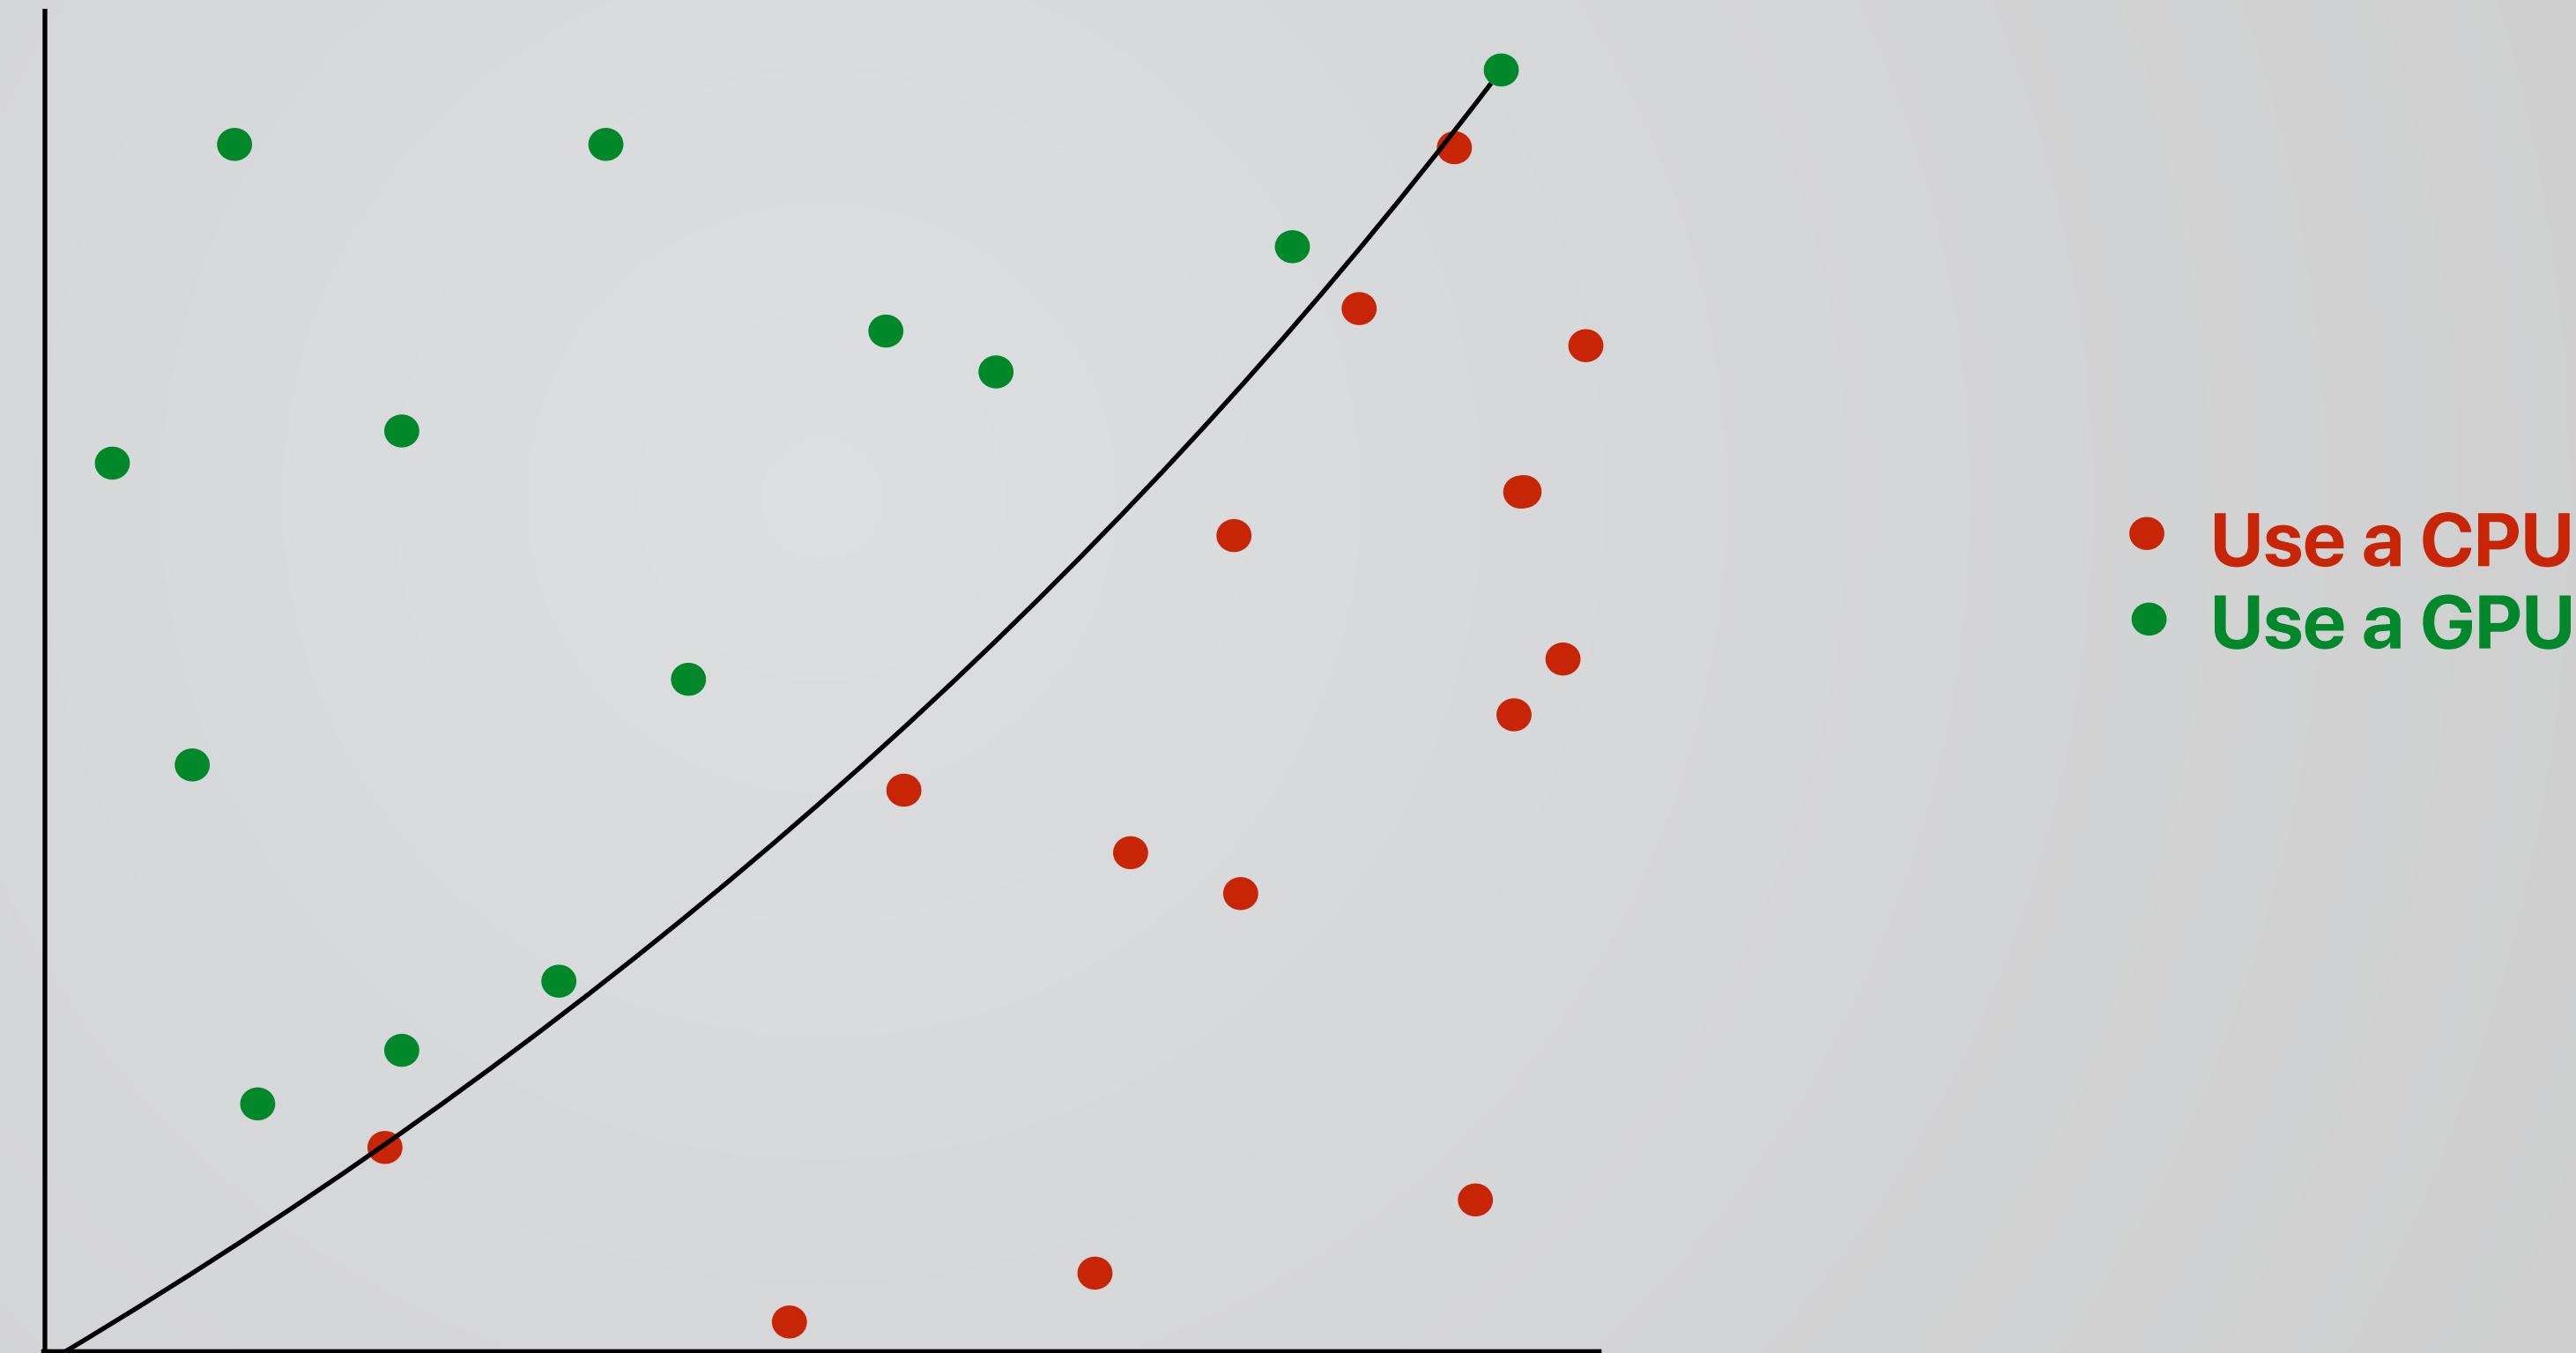
Estimate  $y = f(x)$



The idea.

# Machine Learning

Estimate  $y = f(x)$

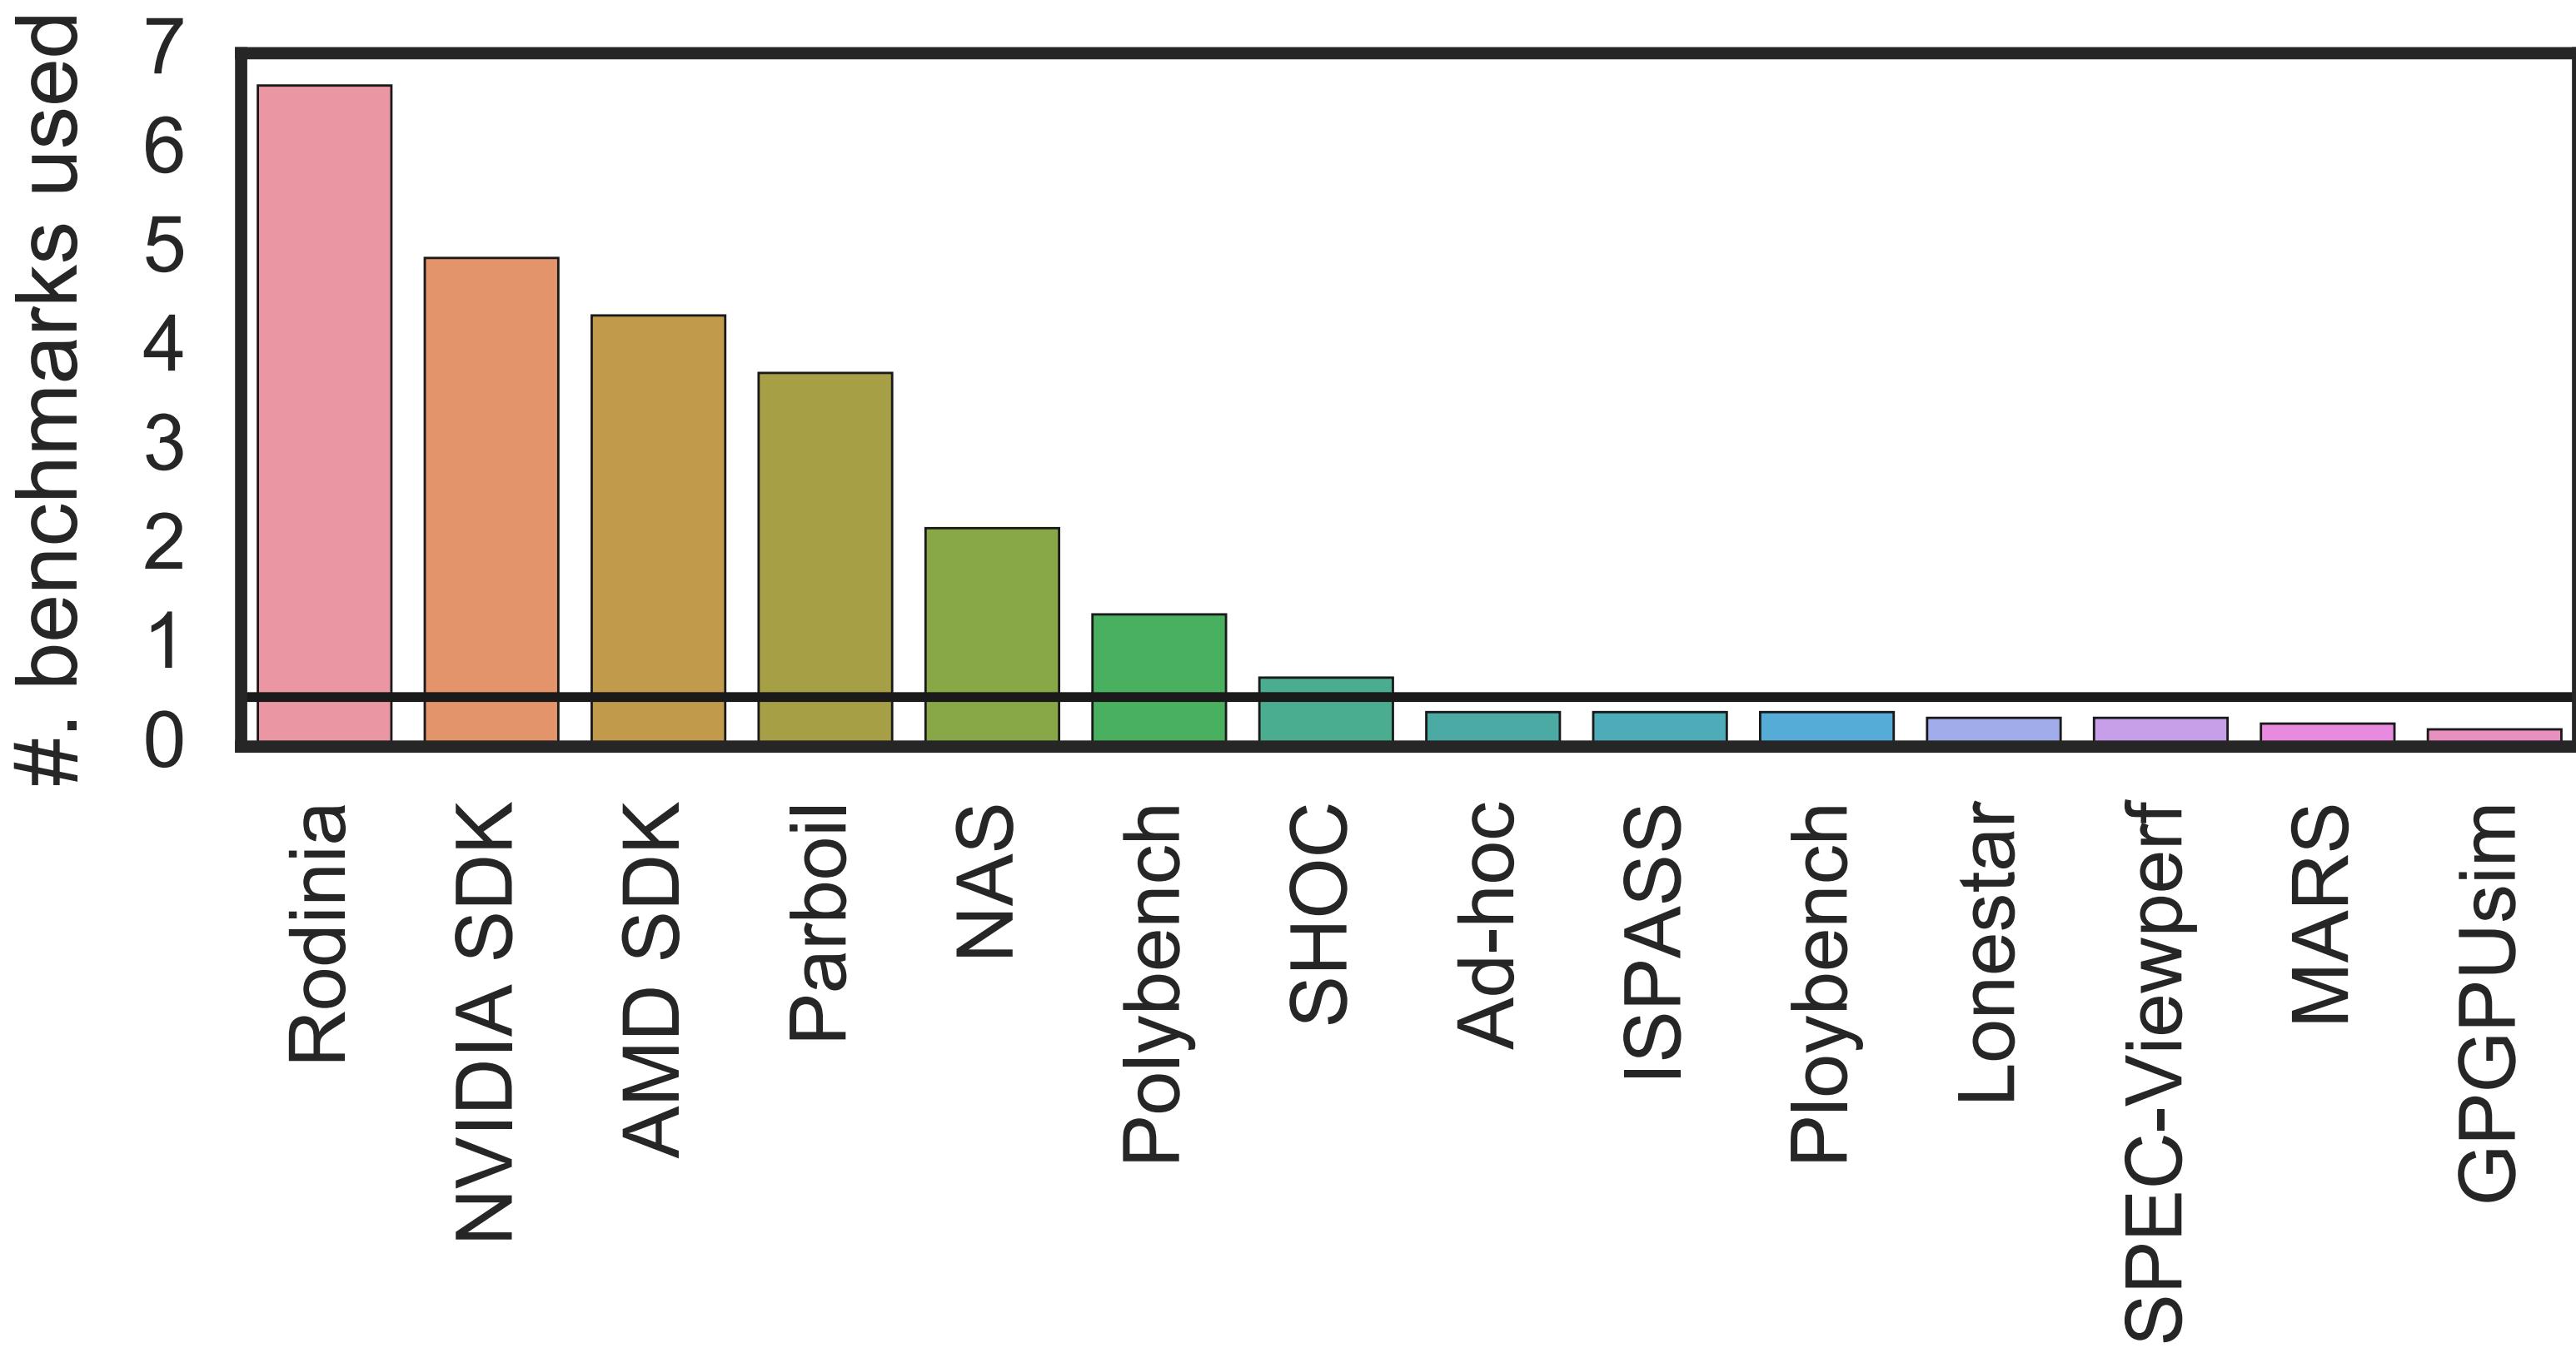


**The reality.**  
(same data, wrong conclusions)

# Synthesizing Benchmarks for Predictive Modeling

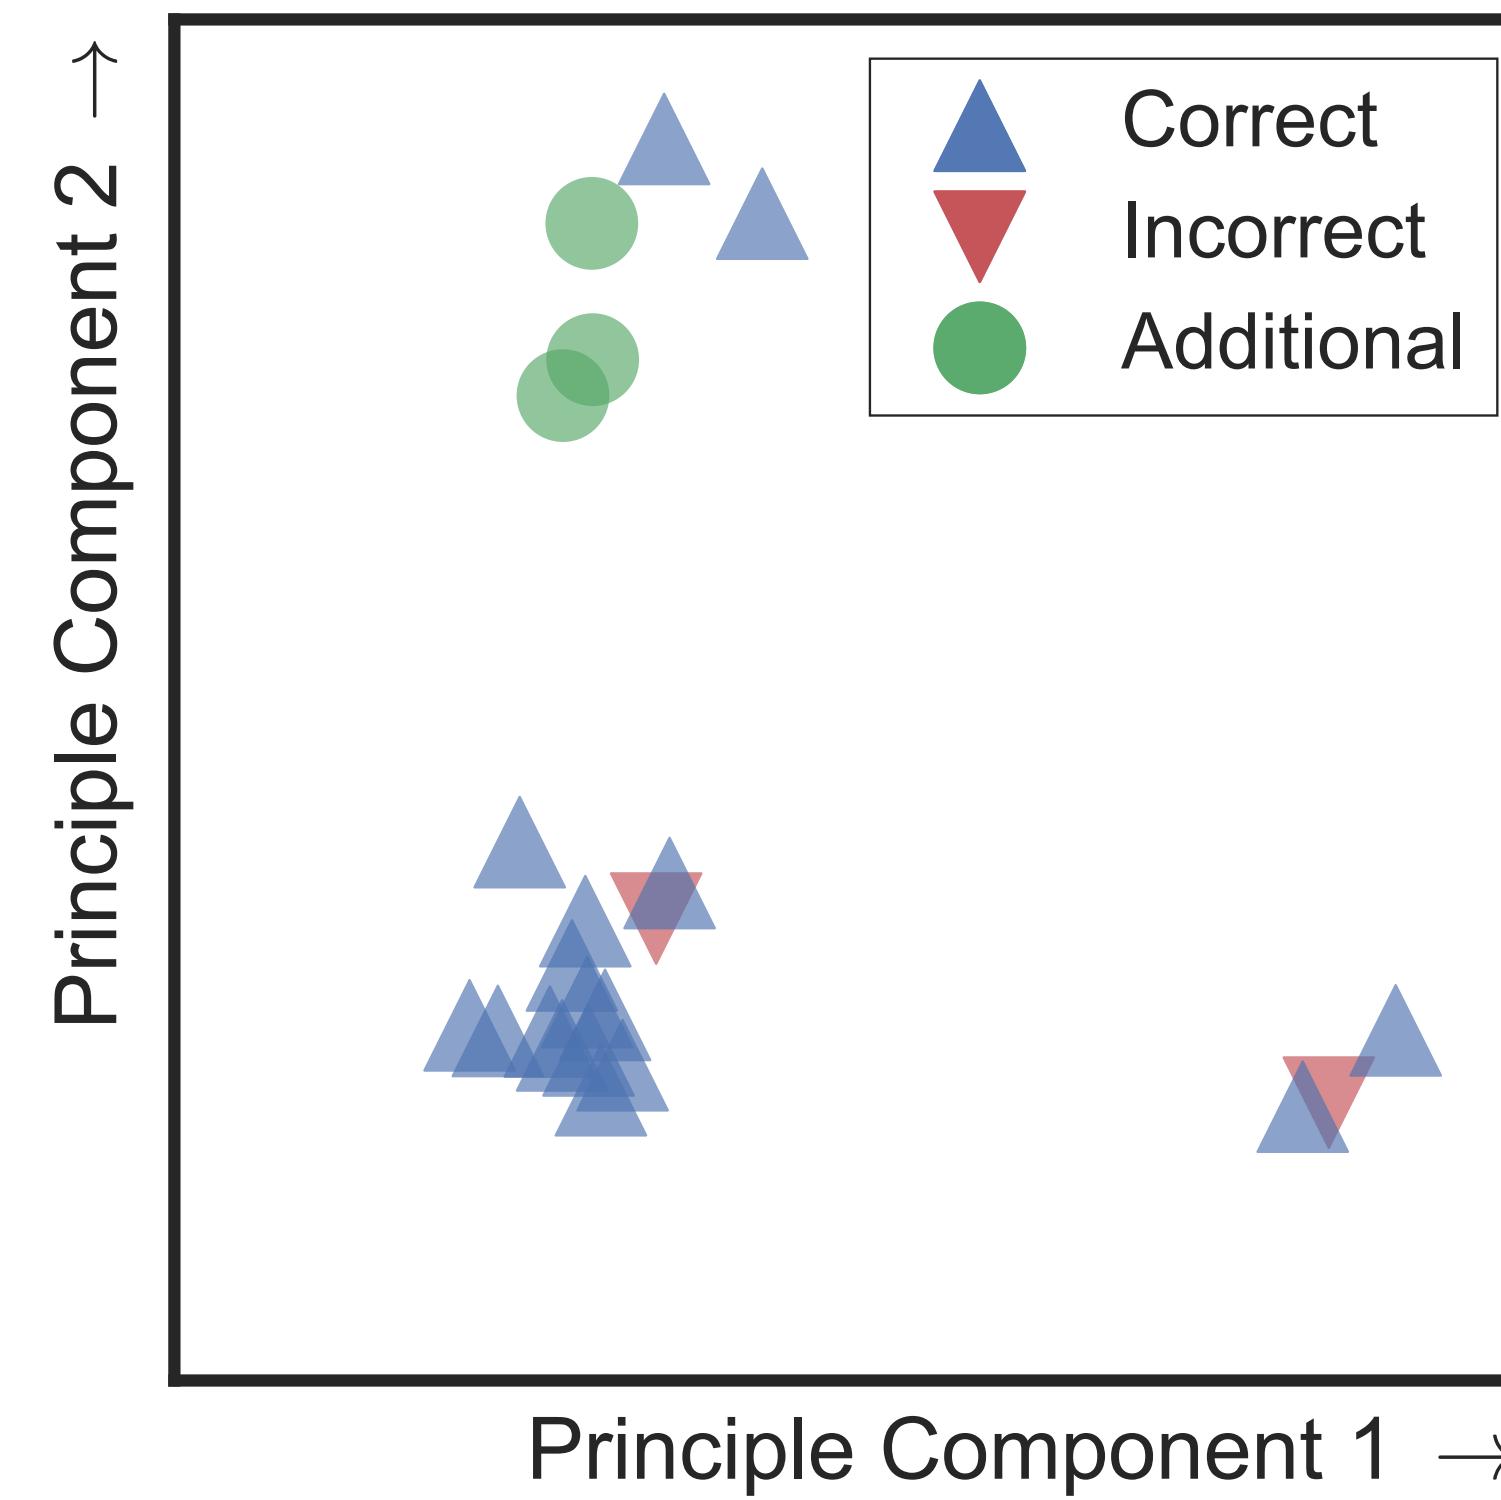
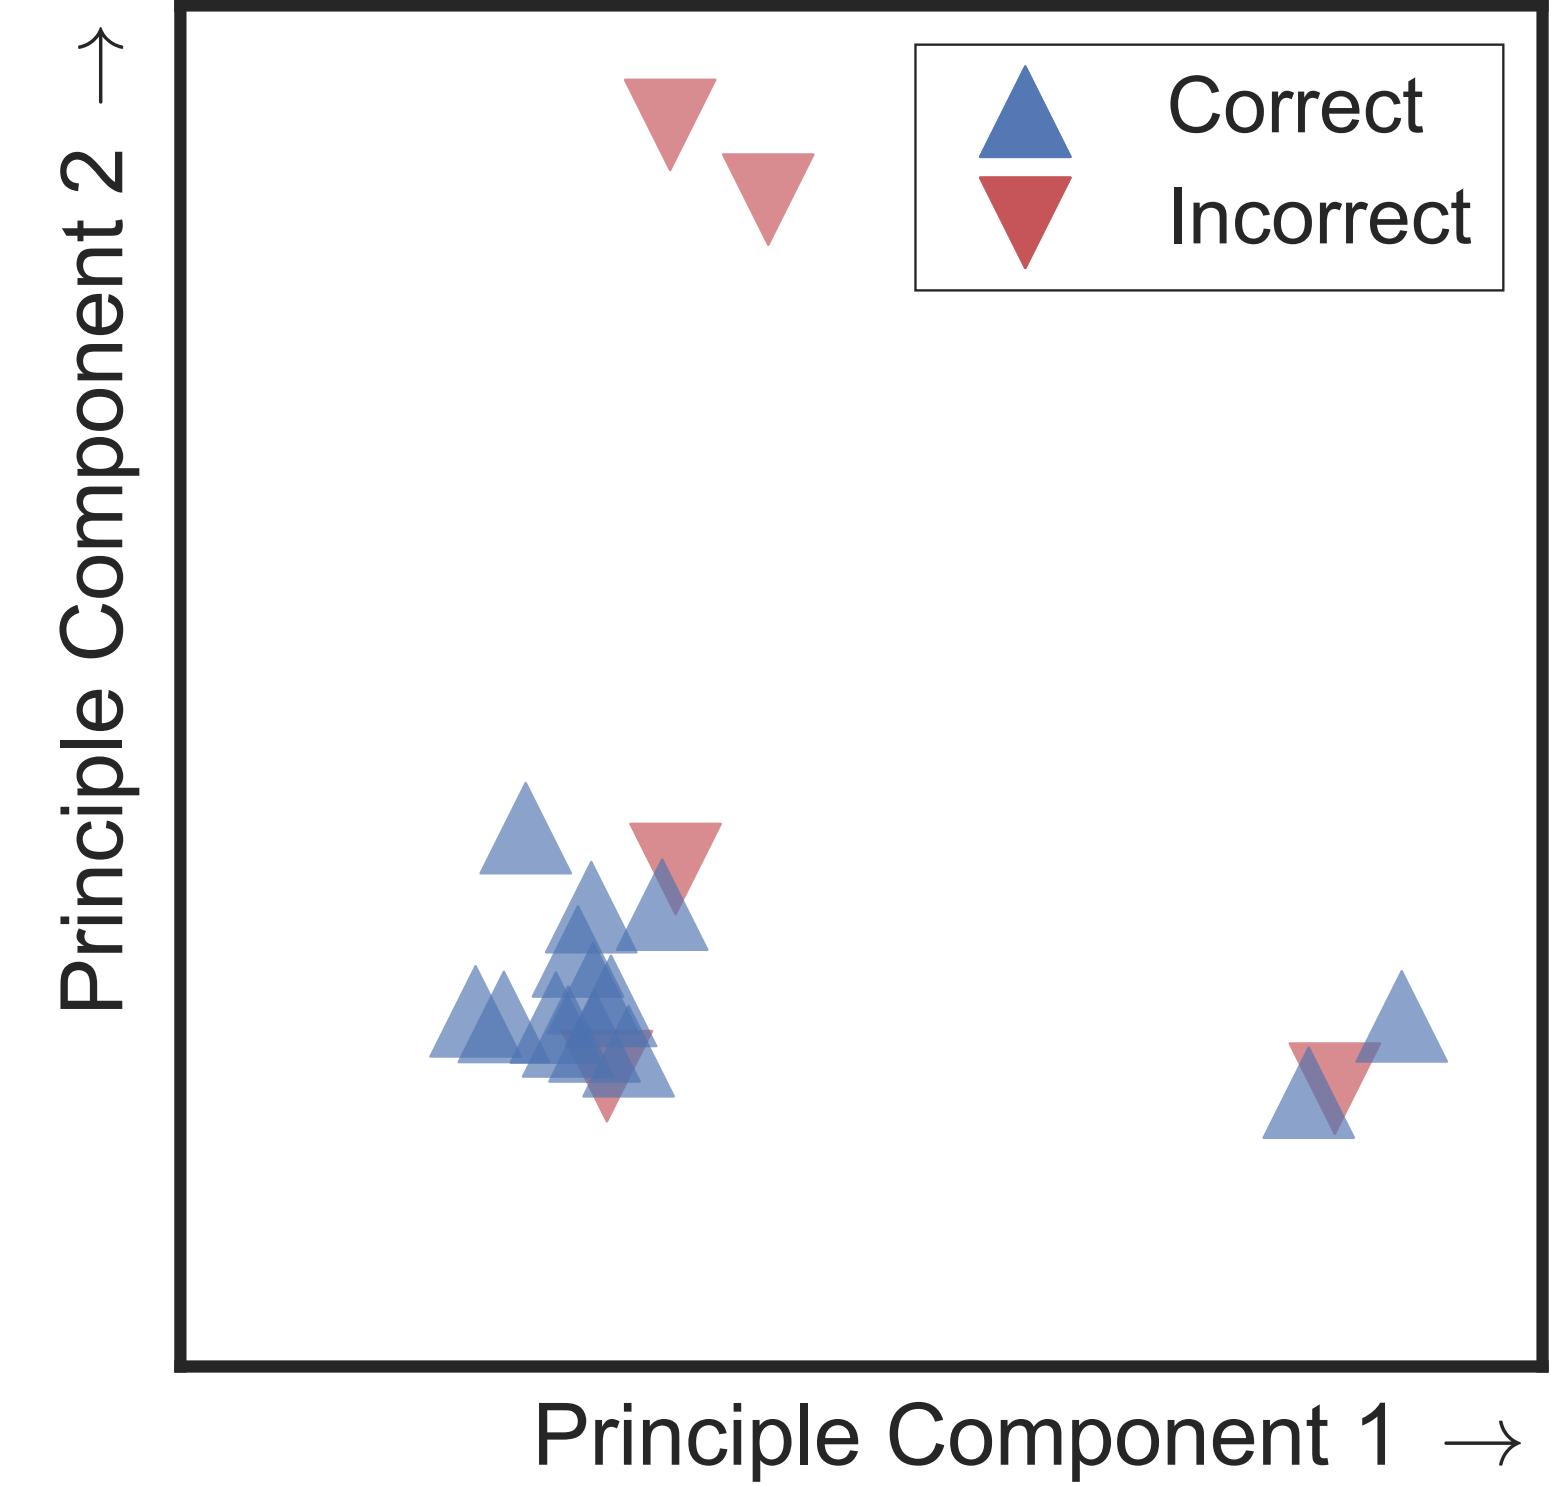
# Why?

There aren't enough benchmarks



# Why?

More benchmarks = better models



# Why?

No adequate solution

into the instance program. Otherwise, nothing is processed or generated by this code.

- The **genloop** construct facilitates repetitive generation. Consider the following example:

```
genloop loopvar:1:5
    ${access(${loopvar})}
end
```

This code produces 5 references to the feature `access`, each with a different value from 1 to 5 passed in as an argument. Note that this does not produce a loop in the instance program, but instead 5 consecutive versions of the code produced when the `access` feature is processed.

- The **geninclude** construct allows information in another file to be used in the current Genesis program. Genesis provides libraries to implement features that may be useful across different Genesis programs of the same target language. For example:

```
geninclude gen_c.glb
```

makes available those features defined in `gen_c.glb`, a library containing features that declare and initialize variables in C programs. For example, `varlistdeclare` is defined in `gen_c.glb`, which initializes C variables in an indicated `varlist`.

## 4. CASE STUDIES

In this section, we present three case studies to show the utility of Genesis. The case studies come from different application domains: auto-tuning of image filtering applications, generation of static program characteristics, and auto-tuning of stencil computations.

### 4.1 Image Filtering

The first case study deals with the generation of image filtering applications for training in performance auto-tuning on GPUs. These applications typically have two perfectly nested loops that sweep over two-dimensional images. Each element of an output image is computed as a function of a subset of the pixels in an input image. Specific image filtering applications differ in the subset and the function used to compute the output.

This case study focuses on memory performance, which is affected by the number and pattern of image accesses and the pixel computations in the loop nest. Thus, we model the body of the loop nest as one or more *read epochs* followed by a *write epoch*, where an *epoch* is a sequence of computations followed by a memory access. We wish to generate a number of such programs where the number of read epochs,

```
1 begin genesis
2
3     geninclude gen_c.glb
4
5     global
6         distribution epochdist = {1:10}
7         distribution numvardist = {8;16;32}
8         distribution compdist = {1:20}
9         distribution coefdist = {0:7}
10        distribution offsdist = {0:15}
11    end
12
13    program
14        value numepochs sample epochdist
15        value numvars sample numvardist
16        varlist temp[${numvars}]
17    end
18
19    feature computation
20        variable dest,src1,src2,src3 from temp
21        ${dest} = ${src1} * ${src2} + ${src3};
22    end
23
24    feature read_access
25        variable dest from temp
26        value coef1,coef2 sample coefdist
27        value offs sample offsdist
28        ${dest} = arr_in[${coef1}*it0 + ${coef2}*
29                    it1 +${offs }];
30    end
31
32    feature write_access
33        variable src from temp
34        value offs sample offsdist
35        arr_out[inner_tc*it0 + it1 +${offs }] = ${
36            src};
37
38    feature epoch (epoch type)
39        value numcomps sample compdist
40        genloop i:1:${numcomps}
41            ${computation}
42        end
43        genif ${epoch_type} eq "read"
44            ${read_access}
45        genelse
46            ${write_access}
47        end
48    end
49
50    feature epochs
51        genloop i:1:${numepochs}
52            ${epoch("read")}
53        end
54        ${epoch("write")}
55    end
56
57    generate 1000
end genesis
```

Type: Conference Proceedings

## Genesis: A Language for Generating Synthetic Training Programs for Machine Learning

Authors: A. Chiu, J. Garvey, T. Abdelrahman



[View research catalog entry for this paper](#)

Proc. Title: **CF**

Year: 2015

Pages: 8

### Abstract:

We describe Genesis, a language for the generation of synthetic programs for use in machine learning-based performance auto-tuning. The language allows users to annotate a template program to customize its code using statistical distributions and to generate program instances based on those distributions. This effectively allows users to generate training programs whose characteristics or features vary in a statistically controlled fashion. We describe the language constructs, a prototype preprocessor for the language, and three case studies that show the ability of Genesis to express a range of training programs in different domains. We evaluate the preprocessor's performance and the statistical quality of the sa...

### Tags:

### Author Keywords:

macro languages; synthetic program generation

### Citation Key:

Chiu2015

### City:

### Editors:

### Publisher:

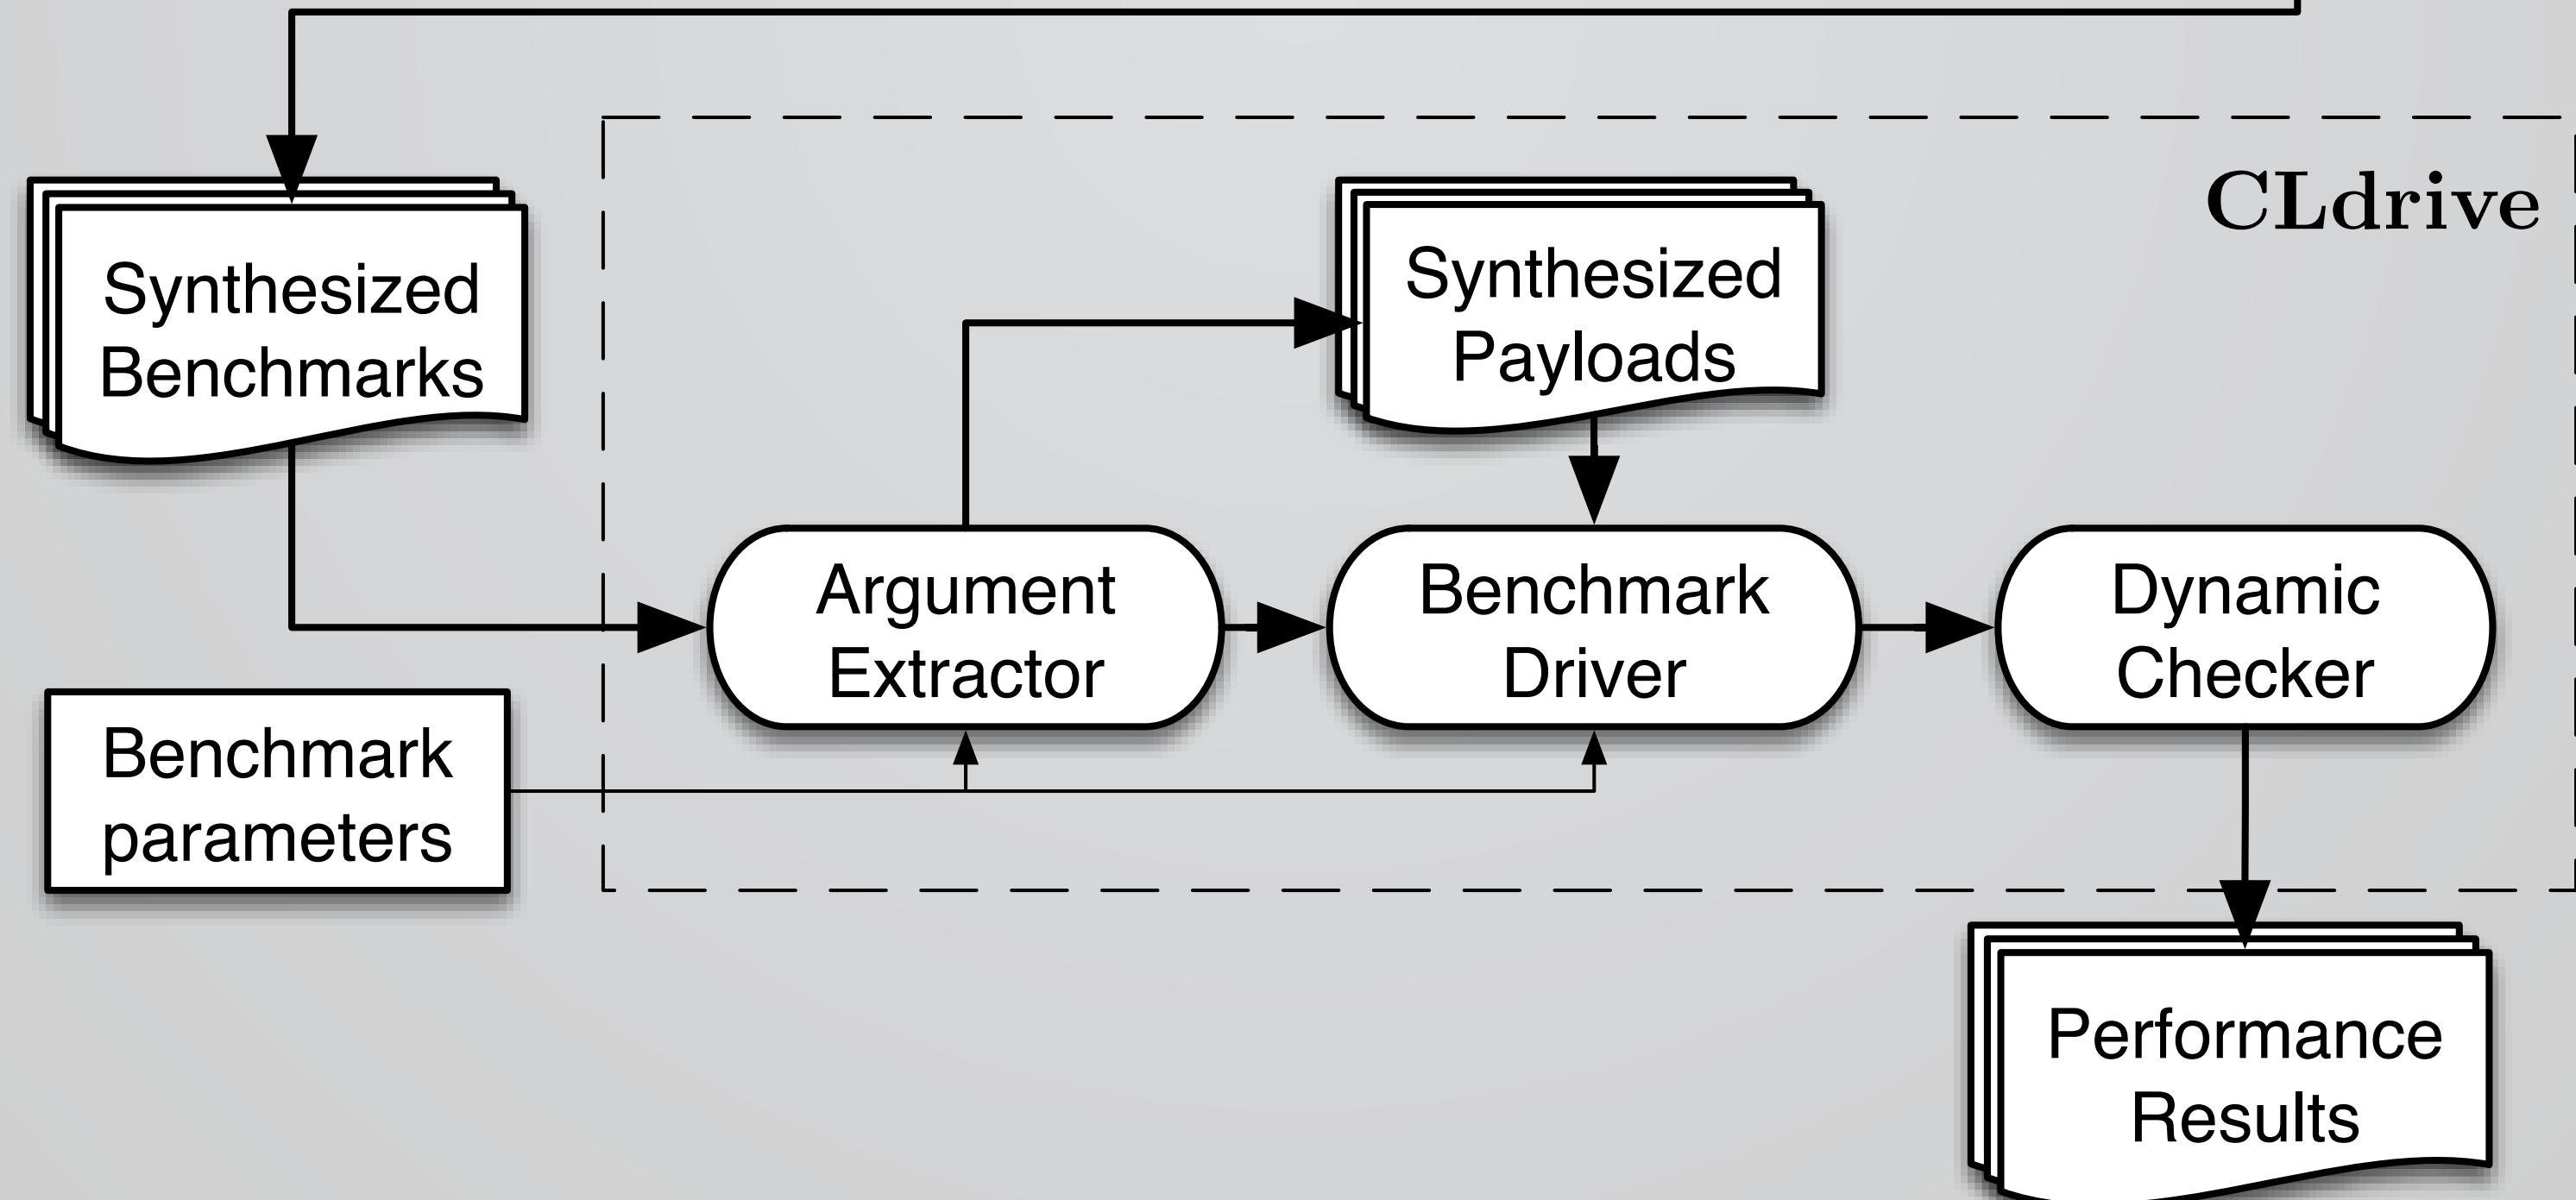
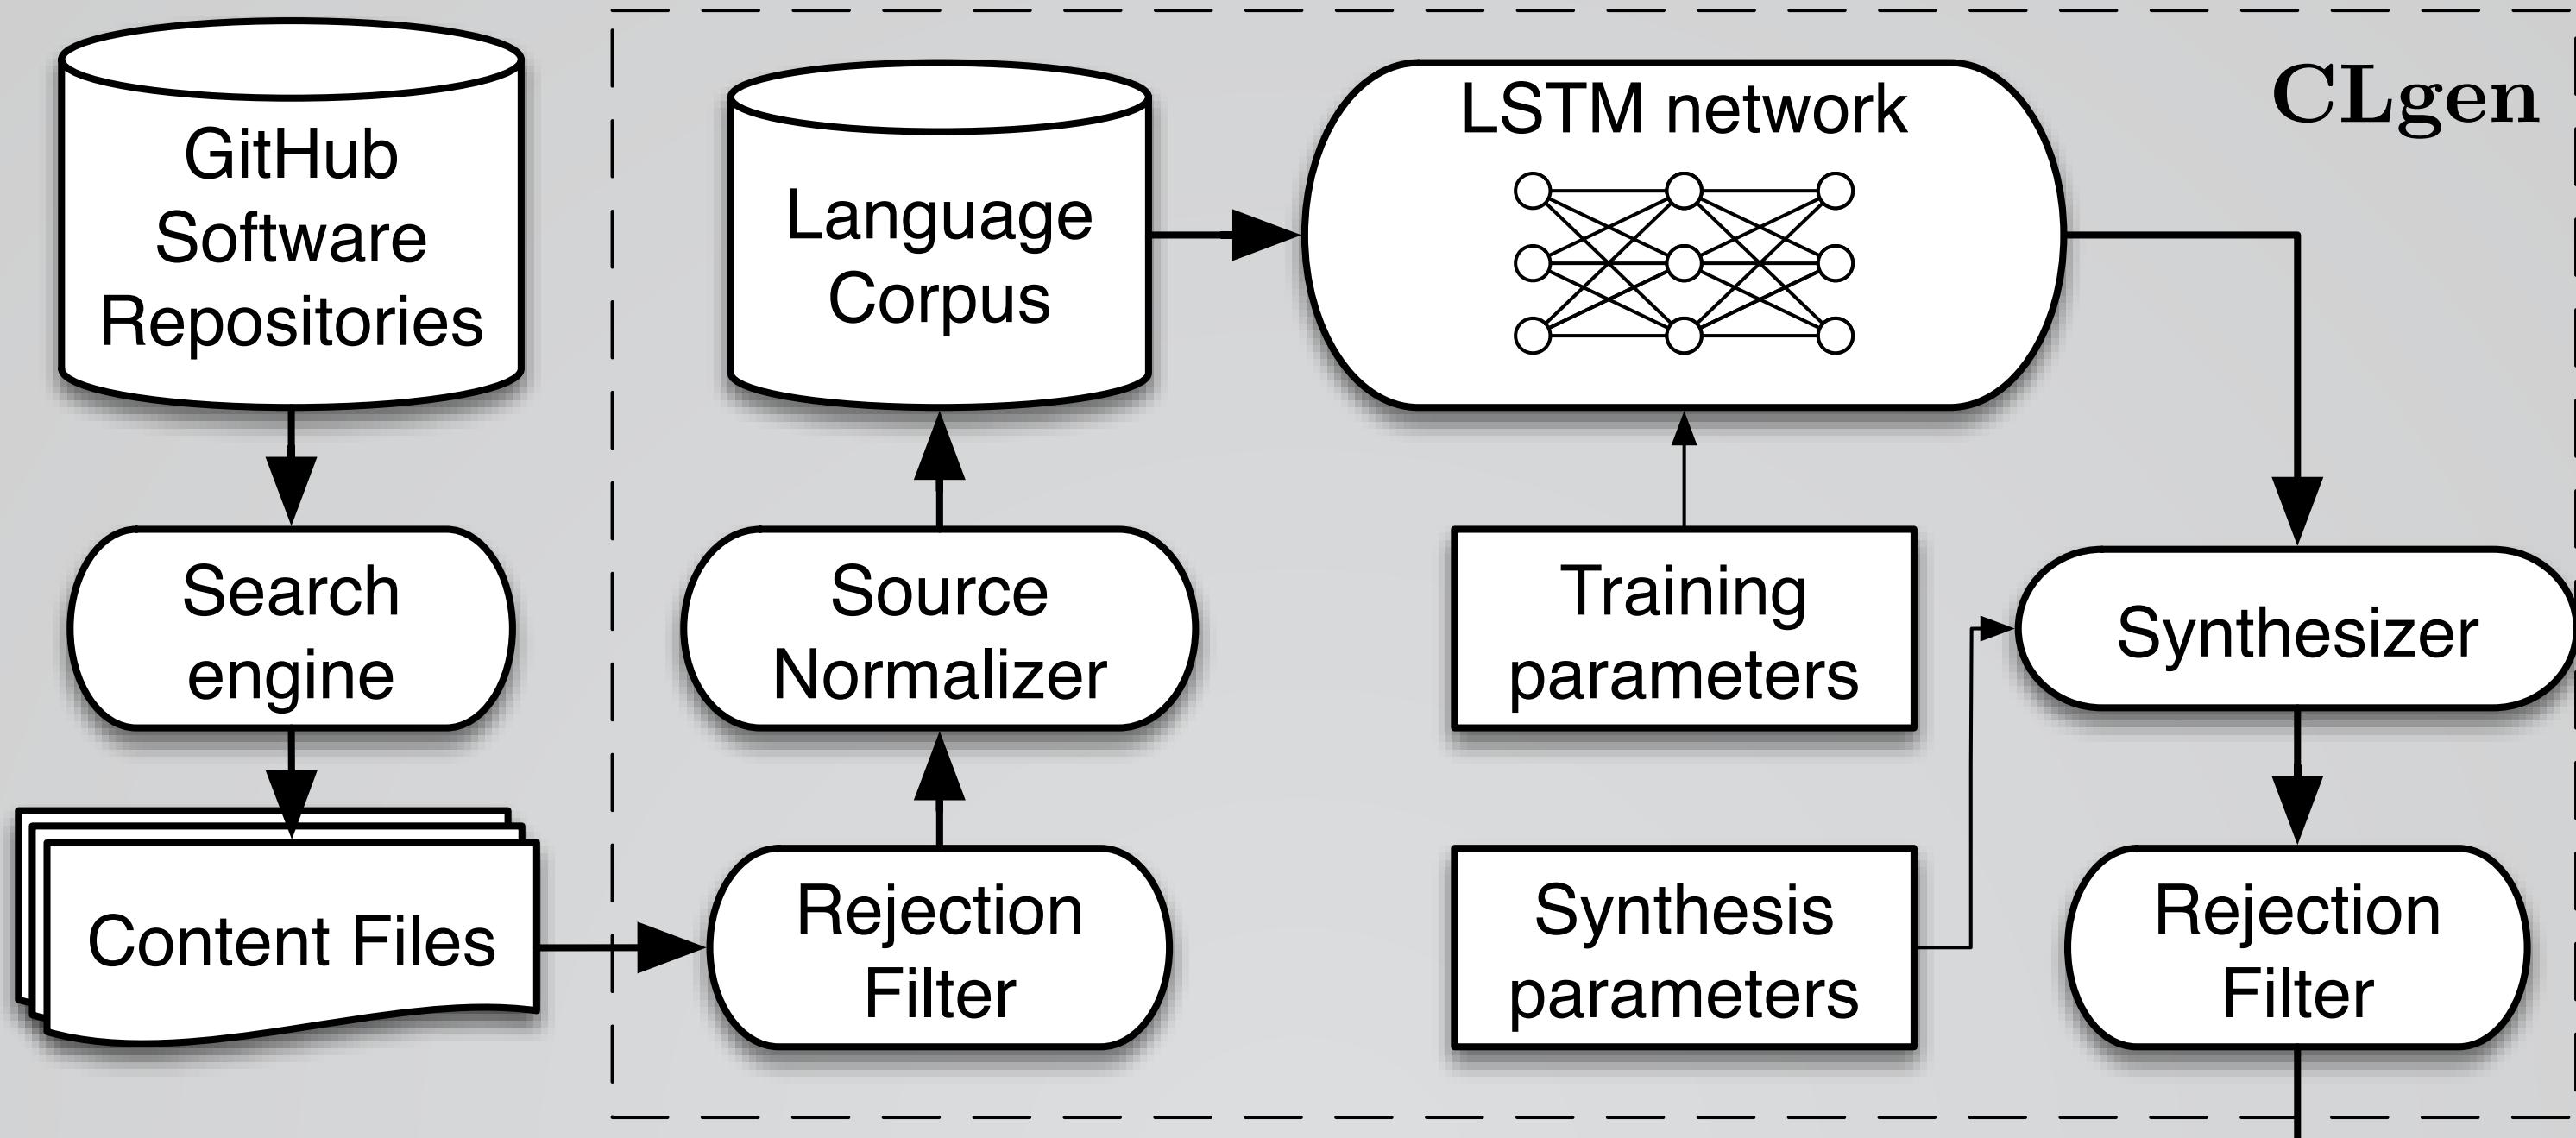
ACM

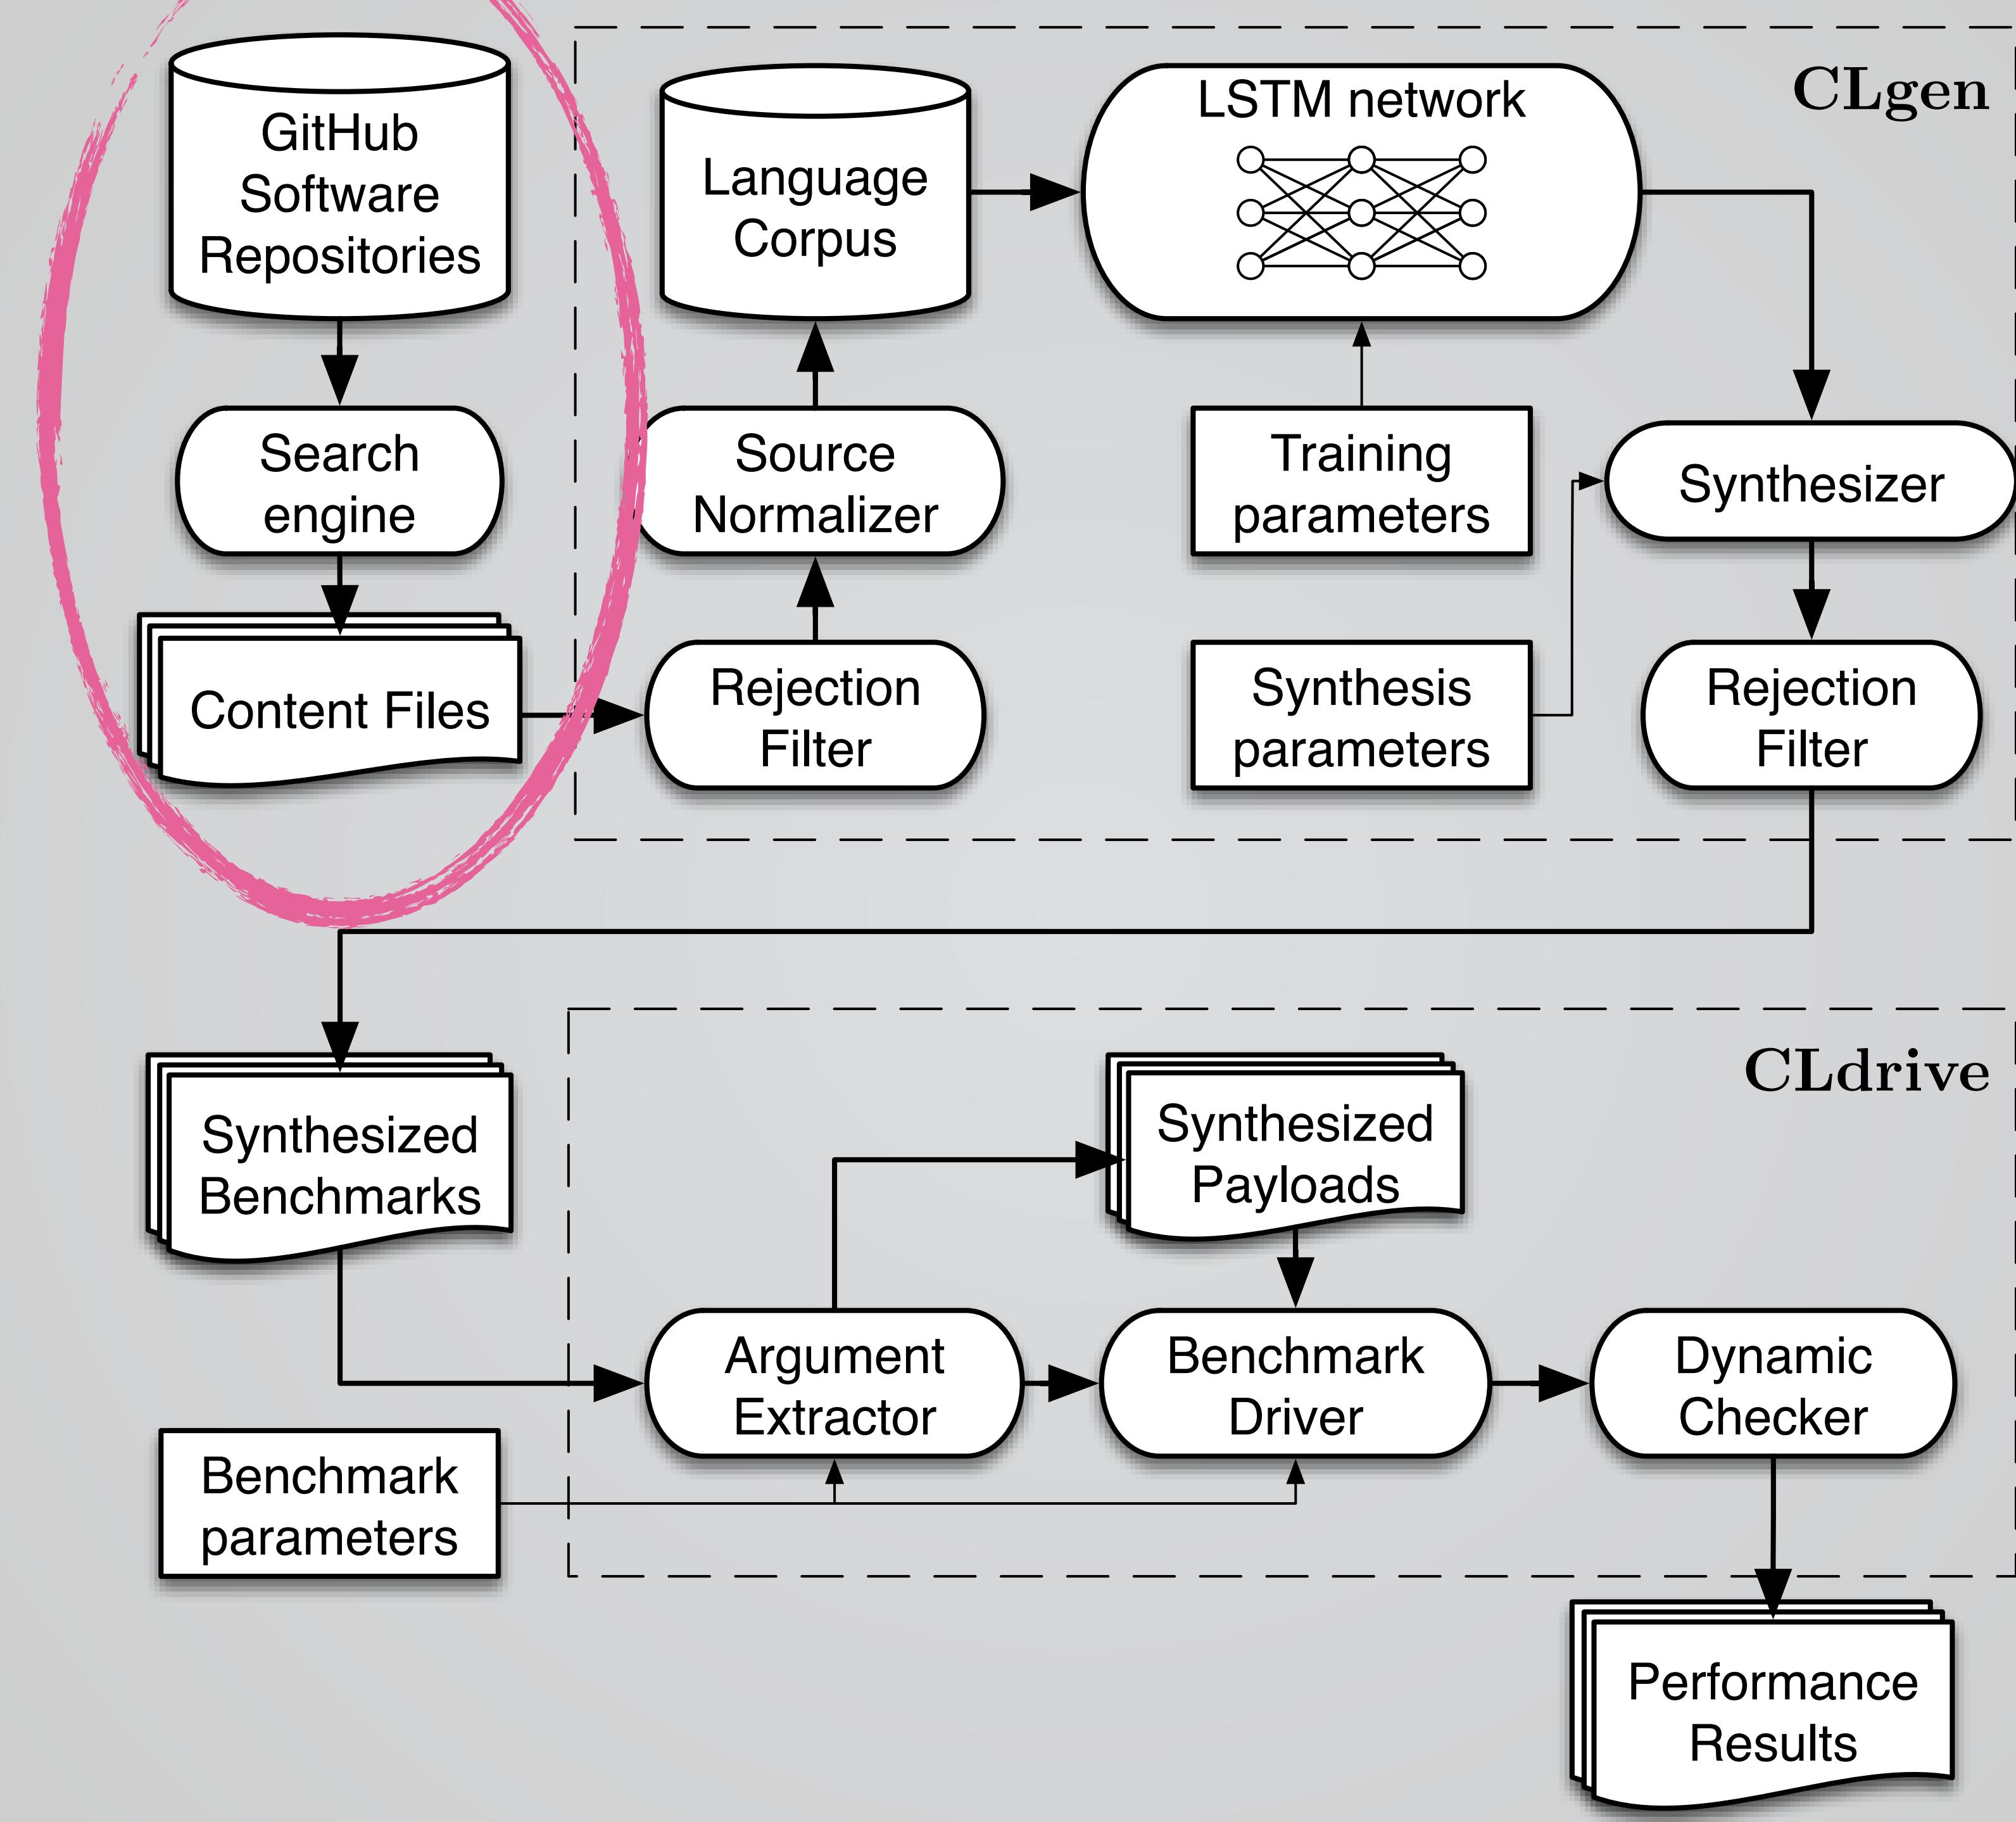
Figure 2: Genesis Program for Image Filtering

# How?

Teach an AI to program from GitHub







We teach an AI to code by showing it lots of code.

Huge repository of public knowledge:



And they have an API :-)

```
$ curl https://api.github.com/search/repositories\?
q\=opengl\&sort\=stars\&order\=desc
{
  "total_count": 3155,
  "incomplete_results": false,
  "items": [
    {
      "id": 7296244,
      "name": "lwjgl3",
      "full_name": "LWJGL/lwjgl3",
```

OpenCL is not a first-class language.

Search repositories using loose keyword terms.

e.g. `opencl`, `nvidia`, `gpu`, `cl`, `amd`.

Recursively iterate over git trees to get `.cl` files.

`Foo:MyOpenCLRepo`

  ↳ `/src/gaussian.cl`

    ↳ `#include <common.h>`

*(0.6% miss rate)*    ↳ `/include/common.h`

          ↳ `#include "detail/math.cl"`

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */
//
//          DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION
//
// 0. You just DO WHAT THE FUCK YOU WANT TO.

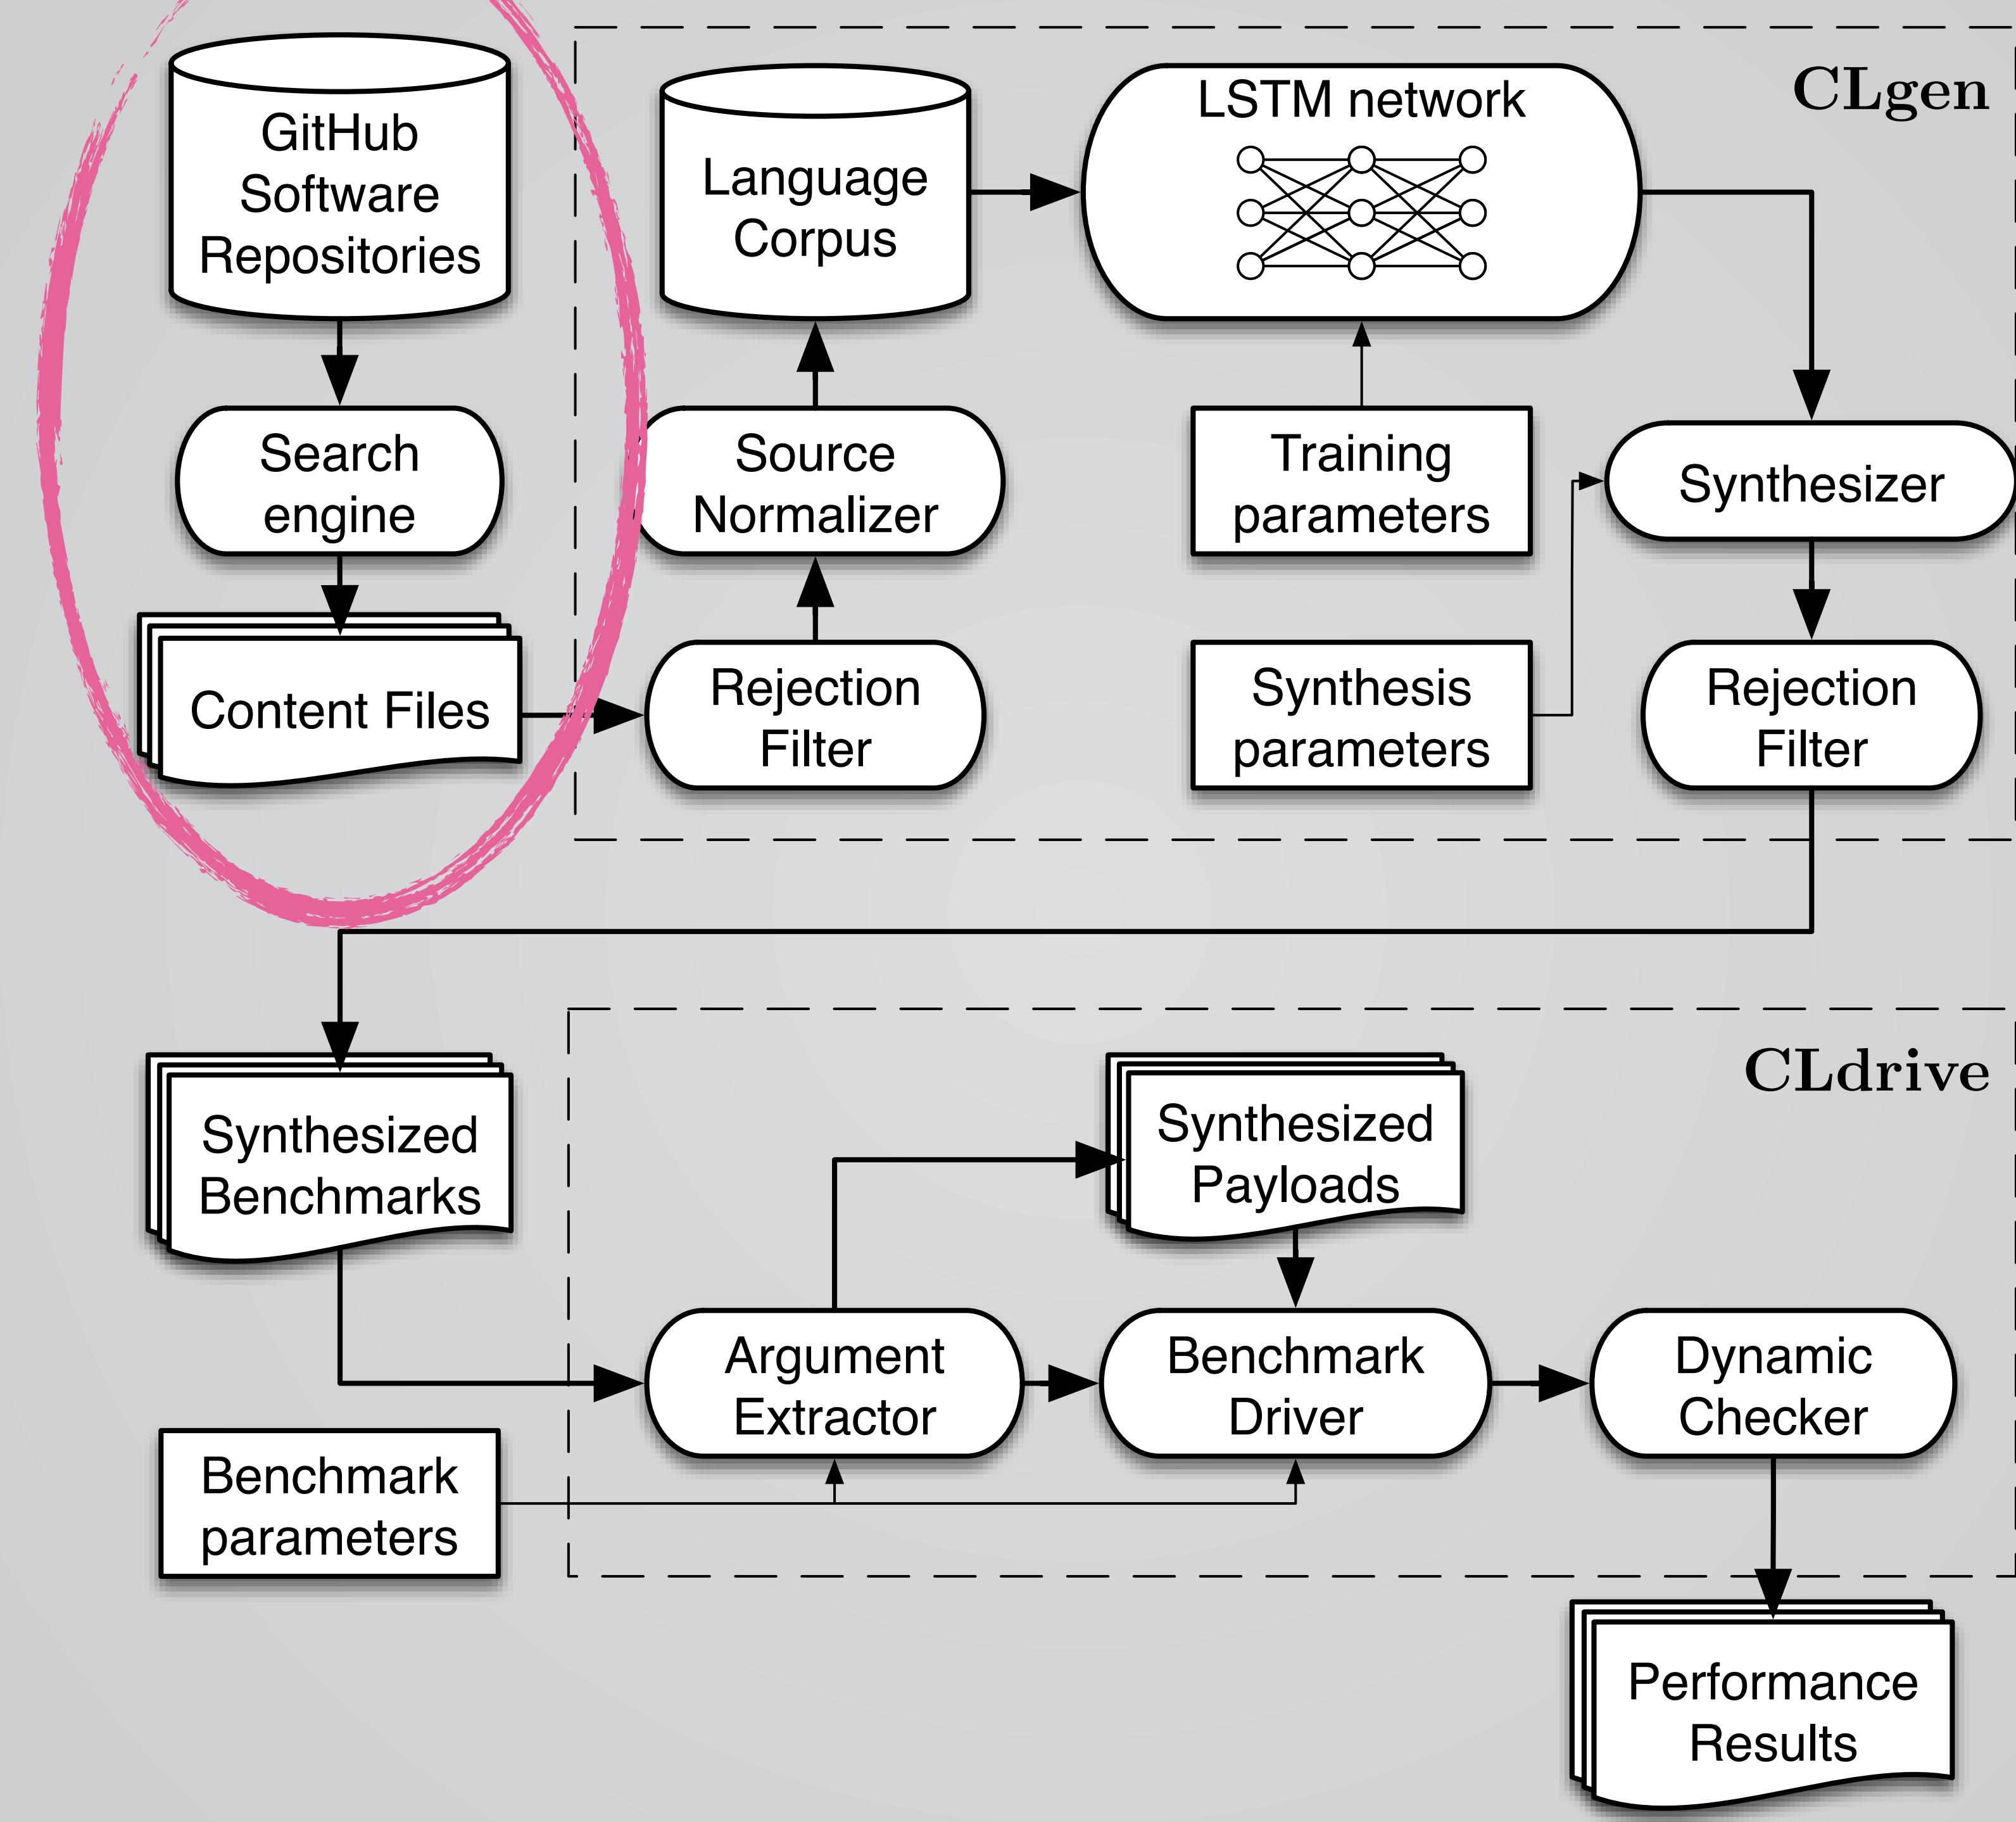
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

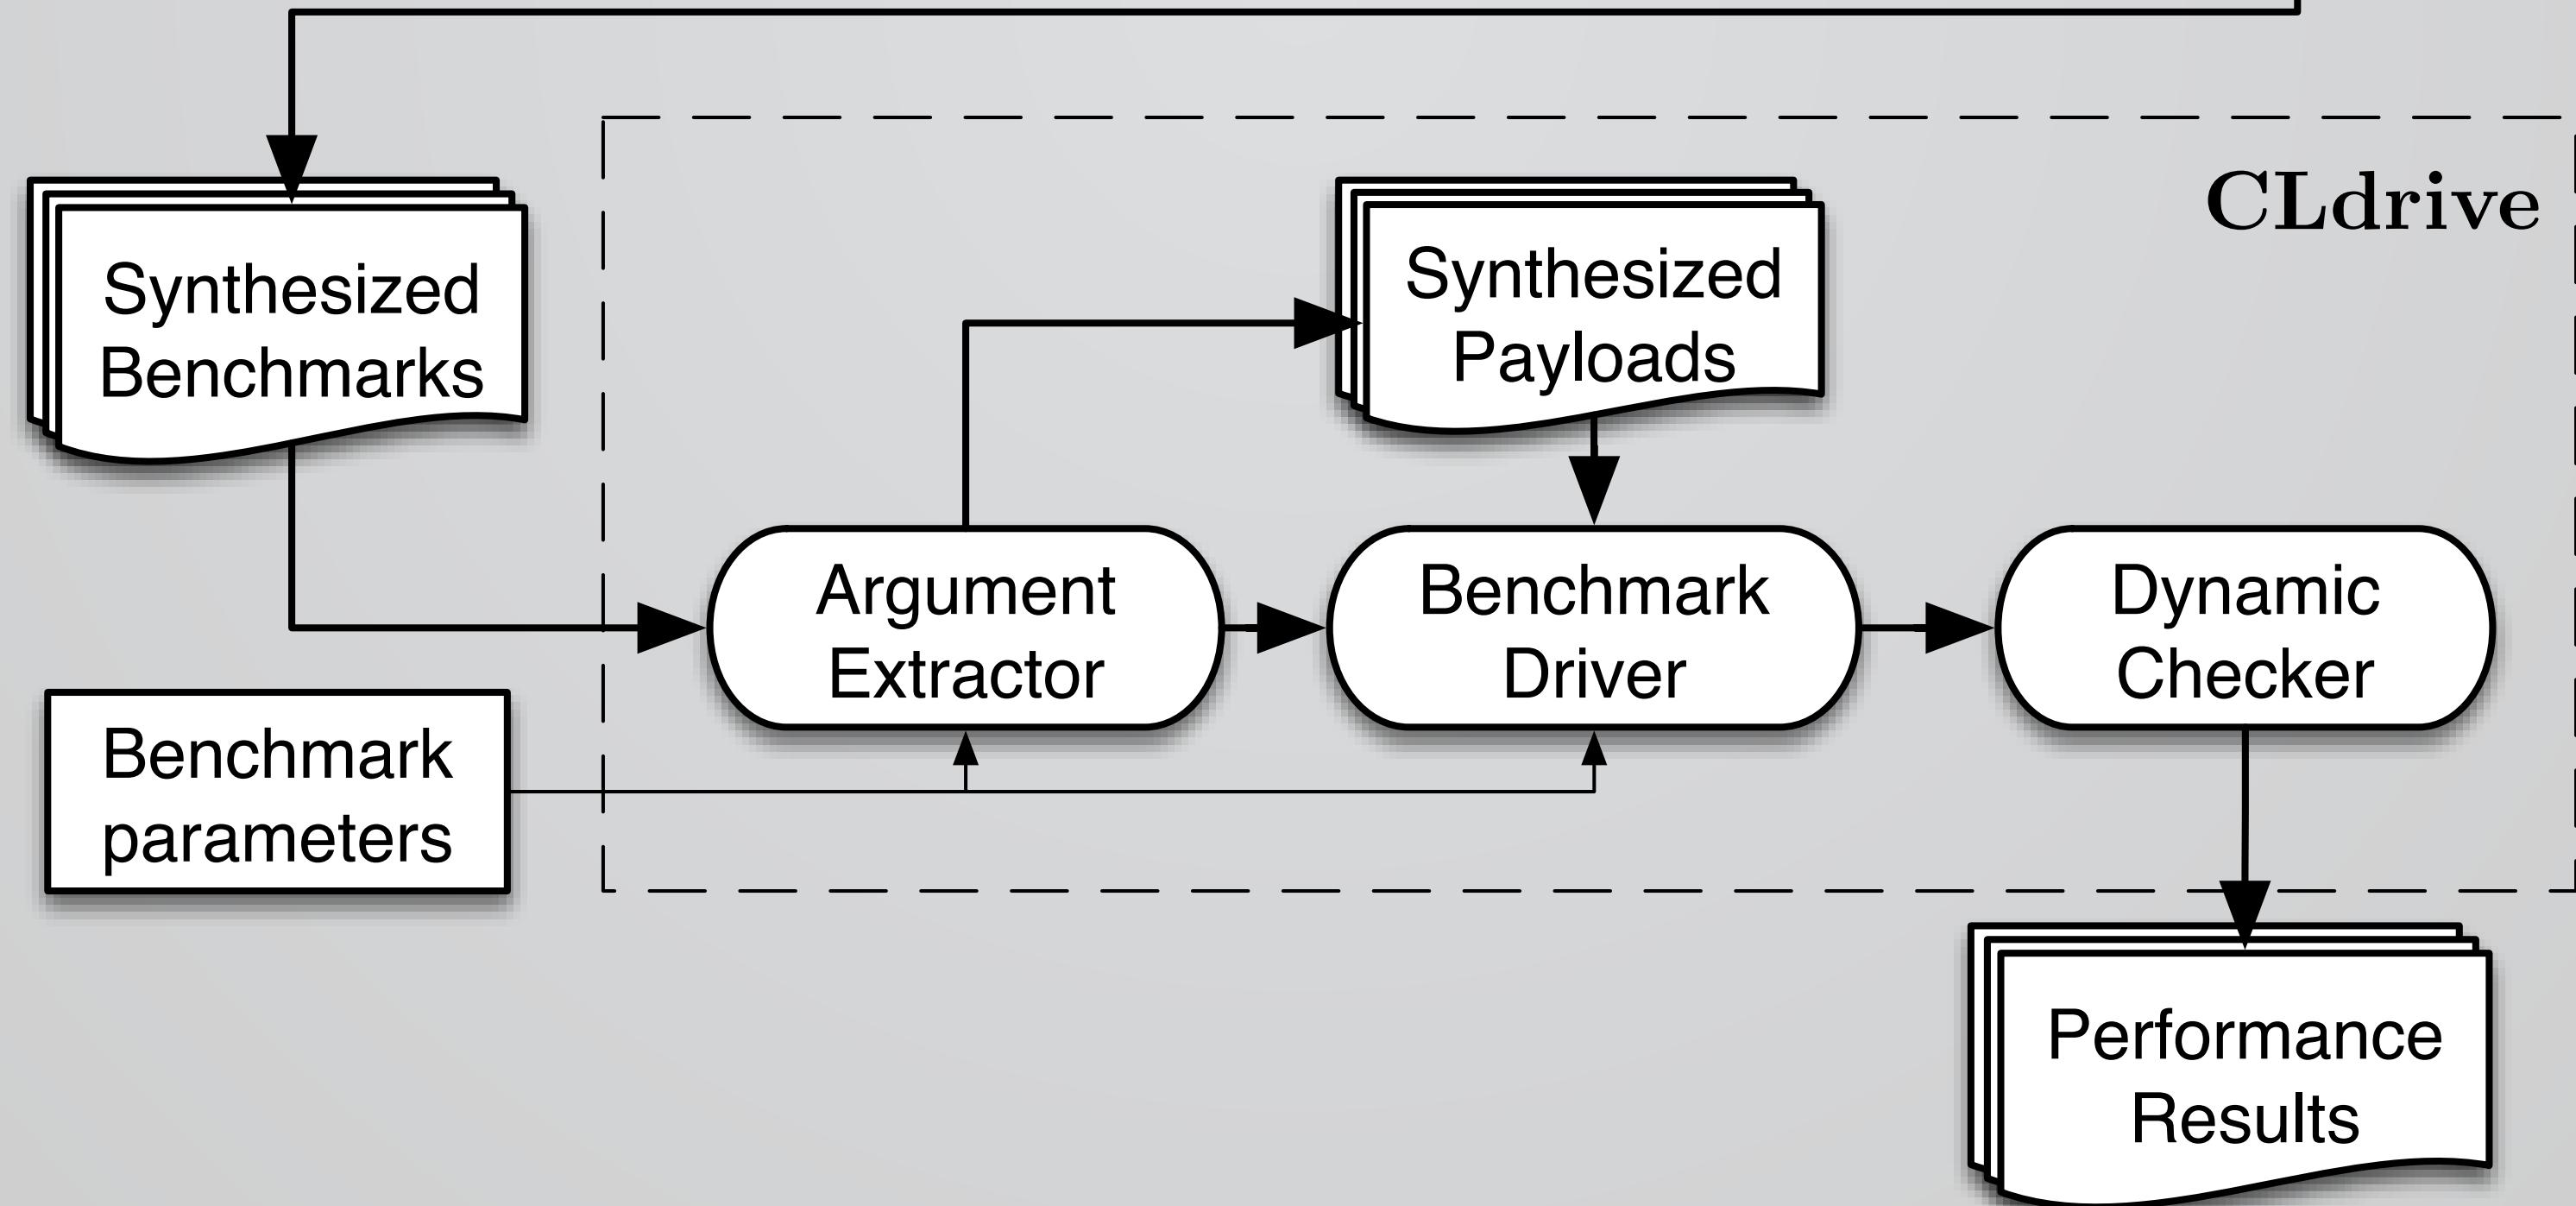
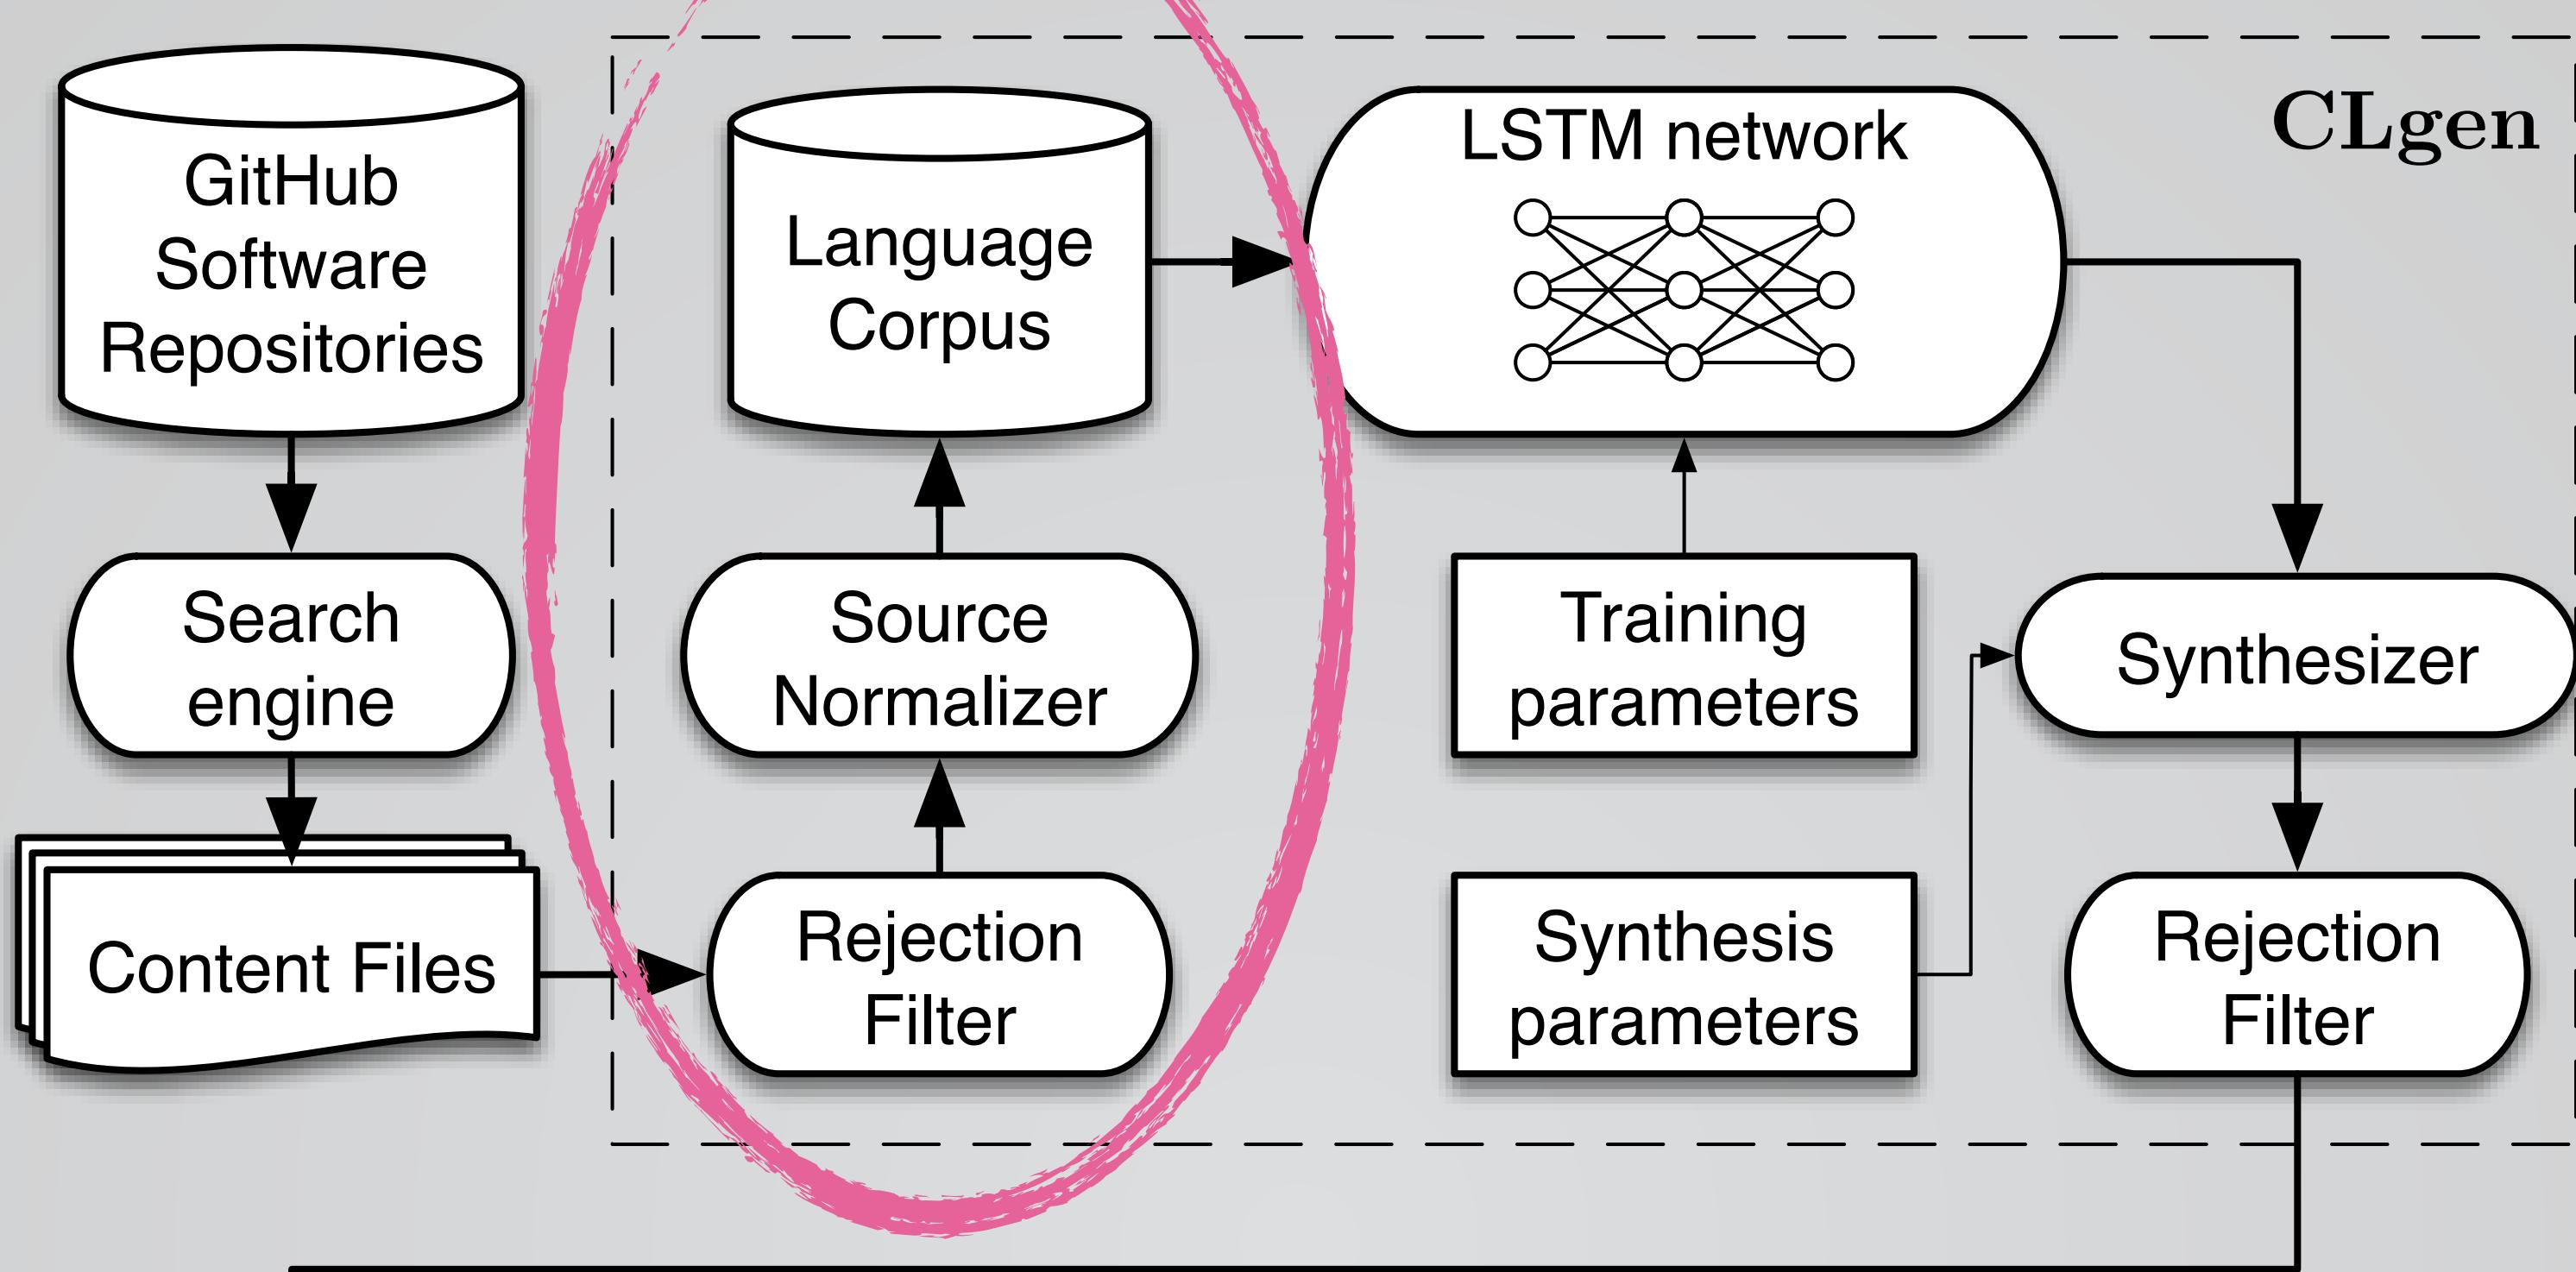
float myclamp(float in) {
#define CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#undef CLAMPING
}

// Do something really flipping cool
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    //
    // 
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

x 8078 files  
2.8 million lines





```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */
//
//          DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION
//
// 0. You just DO WHAT THE FUCK YOU WANT TO.

#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#define CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#undef CLAMPING
#define else
    return in;
#undef endif // CLAMPING
}

// Do something really flipping cool
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    //
    //
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */
//
//          DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION
//
// 0. You just DO WHAT THE FUCK YOU WANT TO.

#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#define CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#undef CLAMPING
#undef CLAMPING
}

// Do something really flipping cool
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    //
    // 
    int id = get_global_id(0);
    if (id < num_elems)
    {
        out[id] = myclamp(in[id]);
    }
}
```

Is this real, valid OpenCL?  
Can we minimise non-functional variance?

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */  
//  
//      DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE  
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION  
//  
// 0. You just DO WHAT THE FUCK YOU WANT TO.  
  
#define CLAMPING  
#define THRESHOLD_MIN 1.0f  
#define THRESHOLD_MAX 1.0f  
  
float myclamp(float in) {  
#ifdef CLAMPING  
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;  
#else  
    return in;  
#endif // CLAMPING  
}  
  
// Do something really flipping cool  
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)  
{  
    //  
    //  
    int id = get_global_id(0);  
    if (id < num_elems)  
    {  
        out[id] = myclamp(in[id]);  
    }  
}
```

Strip comments

Is this real, valid OpenCL?  
Can we minimise non-functional variance?

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */  
//  
// DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE  
// TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION  
//  
// o. You just DO WHAT THE FUCK YOU WANT TO.  
  
#define CLAMPING  
#define THRESHOLD_MIN 1.0f  
#define THRESHOLD_MAX 1.0f  
  
float myclamp(float in) {  
#ifdef CLAMPING  
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;  
#else  
    return in;  
#endif // CLAMPING  
}  
  
// Do something really flipping cool  
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)  
{  
    //  
    int id = get_global_id(0);  
    if (id < num_elems)  
    {  
        out[id] = myclamp(in[id]);  
    }  
}
```

Strip comments

Is this real, valid OpenCL?  
Can we minimise non-functional variance?

~~Strip comments~~

```
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#ifndef CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#else
    return in;
#endif
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~  
Preprocess

```
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#ifndef CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#else
    return in;
#endif
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~  
Preprocess

```
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
    #ifndef CLAMPING
        return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
    #else
        return in;
    #endif
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments  
Preprocess~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments  
Preprocess~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

Does it compile?  
Does it contain instructions?

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Does it compile?~~

~~Does it contain instructions?~~

~~Strip comments~~  
~~Preprocess~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
Rewrite function names

~~Does it compile?~~  
~~Does it contain instructions?~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
Rewrite function names

~~Does it compile?~~  
~~Does it contain instructions?~~

```
float A(float in) {  
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;  
}
```

```
__kernel void B(__global float* in, __global float* out, int num_elems)  
{  
    int id = get_global_id(0);  
    if (id < num_elems)  
    {  
  
        out[id] = A(in[id]);  
    }  
}
```

~~Strip comments~~  
~~Preprocess~~  
~~Rewrite function names~~

~~Does it compile?~~  
~~Does it contain instructions?~~

```
float A(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void B(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = A(in[id]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
~~Rewrite function names~~  
~~Rewrite variable names~~

~~Does it compile?~~  
~~Does it contain instructions?~~

```
float A(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void B(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {
        out[id] = A(in[id]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
~~Rewrite function names~~  
~~Rewrite variable names~~

~~Does it compile?~~  
~~Does it contain instructions?~~

```
float A(float a) {
    return a > 1.0f ? 1.0f : a < 0.0f ? 0.0f : a;
}

__kernel void B(__global float* a, __global float* b, int c)
{
    int d = get_global_id(0);
    if (d < c)
    {

        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
~~Rewrite function names~~  
~~Rewrite variable names~~

~~Does it compile?~~  
~~Does it contain instructions?~~

```
float A(float a) {
    return a > 1.0f ? 1.0f : a < 0.0f ? 0.0f : a;
}

__kernel void B(__global float* a, __global float* b, int c)
{
    int d = get_global_id(0);
    if (d < c)
    {

        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
~~Rewrite function names~~  
~~Rewrite variable names~~  
~~Enforce code style~~

~~Does it compile?~~  
~~Does it contain instructions?~~

```
float A(float a) {
    return a > 1.0f ? 1.0f : a < 0.0f ? 0.0f : a;
}

__kernel void B(__global float* a, __global float* b, int c)
{
    int d = get_global_id(0);
    if (d < c)
    {
        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
~~Rewrite function names~~  
~~Rewrite variable names~~  
~~Enforce code style~~

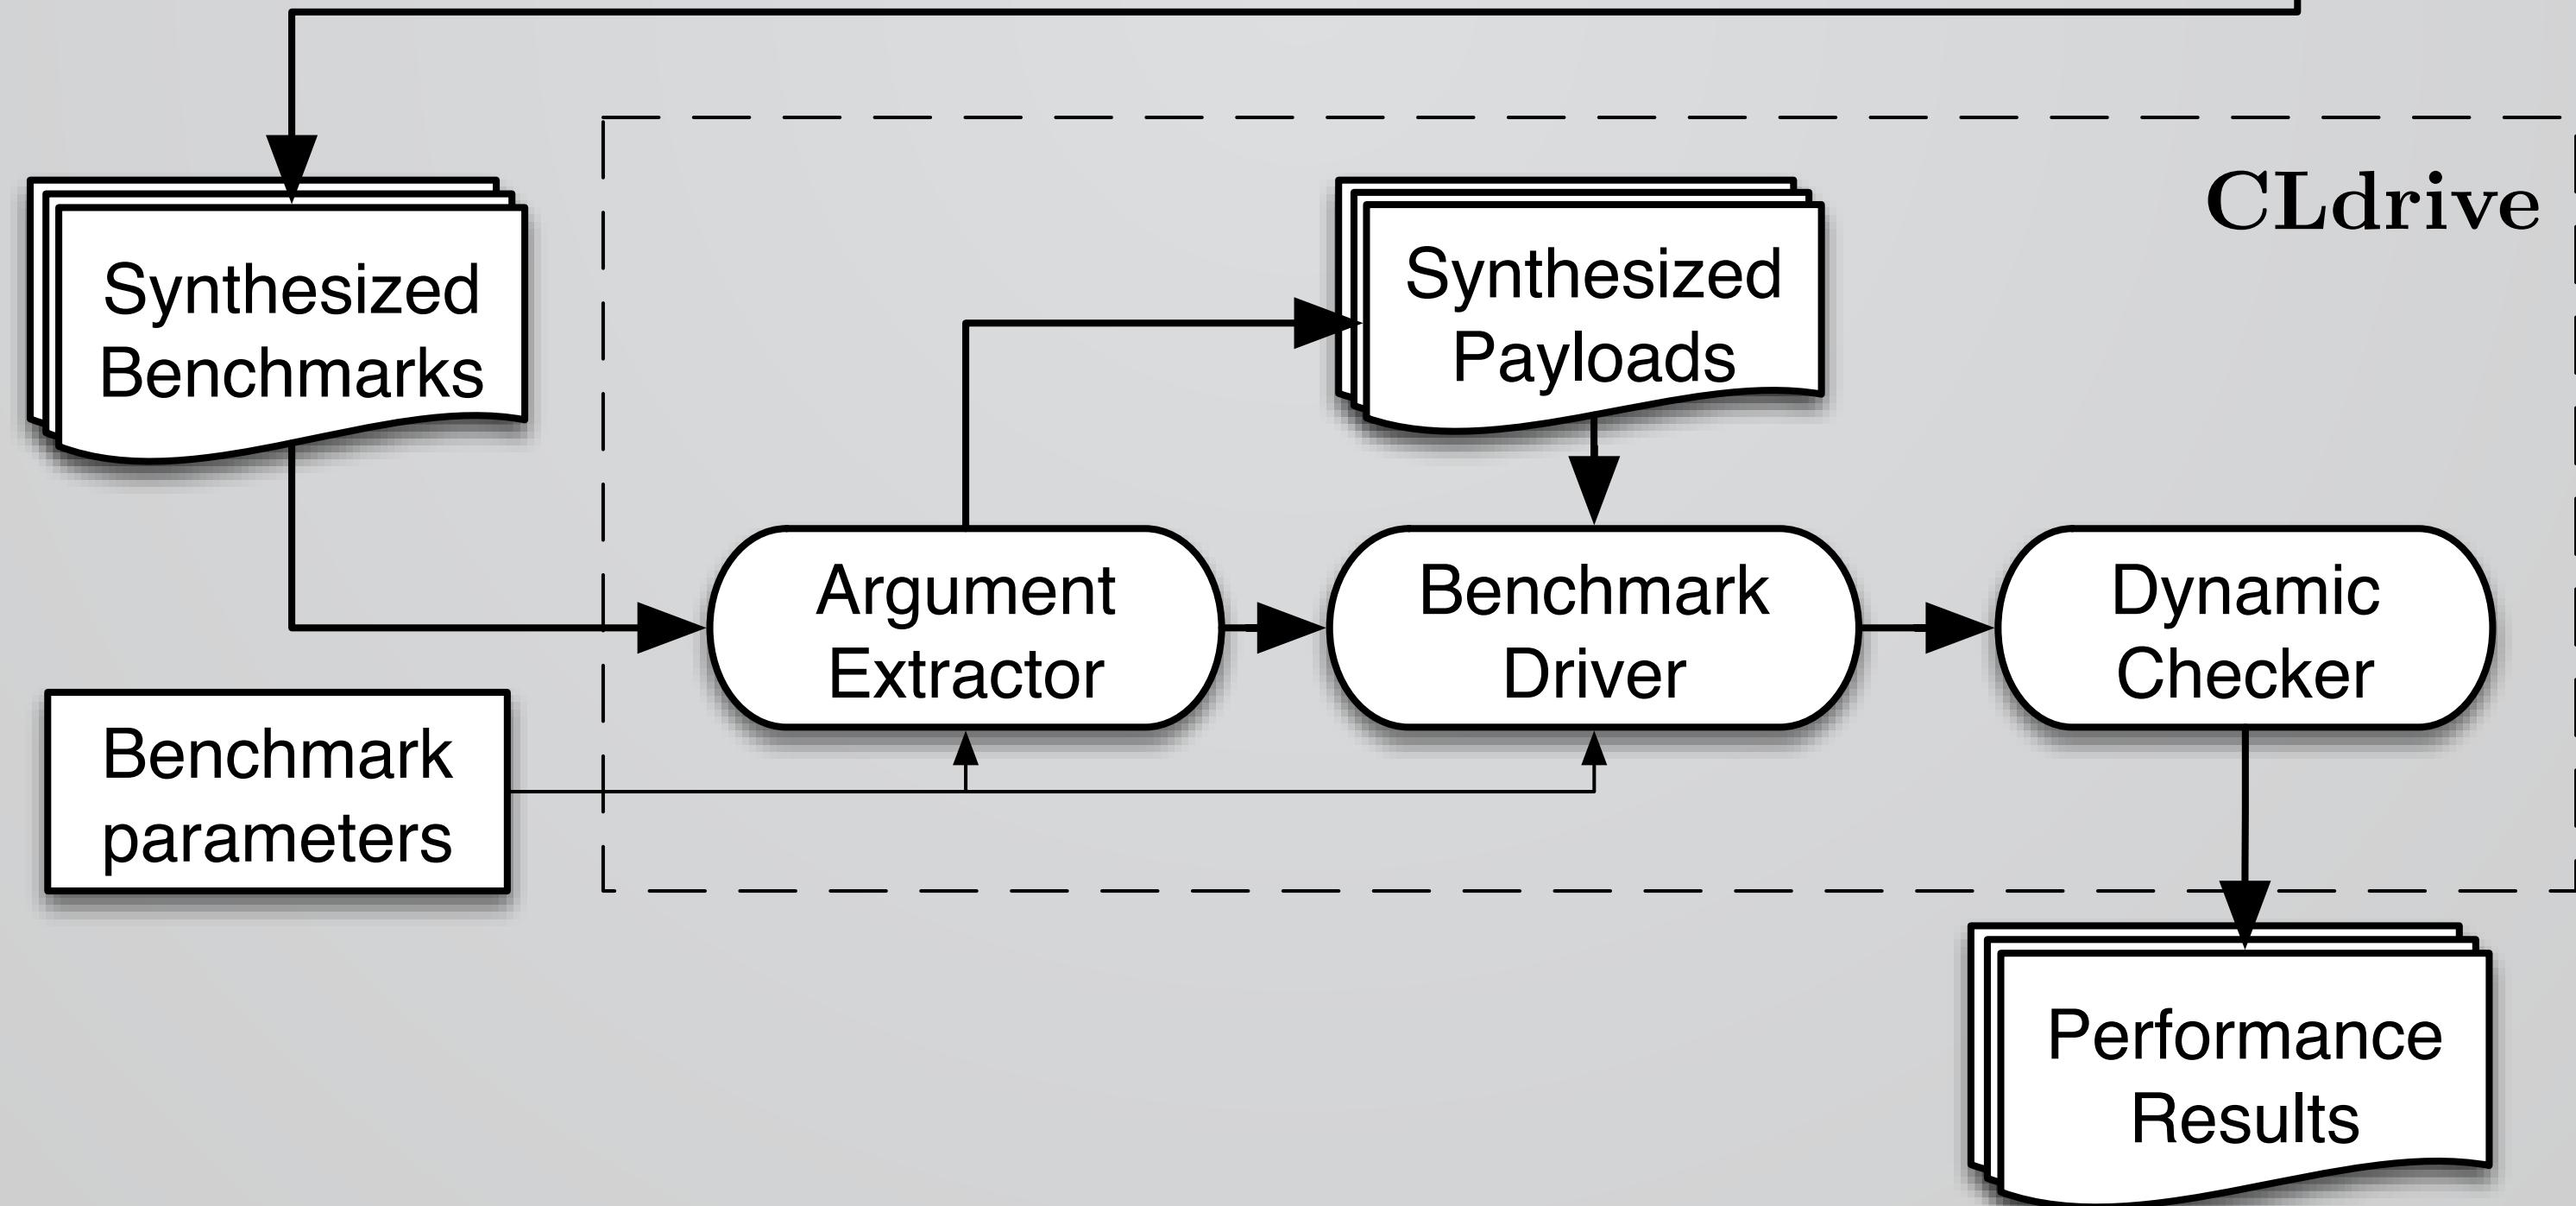
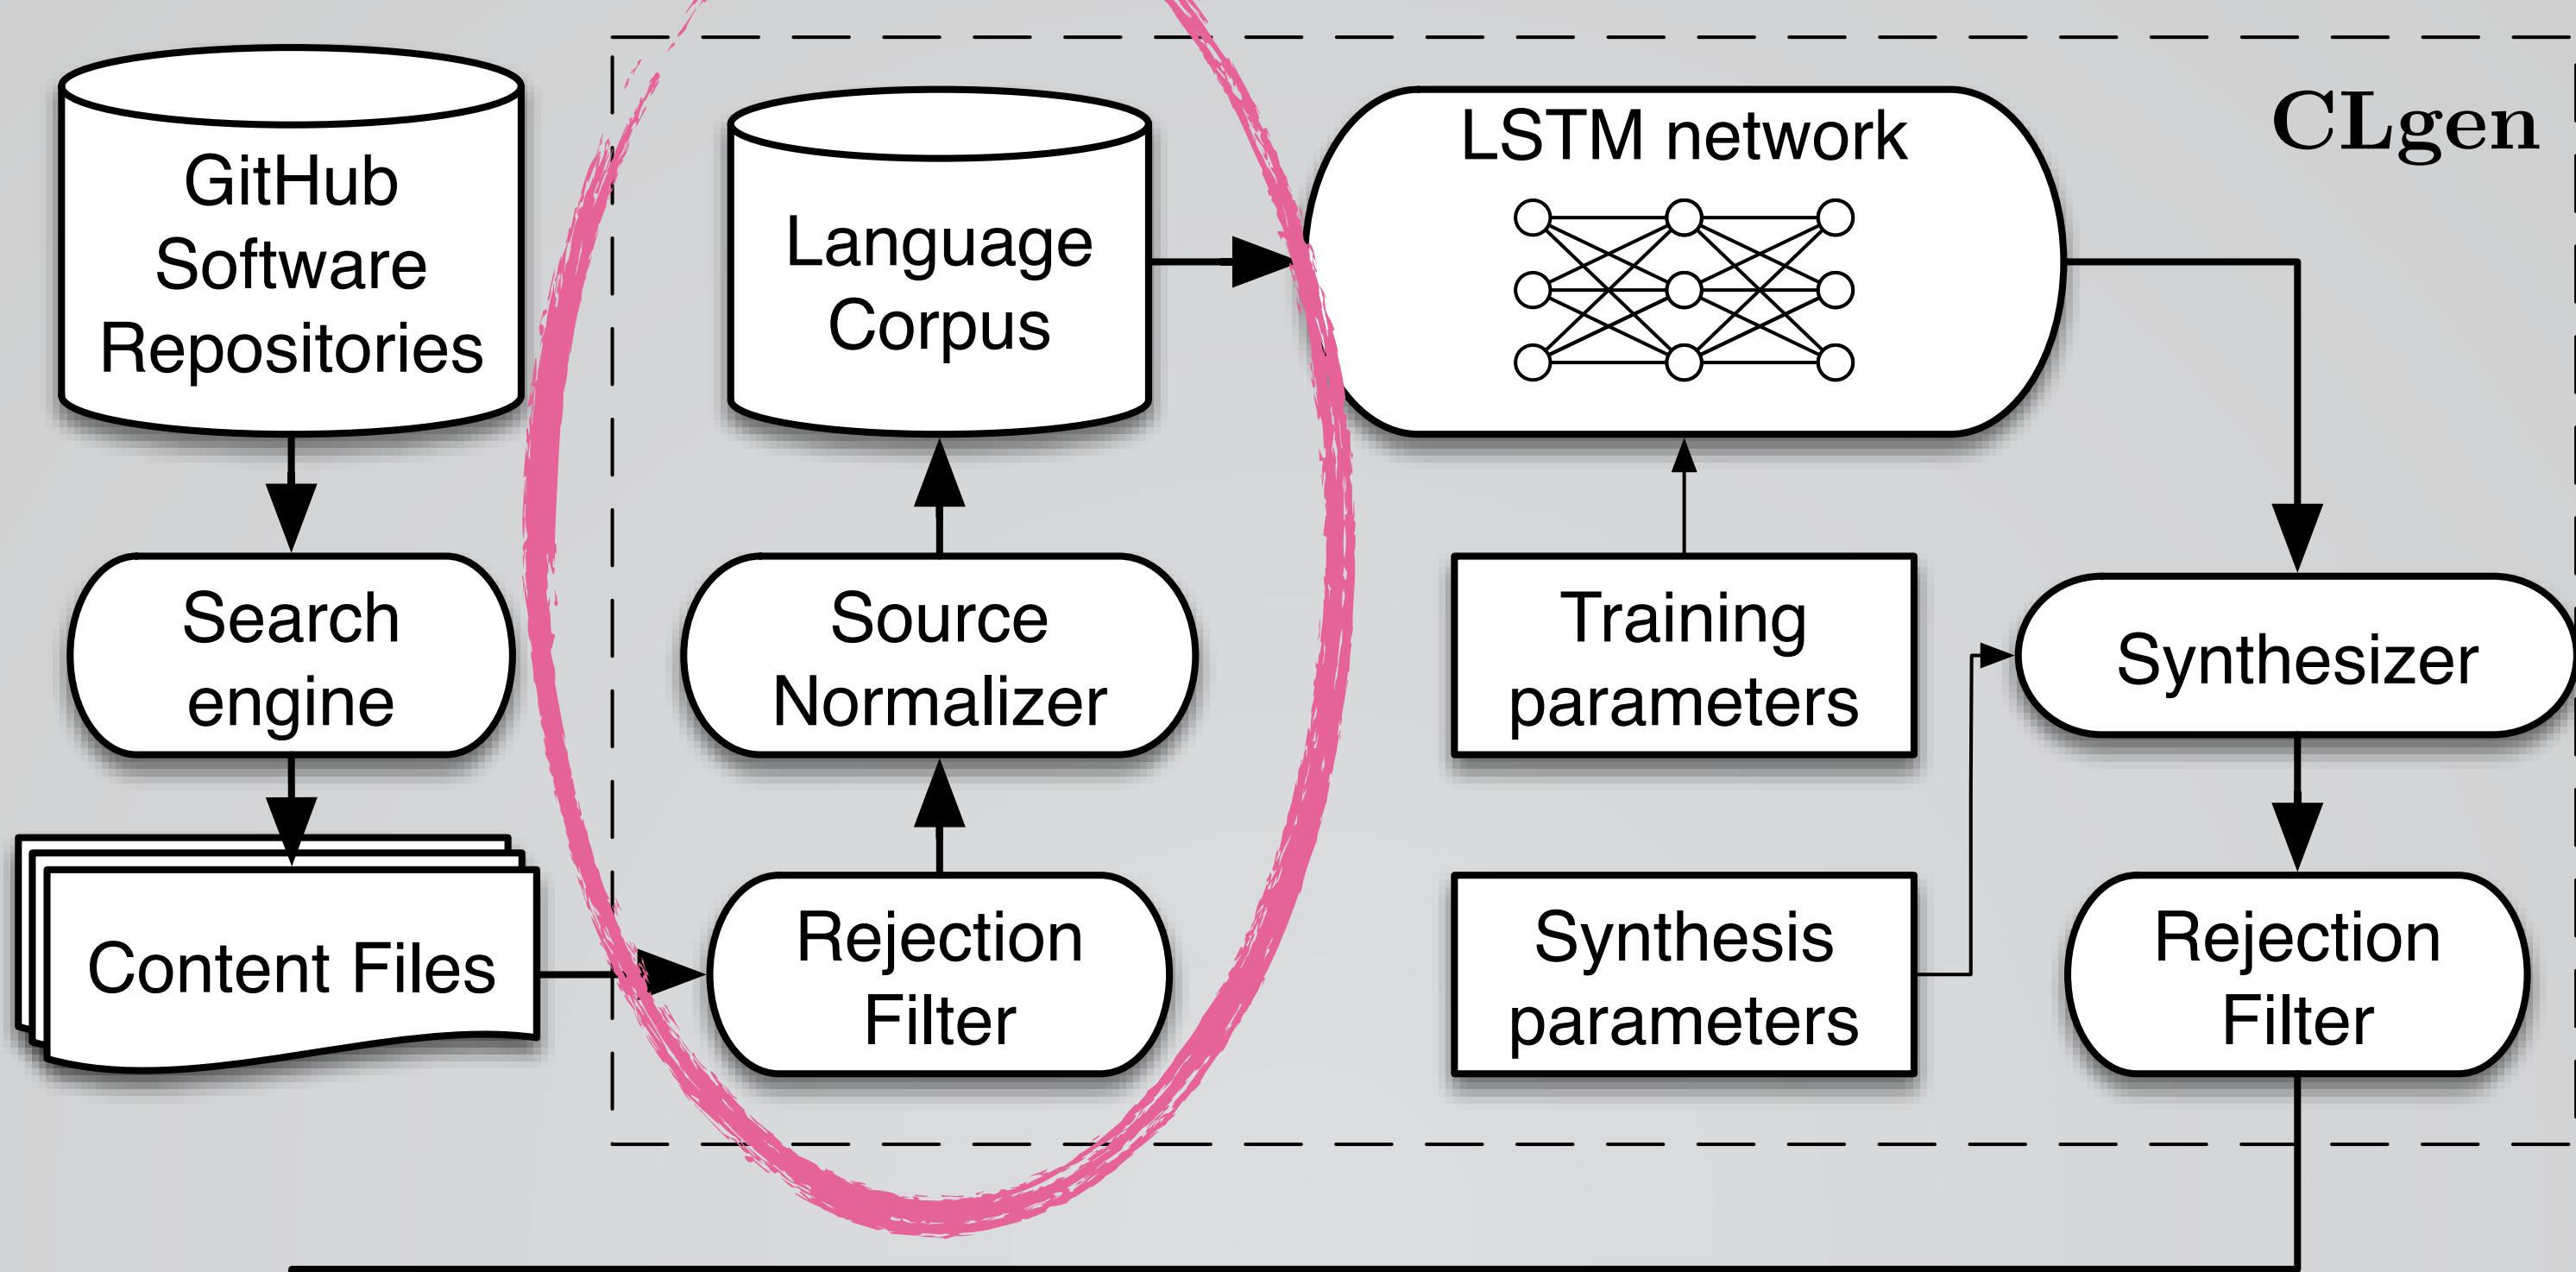
~~Does it compile?~~  
~~Does it contain instructions?~~

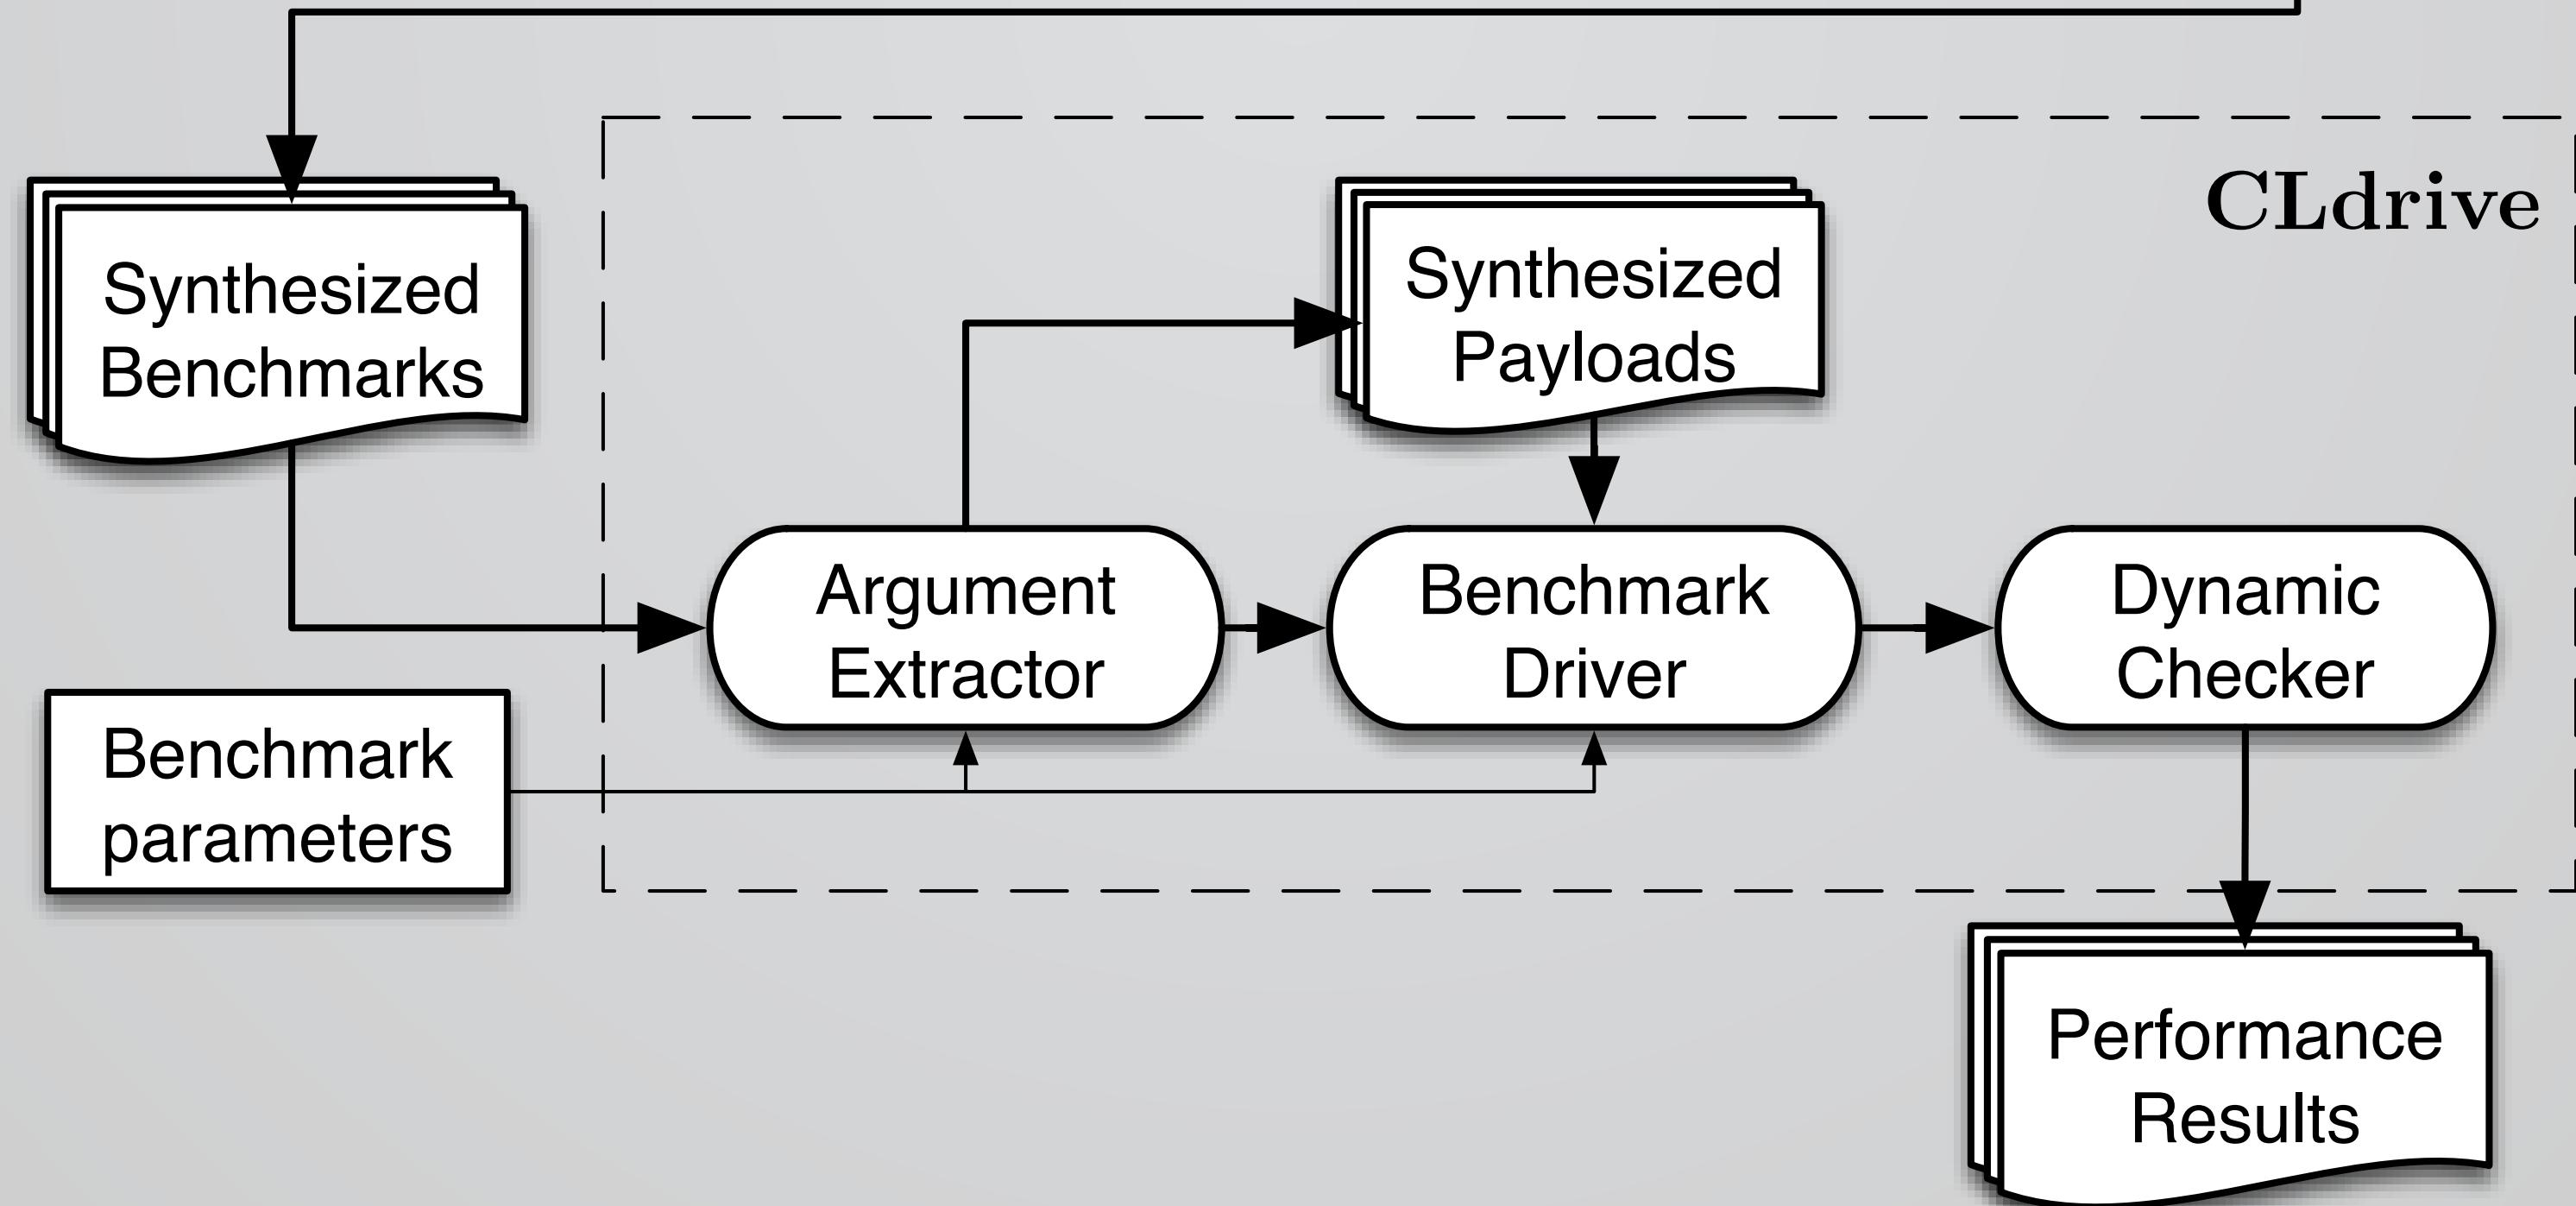
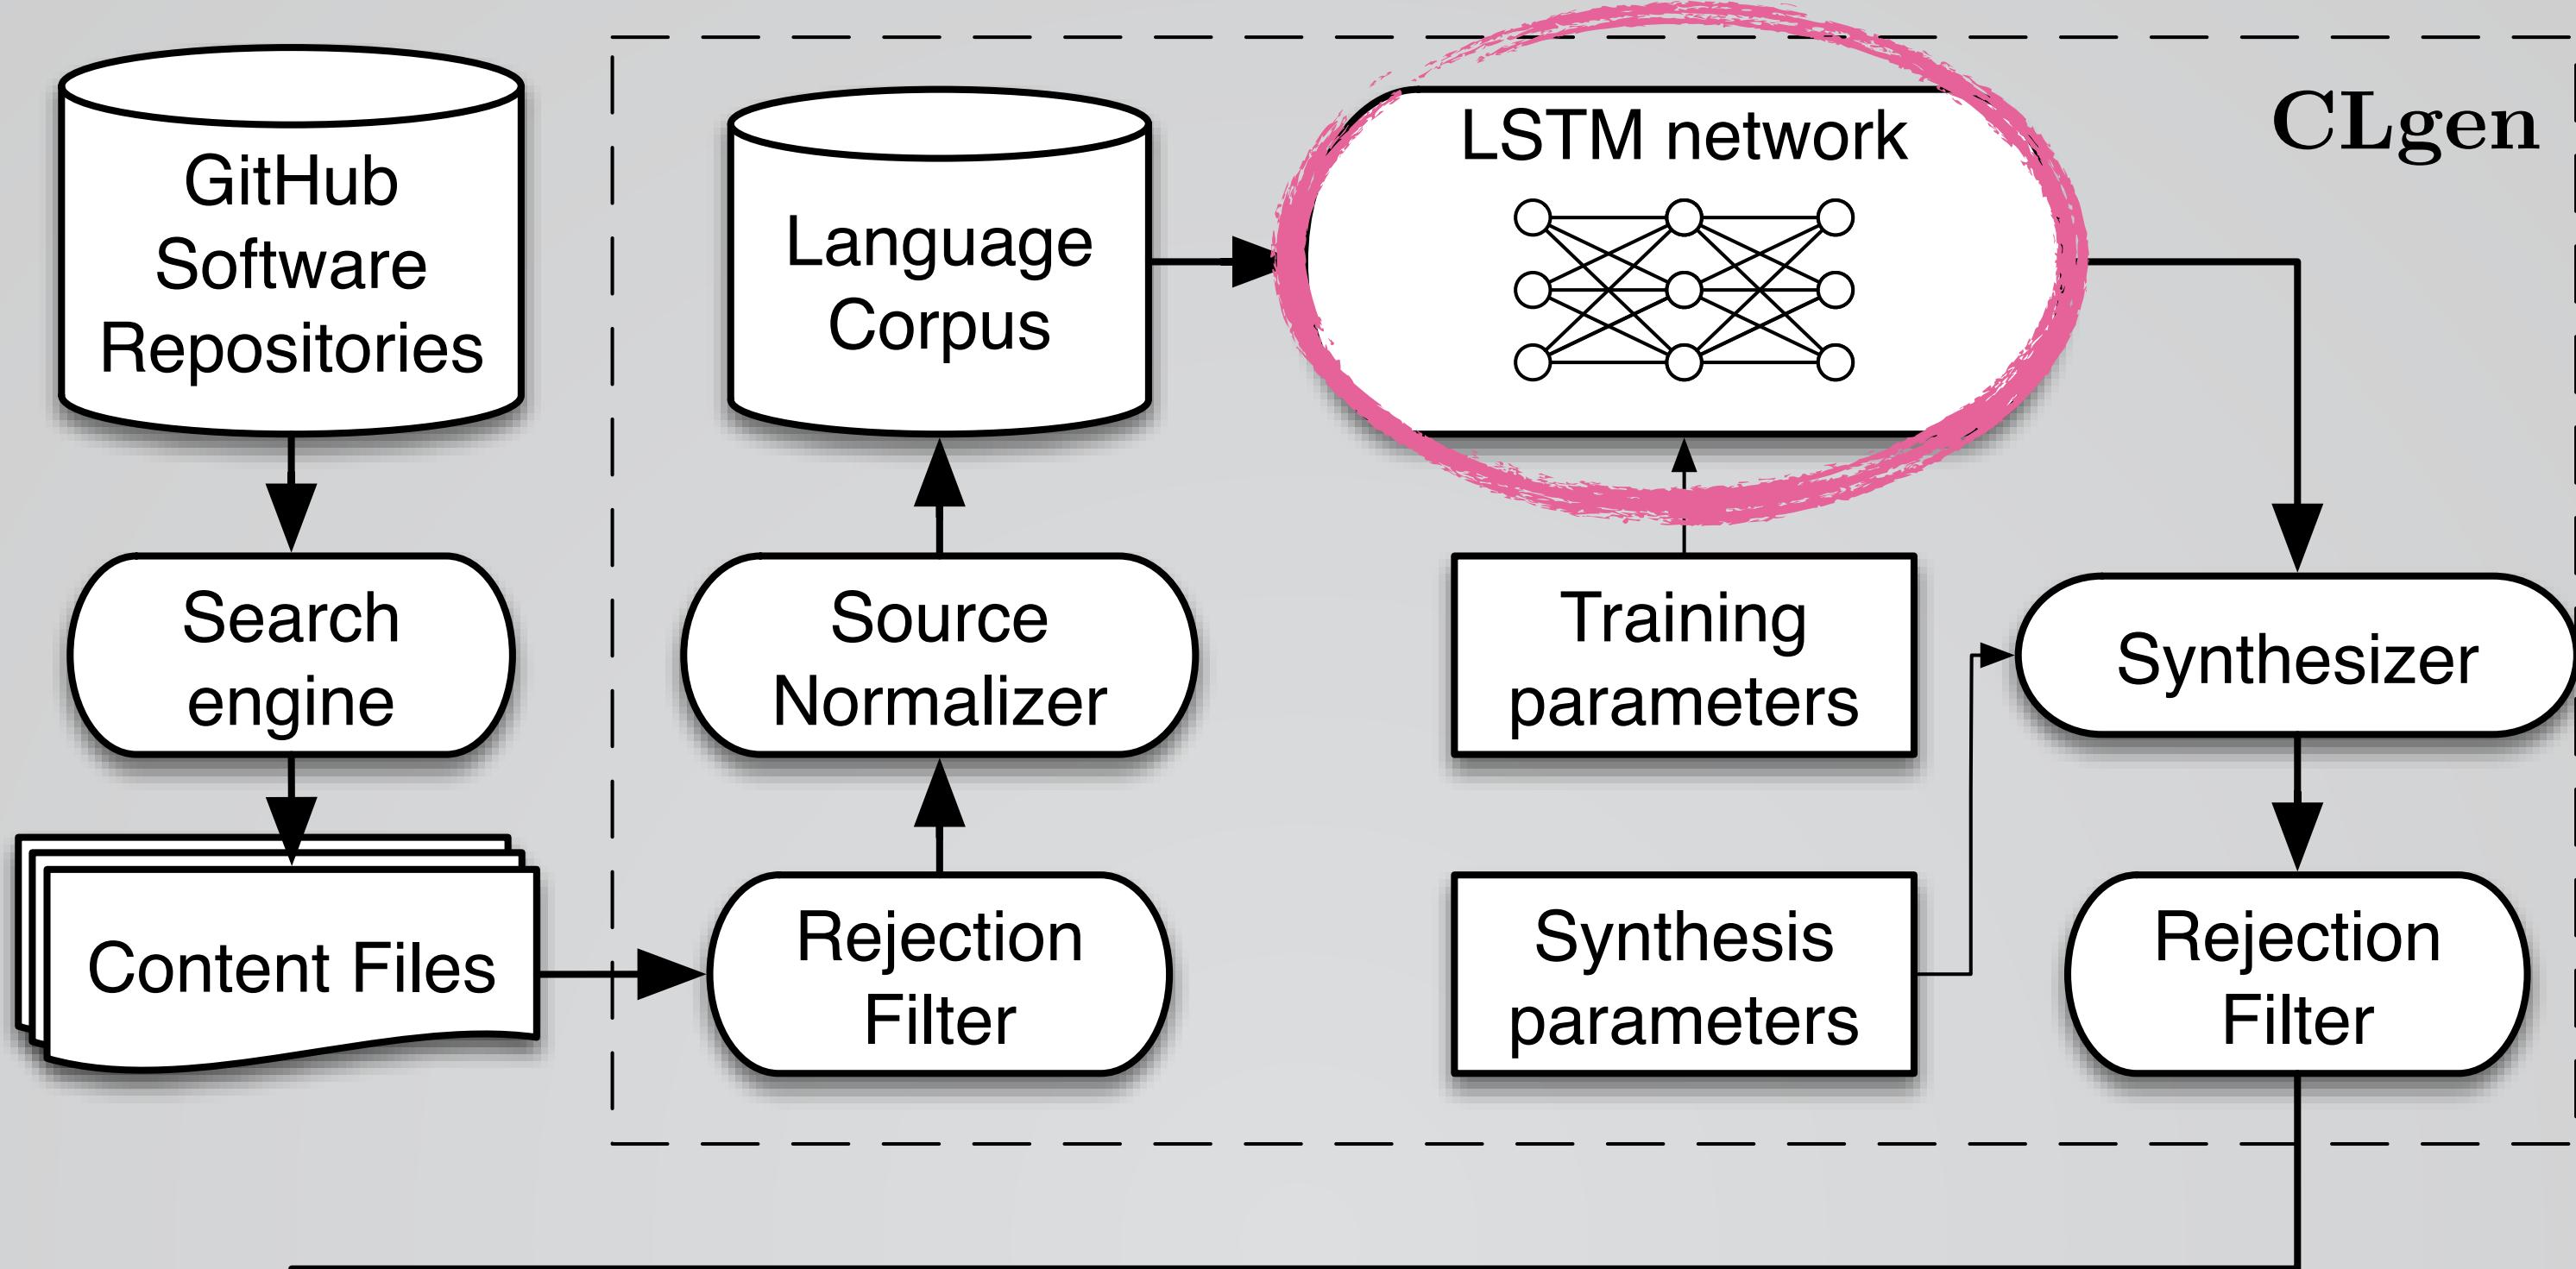
```
float A(float a) {
    return a > 1.0f ? 1.0f : in < 0.0f ? 0.0f : a;
}

__kernel void B(__global float* a, __global float* b, int c) {
    int d = get_global_id(0);
    if (d < c) {
        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~  
~~Preprocess~~  
~~Rewrite function names~~  
~~Rewrite variable names~~  
~~Enforce code style~~

~~Does it compile?~~  
~~Does it contain instructions?~~





## Forward Pass

*Input Gates:*

$$x_t = \sum_{j \in N} w_{tj} y_j(\tau - 1) + \sum_{c \in C} w_{tc} s_c(\tau - 1)$$

$$y_t = f(x_t)$$

*Forget Gates:*

$$x_\phi = \sum_{j \in N} w_{\phi j} y_j(\tau - 1) + \sum_{c \in C} w_{\phi c} s_c(\tau - 1)$$

$$y_\phi = f(x_\phi)$$

*Cells:*

$$\forall c \in C, x_c = \sum_{j \in N} w_{cj} y_j(\tau - 1)$$

$$s_c = y_\phi s_c(\tau - 1) + y_t g(x_c)$$

*Output Gates:*

$$x_\omega = \sum_{j \in N} w_{\omega j} y_j(\tau - 1) + \sum_{c \in C} w_{\omega c} s_c(\tau)$$

$$y_\omega = f(x_\omega)$$

*Forget Gates:*

## Backward Pass

$$\text{define } \delta_k(\tau) = \frac{\partial E(\tau)}{\partial x_k}$$

$$\delta_k(\tau) = y_k(\tau) - t_k(\tau) \quad k \in \text{output units}$$

*Cell Outputs:*

$$\forall c \in C, \text{ define } \epsilon_c = \sum_{j \in N} w_{jc} \delta_j(\tau + 1)$$

*Output Gates:*

$$\delta_\omega = f'(x_\omega) \sum_{c \in C} \epsilon_c h(s_c)$$

*States:*

$$\frac{\partial E}{\partial s_c}(\tau) = \epsilon_c y_\omega h'(s_c) + \frac{\partial E}{\partial s_c}(\tau + 1) y_\phi(\tau + 1)$$

$$+ \delta_t(\tau + 1) w_{tc} + \delta_\phi(\tau + 1) w_{\phi c} + \delta_\omega w_{\omega c}$$

*Cells:*

$$\forall c \in C, \delta_c = y_t g'(x_c) \frac{\partial E}{\partial s_c}$$

*Input Gates:*

$$\delta_\phi = f'(x_\phi) \sum_{c \in C} \frac{\partial E}{\partial s_c} s_c(\tau - 1)$$

*Input Gates:*

$$\delta_t = f'(x_t) \sum_{c \in C} \frac{\partial E}{\partial s_c} g(x_c)$$

## Update Weights

$$\text{define } E_{total}(S) = \sum_{\tau=\tau_0}^{\tau_1} E(\tau)$$

$$\text{define } \nabla_{ij}(S) = \frac{\partial E_{total}(S)}{\partial w_{ij}}$$

$$\implies \nabla_{ij}(S) = \sum_{\tau=\tau_0+1}^{\tau_1} \delta_i(\tau) y_j(\tau - 1)$$

$$\Delta w_{ij}(S) = -\alpha \nabla_{ij}(S) + m \Delta w_{ij}(S - 1)$$

## Network Design

2048 nodes, 3 layers

Stochastic Gradient Descent

Initial learning rate 0.005

Train for 50 epochs

Learning rate decay every 5 epochs

## Forward Pass

*Input Gates:*

$$x_t = \sum_{j \in N} w_{tj} y_j(\tau - 1) + \sum_{c \in C} w_{tc} s_c(\tau - 1)$$

$$y_t = f(x_t)$$

*Forget Gates:*

$$x_\phi = \sum_{j \in N} w_{\phi j} y_j(\tau - 1) + \sum_{c \in C} w_{\phi c} s_c(\tau - 1)$$

$$y_\phi = f(x_\phi)$$

*Cells:*

$$\forall c \in C, x_c = \sum_{j \in N} w_{cj} y_j(\tau - 1)$$

$$s_c = y_\phi s_c(\tau - 1) + y_t g(x_c)$$

*Output Gates:*

$$x_\omega = \sum_{j \in N} w_{\omega j} y_j(\tau - 1) + \sum_{c \in C} w_{\omega c} s_c(\tau)$$

$$y_\omega = f(x_\omega)$$

## Backward Pass

$$\text{define } \delta_k(\tau) = \frac{\partial E(\tau)}{\partial x_k}$$

$$\delta_k(\tau) = y_k(\tau) - t_k(\tau) \quad k \in \text{output units}$$

*Cell Outputs:*

$$\forall c \in C, \text{ define } \epsilon_c = \sum_j w_{cj} y_j(\tau + 1)$$

*Output Gates:*

$$\delta_\omega = f'(x_\omega) \sum_{c \in C} \epsilon_c h(s_c)$$

*Gates:*

$$\frac{\partial E}{\partial s_c}(\tau) = \epsilon_c y_\omega h'(s_c) + \frac{\partial E}{\partial s_c}(\tau + 1) y_\phi(\tau + 1)$$

$$+ \delta_t(\tau + 1) w_{tc} + \delta_\phi(\tau + 1) w_{\phi c} + \delta_\omega w_{\omega c}$$

*Cells:*

$$\forall c \in C, \delta_c = y_t g'(x_c) \frac{\partial E}{\partial s_c}$$

*Forget Gates:*

$$\delta_\phi = f'(x_\phi) \sum_{c \in C} \frac{\partial E}{\partial s_c} s_c(\tau - 1)$$

*Input Gates:*

$$\delta_t = f'(x_t) \sum_{c \in C} \frac{\partial E}{\partial s_c} g(x_c)$$

## Update Weights

$$\text{define } E_{total}(S) = \sum_{\tau=\tau_0}^{\tau_1} E(\tau)$$

$$\text{define } \nabla_{ij}(S) = \frac{\partial E_{total}(S)}{\partial w_{ij}}$$

$$\Rightarrow \nabla_{ij}(S) = \sum_{\tau=\tau_0+1}^{\tau_1} \delta_i(\tau) y_j(\tau - 1)$$

$$\Delta w_{ij}(S) = -\alpha \nabla_{ij}(S) + m \Delta w_{ij}(S - 1)$$

## Network Design

2048 nodes, 3 layers

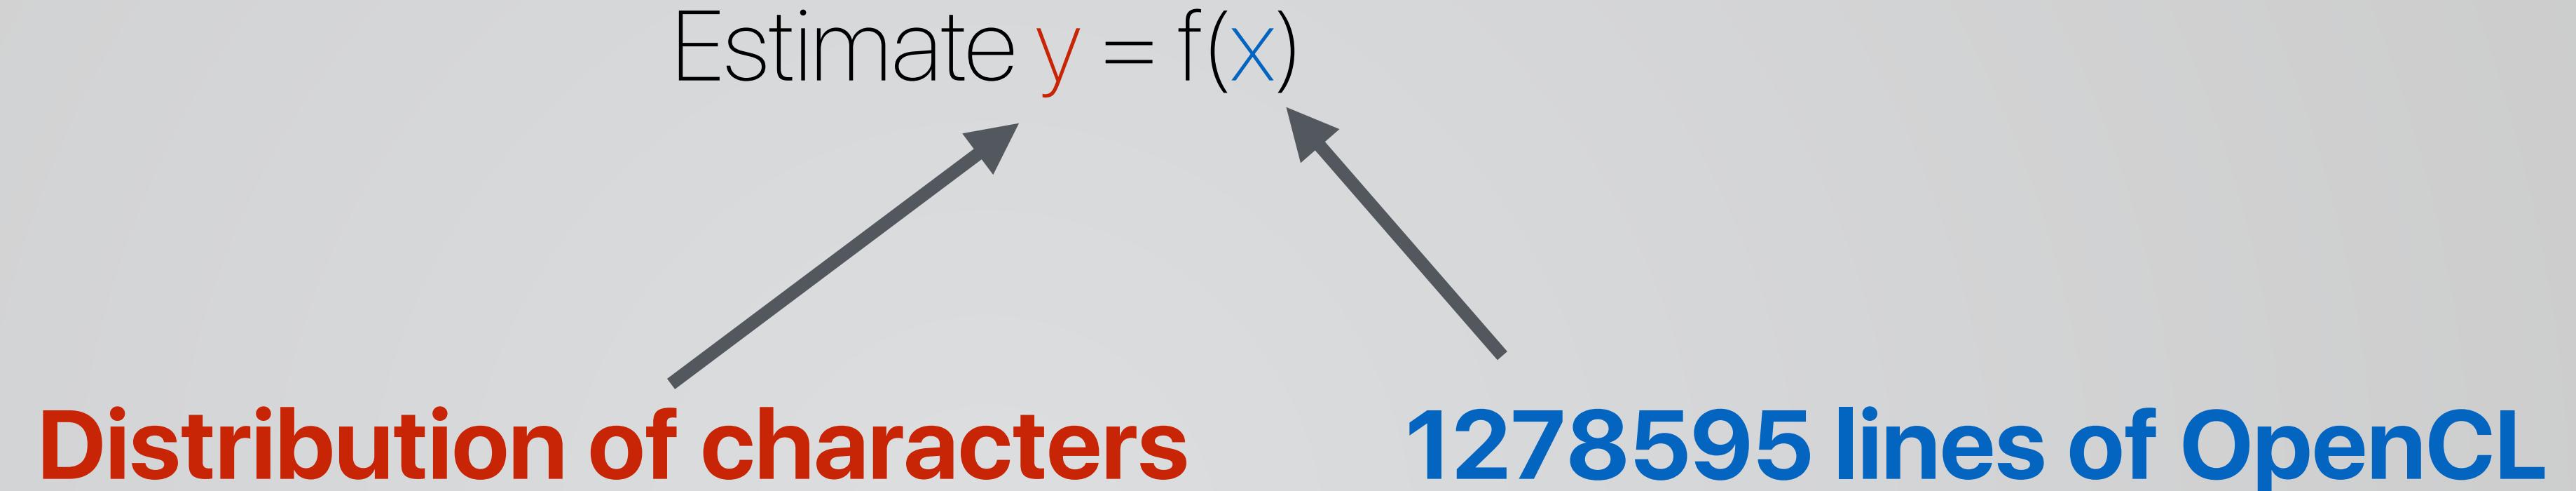
Stochastic Gradient Descent

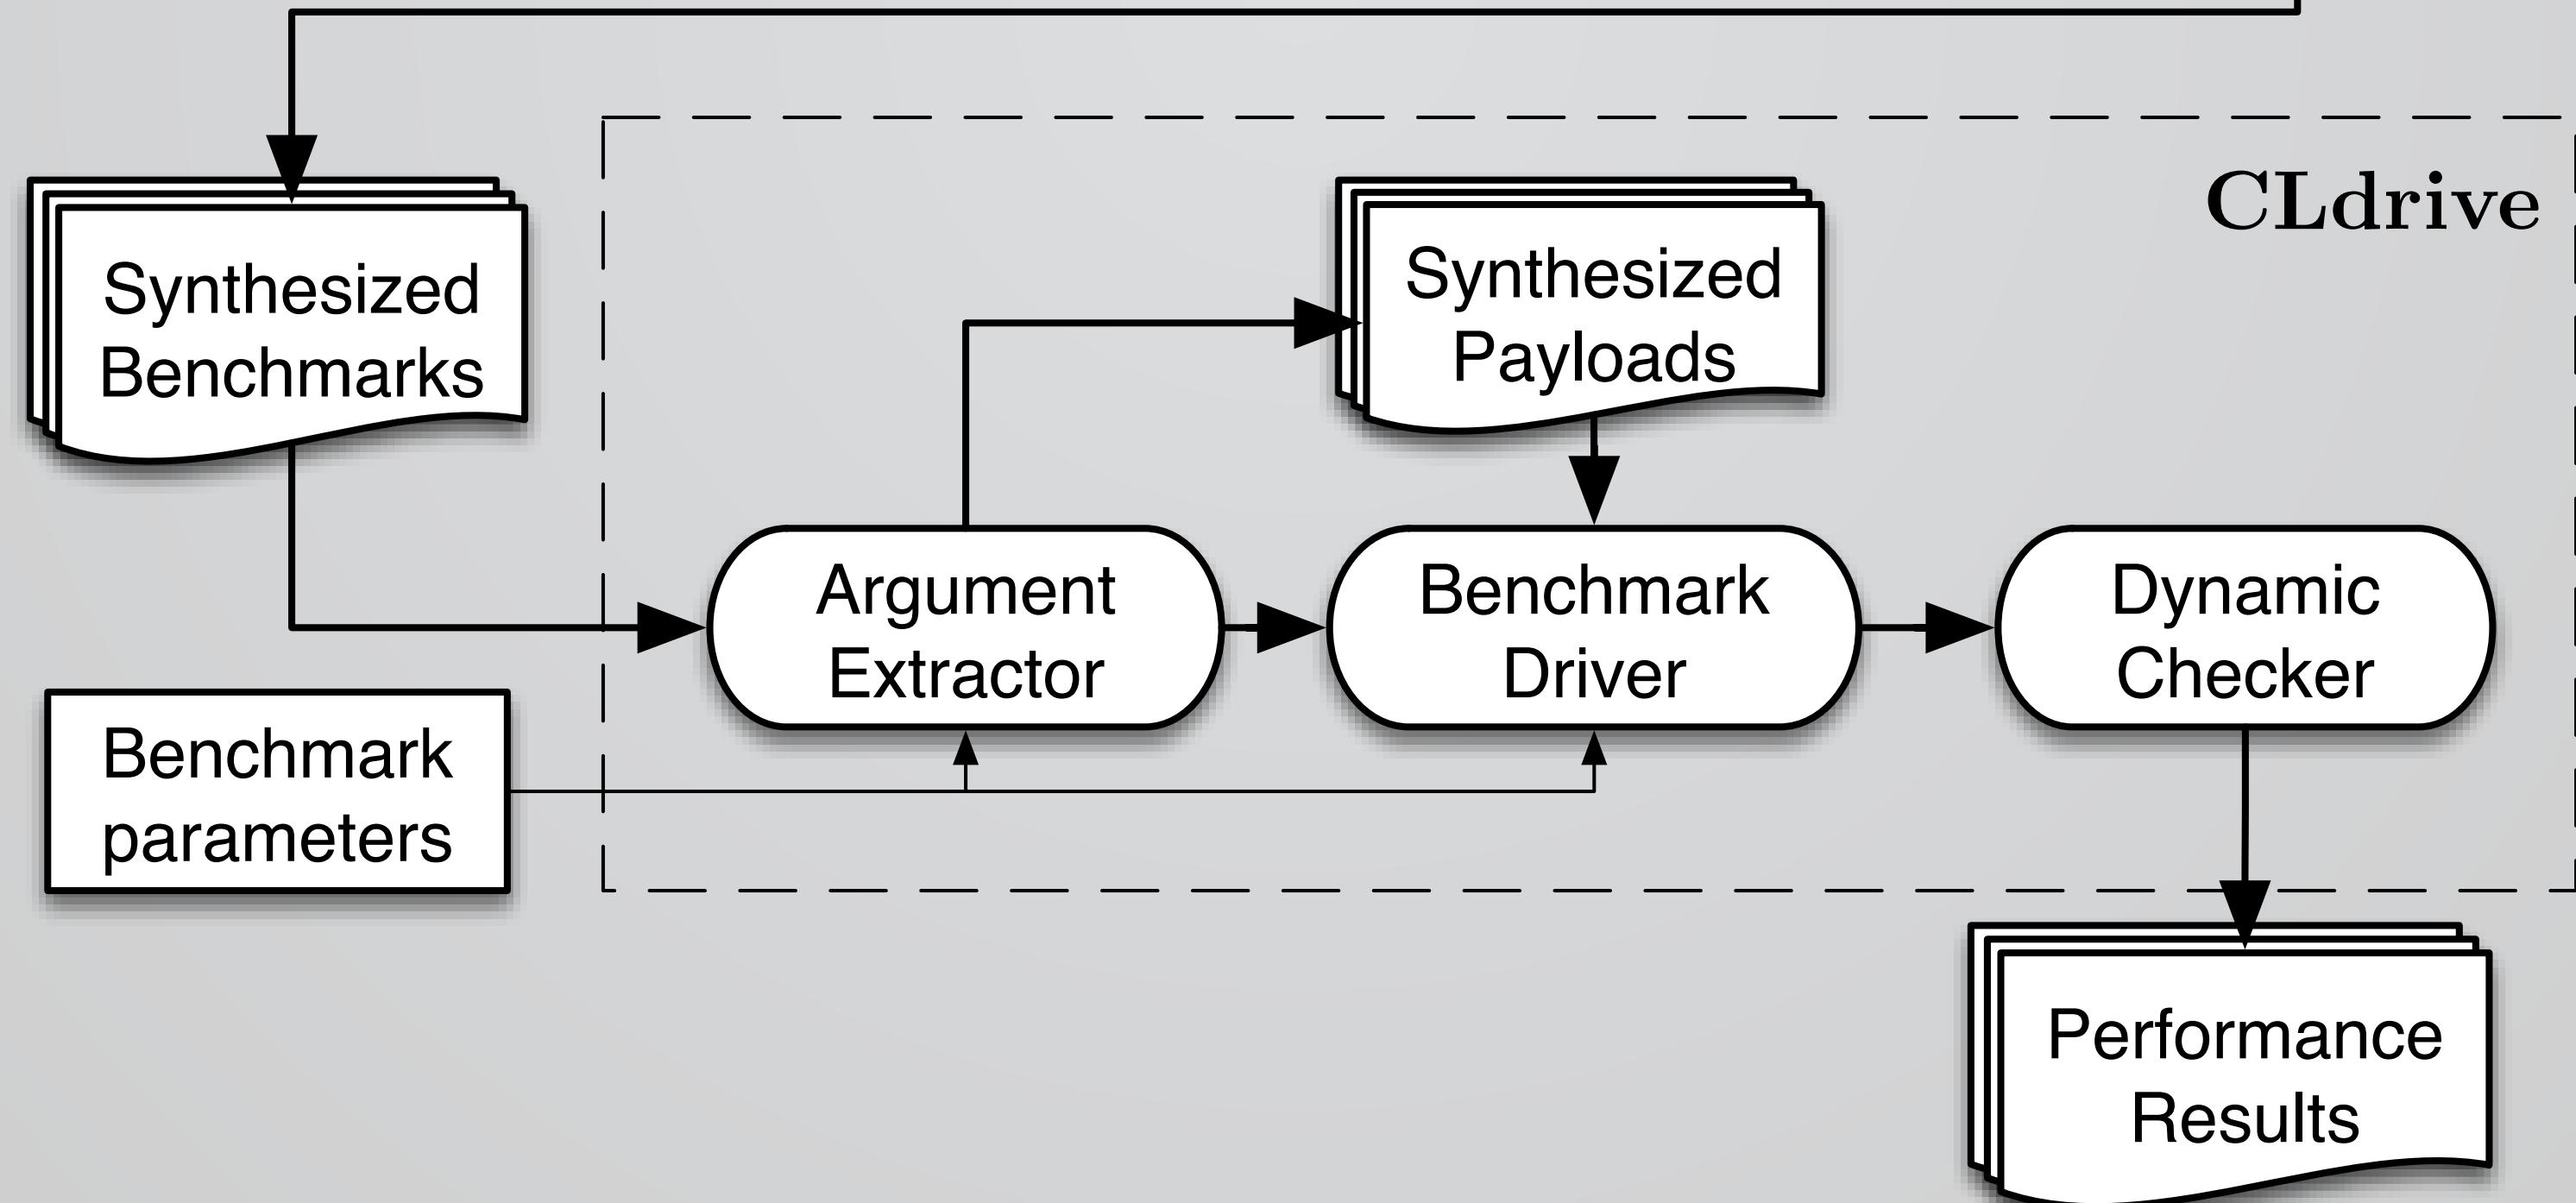
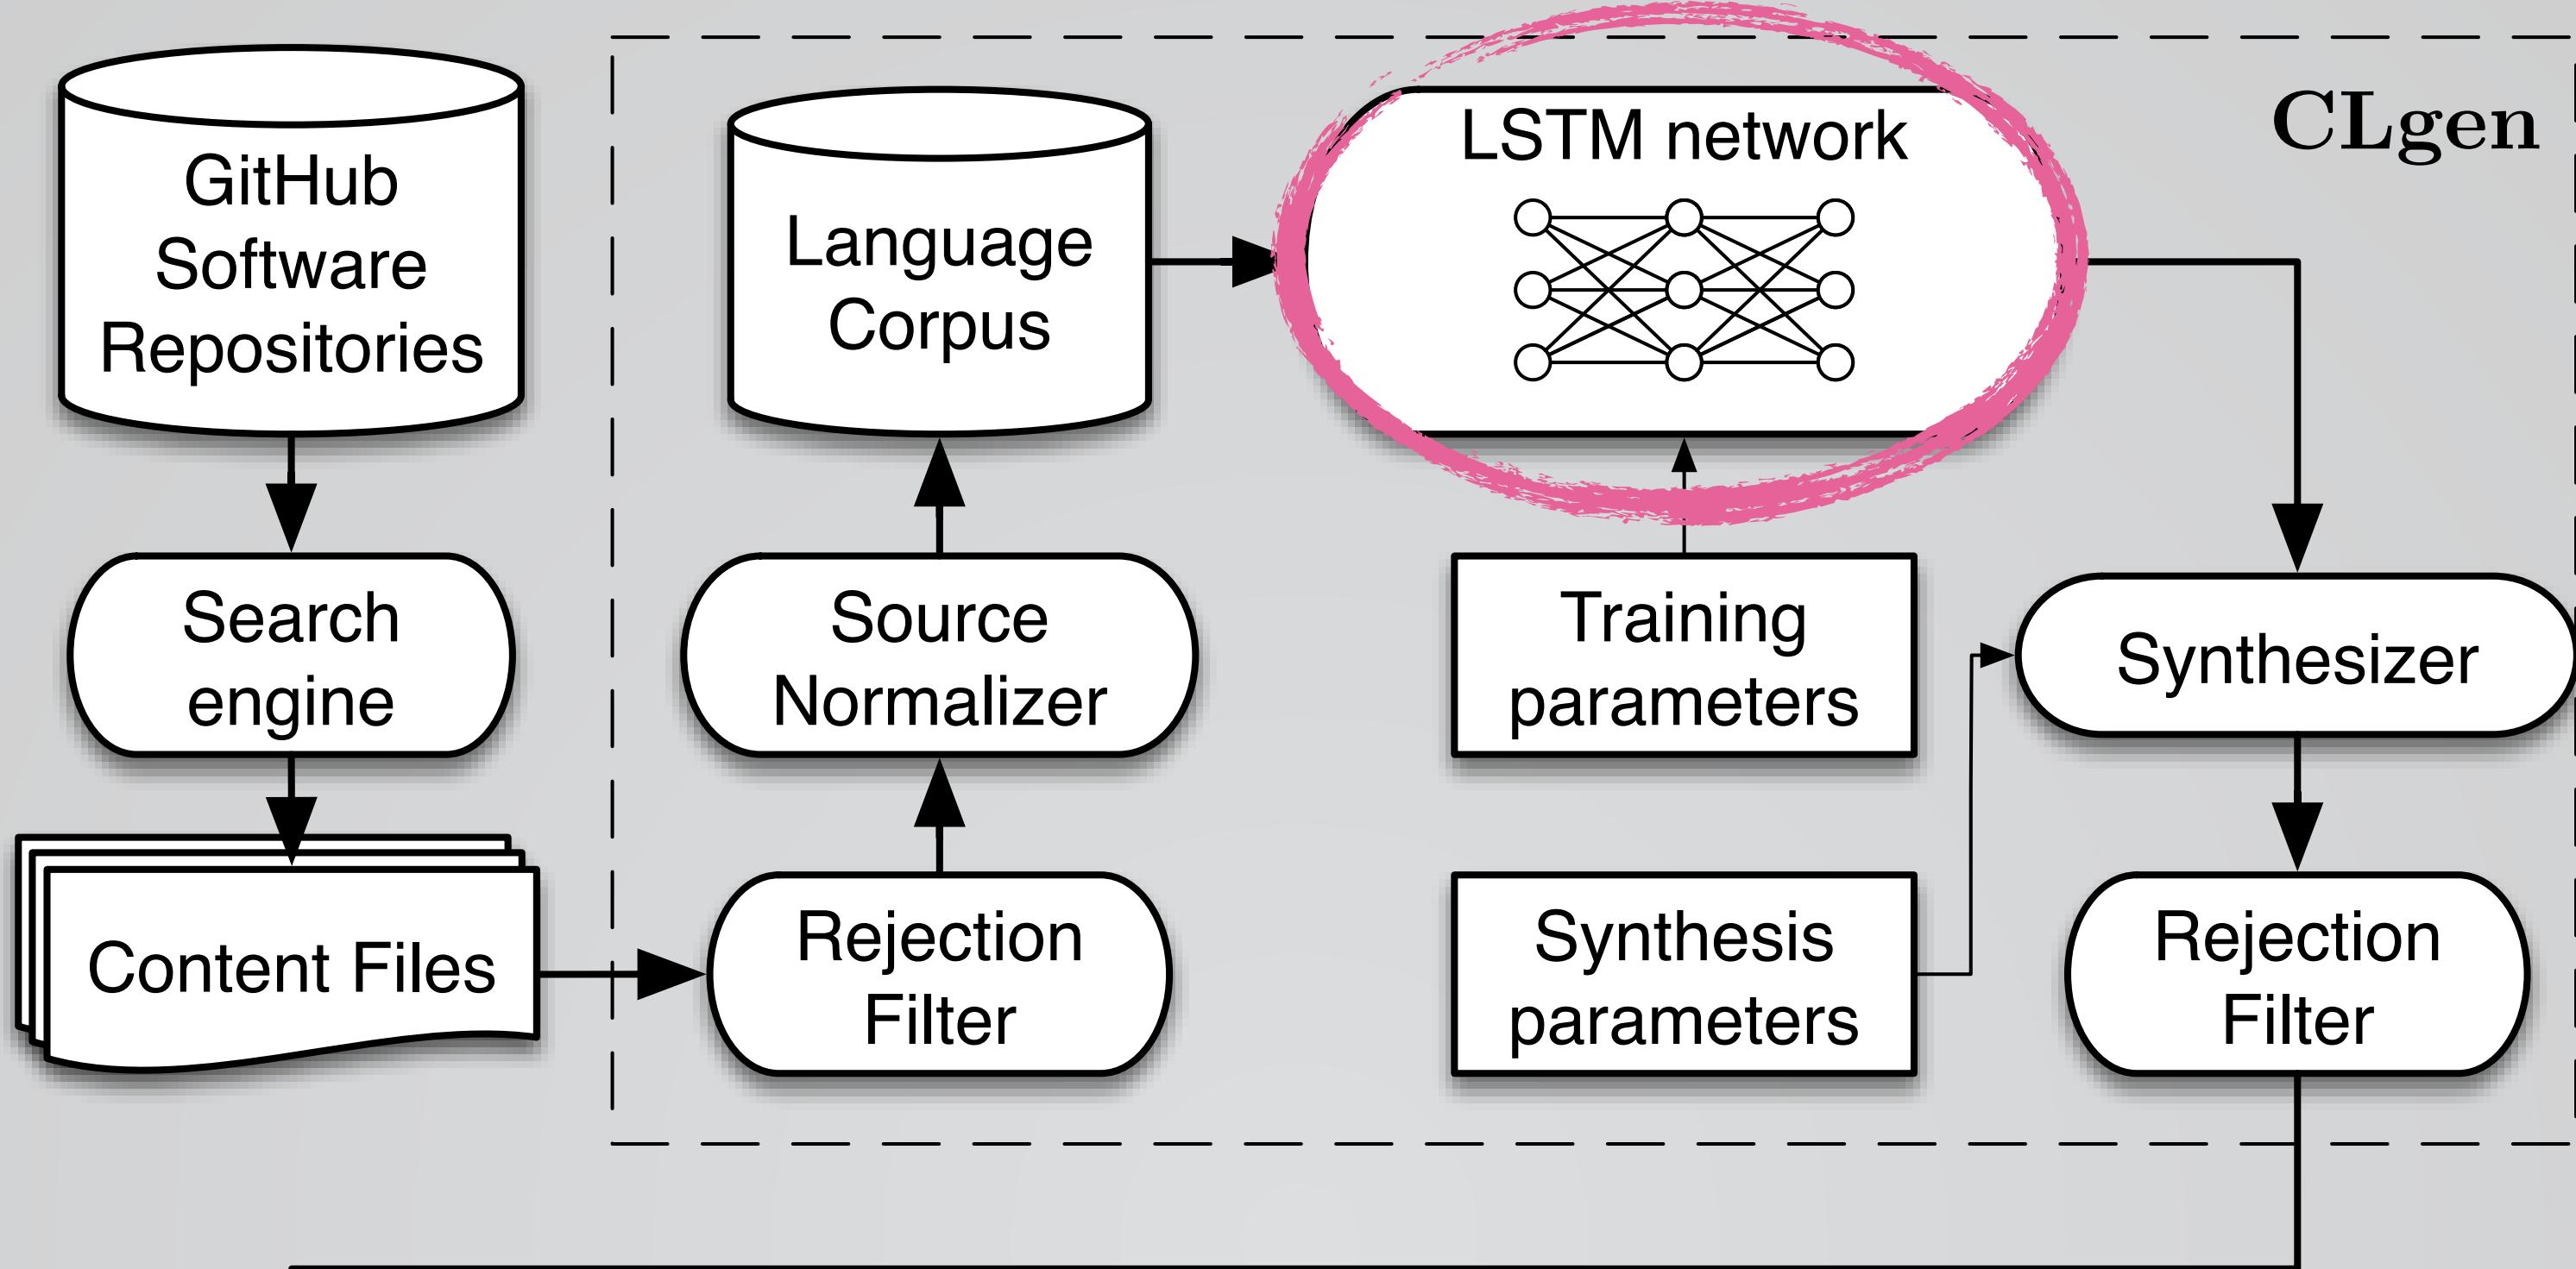
Initial learning rate 0.005

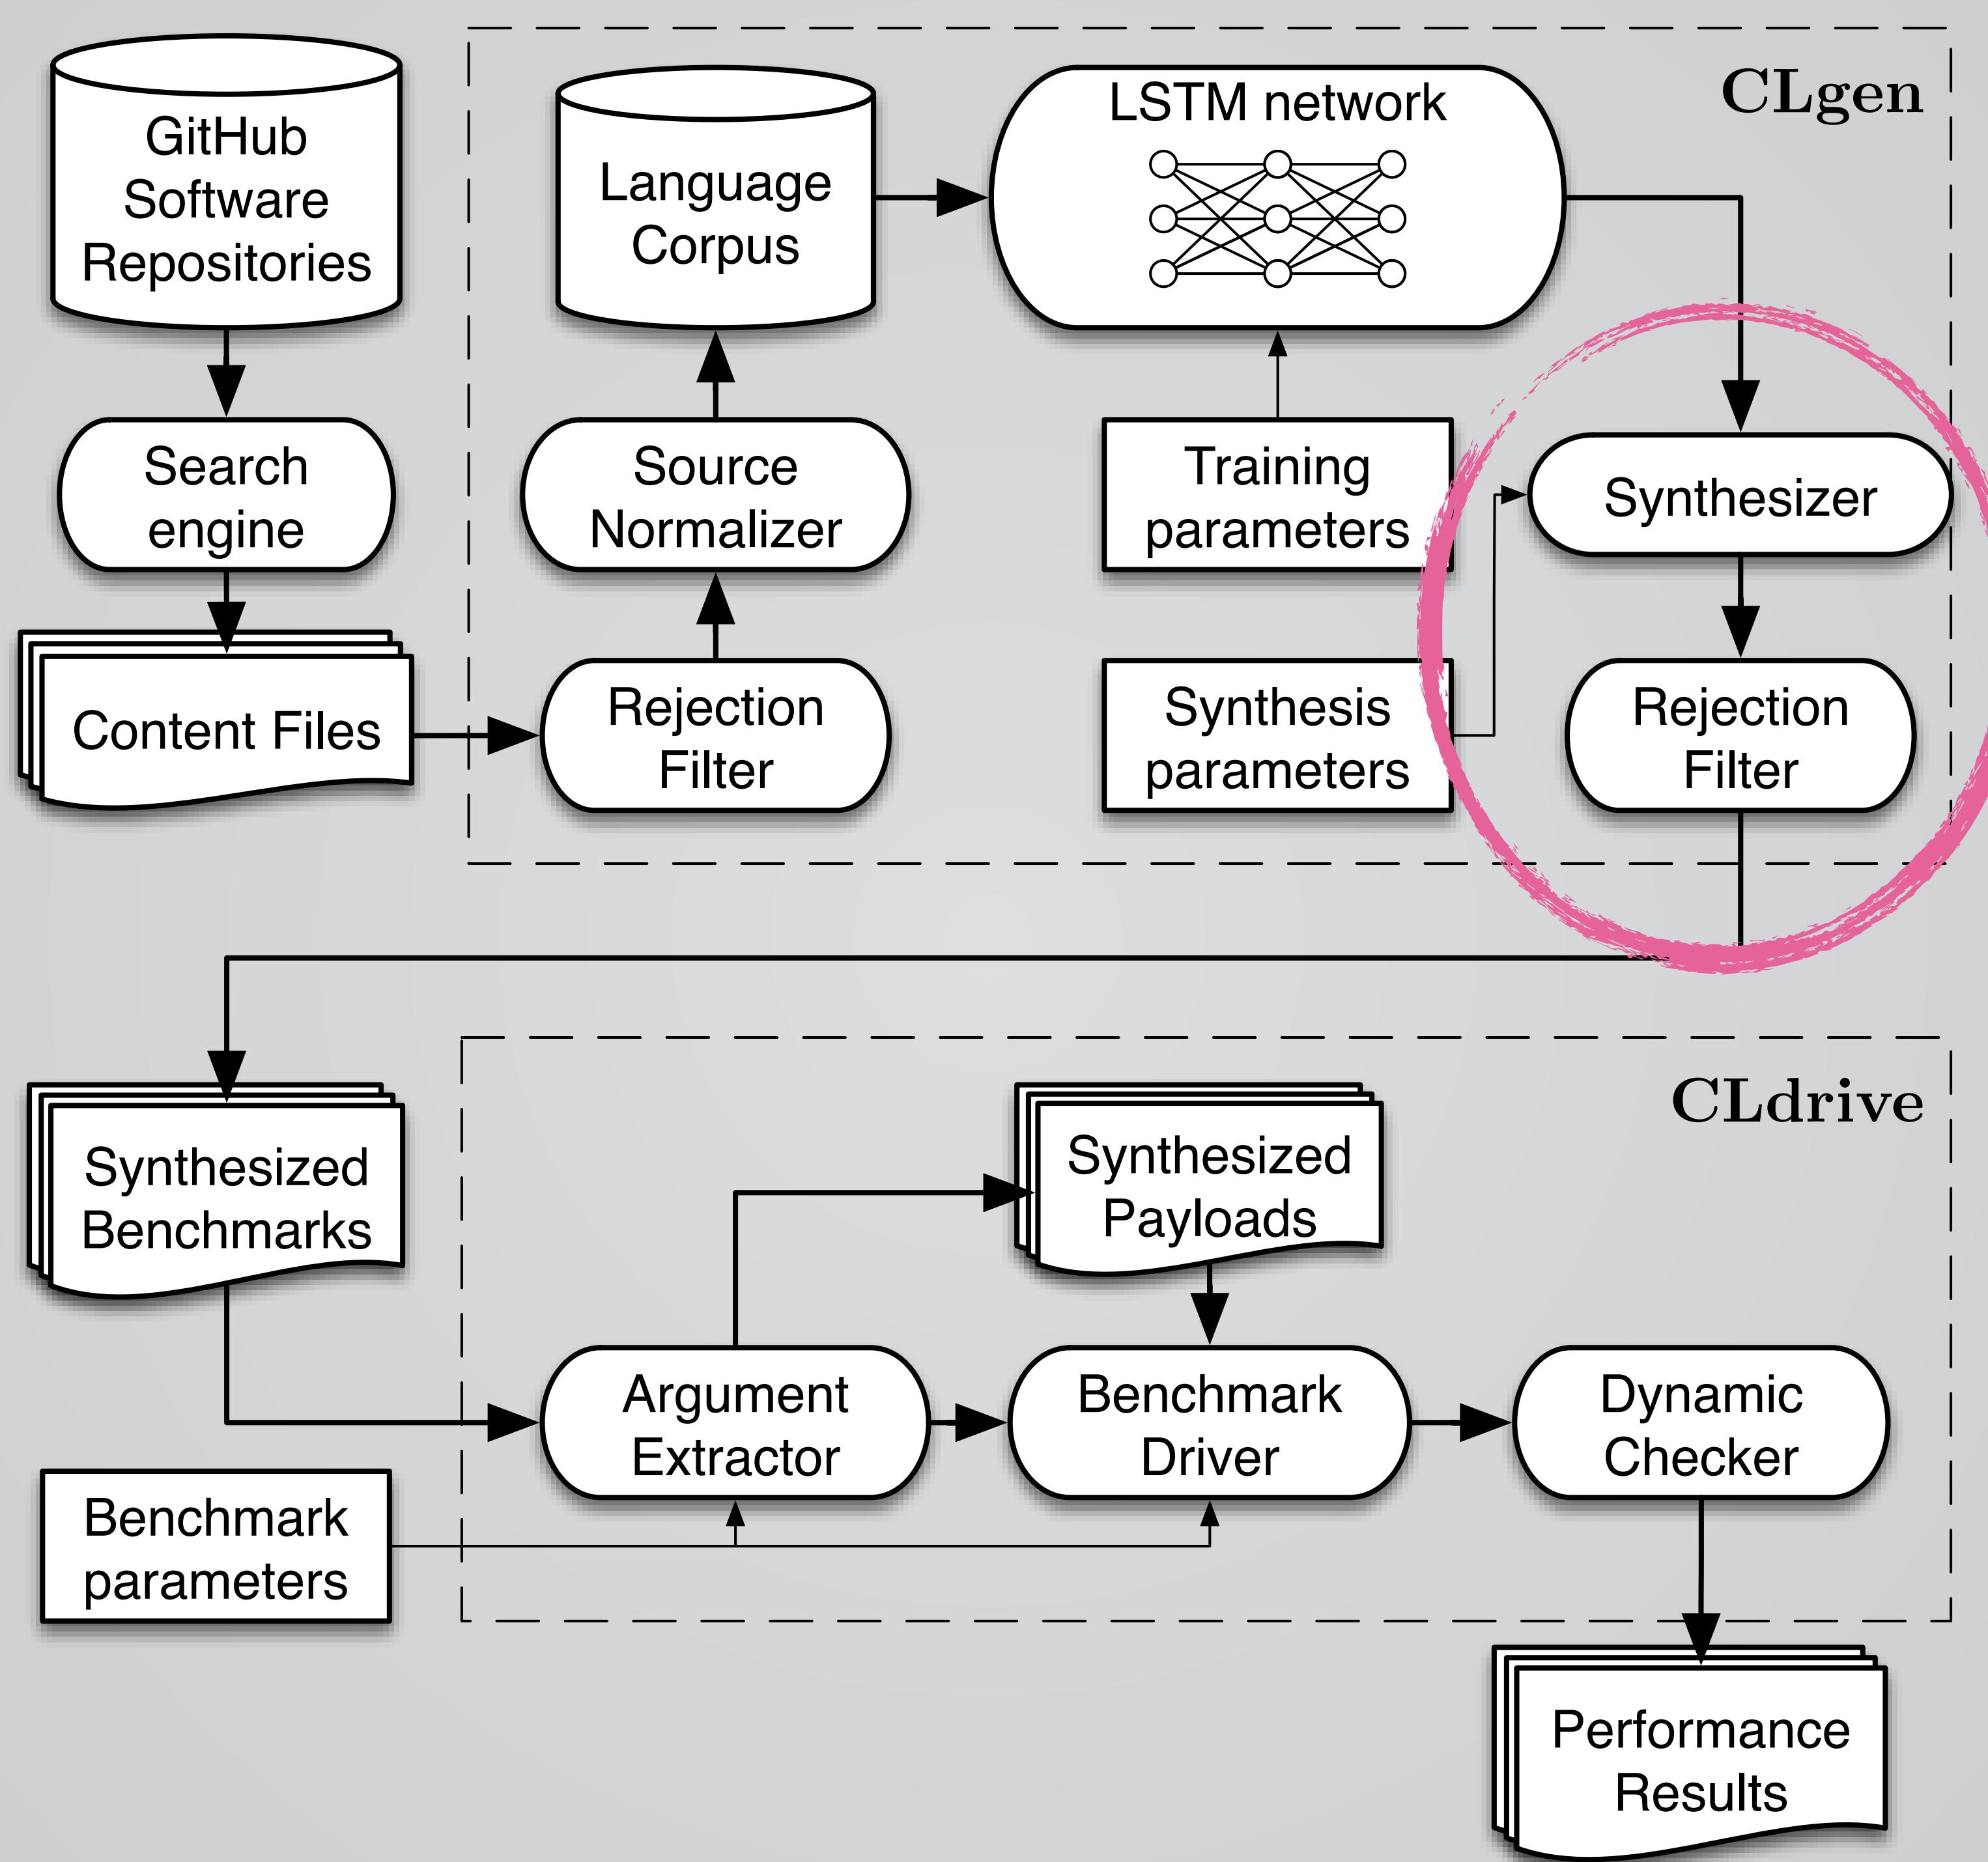
Train for 50 epochs

Learning rate decay every 5 epochs

# Machine Learning







# Kernel Synthesis

```
S = '__kernel void A(__global float* a) {'  
depth = 1  
while depth > 0:  
    c = predict_next_character(S)  
    if c == '{':  
        depth += 1  
    if c == '}':  
        depth -= 1  
    S += c  
  
return S
```

# DEMO

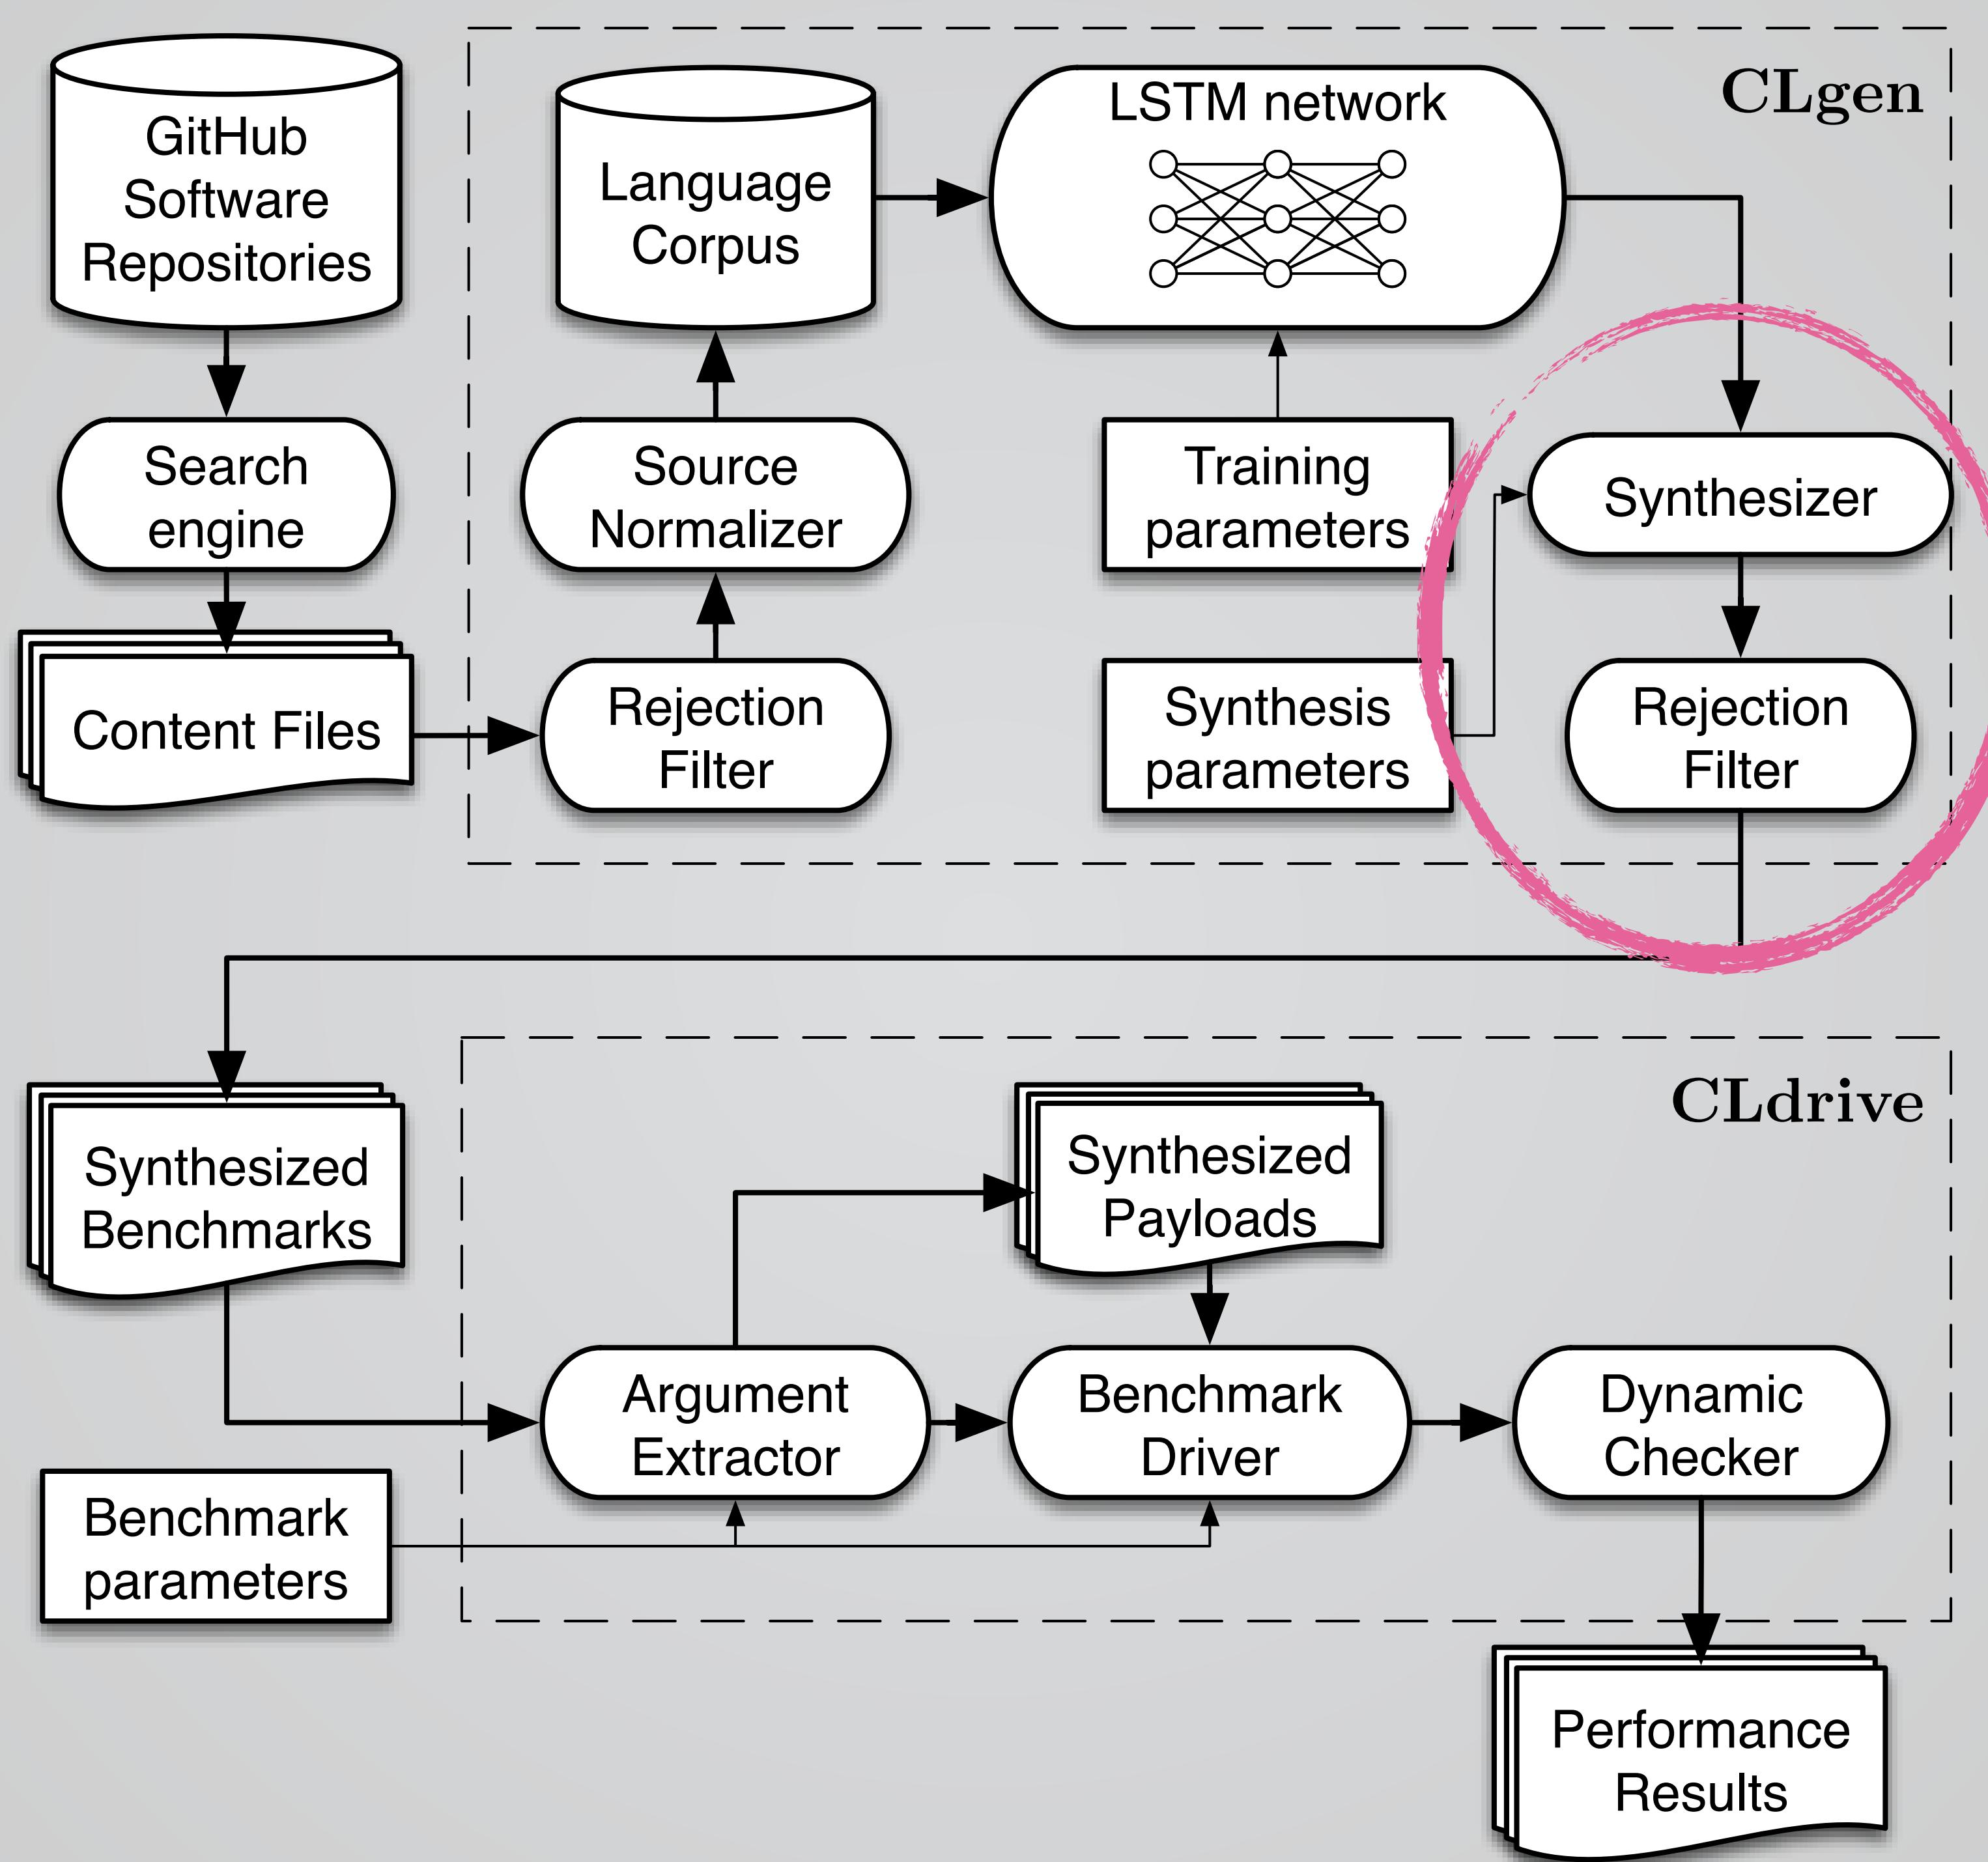
*(you had to be there)*

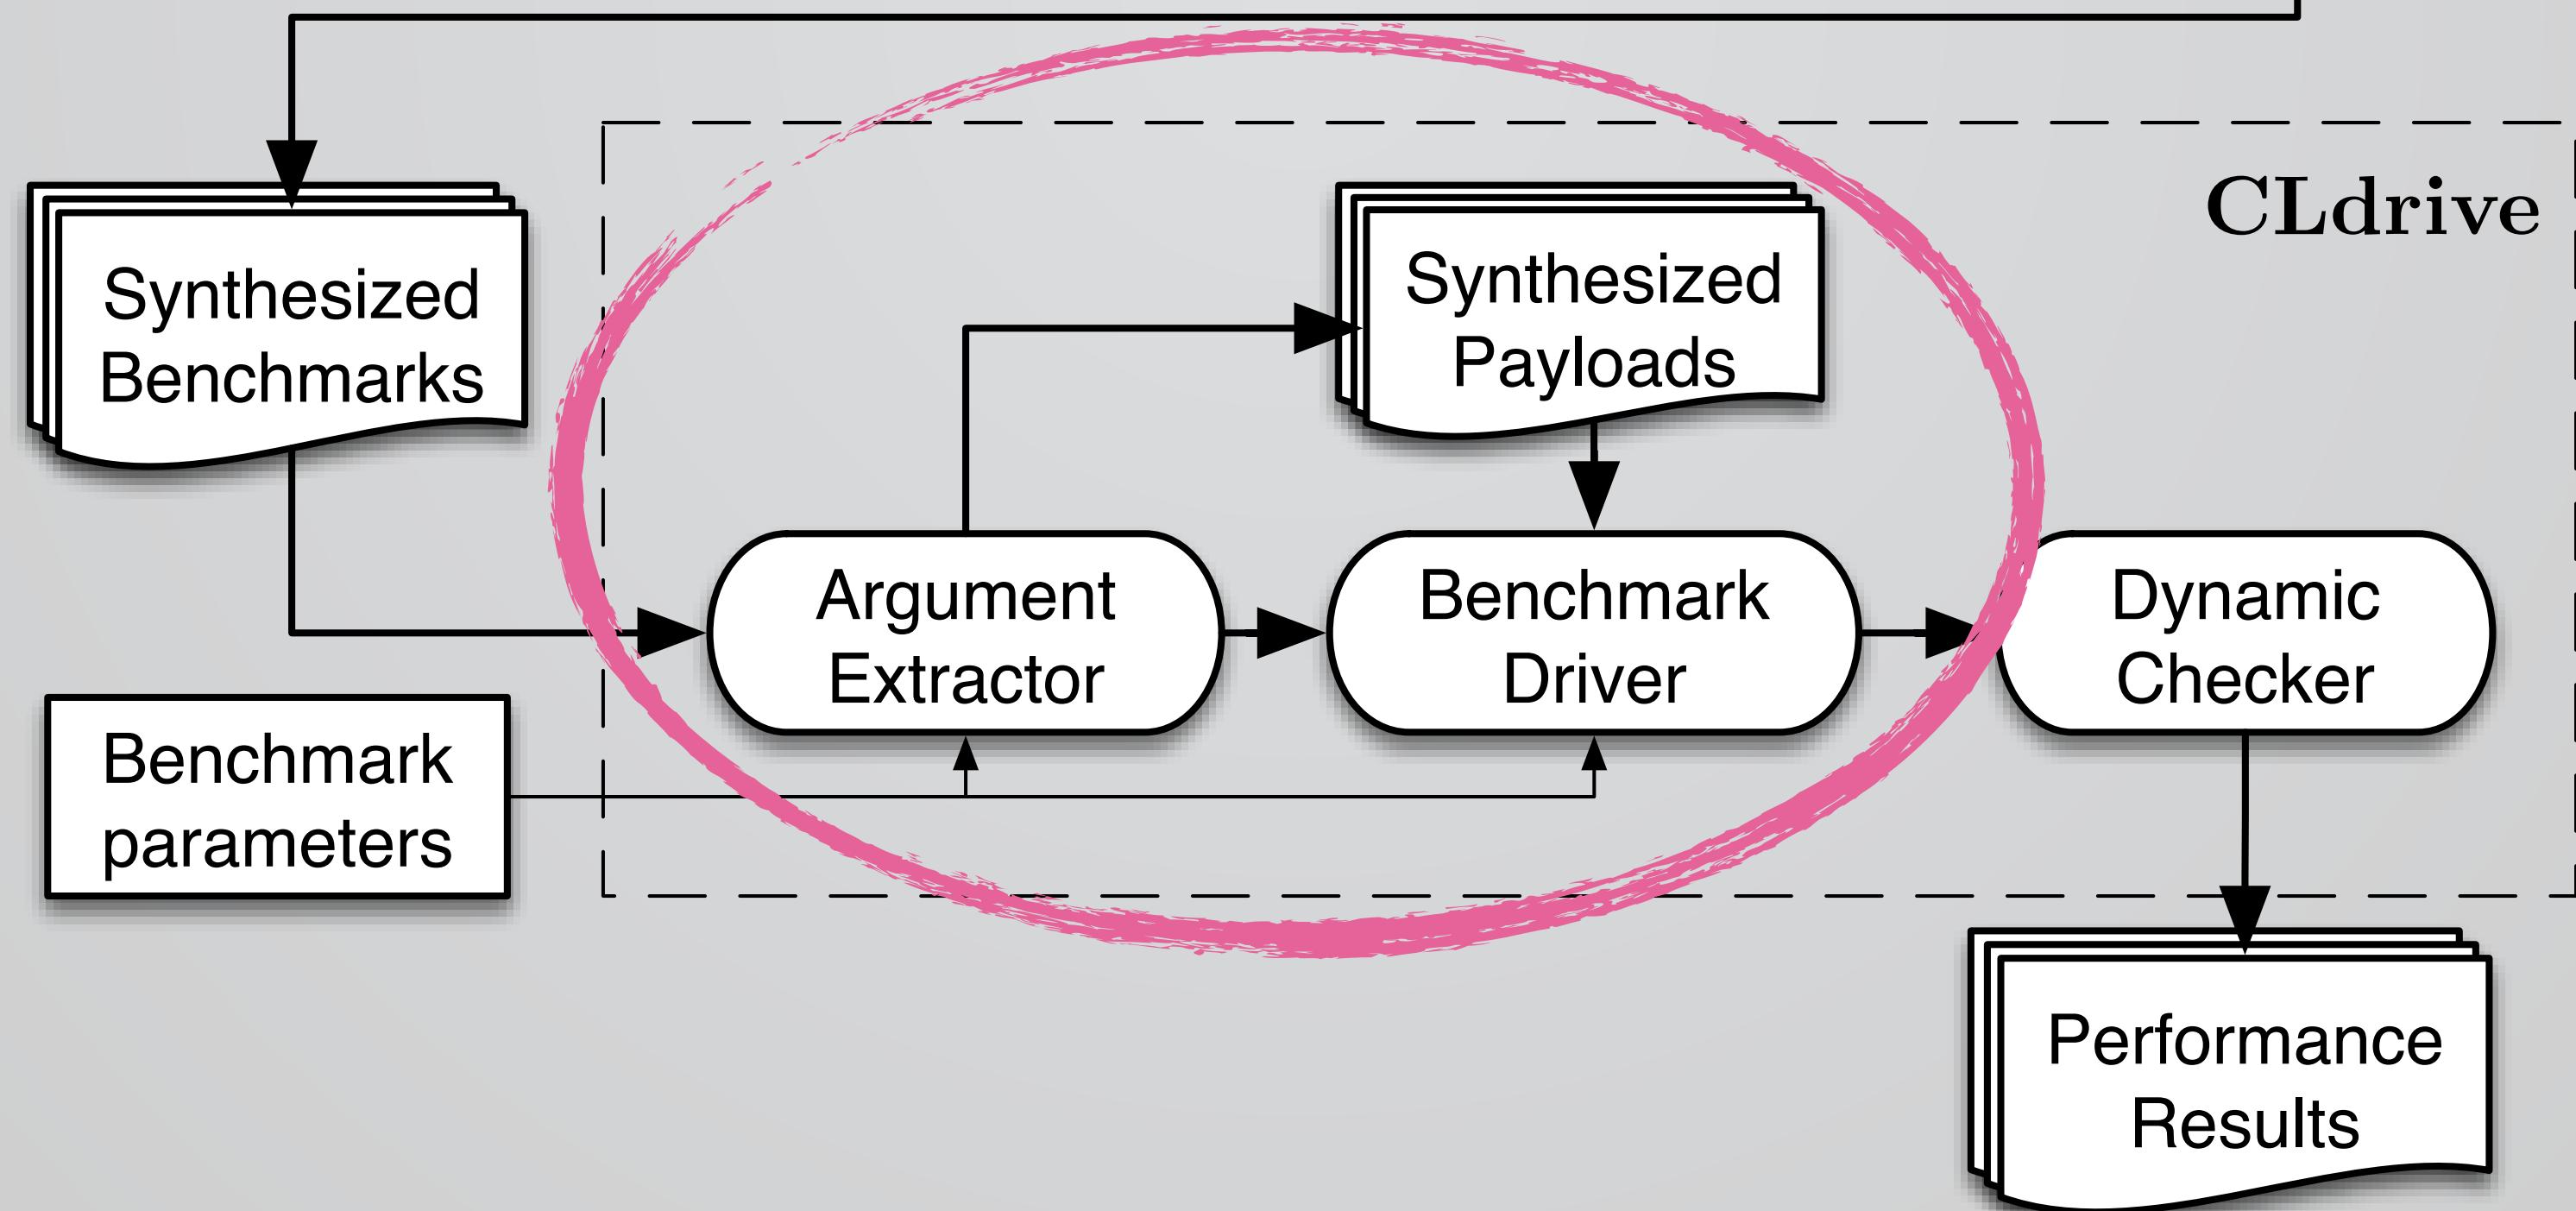
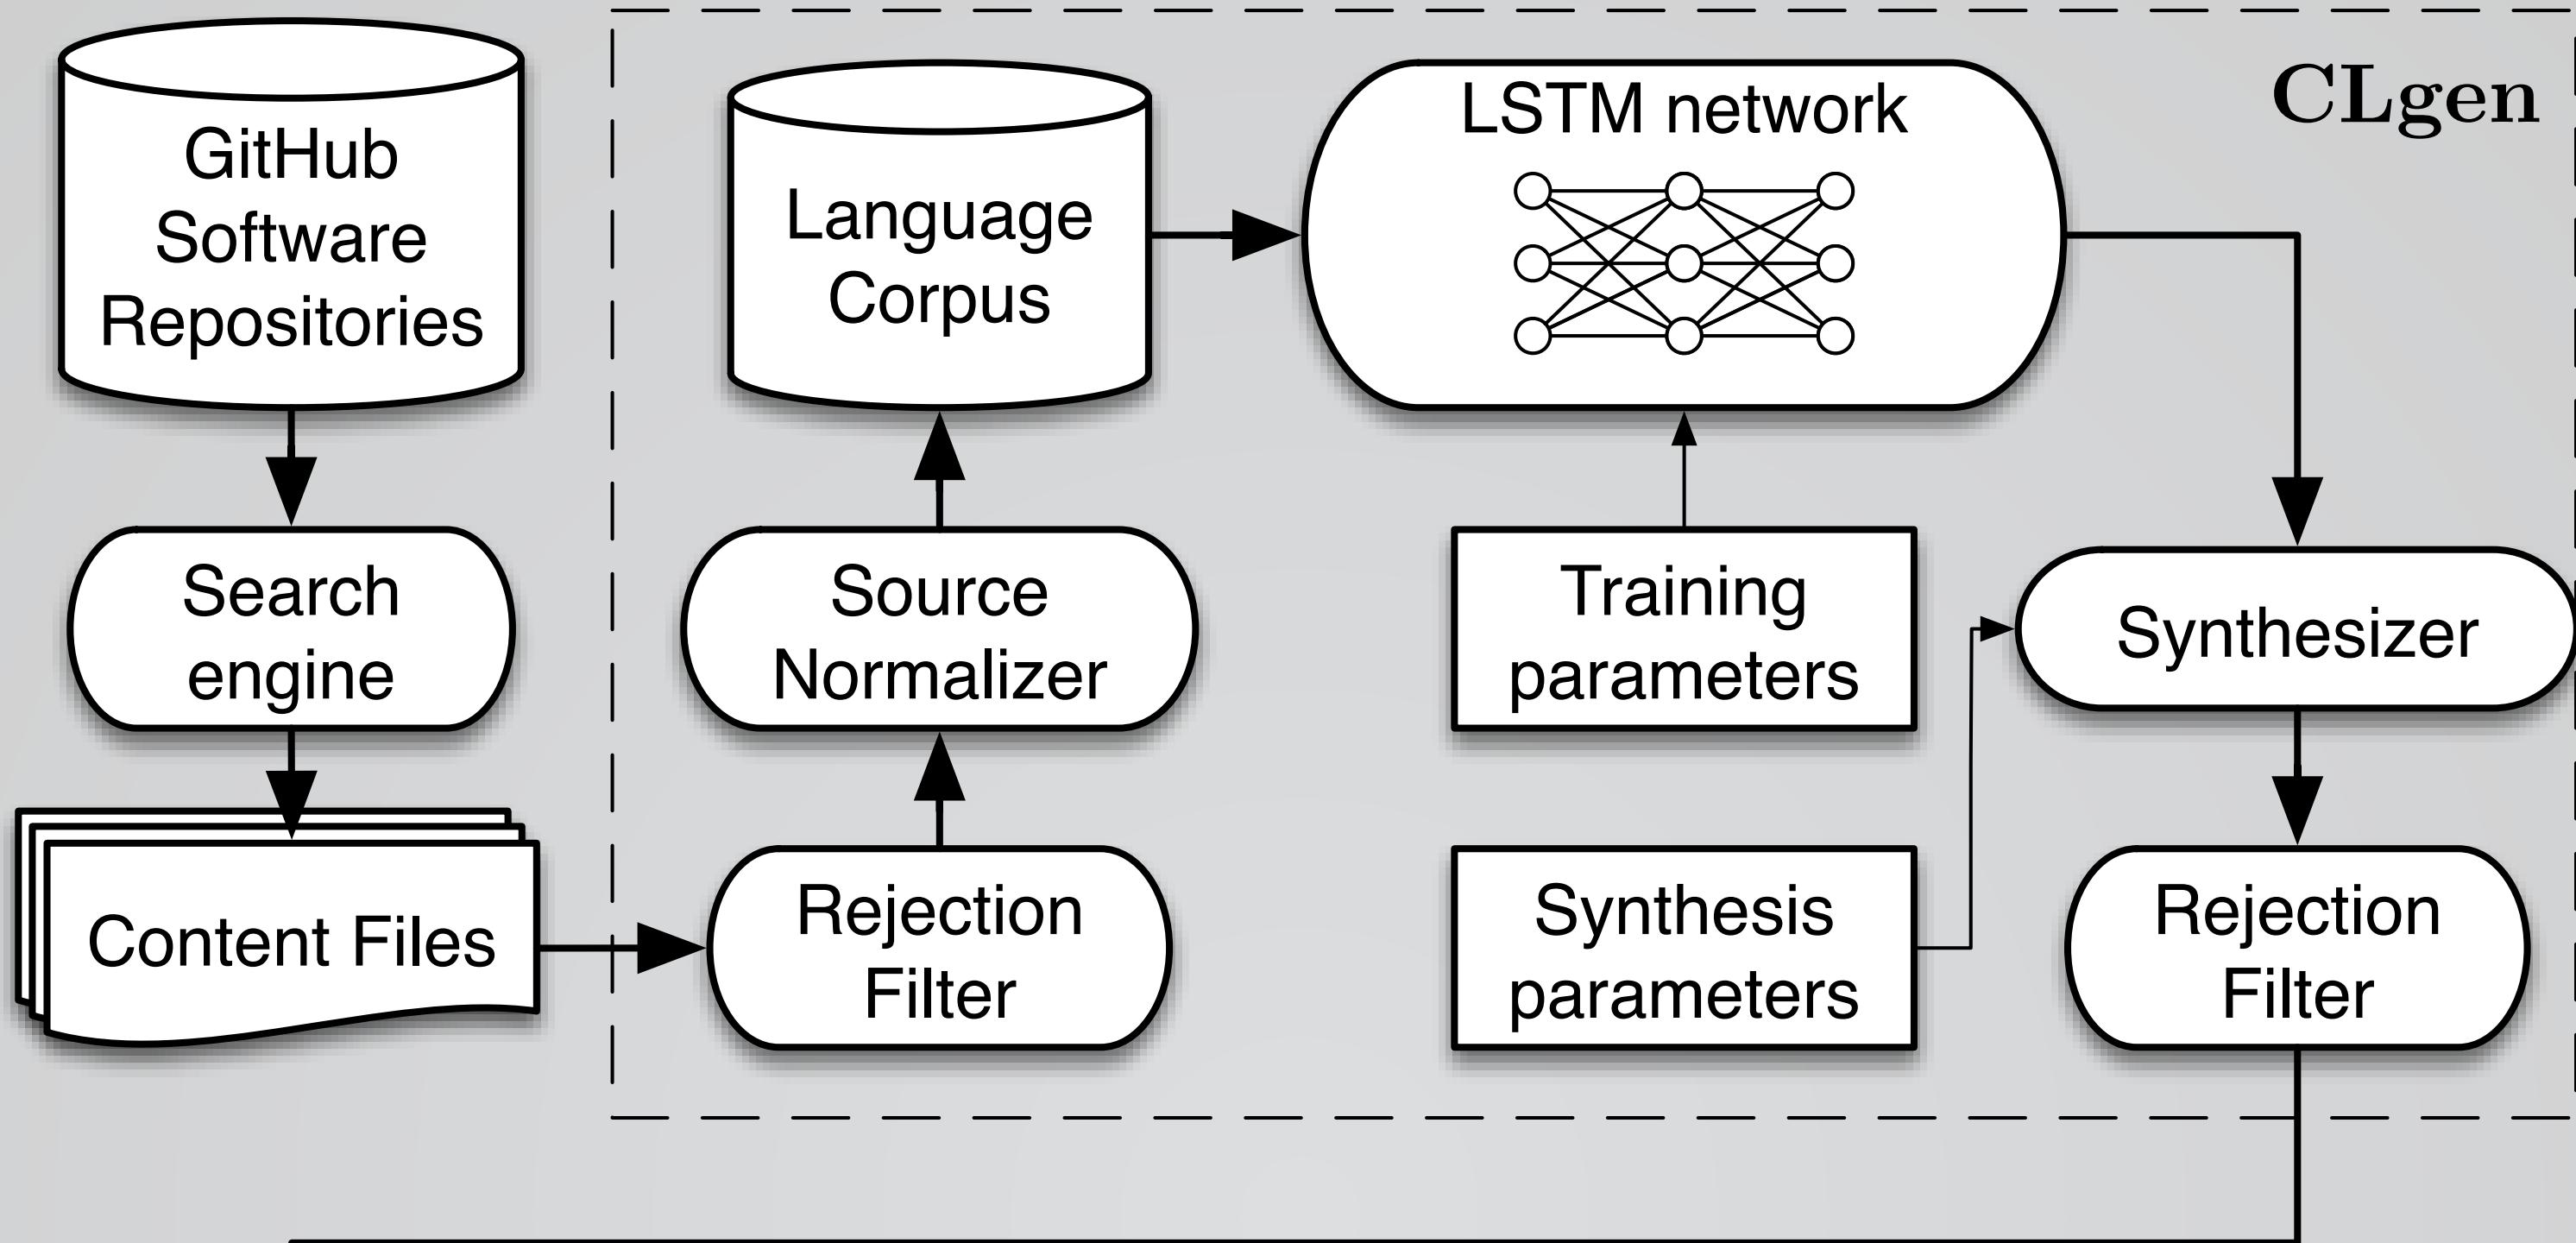
```
__kernel void A(__global float* a,
                __global float* b,
                __global float* c,
                const int d) {
    int e = get_global_id(0);
    float f = 0.0;
    for (int g = 0; g < d; g++) {
        c[g] = 0.0f;
    }
    barrier(1);

    a[get_global_id(0)] = 2*b[get_global_id(0)];
}
```

```
_kernel void A(__global float* a,
           __global float* b,
           __global float* c,
           const int d) {
    int e = get_global_id(0);
    if (e >= d) {
        return;
    }
    c[e] = a[e] + b[e] + 2 * a[e] + b[e] + 4;
}
```

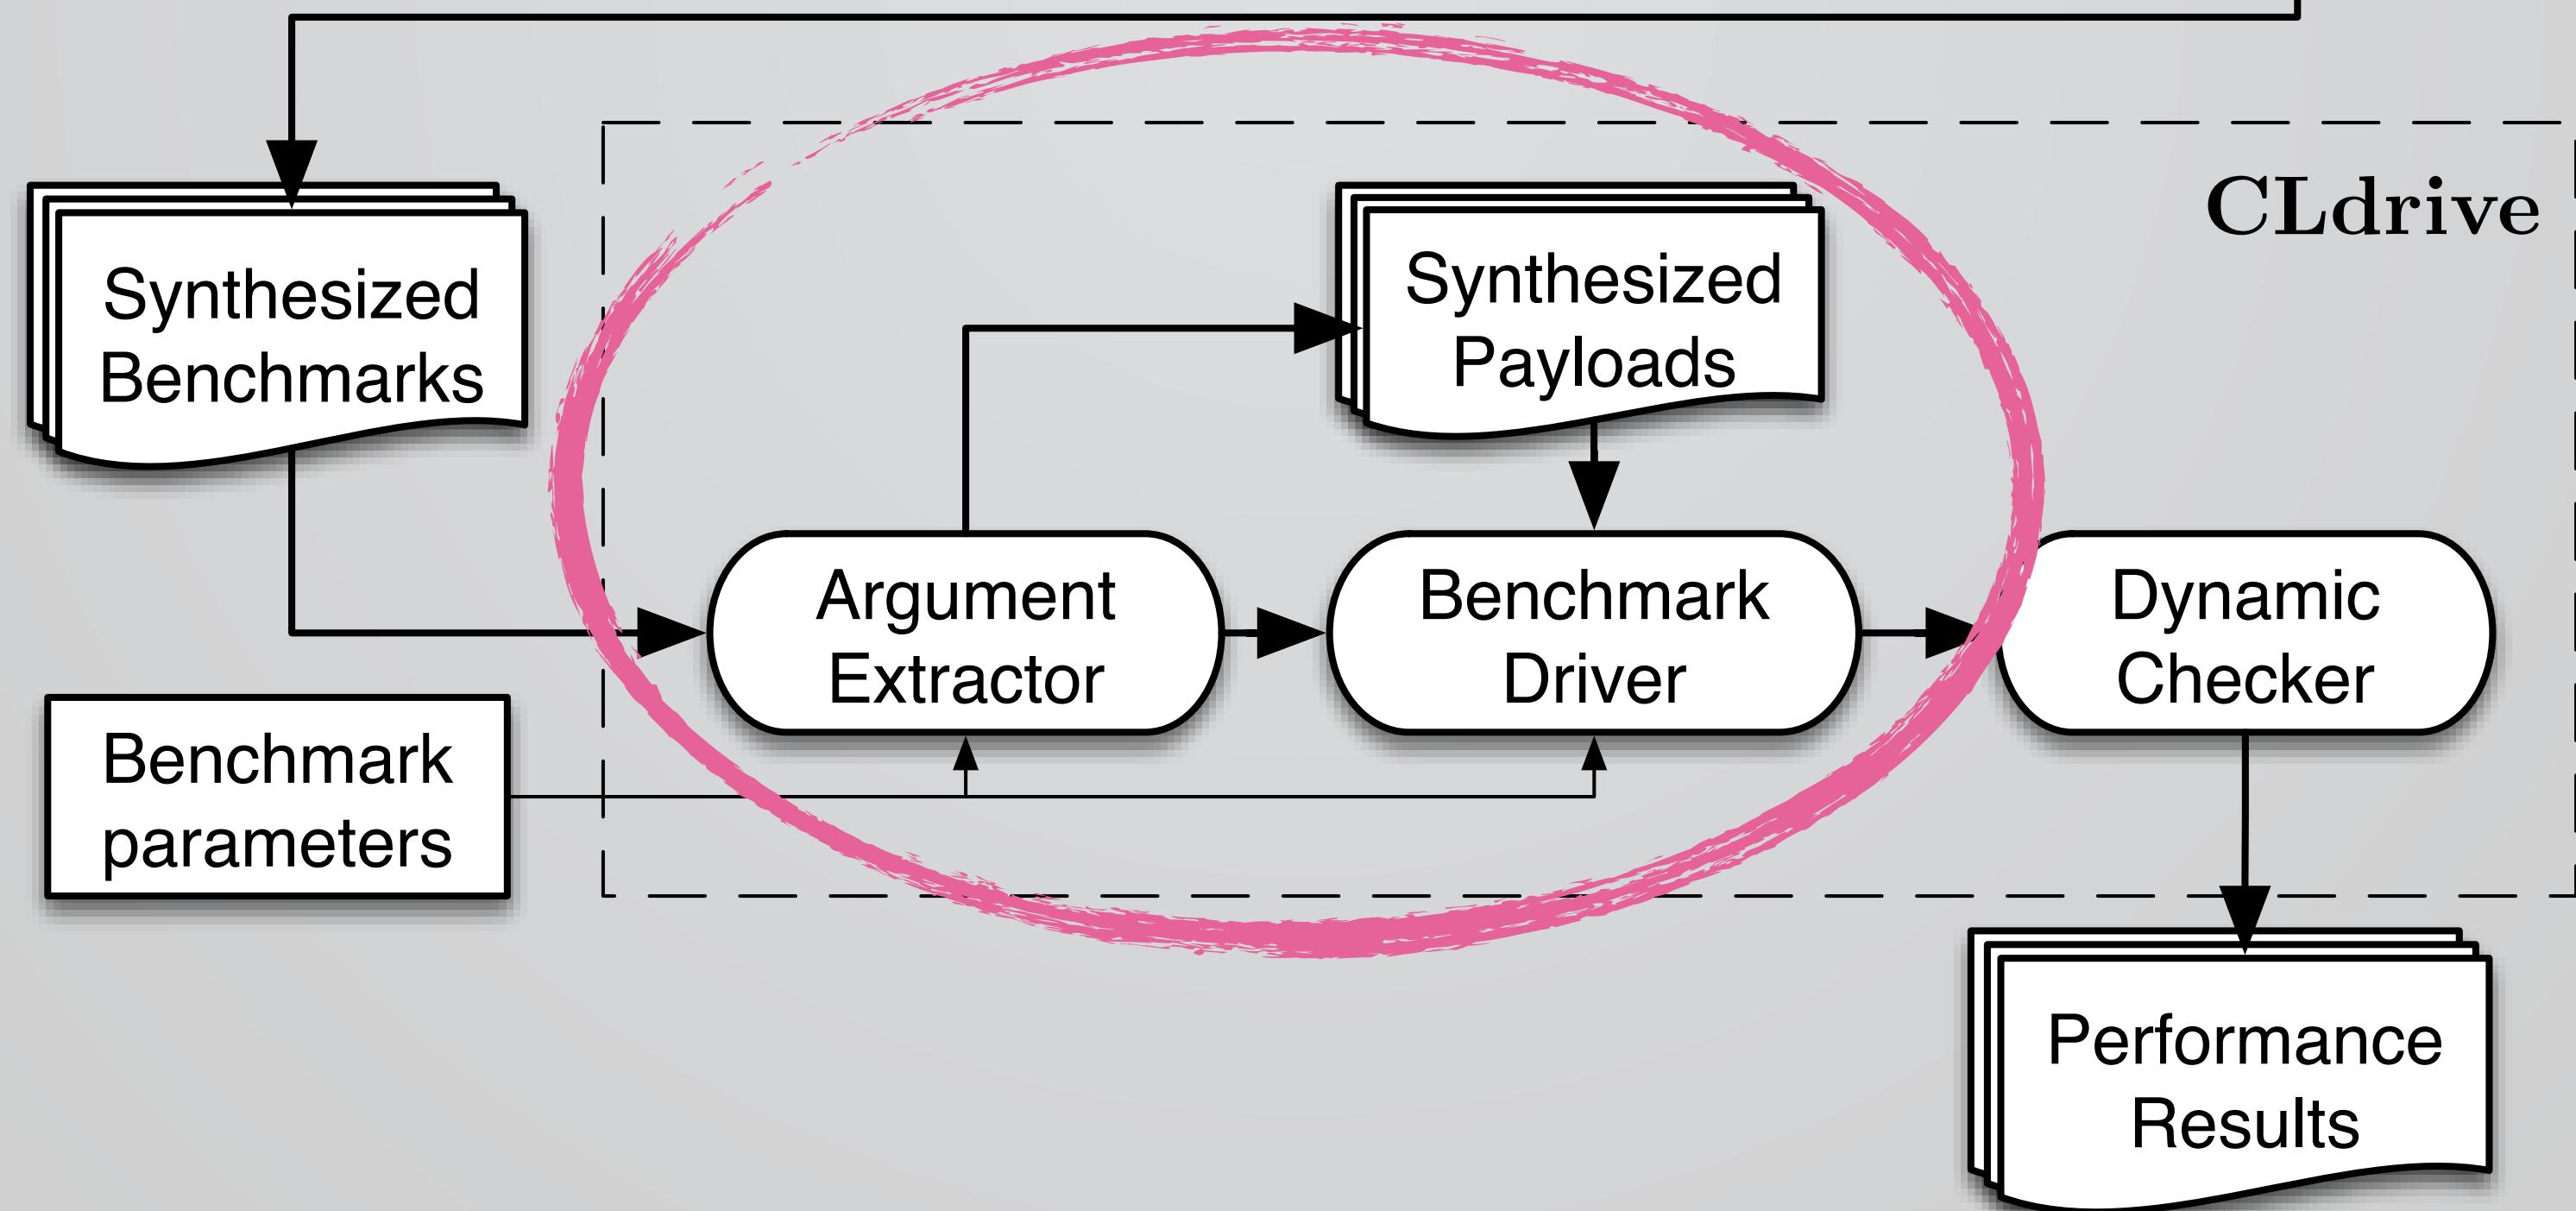
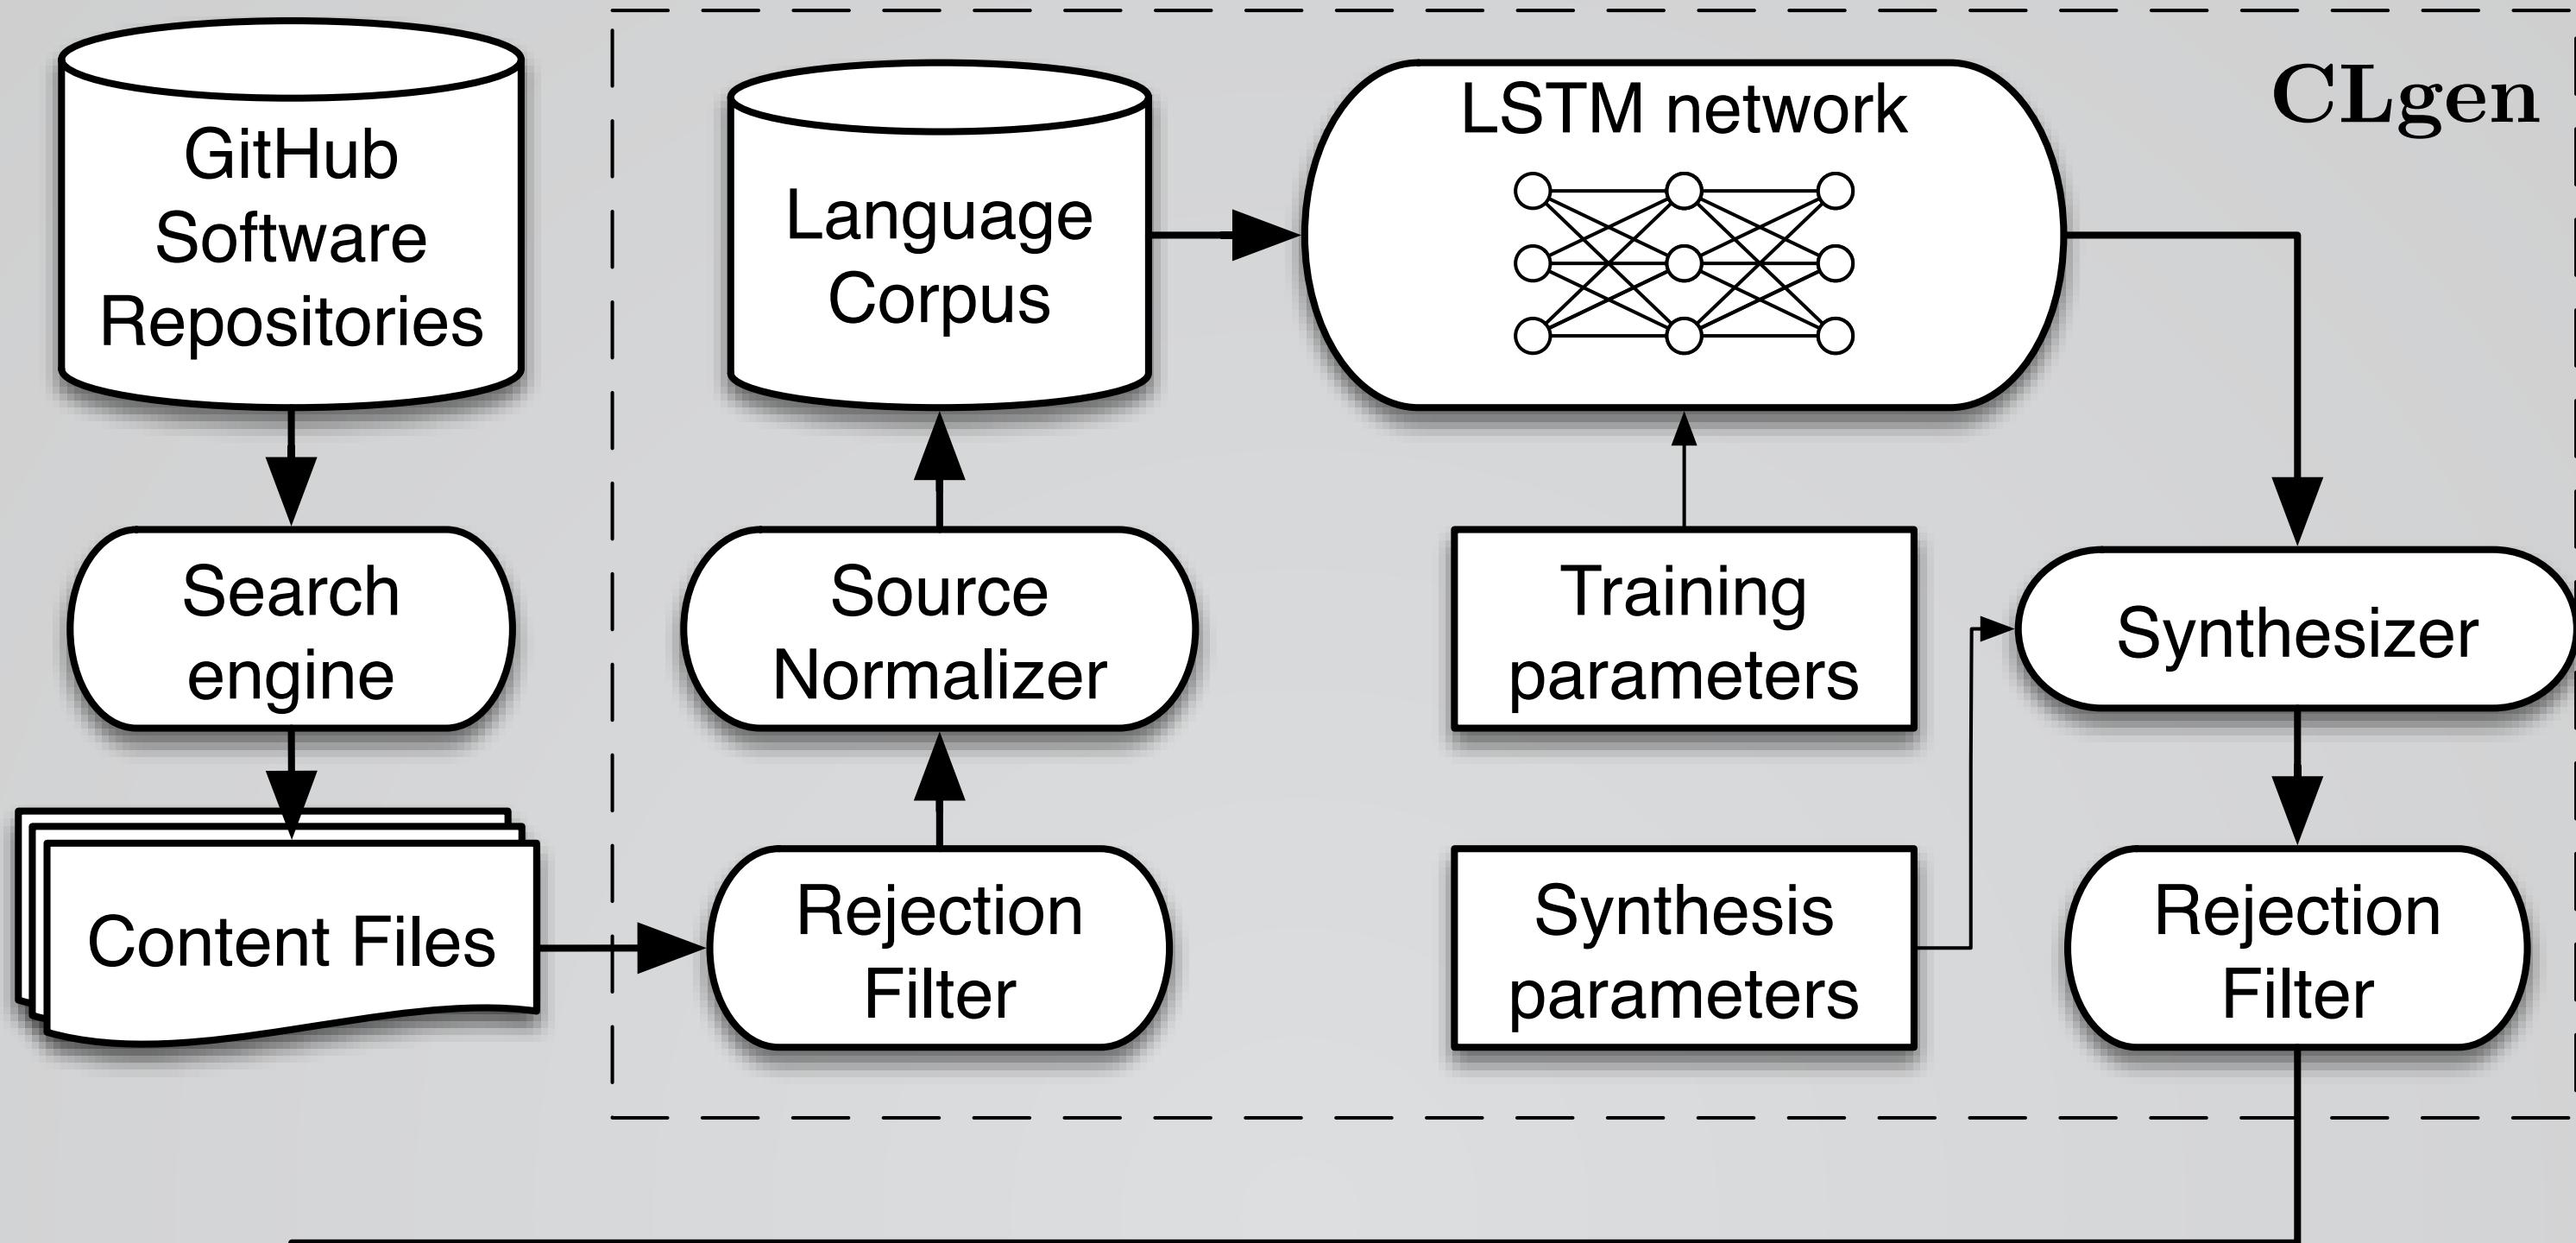
```
__kernel void A(__global float* a,
                __global float* b,
                __global float* c,
                const int d) {
    unsigned int e = get_global_id(0);
    float16 f = (float16)(0.0);
    for (unsigned int g = 0; g < d; g++) {
        float16 h = a[g];
        f.s0 += h.s0;
        f.s1 += h.s1;
        /* snip ... */
        f.sE += h.sE;
        f.sF += h.sF;
    }
    b[e] = f.s0 + f.s1 + f.s2 + f.s3 + f.s4 +
           f.s5 + f.s6 + f.s7 + f.s8 + f.s9 + f.sA +
           f.sB + f.sC + f.sD + f.sE + f.sF;
}
```

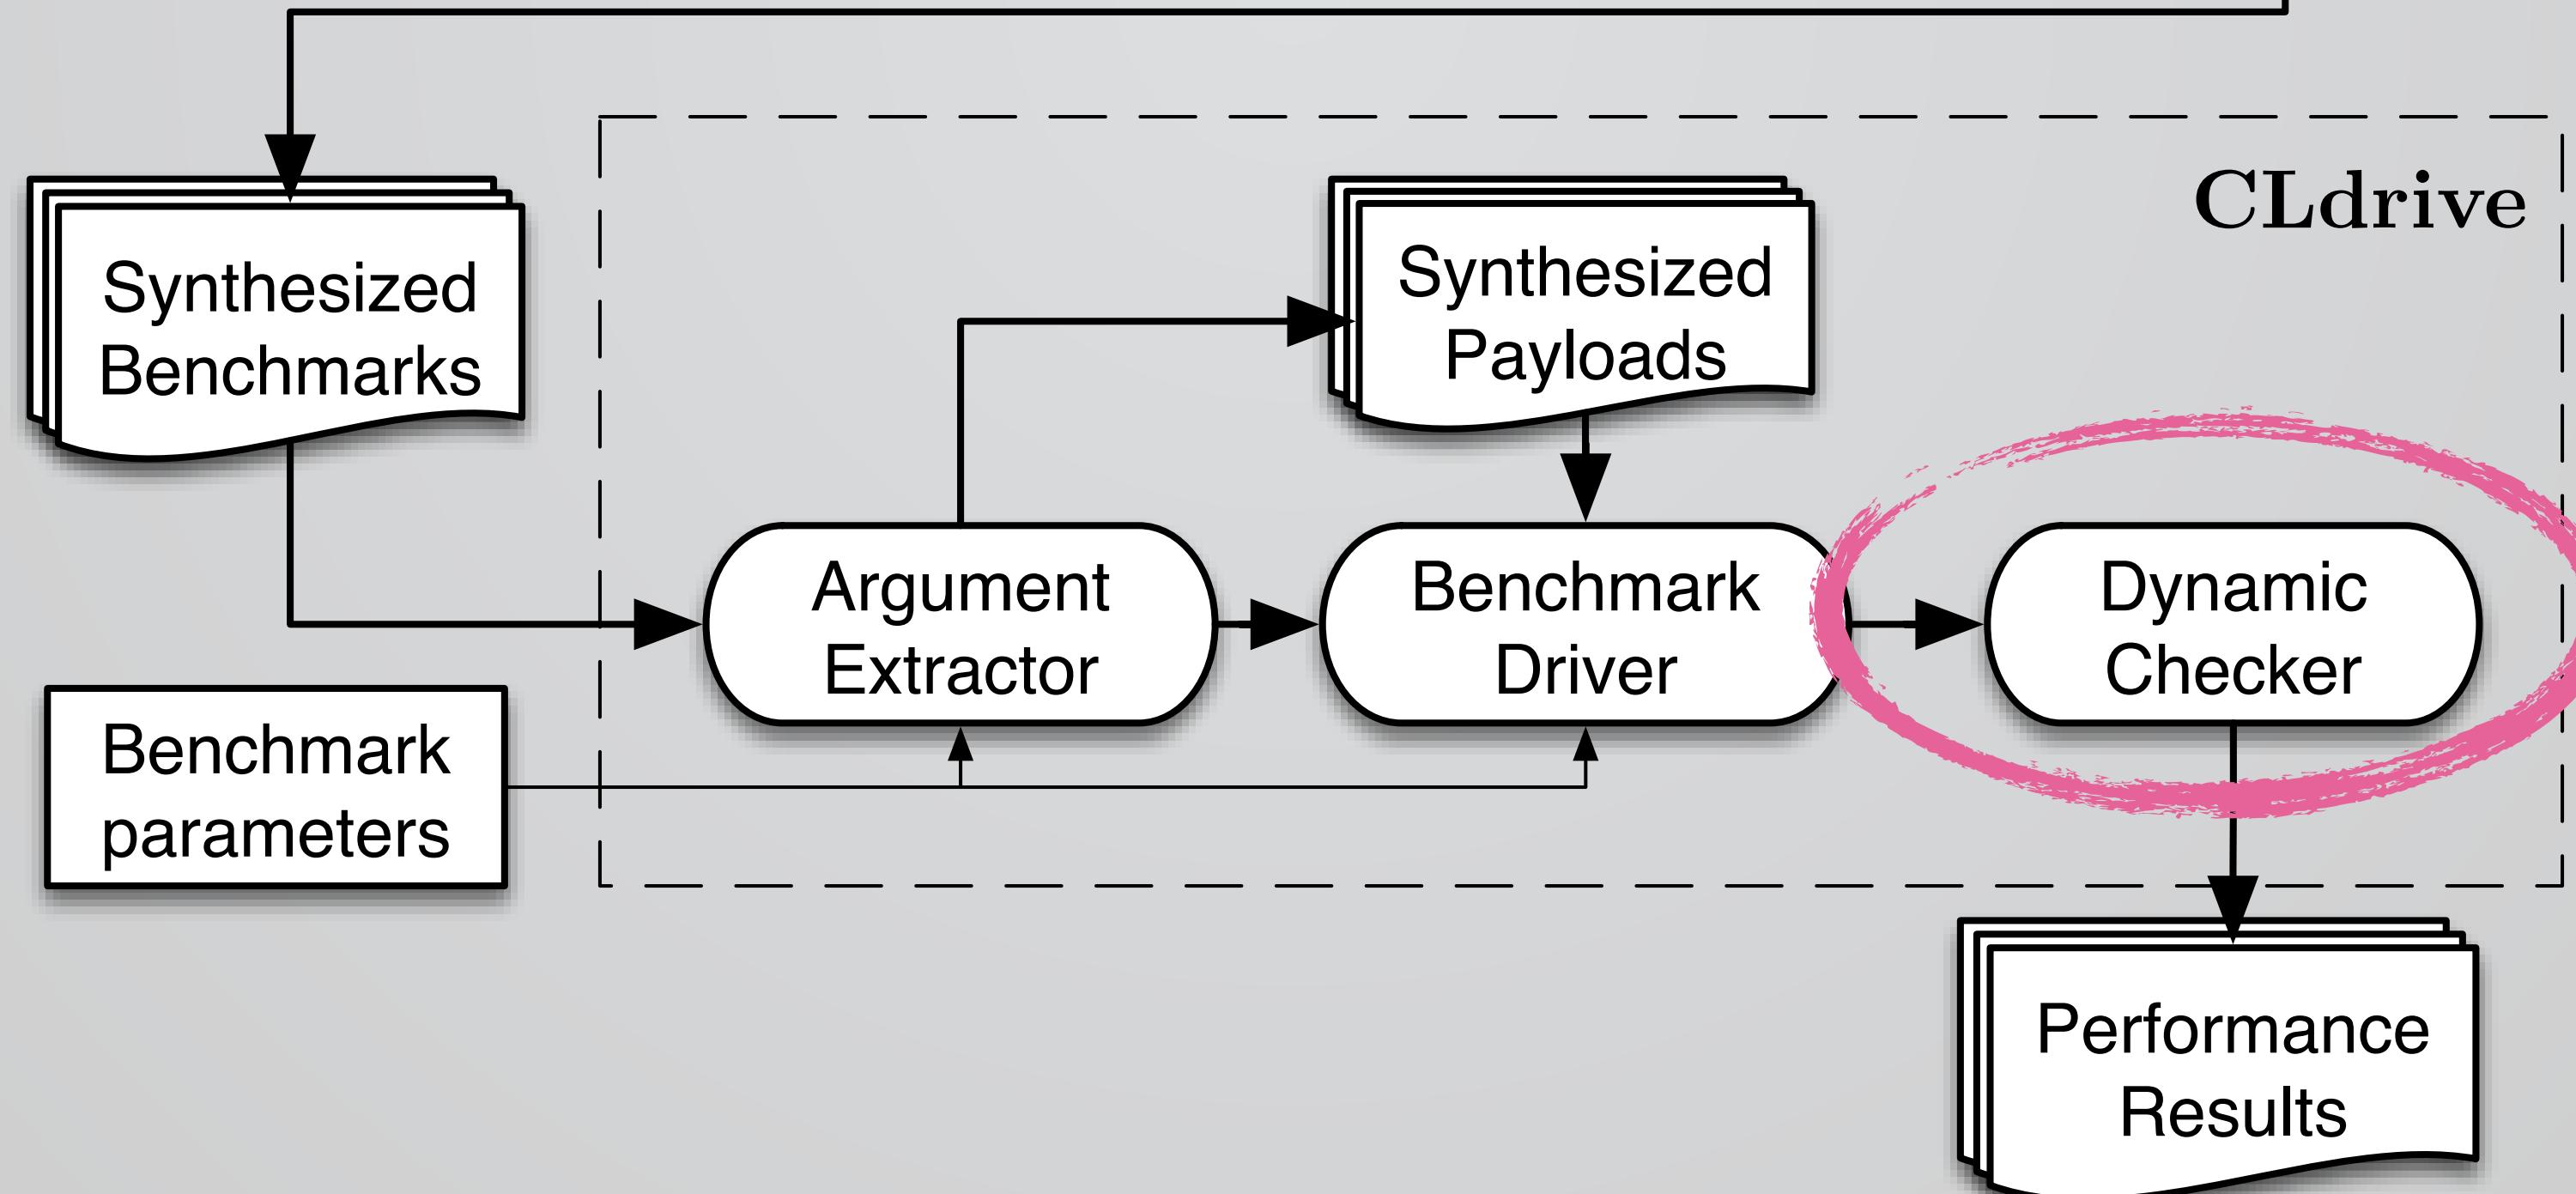
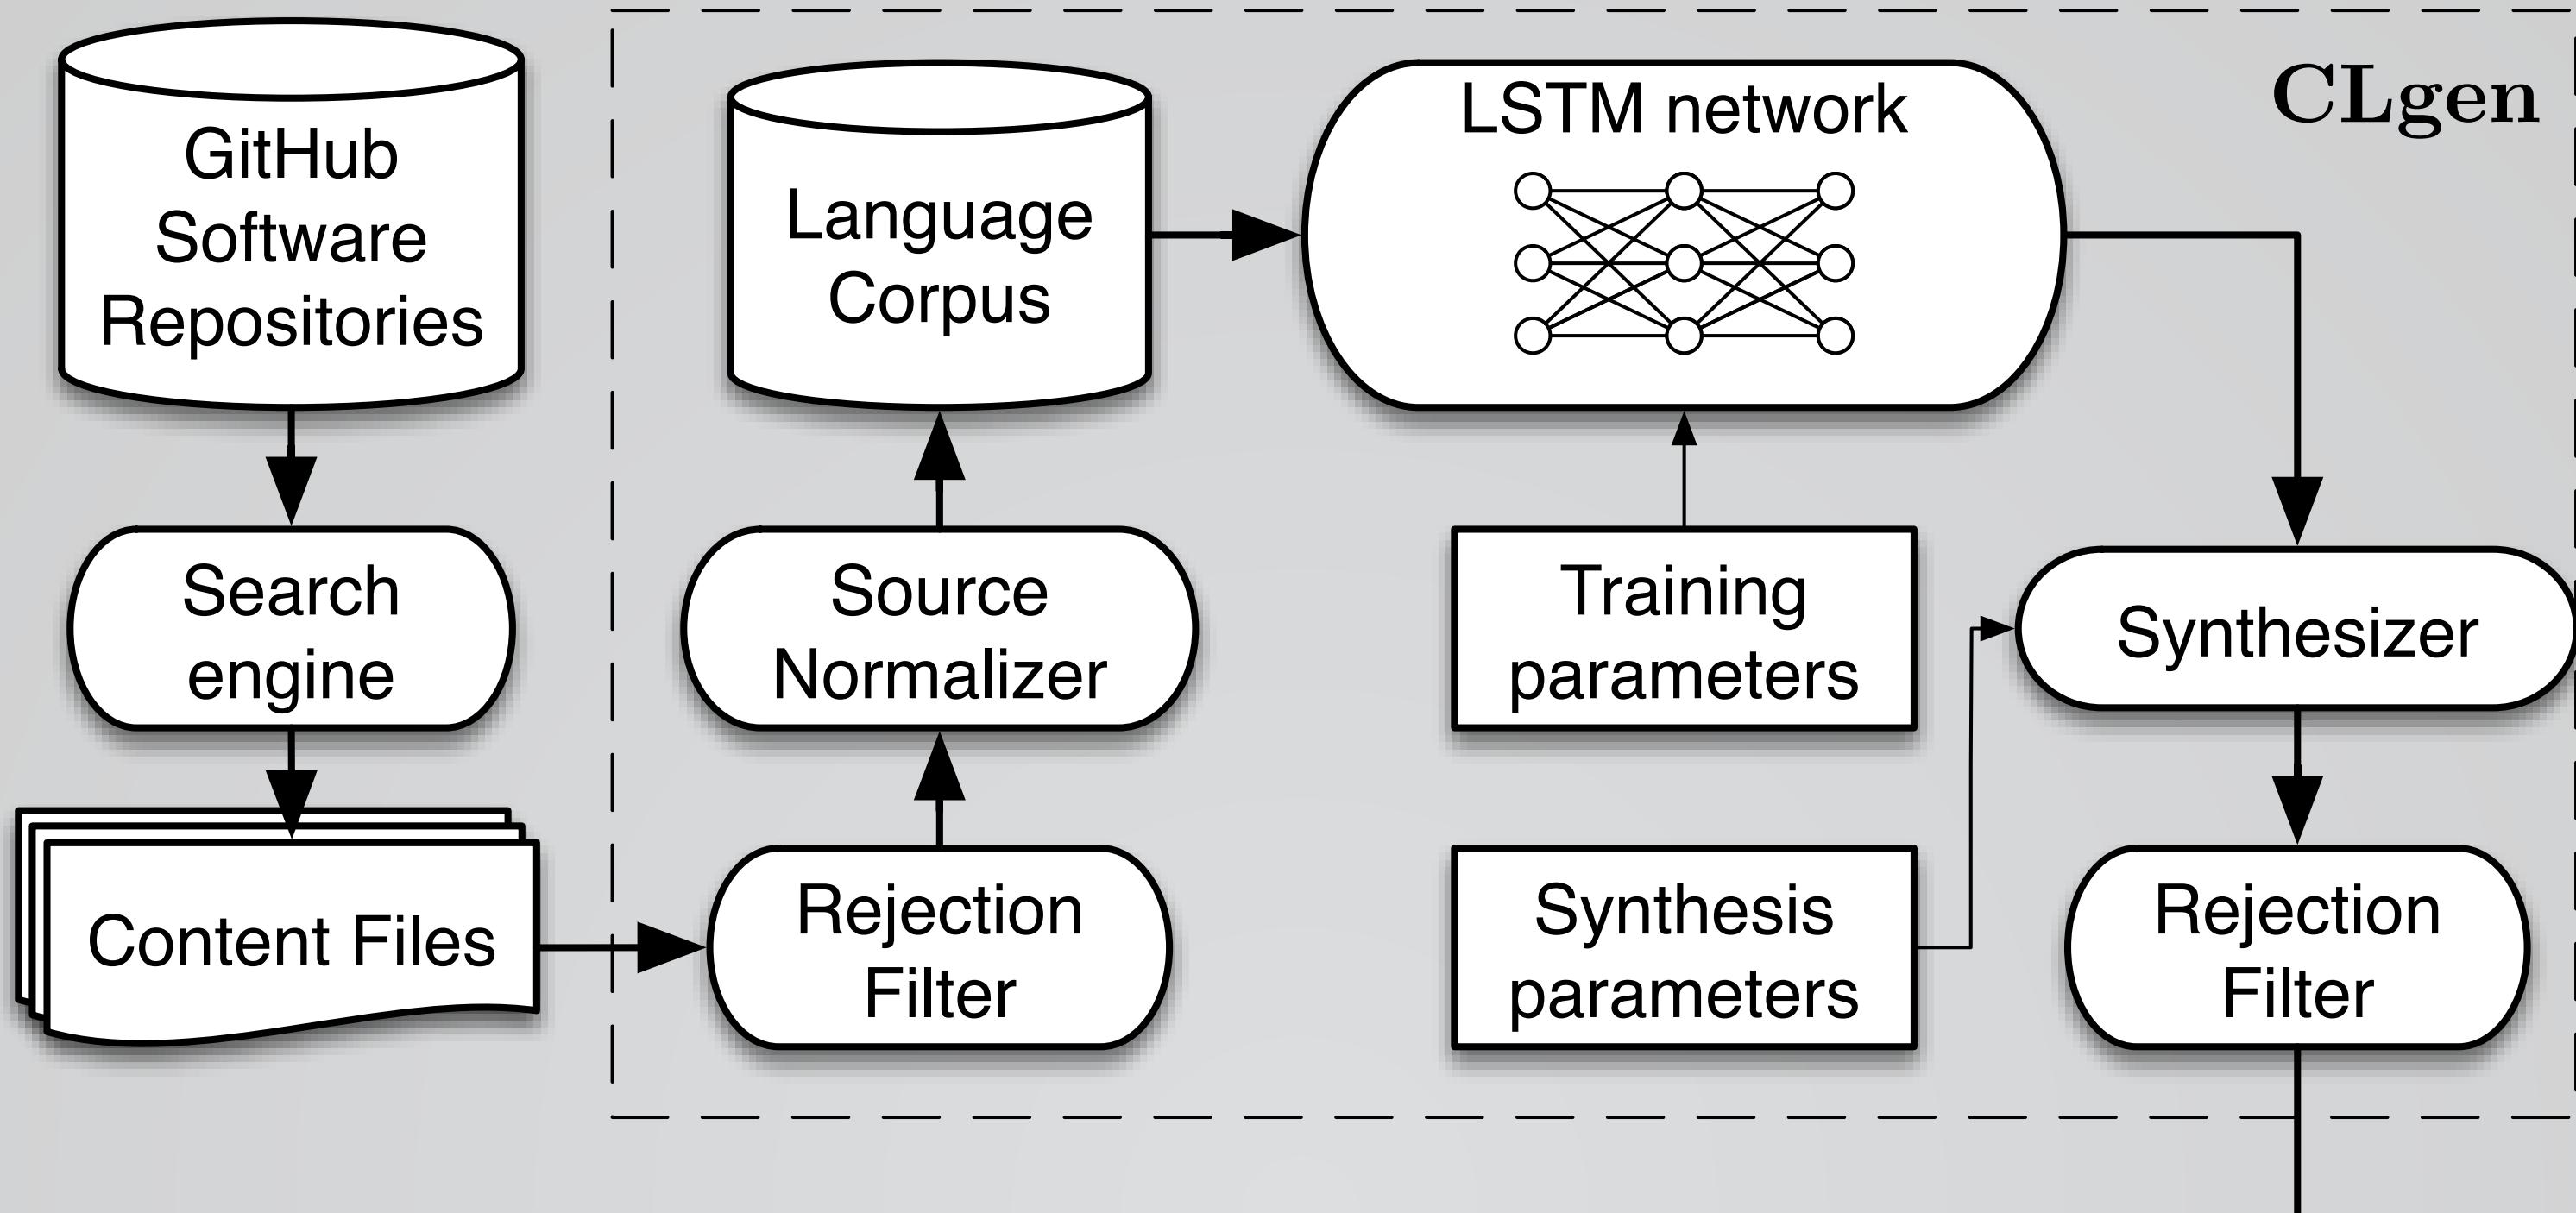




```
_kernel void A(__global float* a,  
          __global float* b,  
          __global float* c,  
          const int d) {  
  
    int e = get_global_id(0);  
    if (e >= d) {  
        return;  
    }  
    c[e] = a[e] + b[e] + 2 * a[e] + b[e] + 4;  
}
```

```
_kernel void A(__global float* a,
          __global float* b,
          __global float* c,
          const int d) {
    int e = get_global_id(0);
    if (e >= d) {
        return;
    }
    c[e] = a[e] + b[e] + 2 * a[e] + b[e] + 4;
}
```

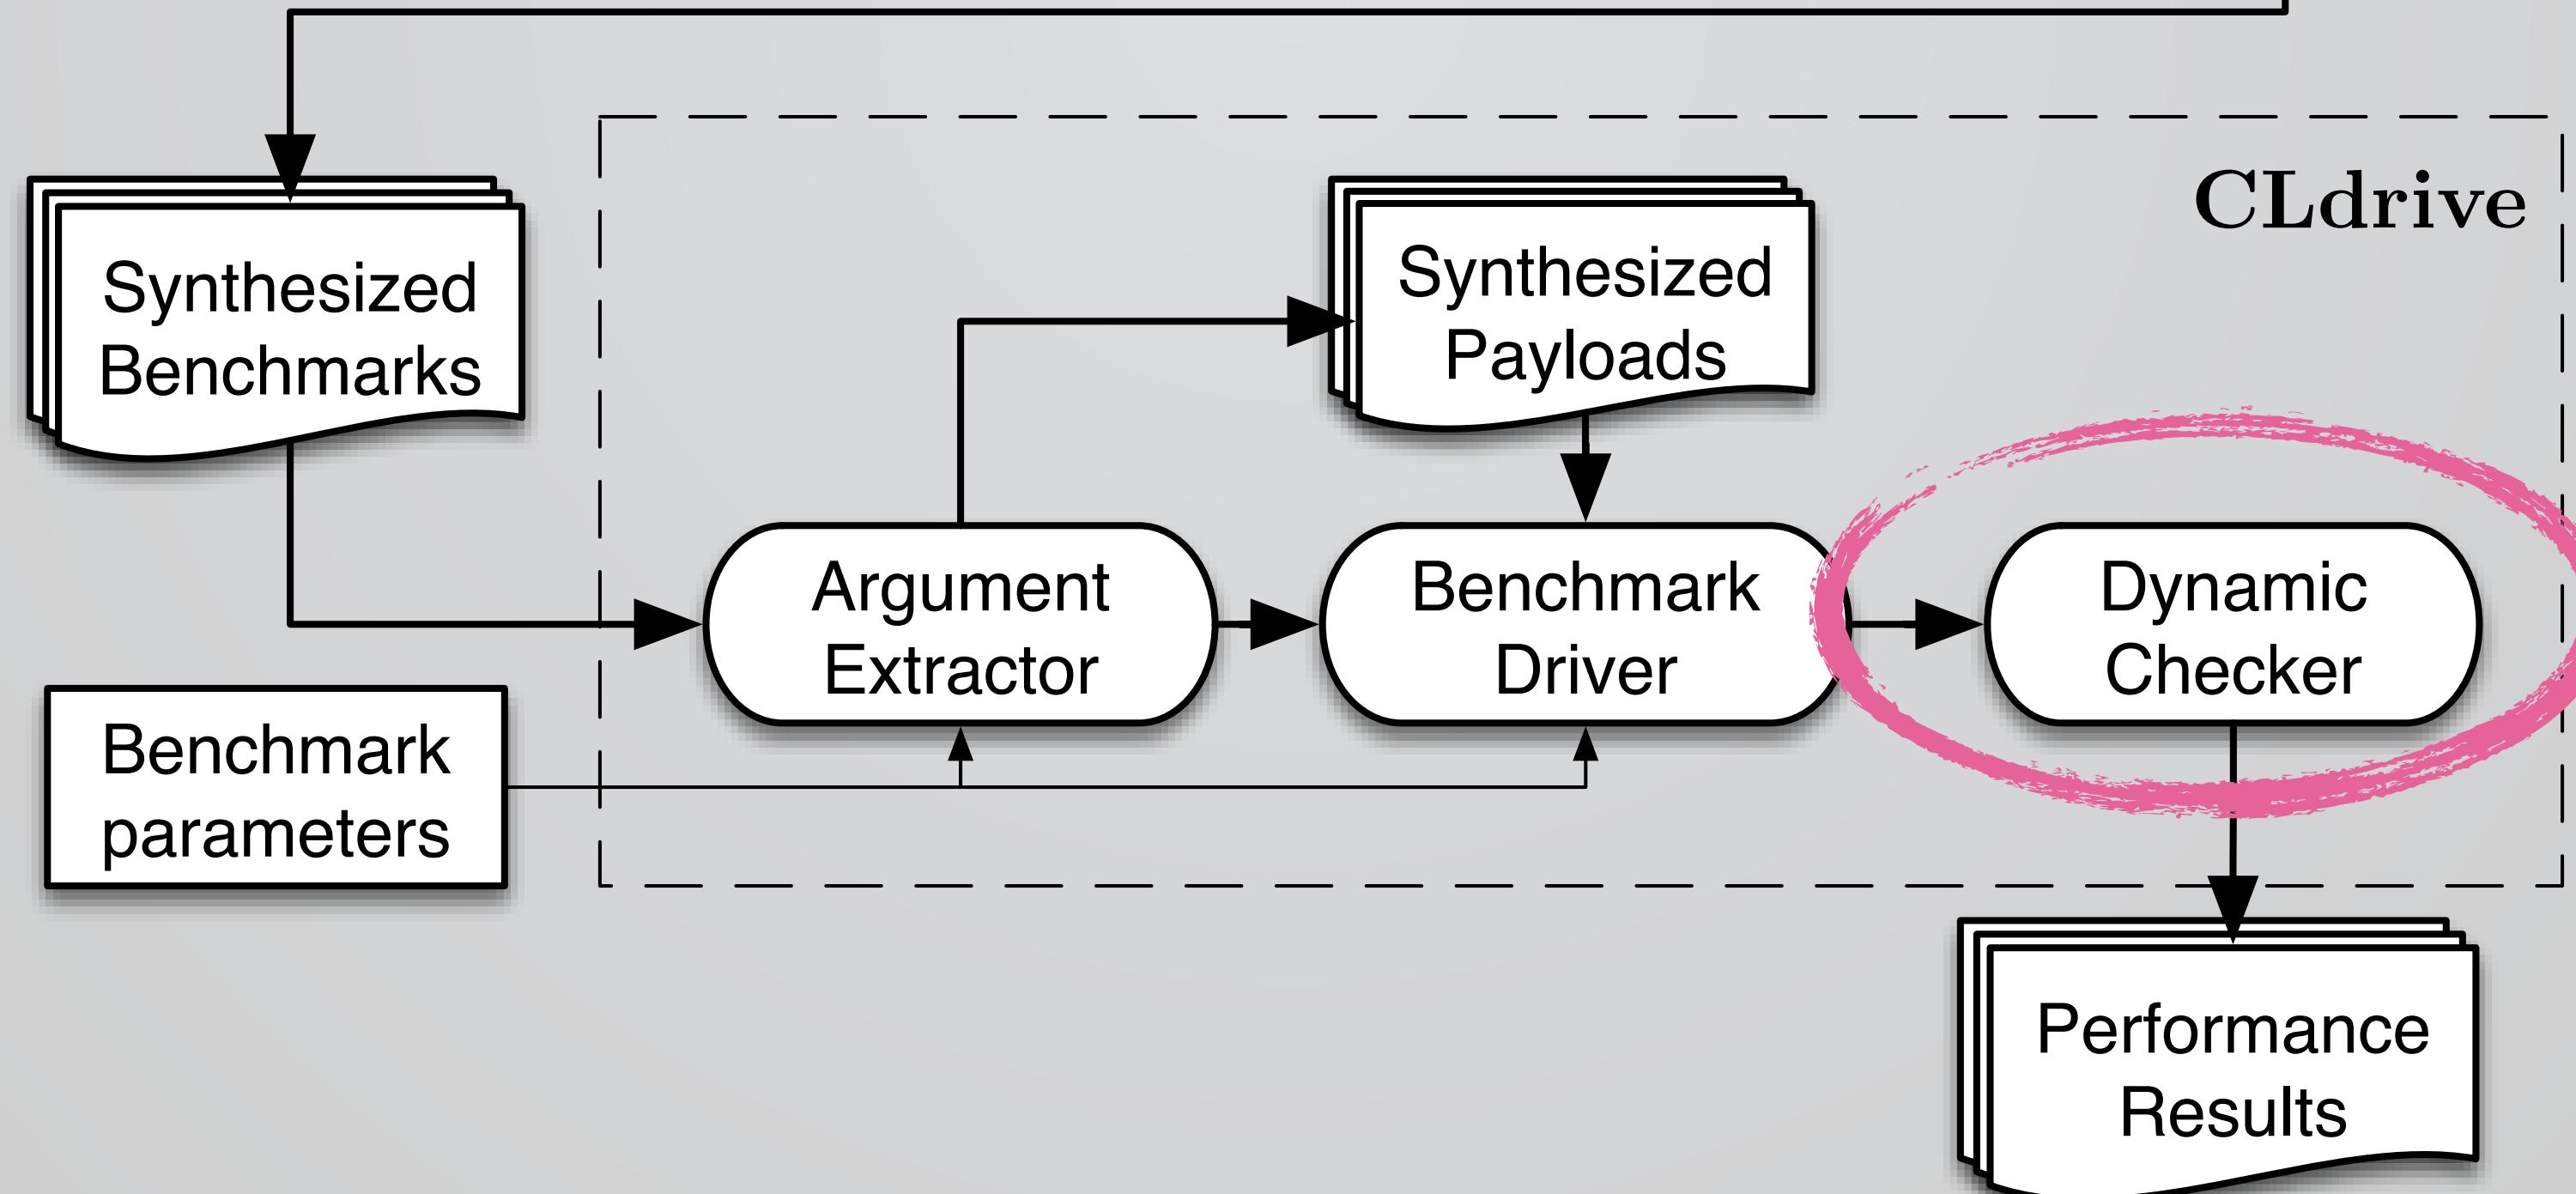
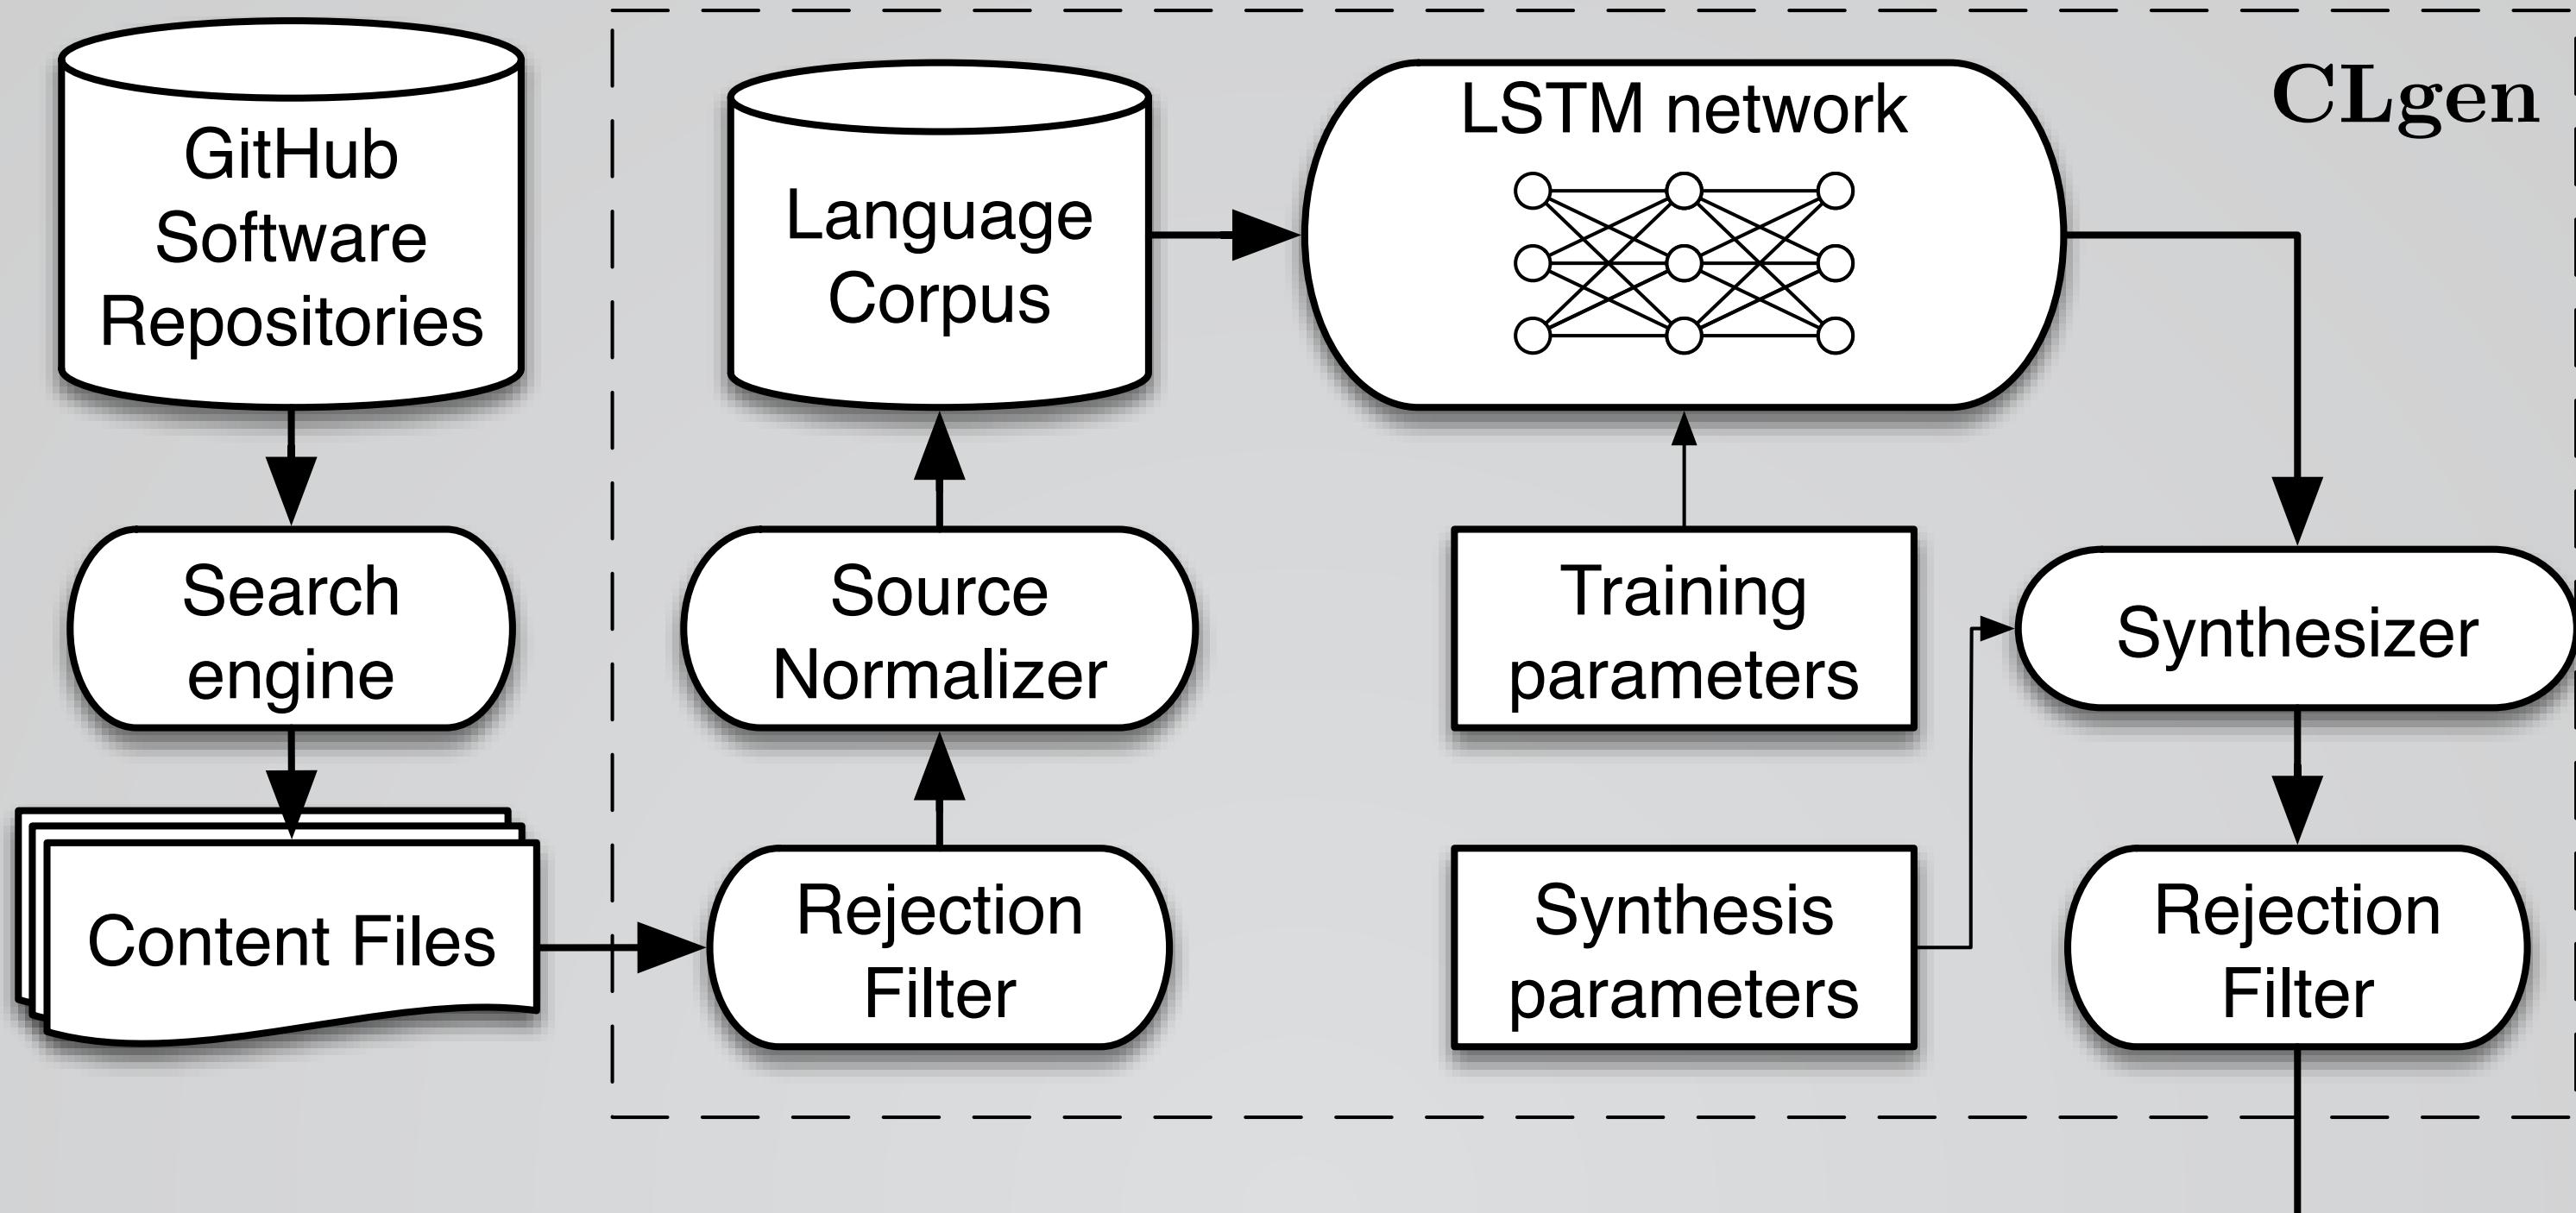




```
_A = random_payload(_A)    # generate inputs
_B = random_payload(_B)
_C = copy(_C)
_D = copy(_D)

A_ = k(_A)    # compute outputs
B_ = k(_B)
C_ = k(_C)
D_ = k(_D)

# differential test
assert (A_ != _A || B_ != _B) else NO_OUTPUTS
assert (A_ != B_ || C_ != D_) else INPUT_INSENSITIVE
assert (A_ == C_ && B_ == D_) else NON_DETERMINISTIC
```





## Round 1

Player: 1002, Robot: 998

**★★★★★ What a piece of junk!**

By George S. Mitchell "gsmitchell"

I checked around Amazon as well as some other sites, and decided these were what I needed. I read the reviews, looked at the pictures and bought two of these. I received them today. I unpacked one of them, read the instructions and immediately filled out the return request for Amazon. These mounts are very, very heavy and require 6, yes SIX, screws to mount to the wall. Are you kidding me? The TVs are only 22 lbs! Two screws in the studs and then more damage to the wall than is necessary. What a joke! And Lag bolts? Ever heard of screws? The other issue is that these are very hard to move. I bought an articulating wall mount for a reason, and I expect them to be able to move without a wrench in my hand.

**★★★★★ Excellent for the price**

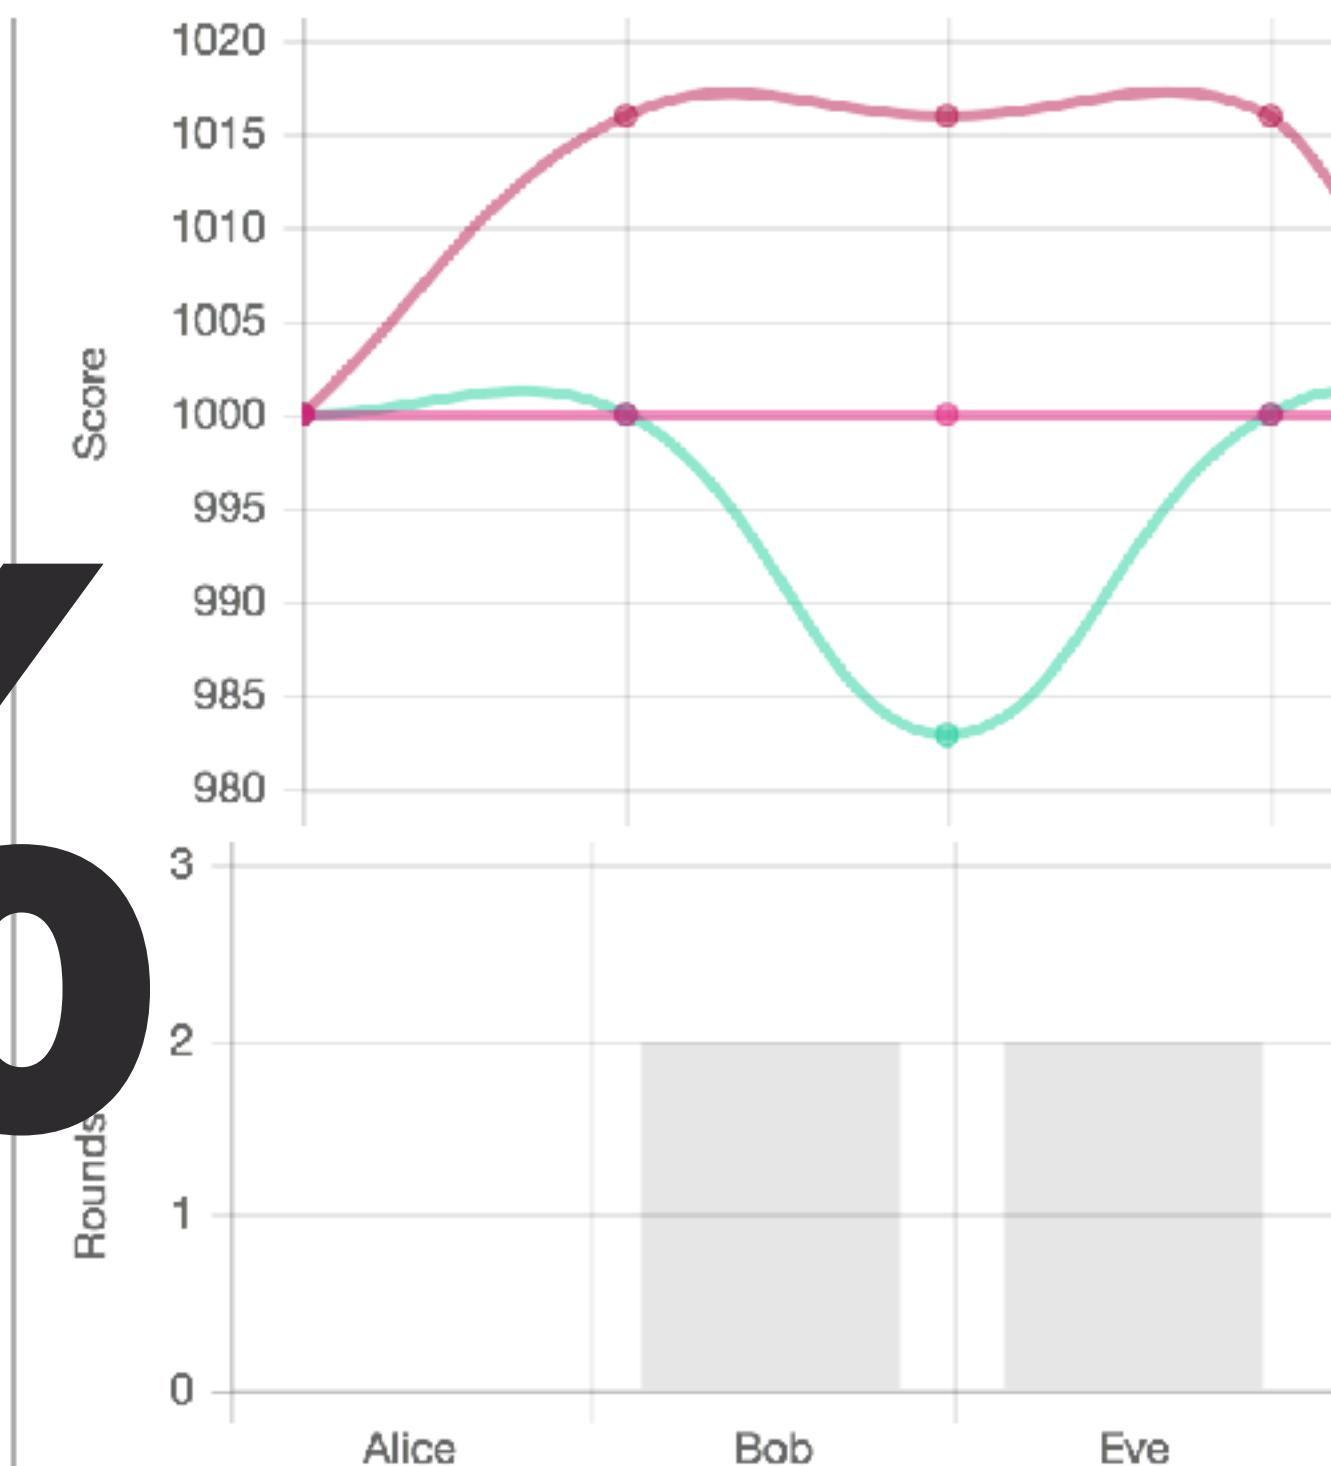
By A. Lopez "David Roberts"

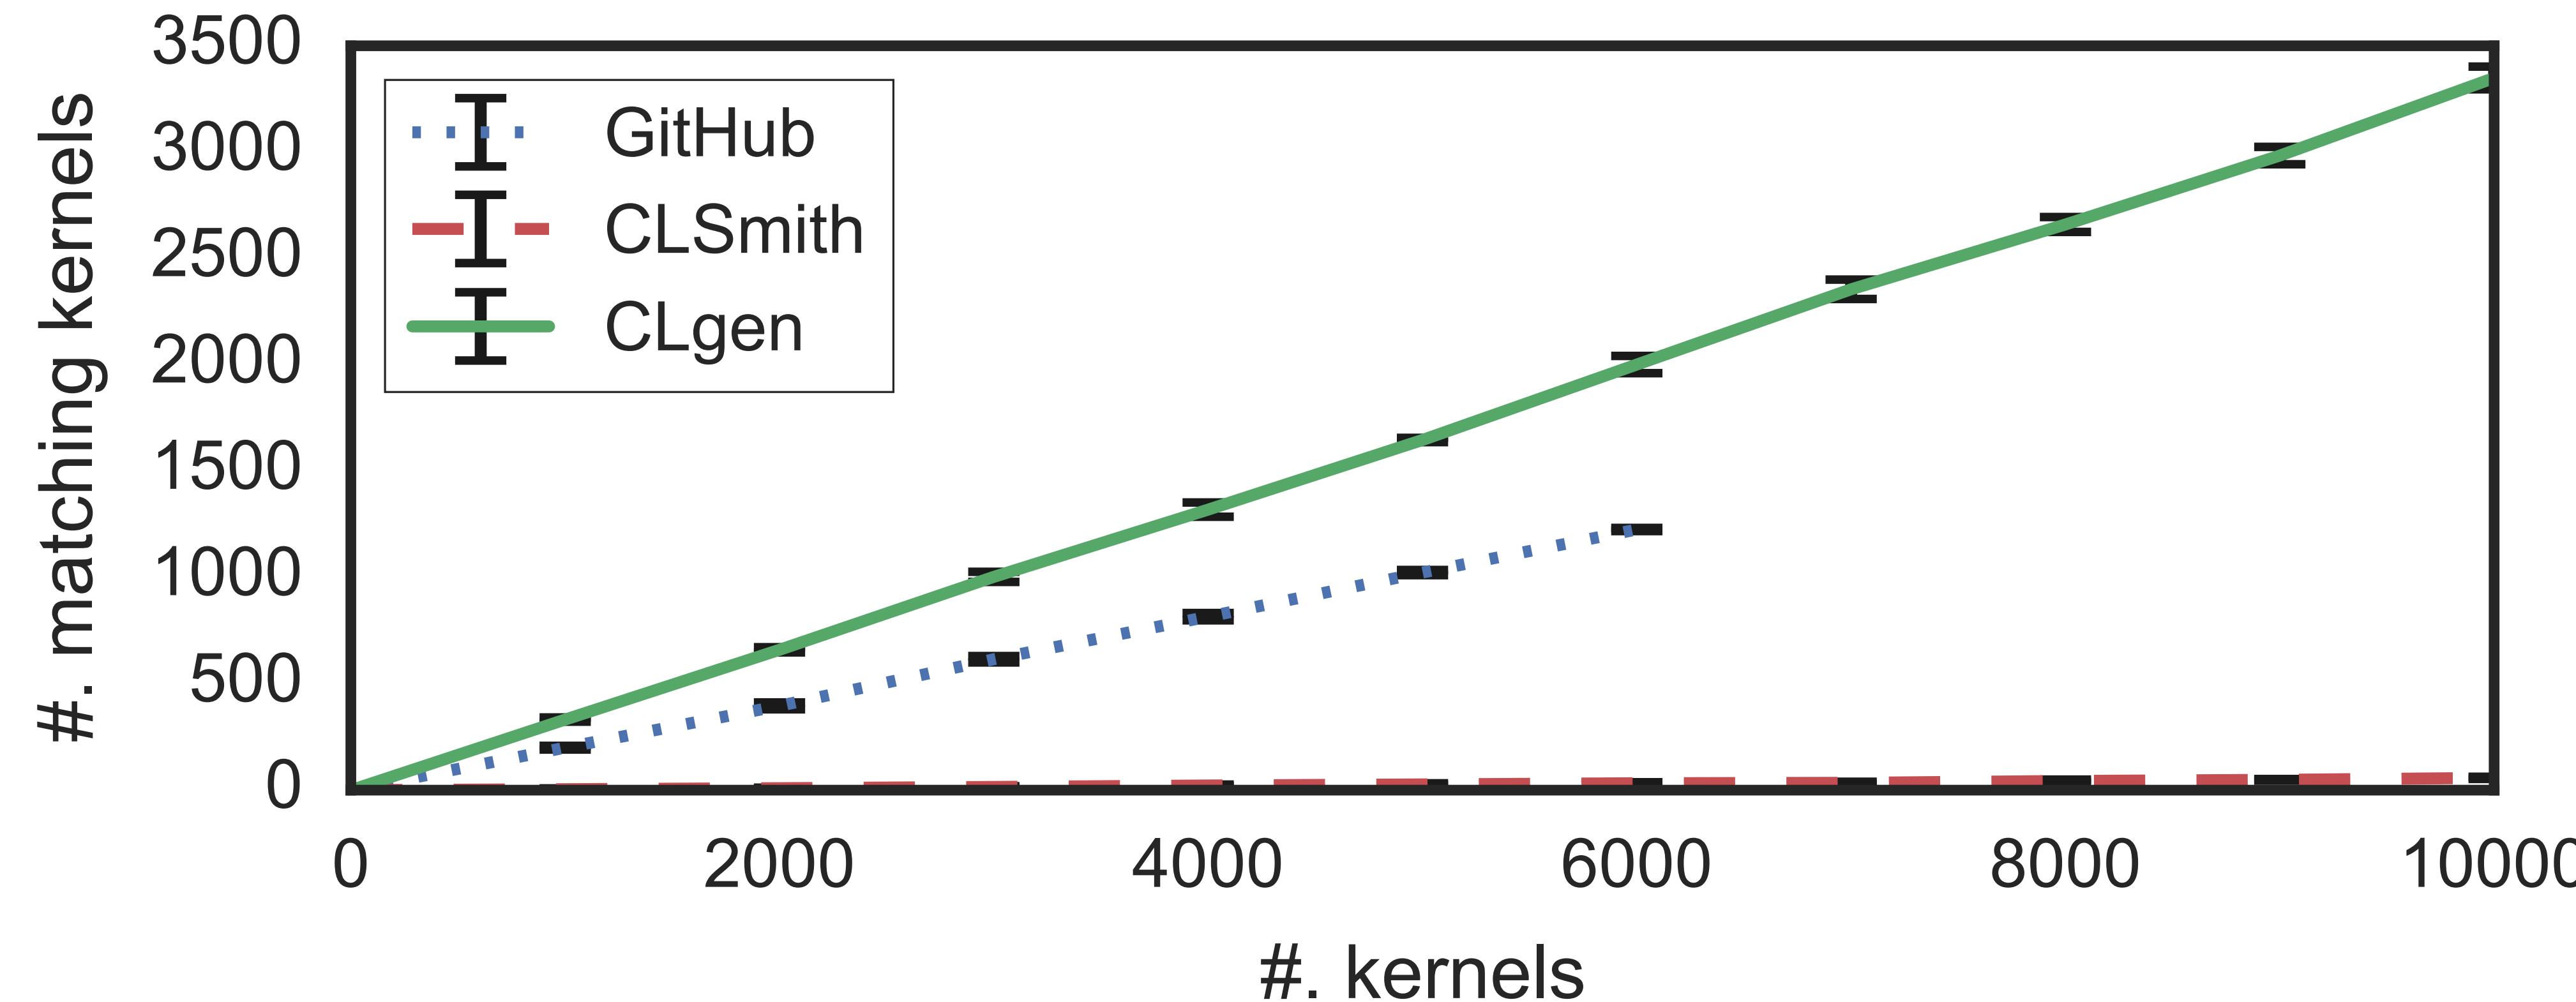
This thing is fantastic. It has the same size and speed and sharp sound. I like the power turntable. I use it every time you connect your computer and trying to figure out you want to get a problem with a ti  
e and a product. I never h  
a problems with it. It's a conv  
tional product that works out. Good price for it and a couple of years ago we w  
hile software installed so far I'm rem  
appy with this product. It works great for what I wanted to do. The way the laptop company you a lot more expensive, you can't tell that any other devic  
on this is a protective product that works well for a reasonable price and works great for the price, and it works.

52%

This side is the Robot

This side is the Robot



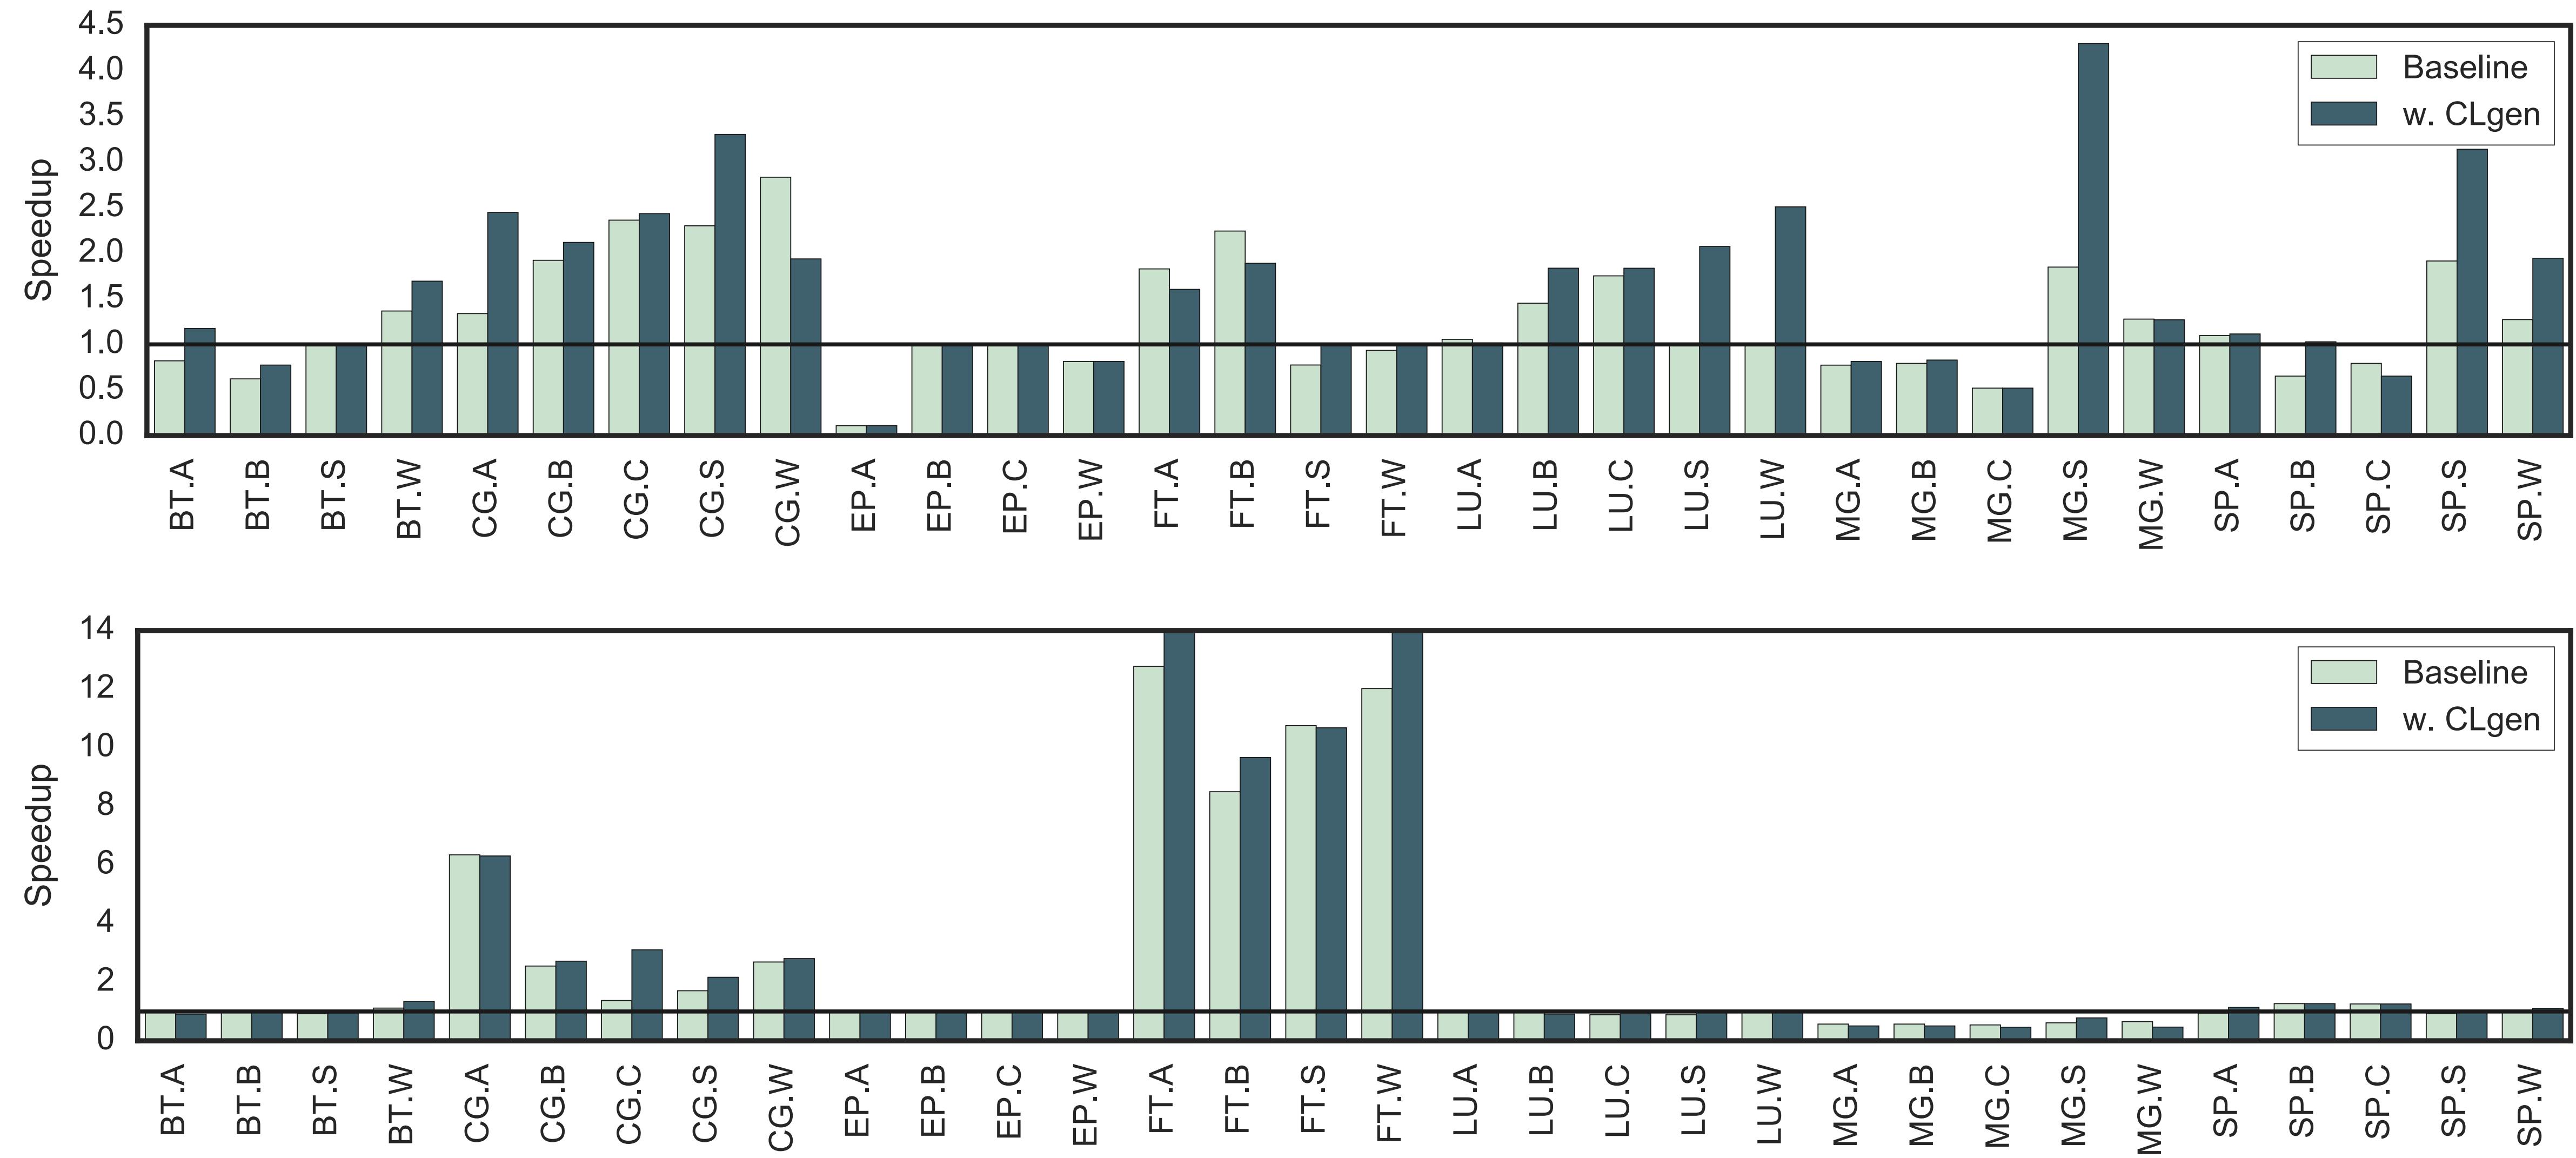


---

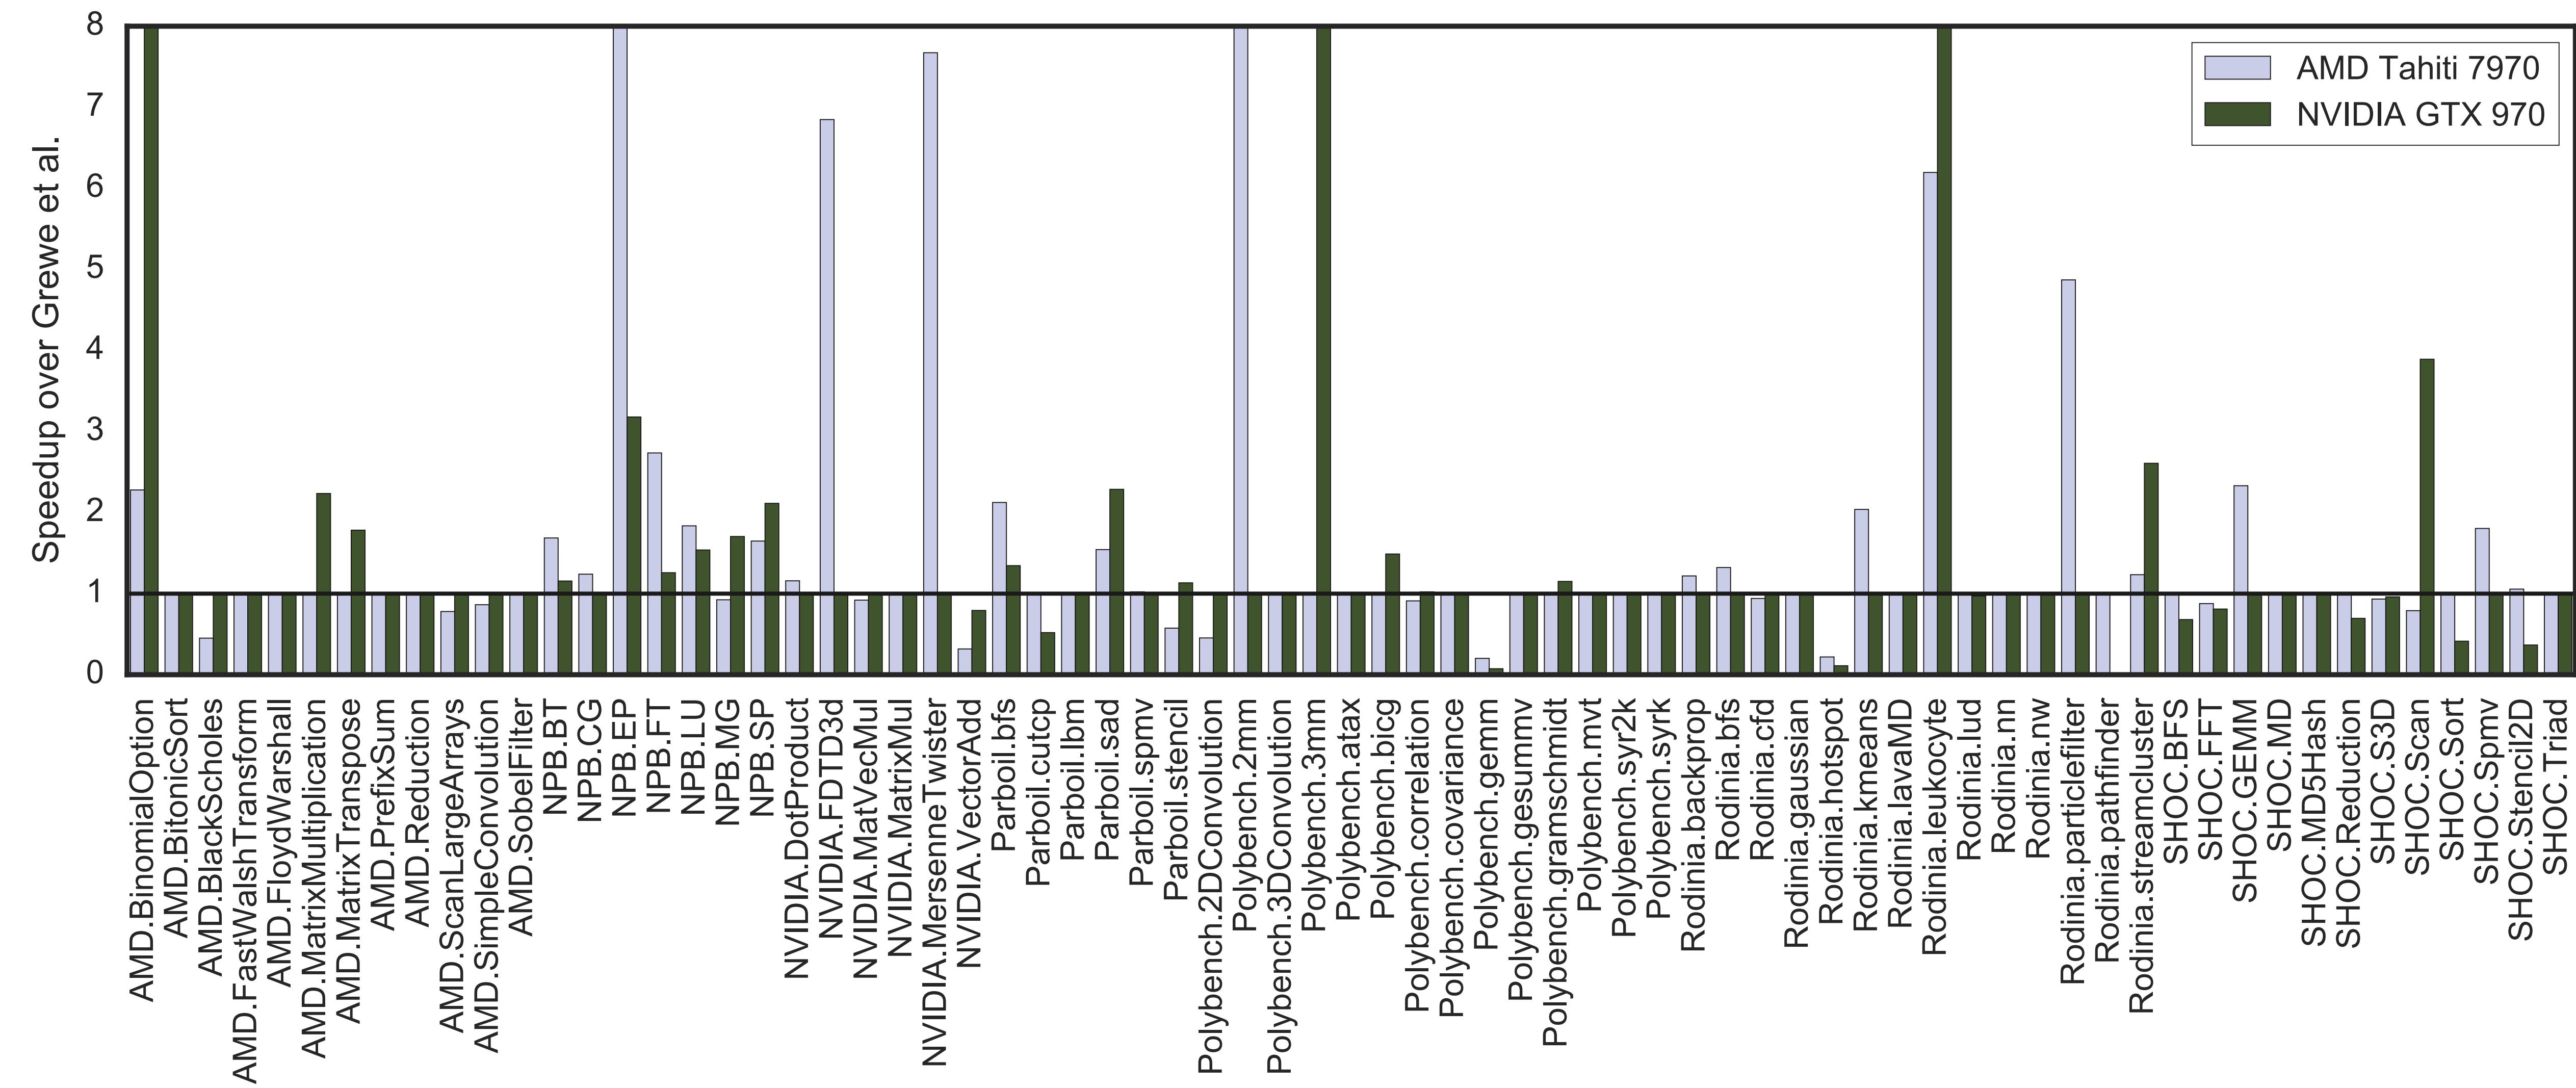
#### Raw Code Features

---

comp	static	#. compute operations
mem	static	#. accesses to global memory
localmem	static	#. accesses to local memory
coalesced	static	#. coalesced memory accesses



7 programs, 1,000 synthetic benchmarks. 1.27x faster



71 programs, 1,000 synthetic benchmarks. 2.66x faster

# Good Things

Basically\* language agnostic.

35 million repos on GitHub. We're using 0.00004%.

Generates 2000 OpenCL benchmarks per machine per day.

# Bad Things

No support for things declared outside of kernel scope.

Undirected almost to a fault.

AMD rage.

# Thanks!

More benchmarks = Better models

No general way of creating benchmarks

I taught a Neural Network to program from GitHub

Improves state of the art by 3.38x