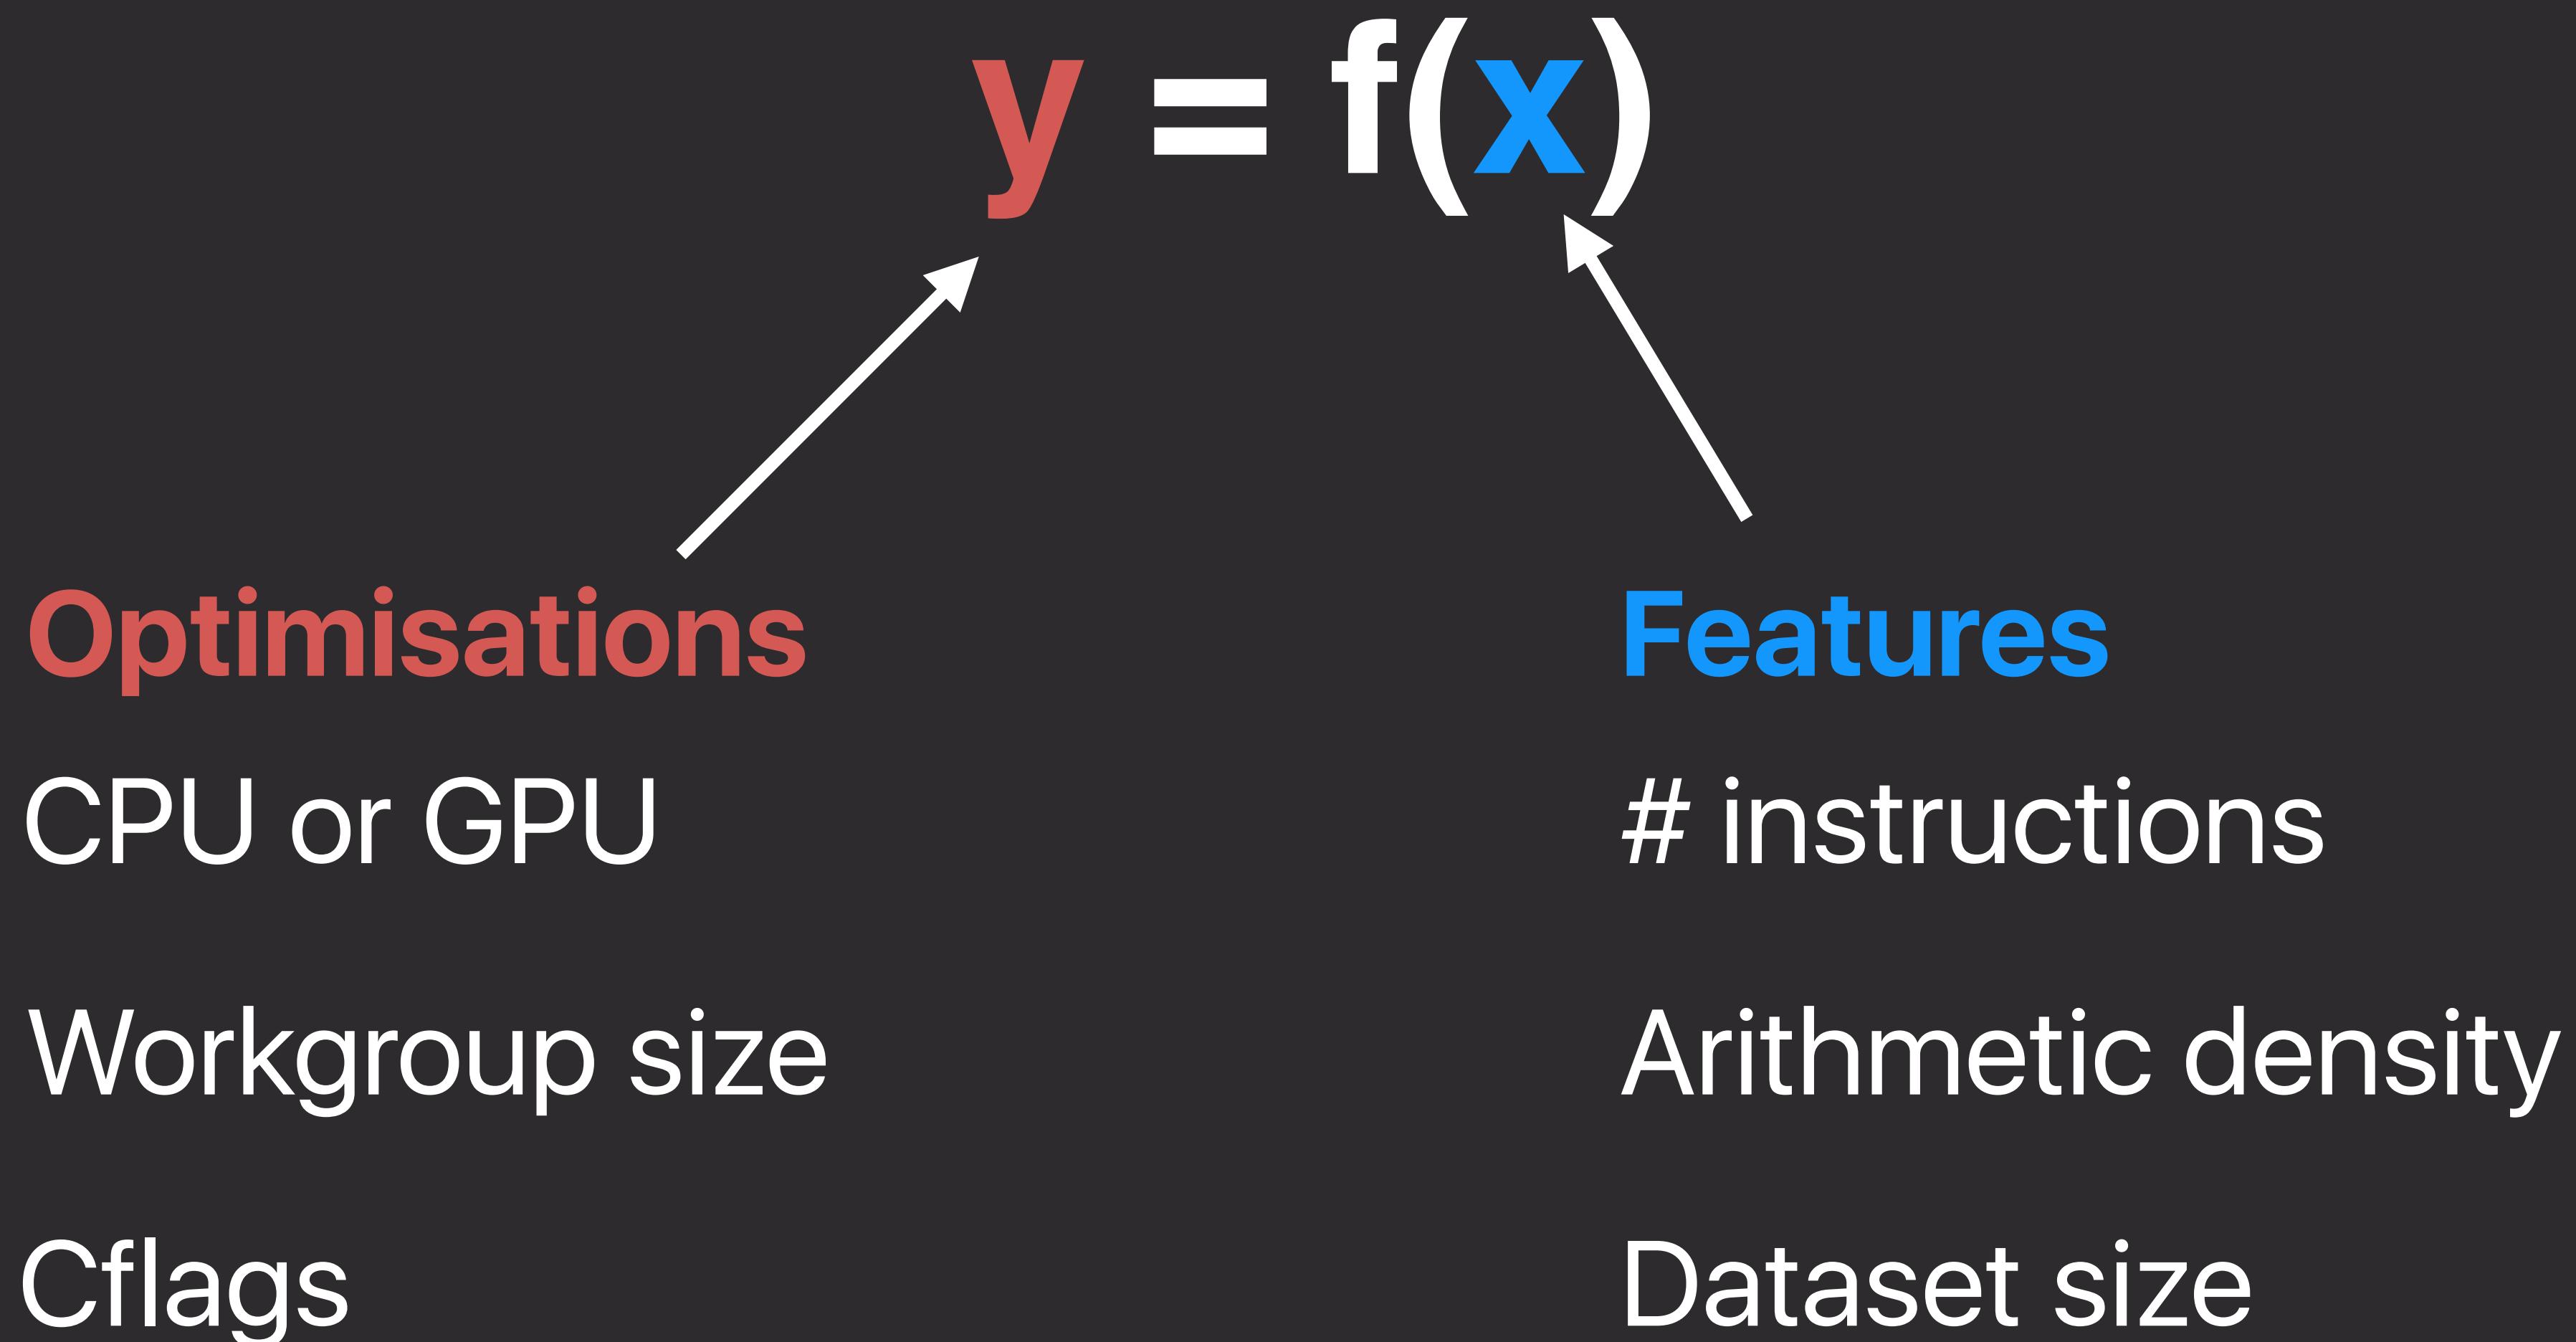


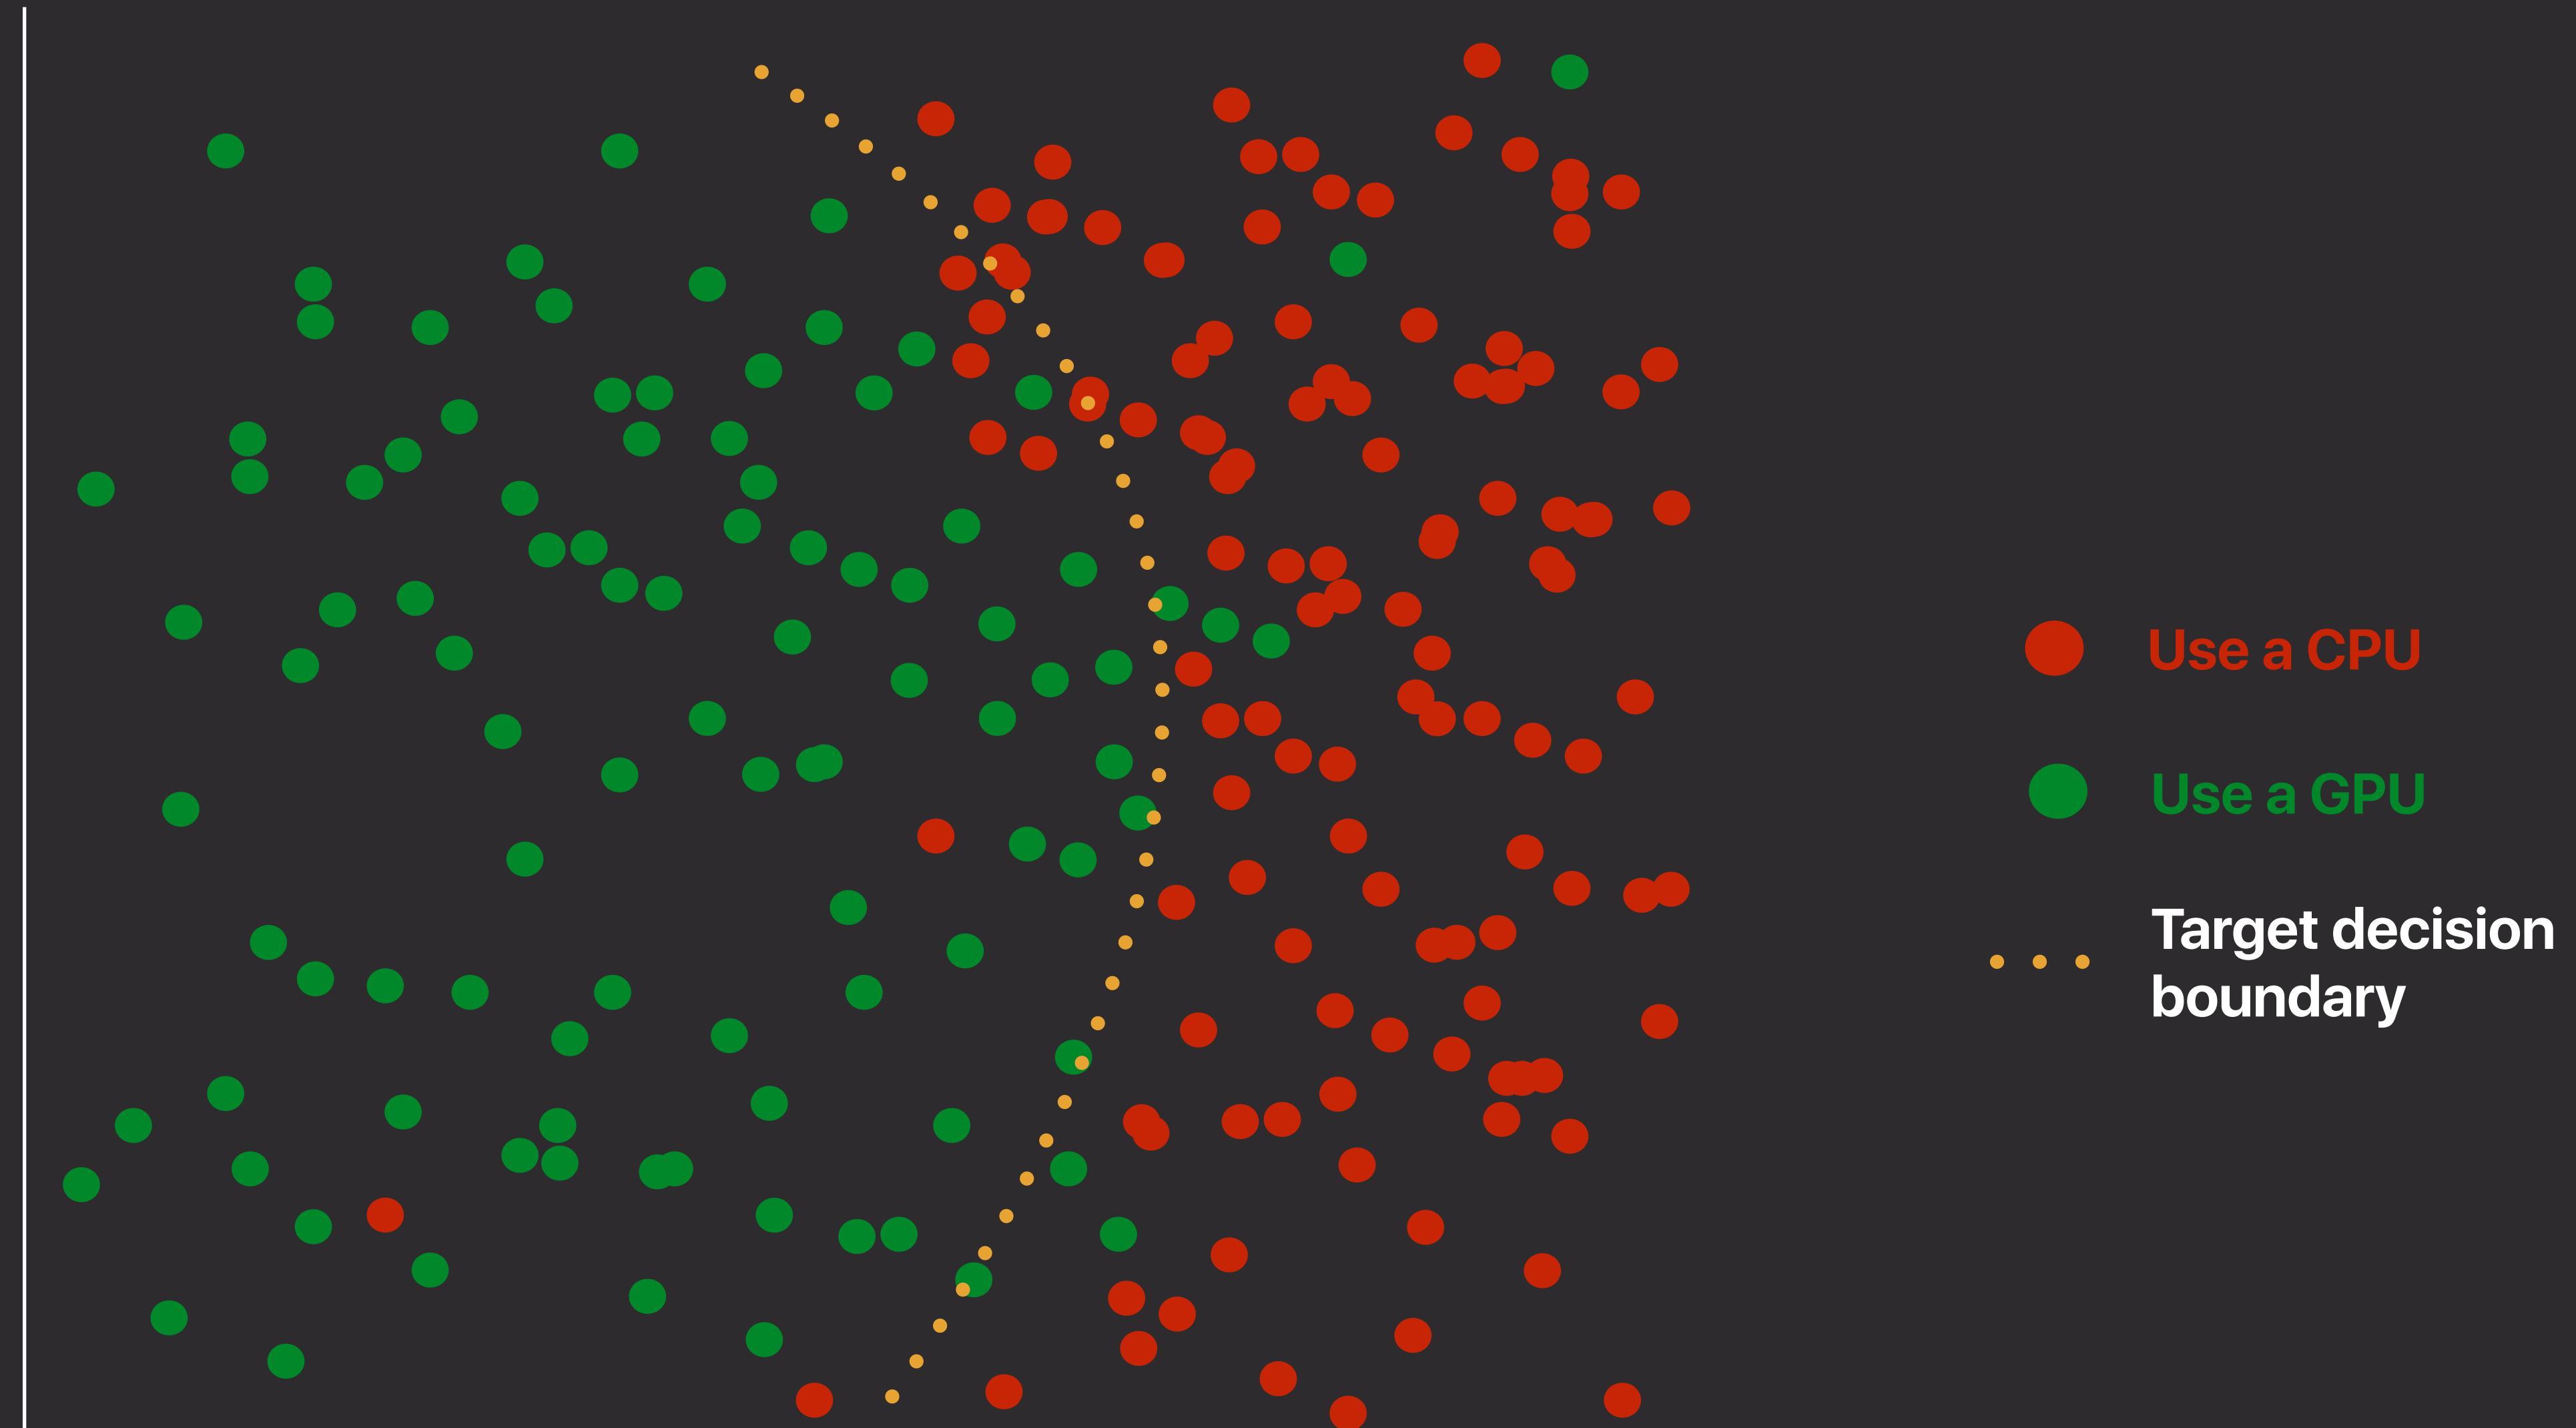
first year Review

Chris Cummins

machine learning for compilers

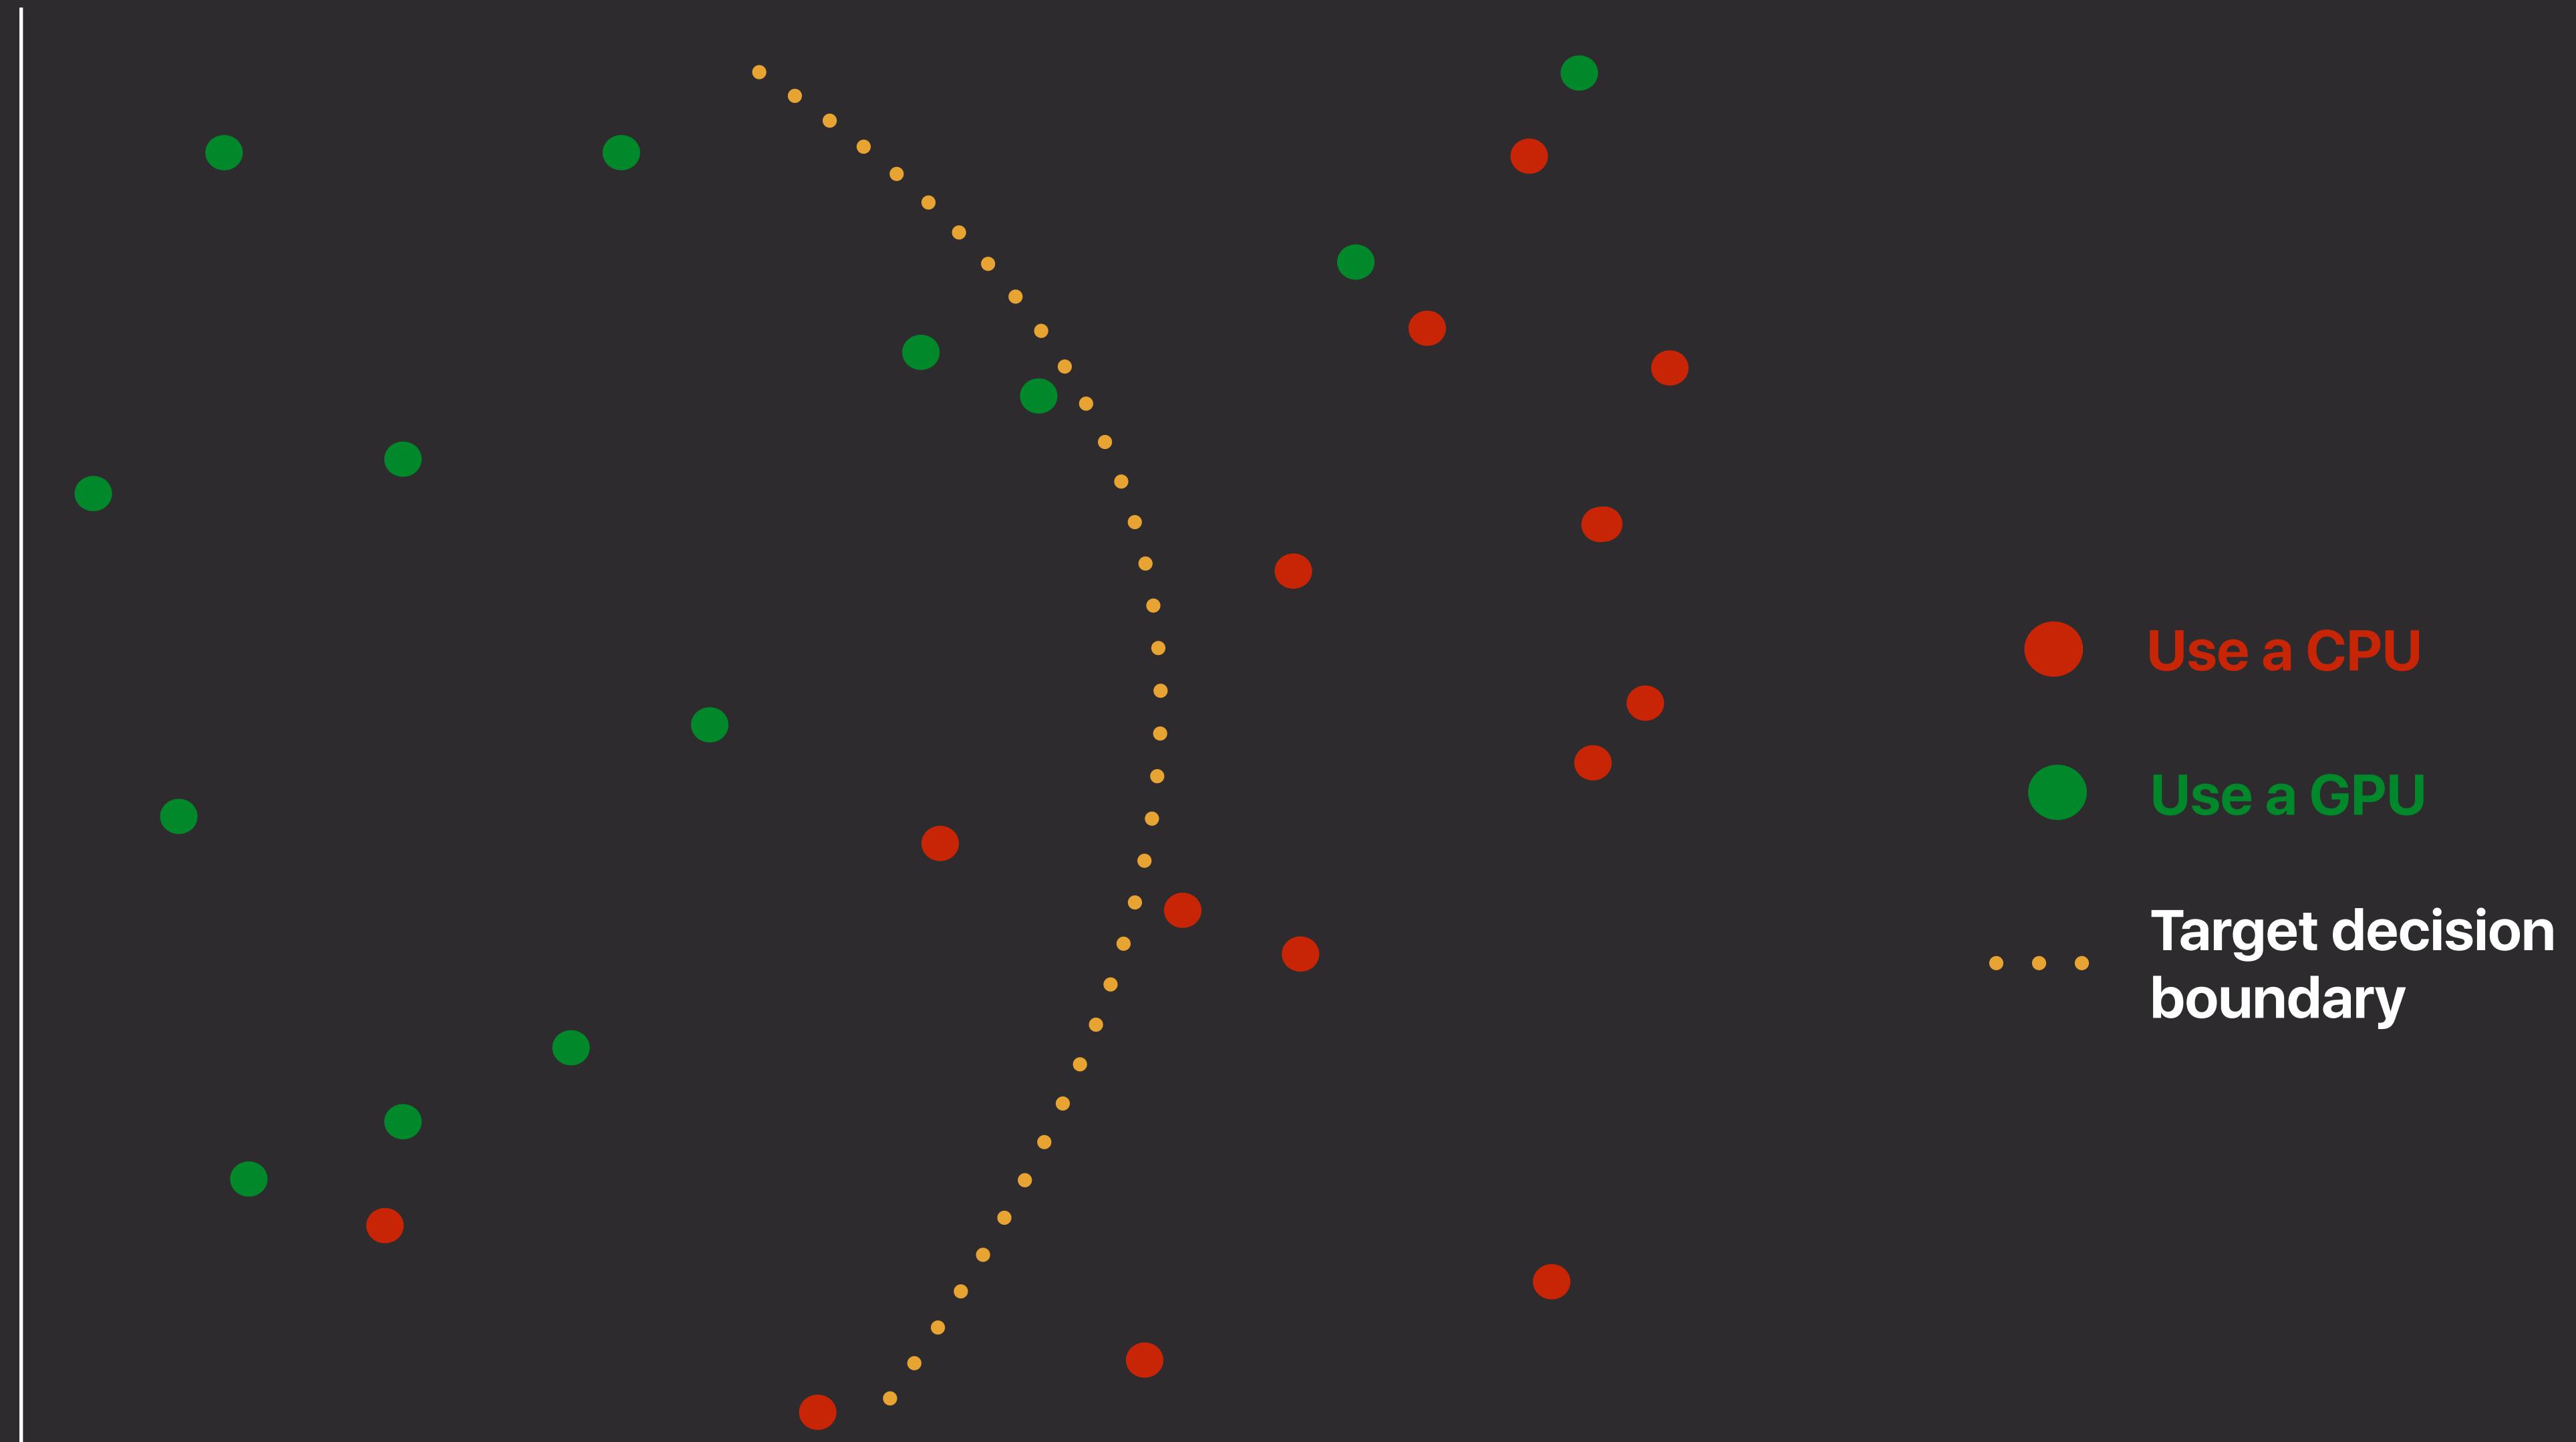


machine learning for compilers



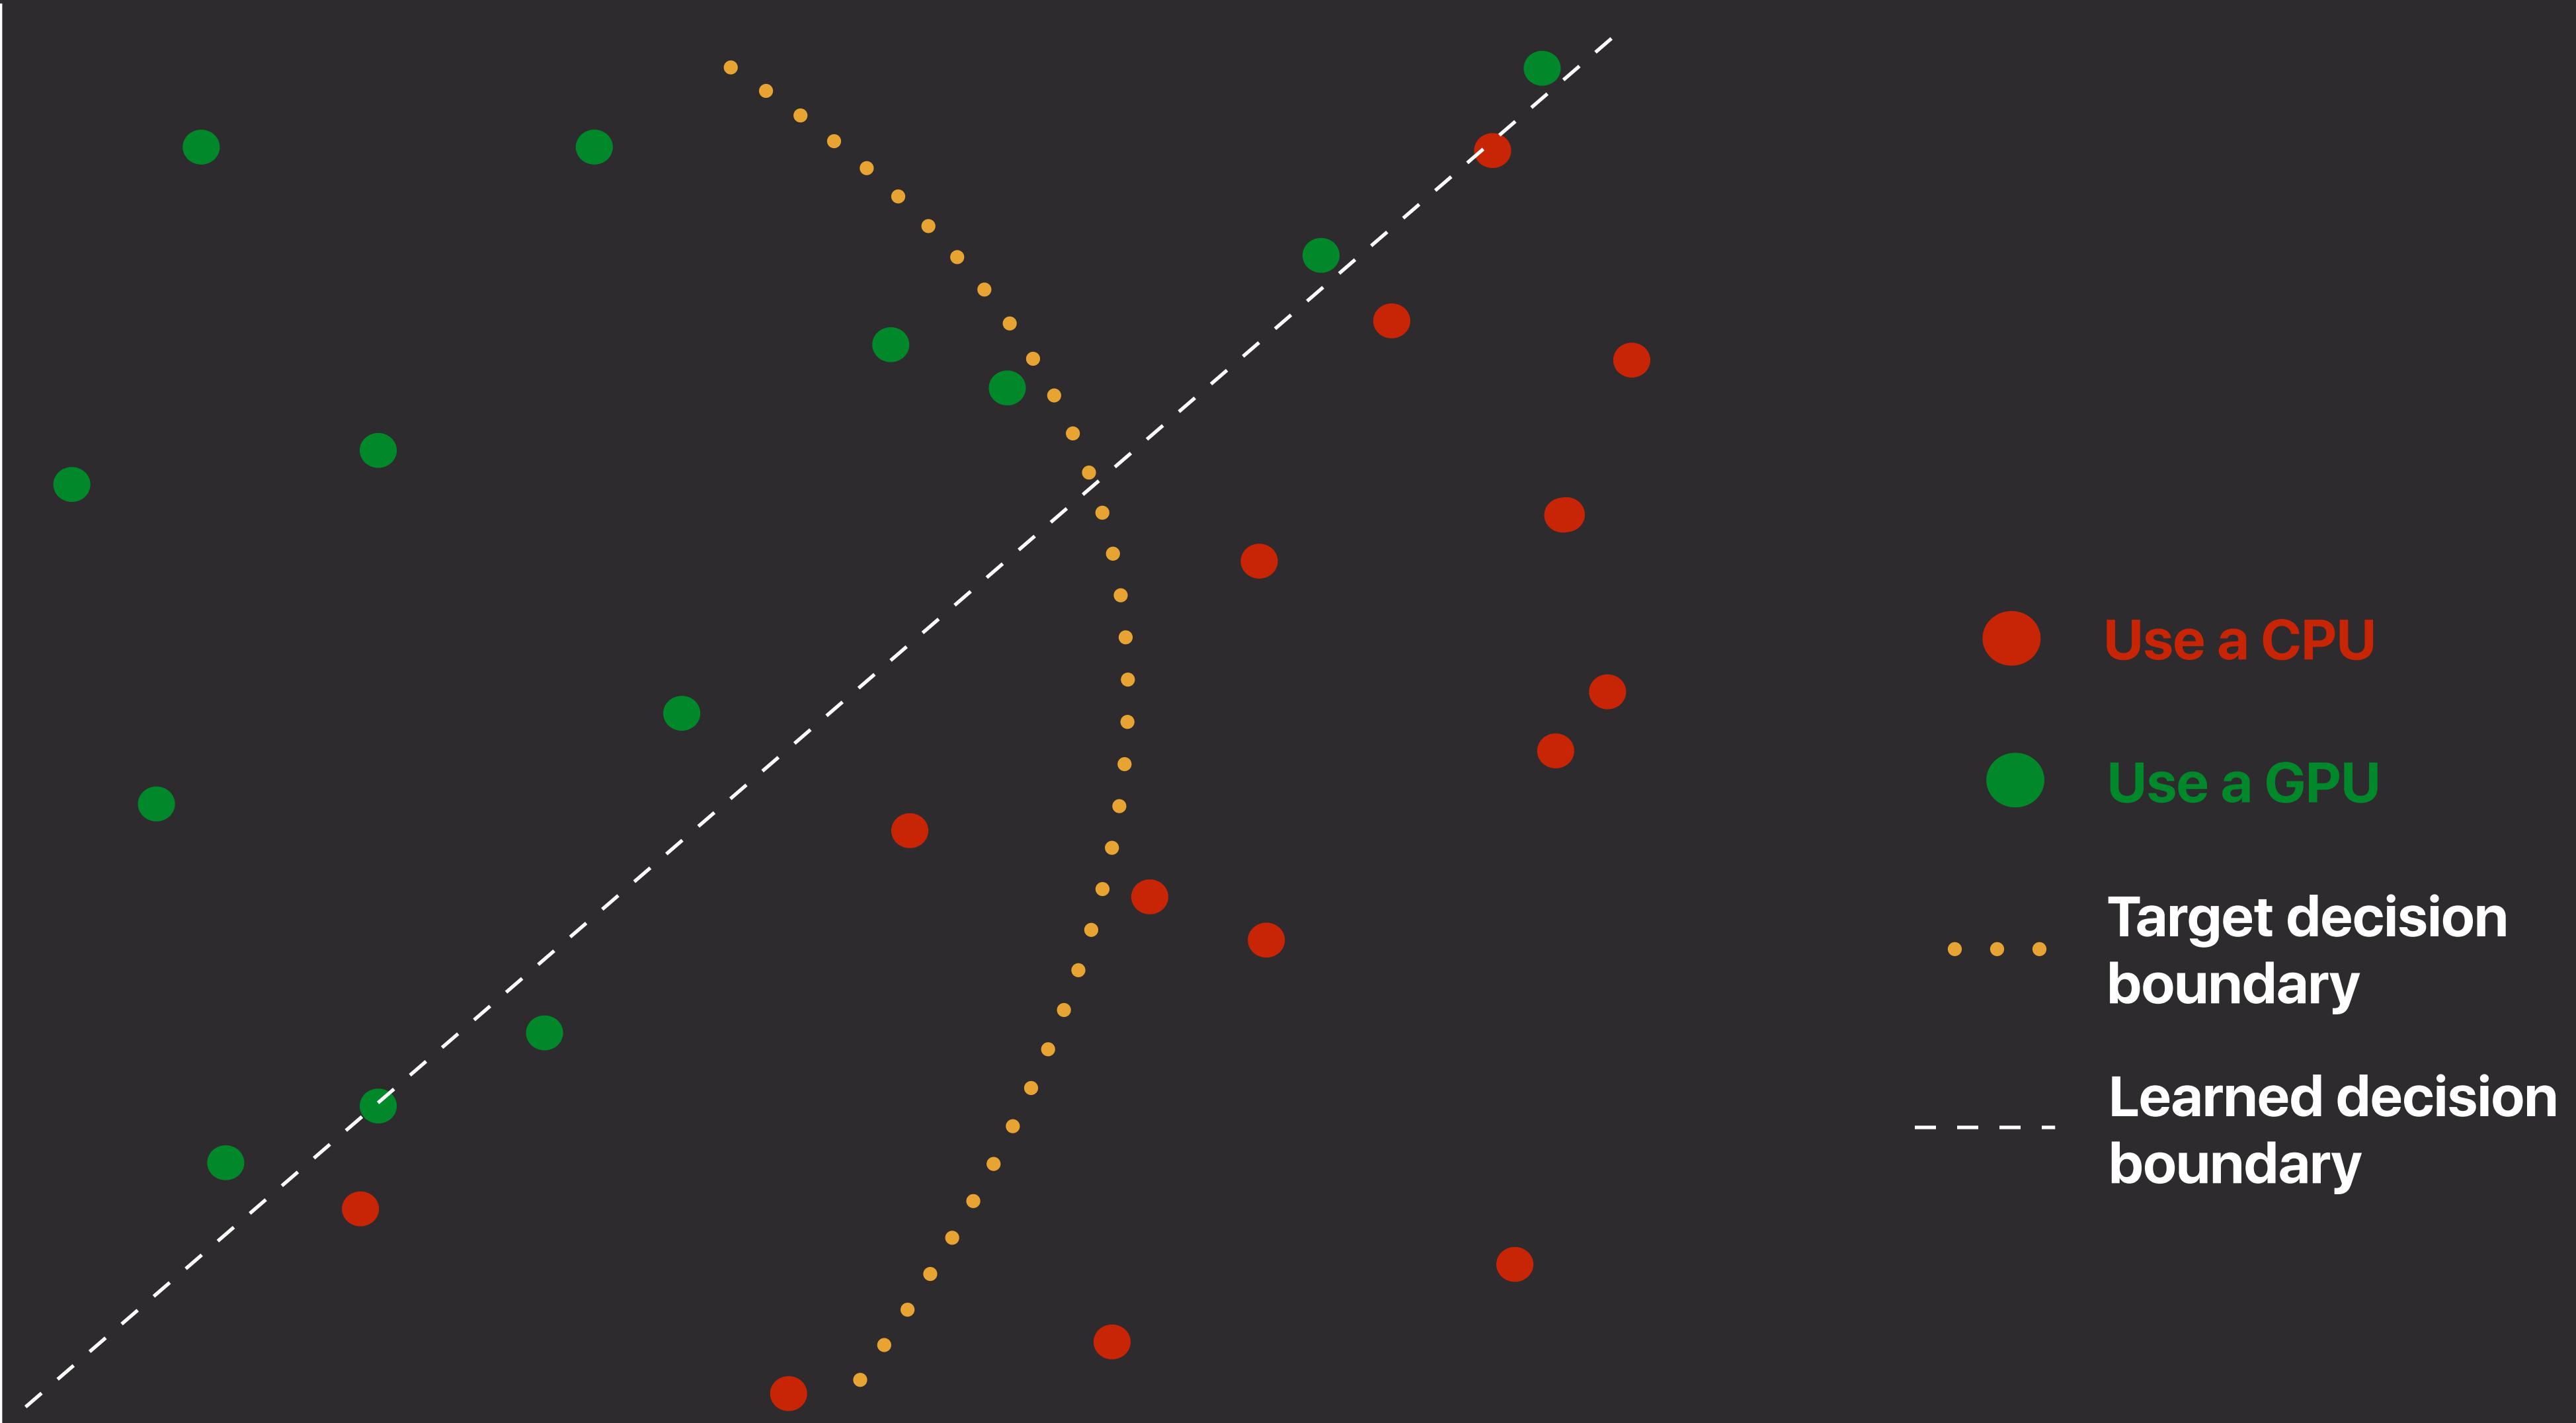
The idea.

machine learning for compilers



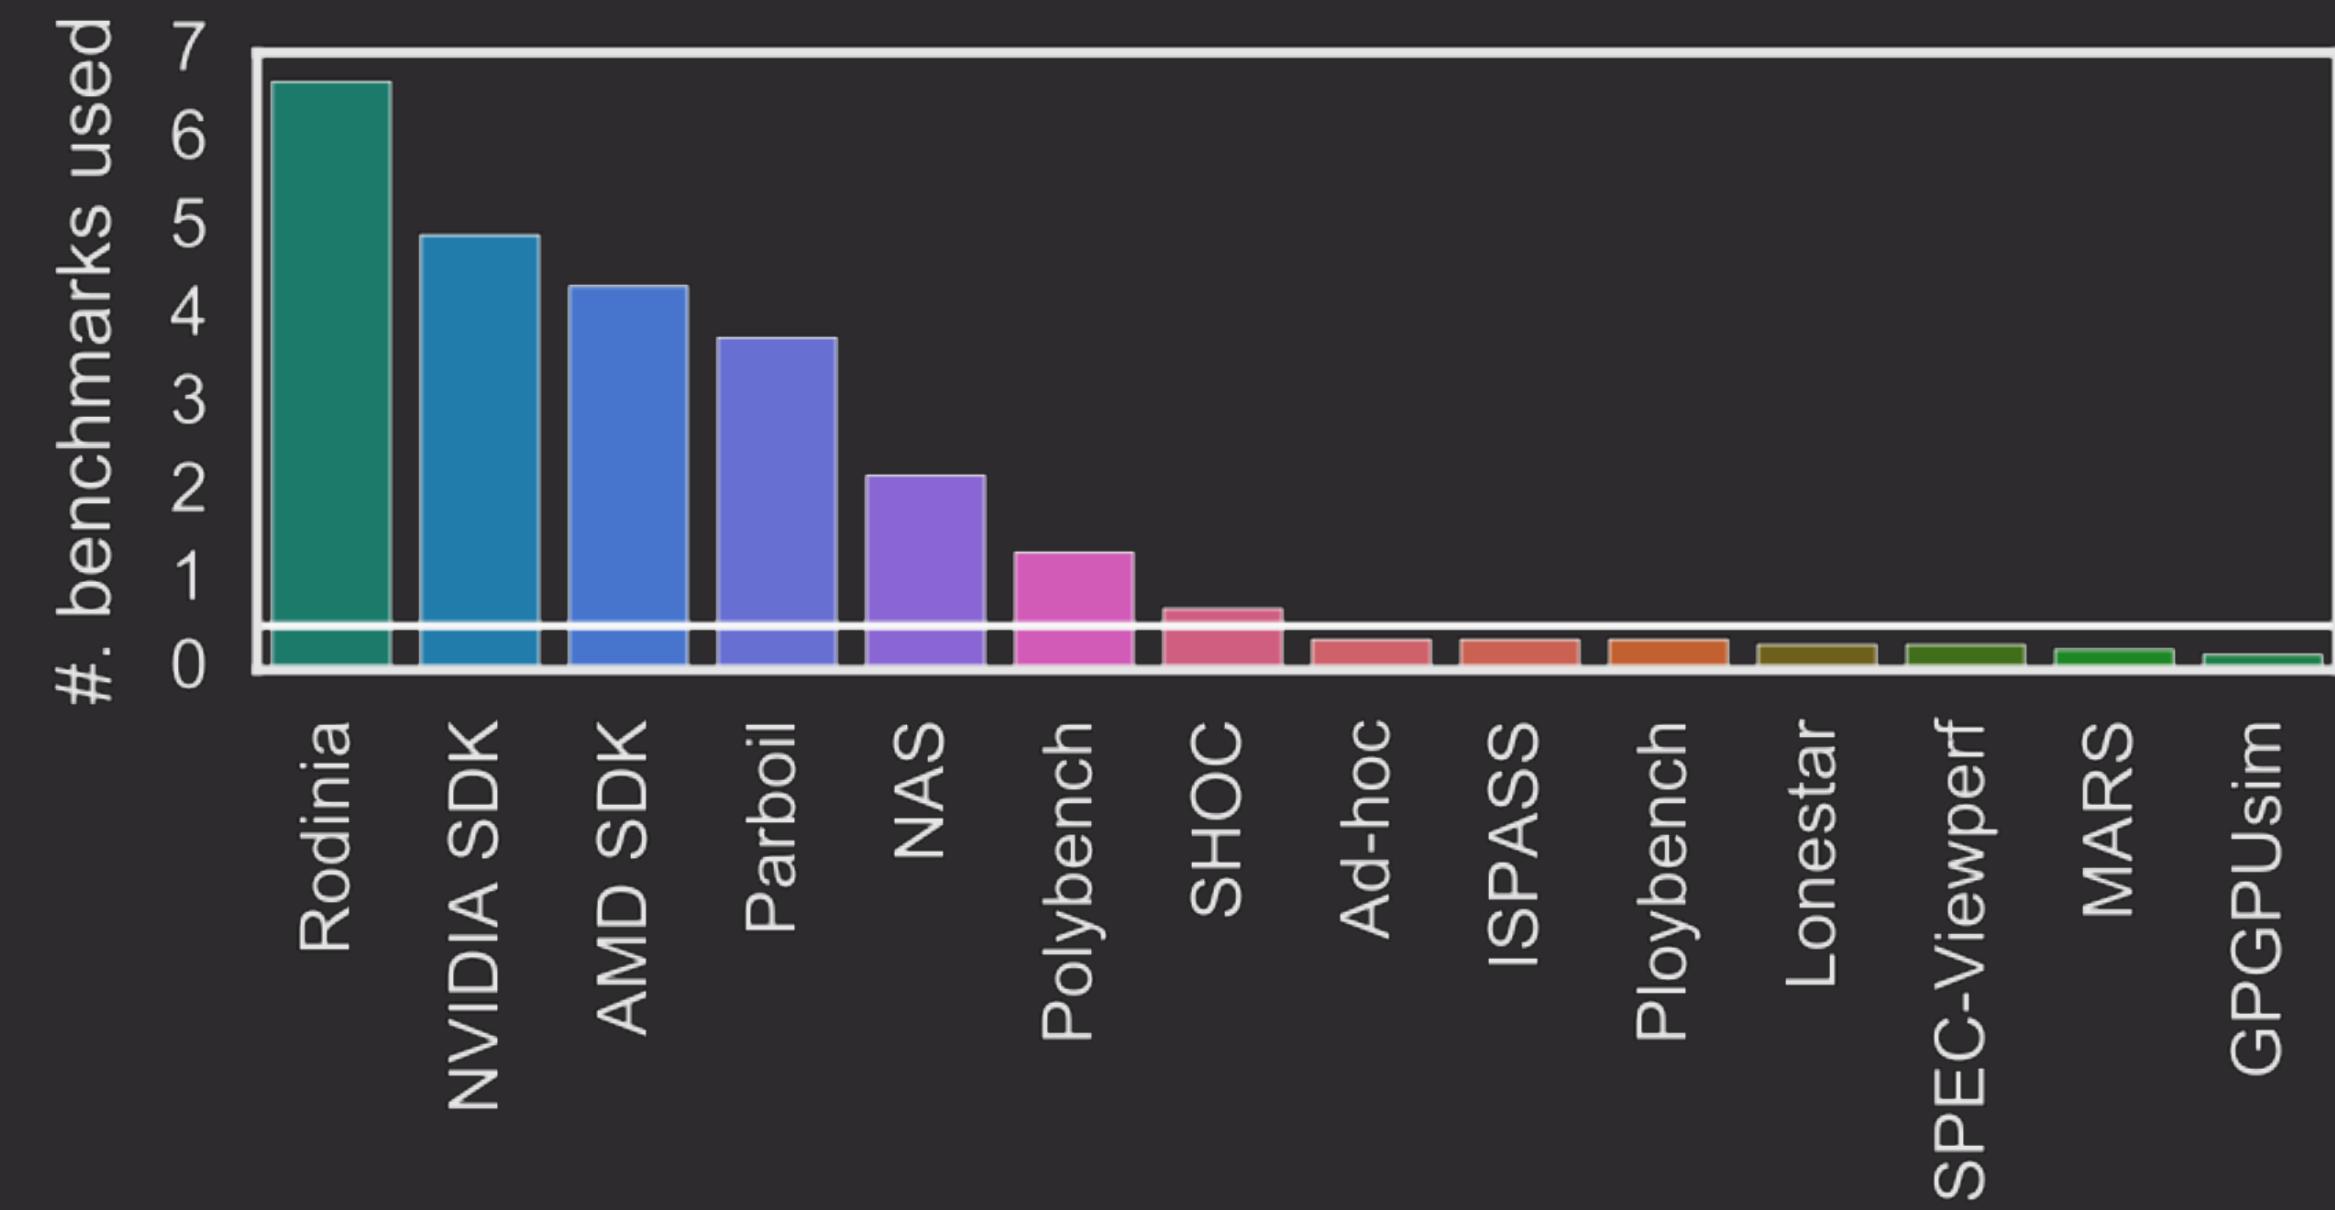
The reality.

machine learning for compilers

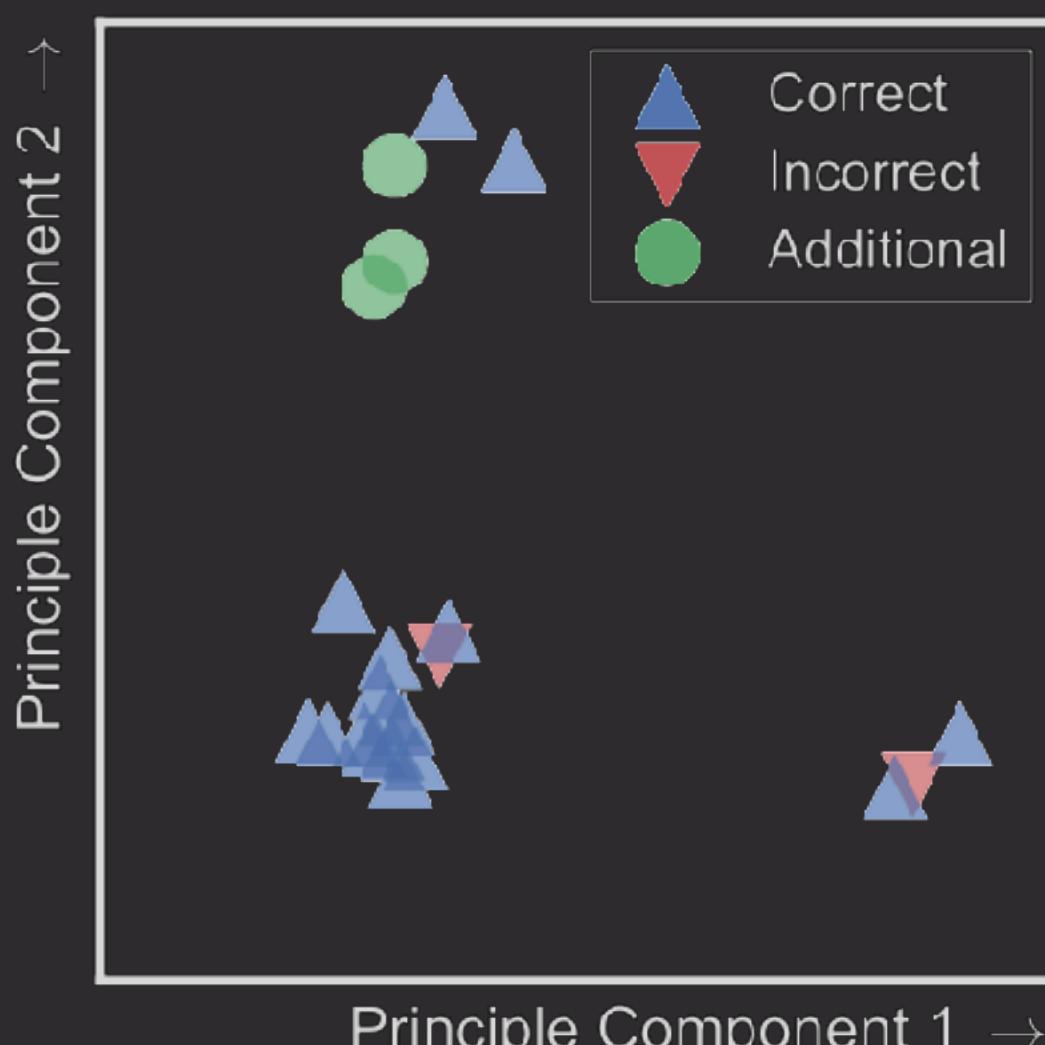
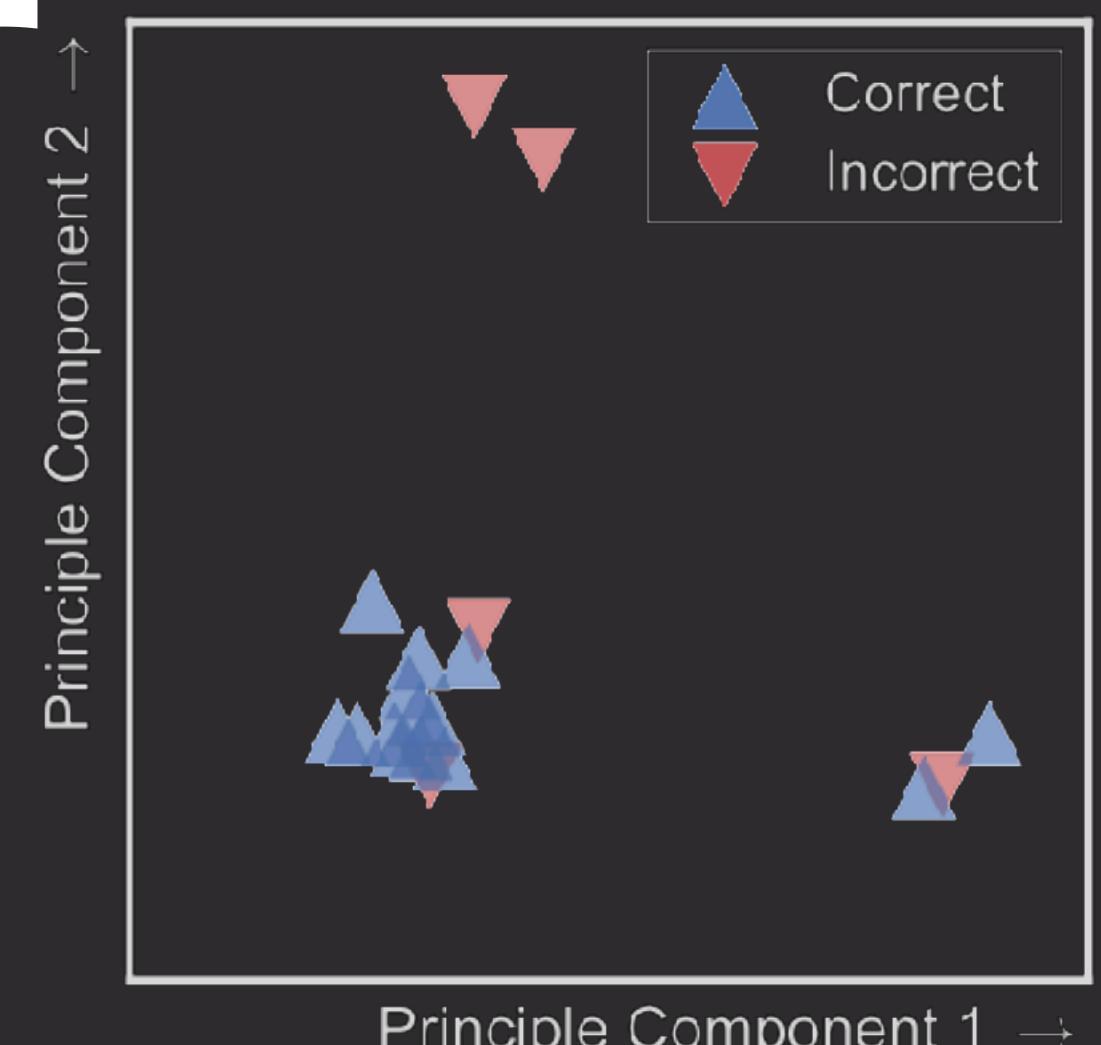


The reality.

1. there aren't enough benchmarks



2. more benchmarks = better models

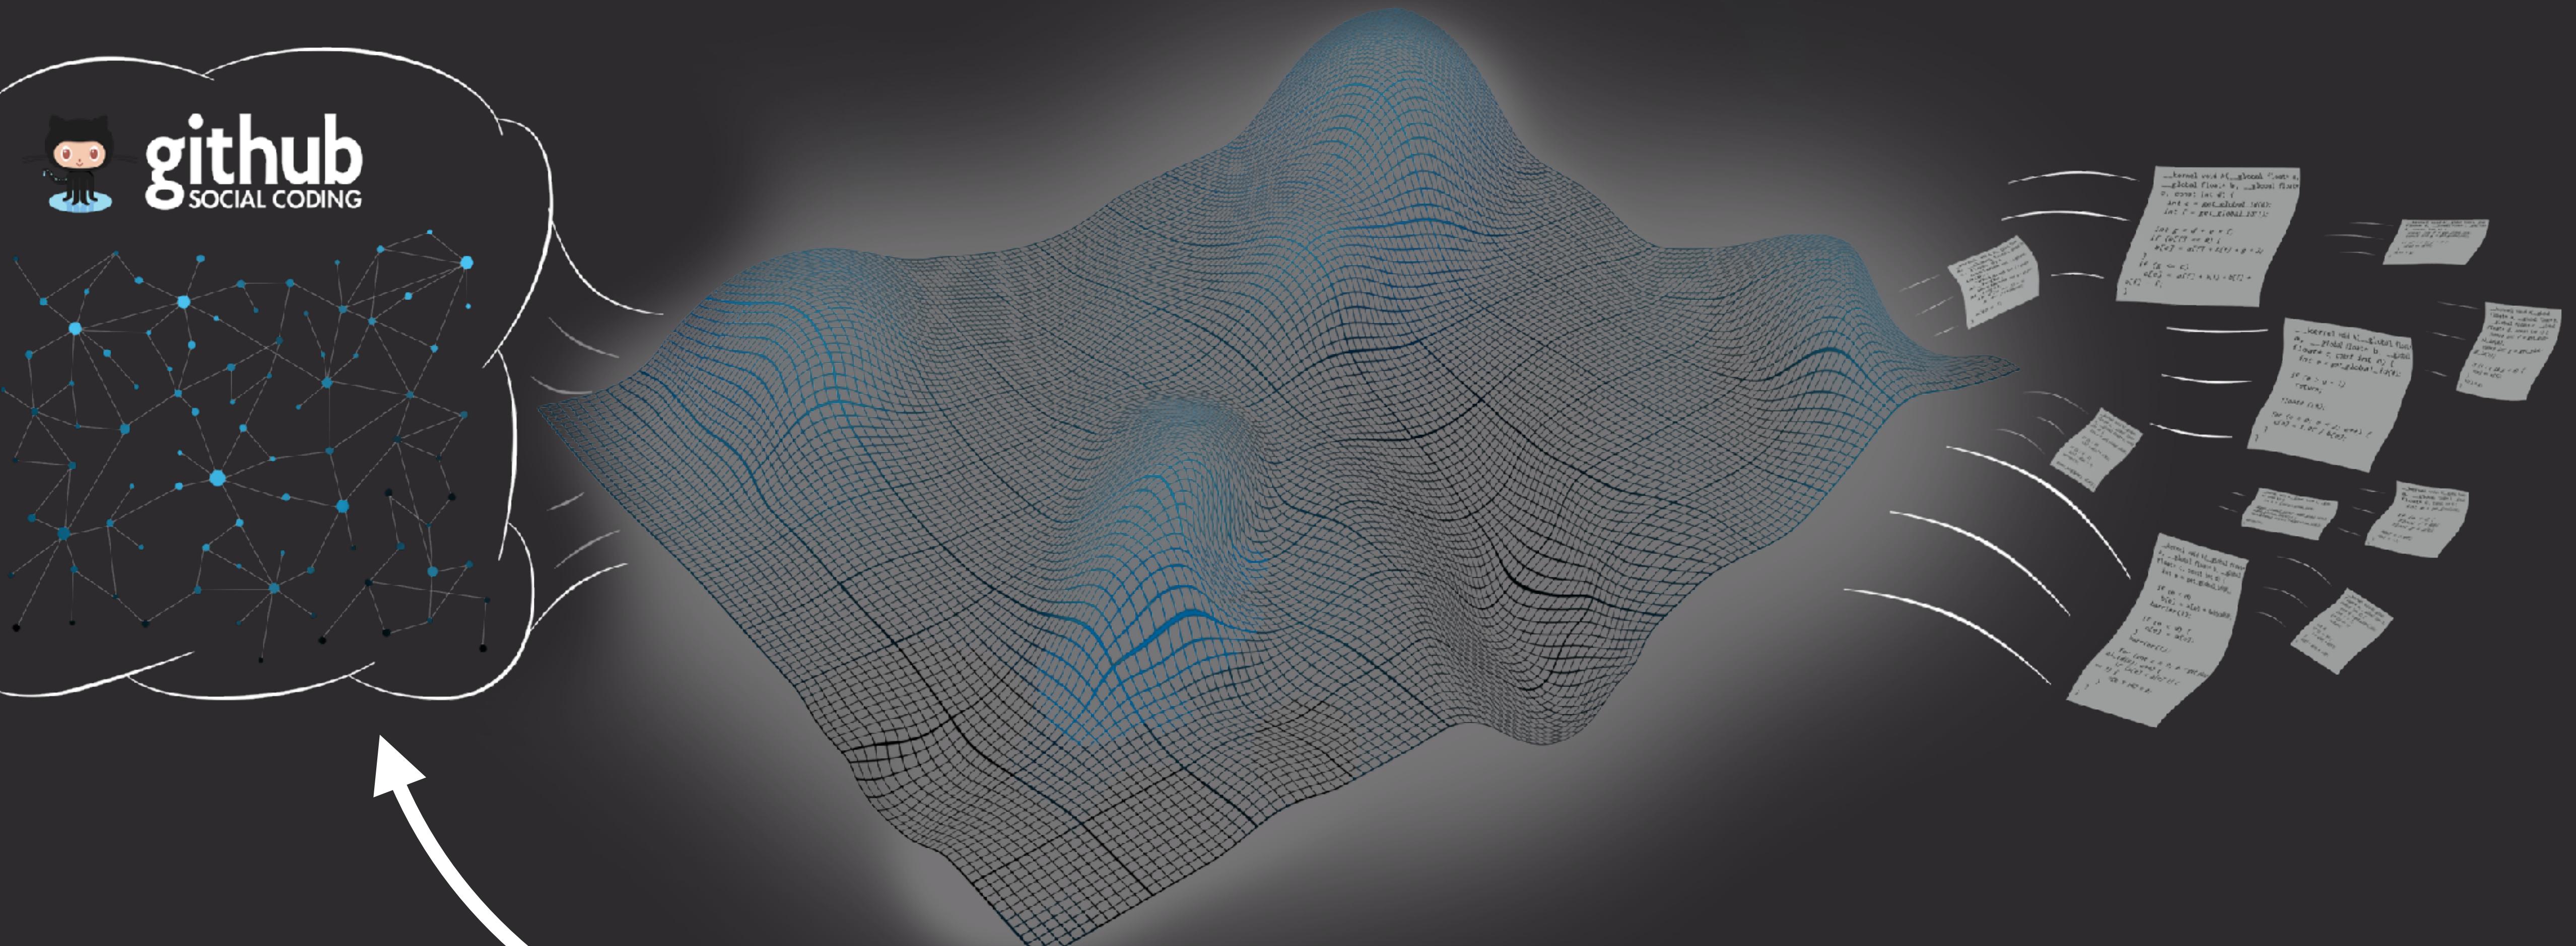


3. no adequate solution

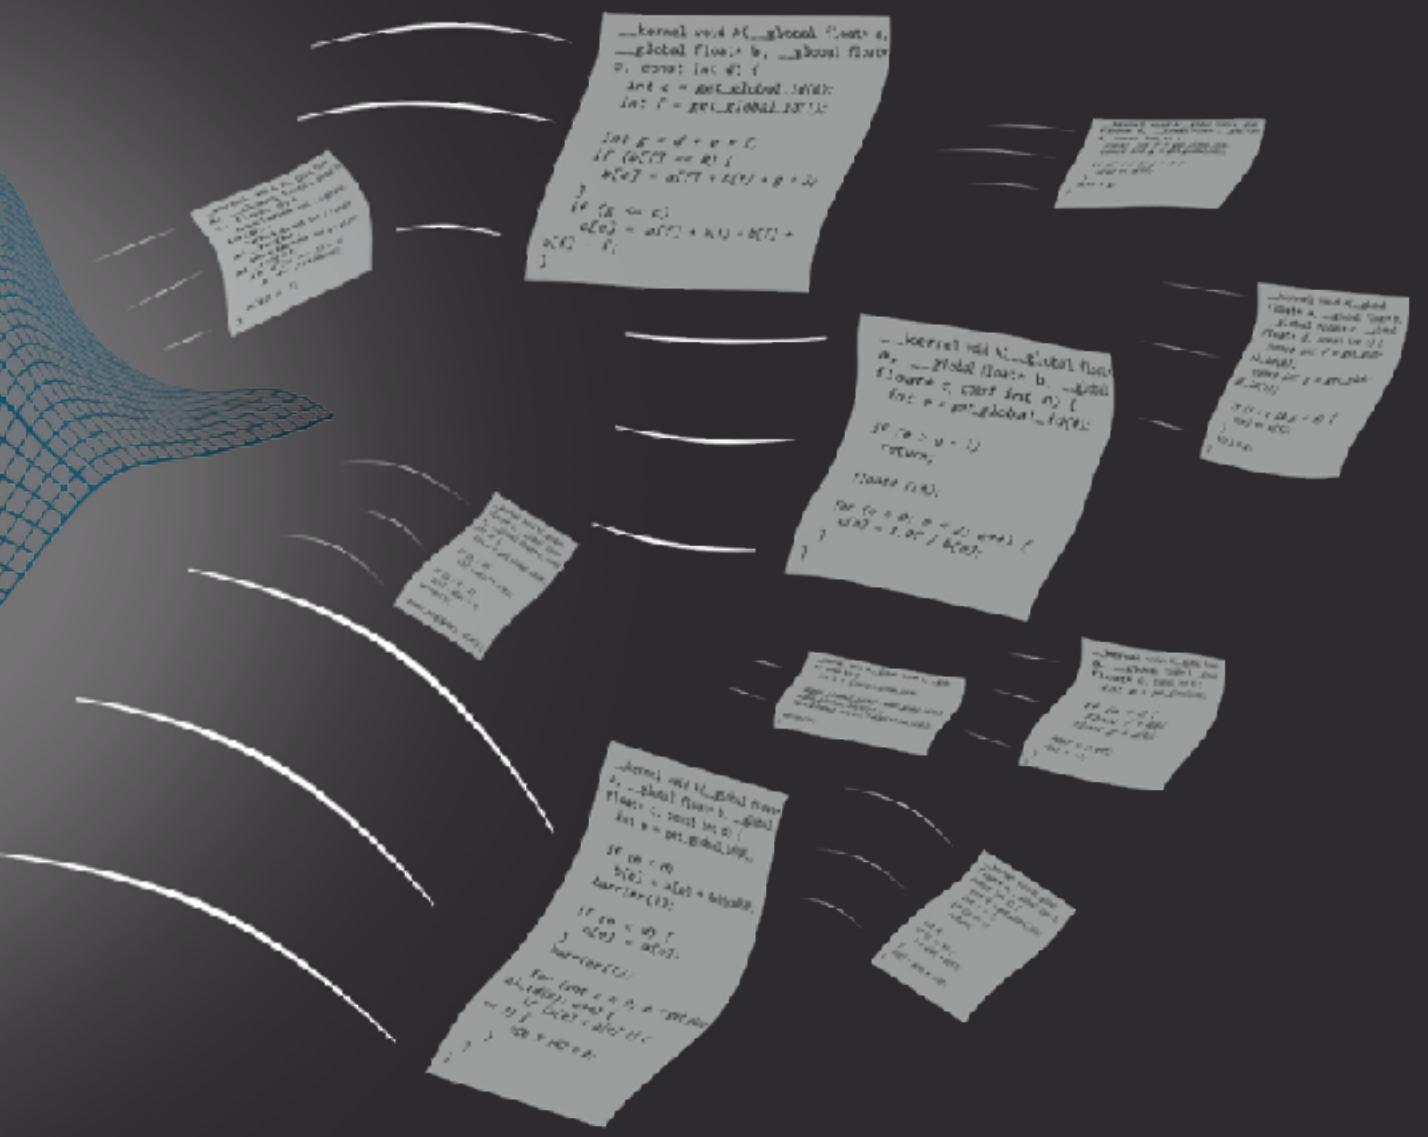
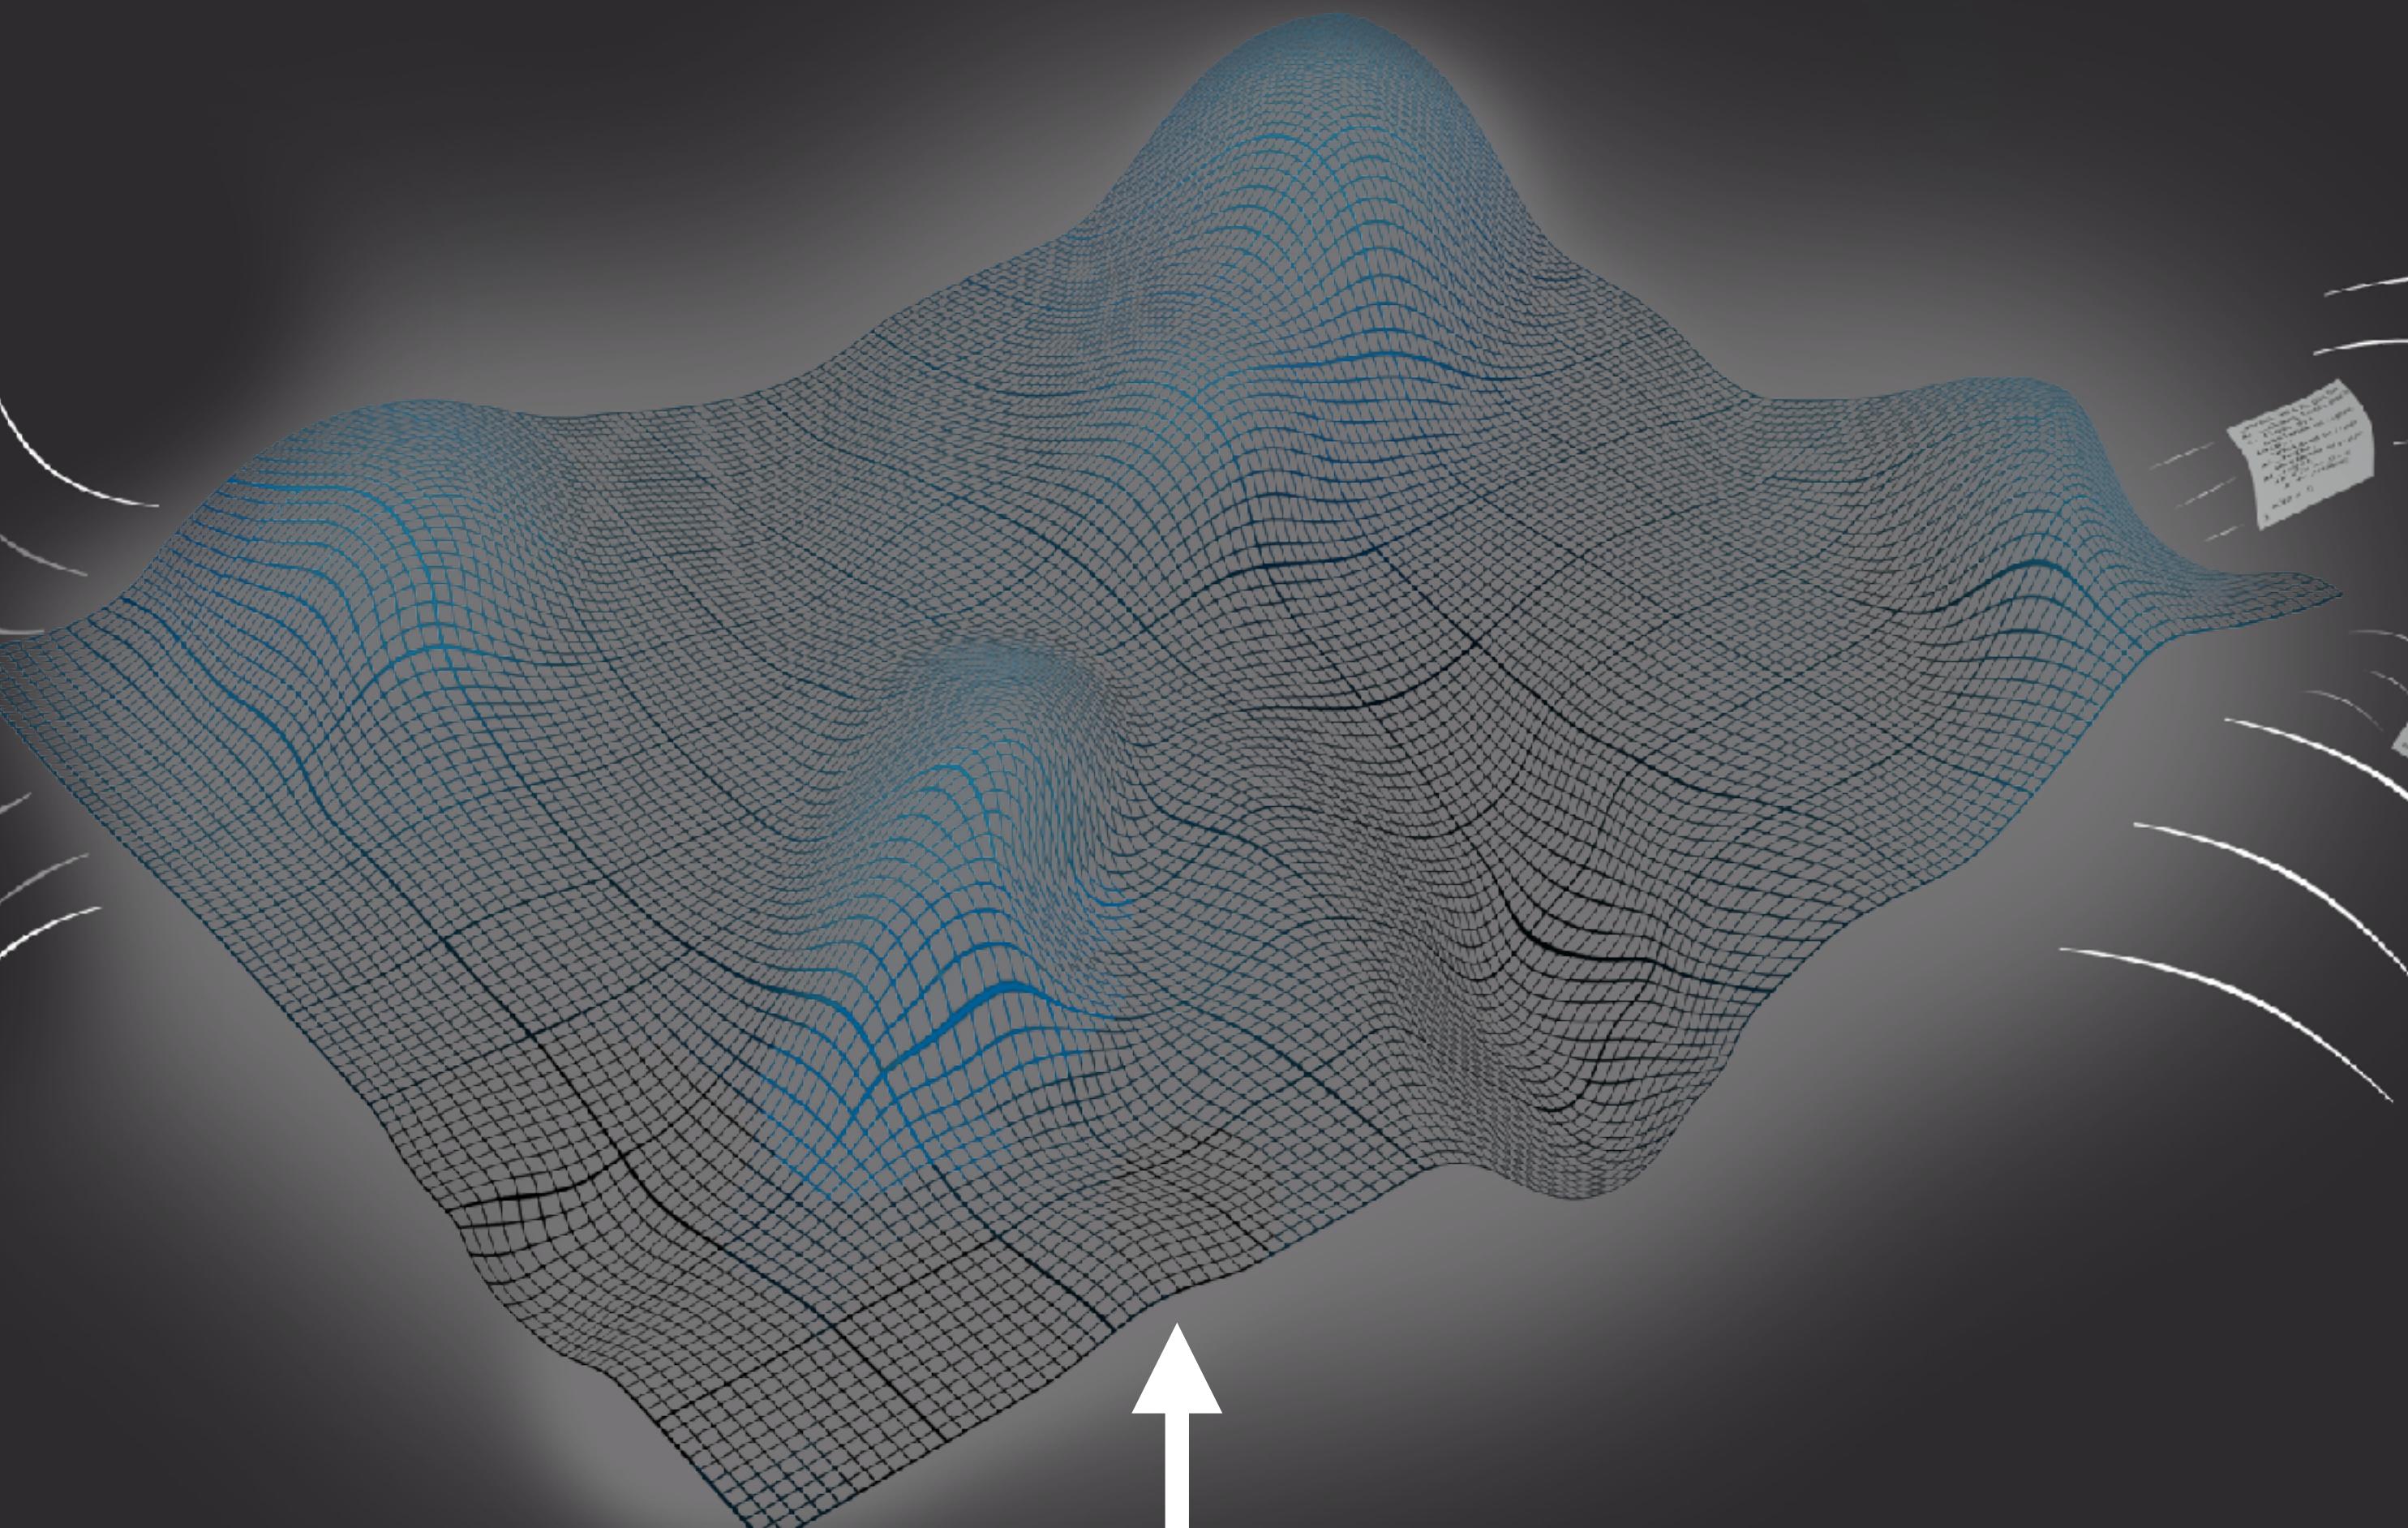
```
1 begin genesis
2
3 geninclude gen_c.glb
4
5 global
6   distribution epochdist = {1:10}
7   distribution numvardist = {8;16;32}
8   distribution compdist = {1:20}
9   distribution coefdist = {0:7}
10  distribution offsdist = {0:15}
11 end
12
13 program
14   value numepochs sample epochdist
15   value numvars sample numvardist
16   varlist temp[$numvars]
17 end
18
19 feature computation
20   variable dest,src1,src2,src3 from temp
21   ${dest} = ${src1} * ${src2} + ${src3};
22 end
23
24 feature read_access
25   variable dest from temp
26   value coef1,coef2 sample coefdist
27   value offs sample offsdist
28   ${dest} = arr_in[ ${coef1}*it0 + ${coef2}*
29           it1 +${offs }];
30 end
31
32 feature write_access
33   variable src from temp
34   value offs sample offsdist
35   arr_out[ inner_tc*it0 + it1 +${offs}] = ${src};
36 end
37
38 feature epoch (epoch_type)
39   value numcomps sample compdist
40   genloop i:1:$numcomps
41     ${computation}
42   end
43   genif ${epoch_type} eq "read"
44     ${read_access}
45   genelse
46     ${write_access}
47   end
48 end
49
50 feature epochs
51   genloop i:1:$numepochs
52     ${epoch("read")}
53   end
54   ${epoch("write")}
55 end
56 generate 1000
57 end genesis
```

our approach:

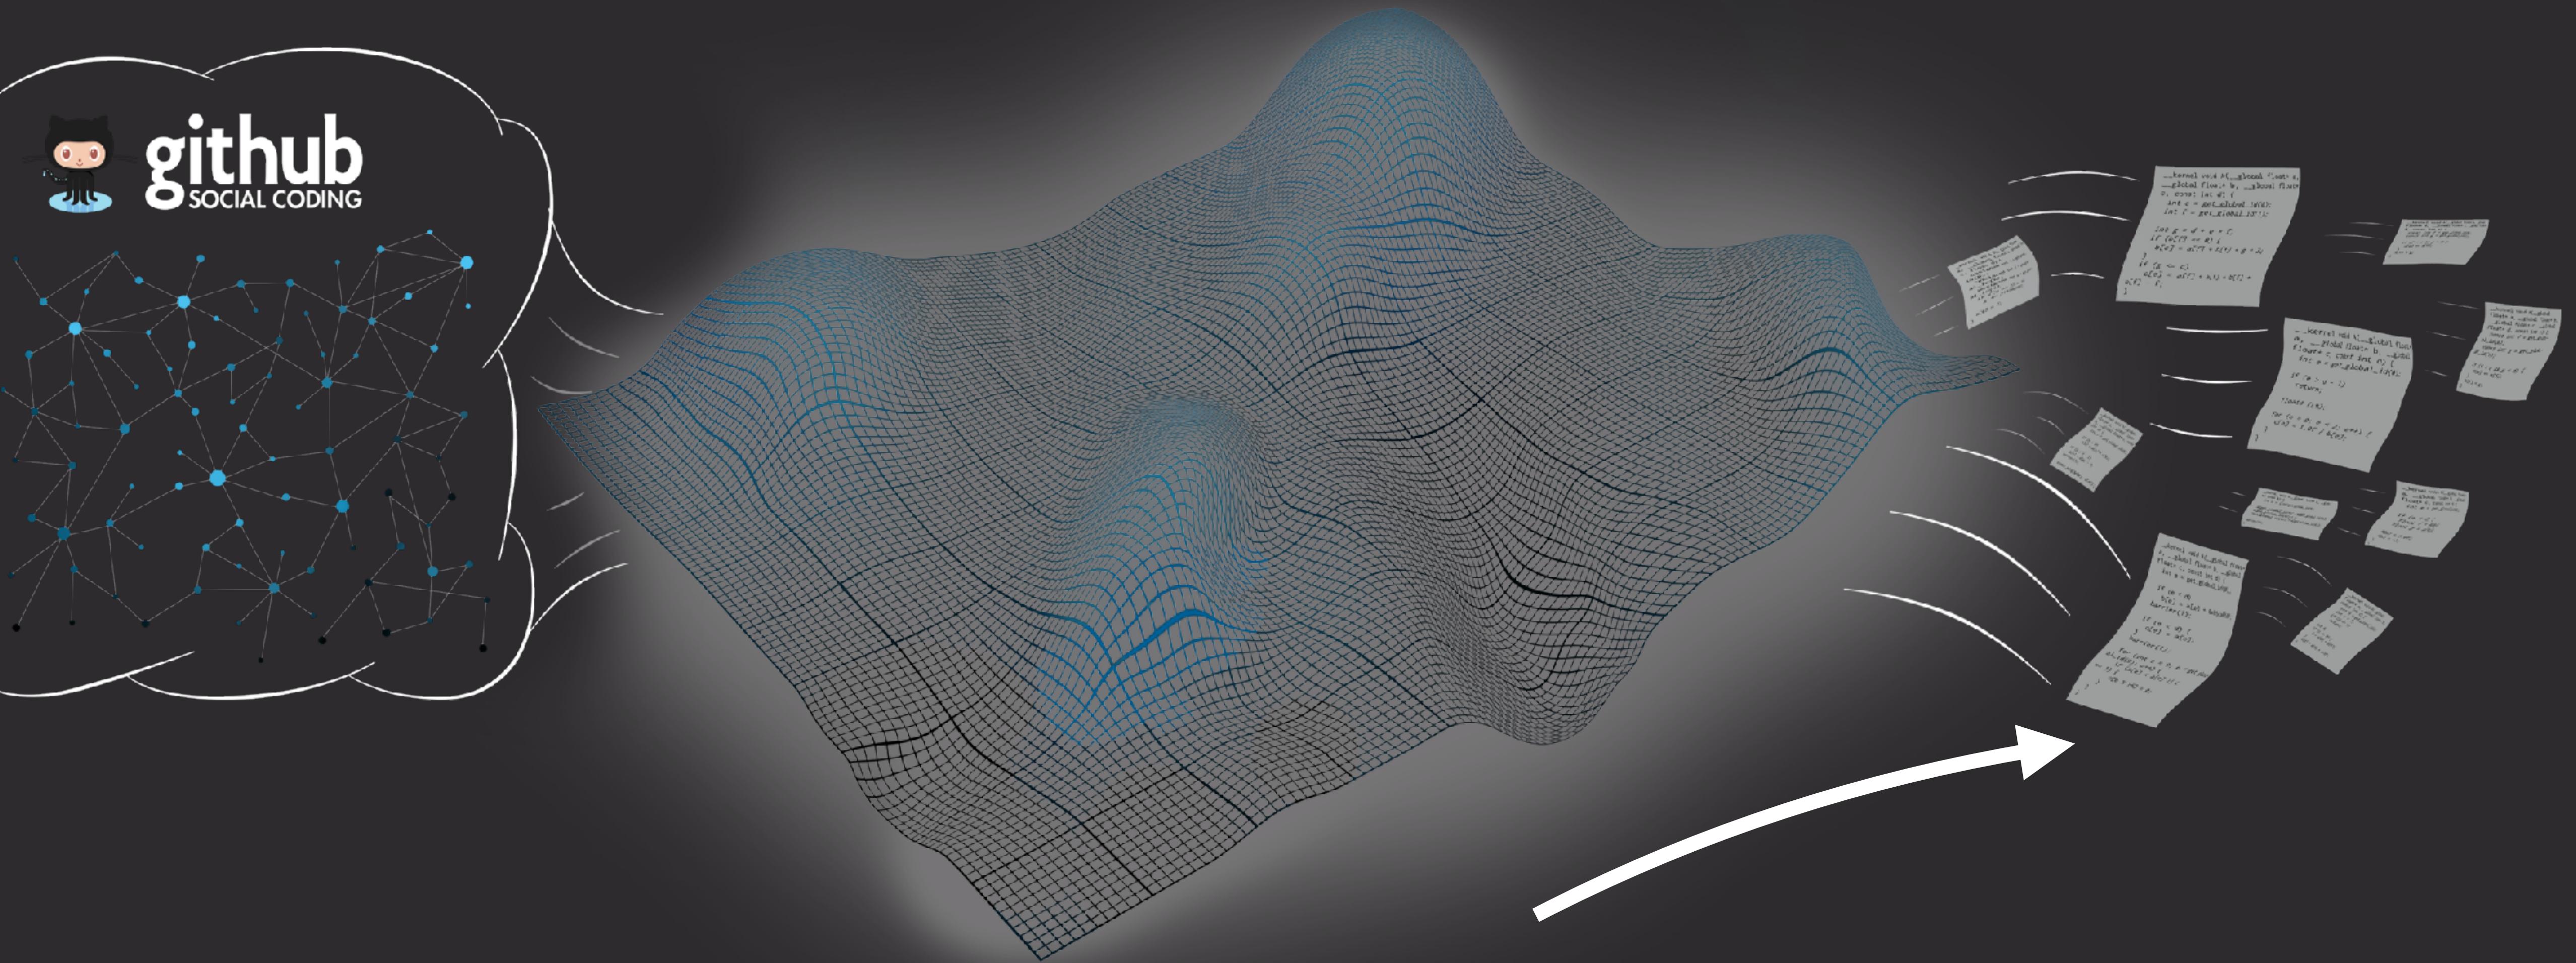
automatic program generation



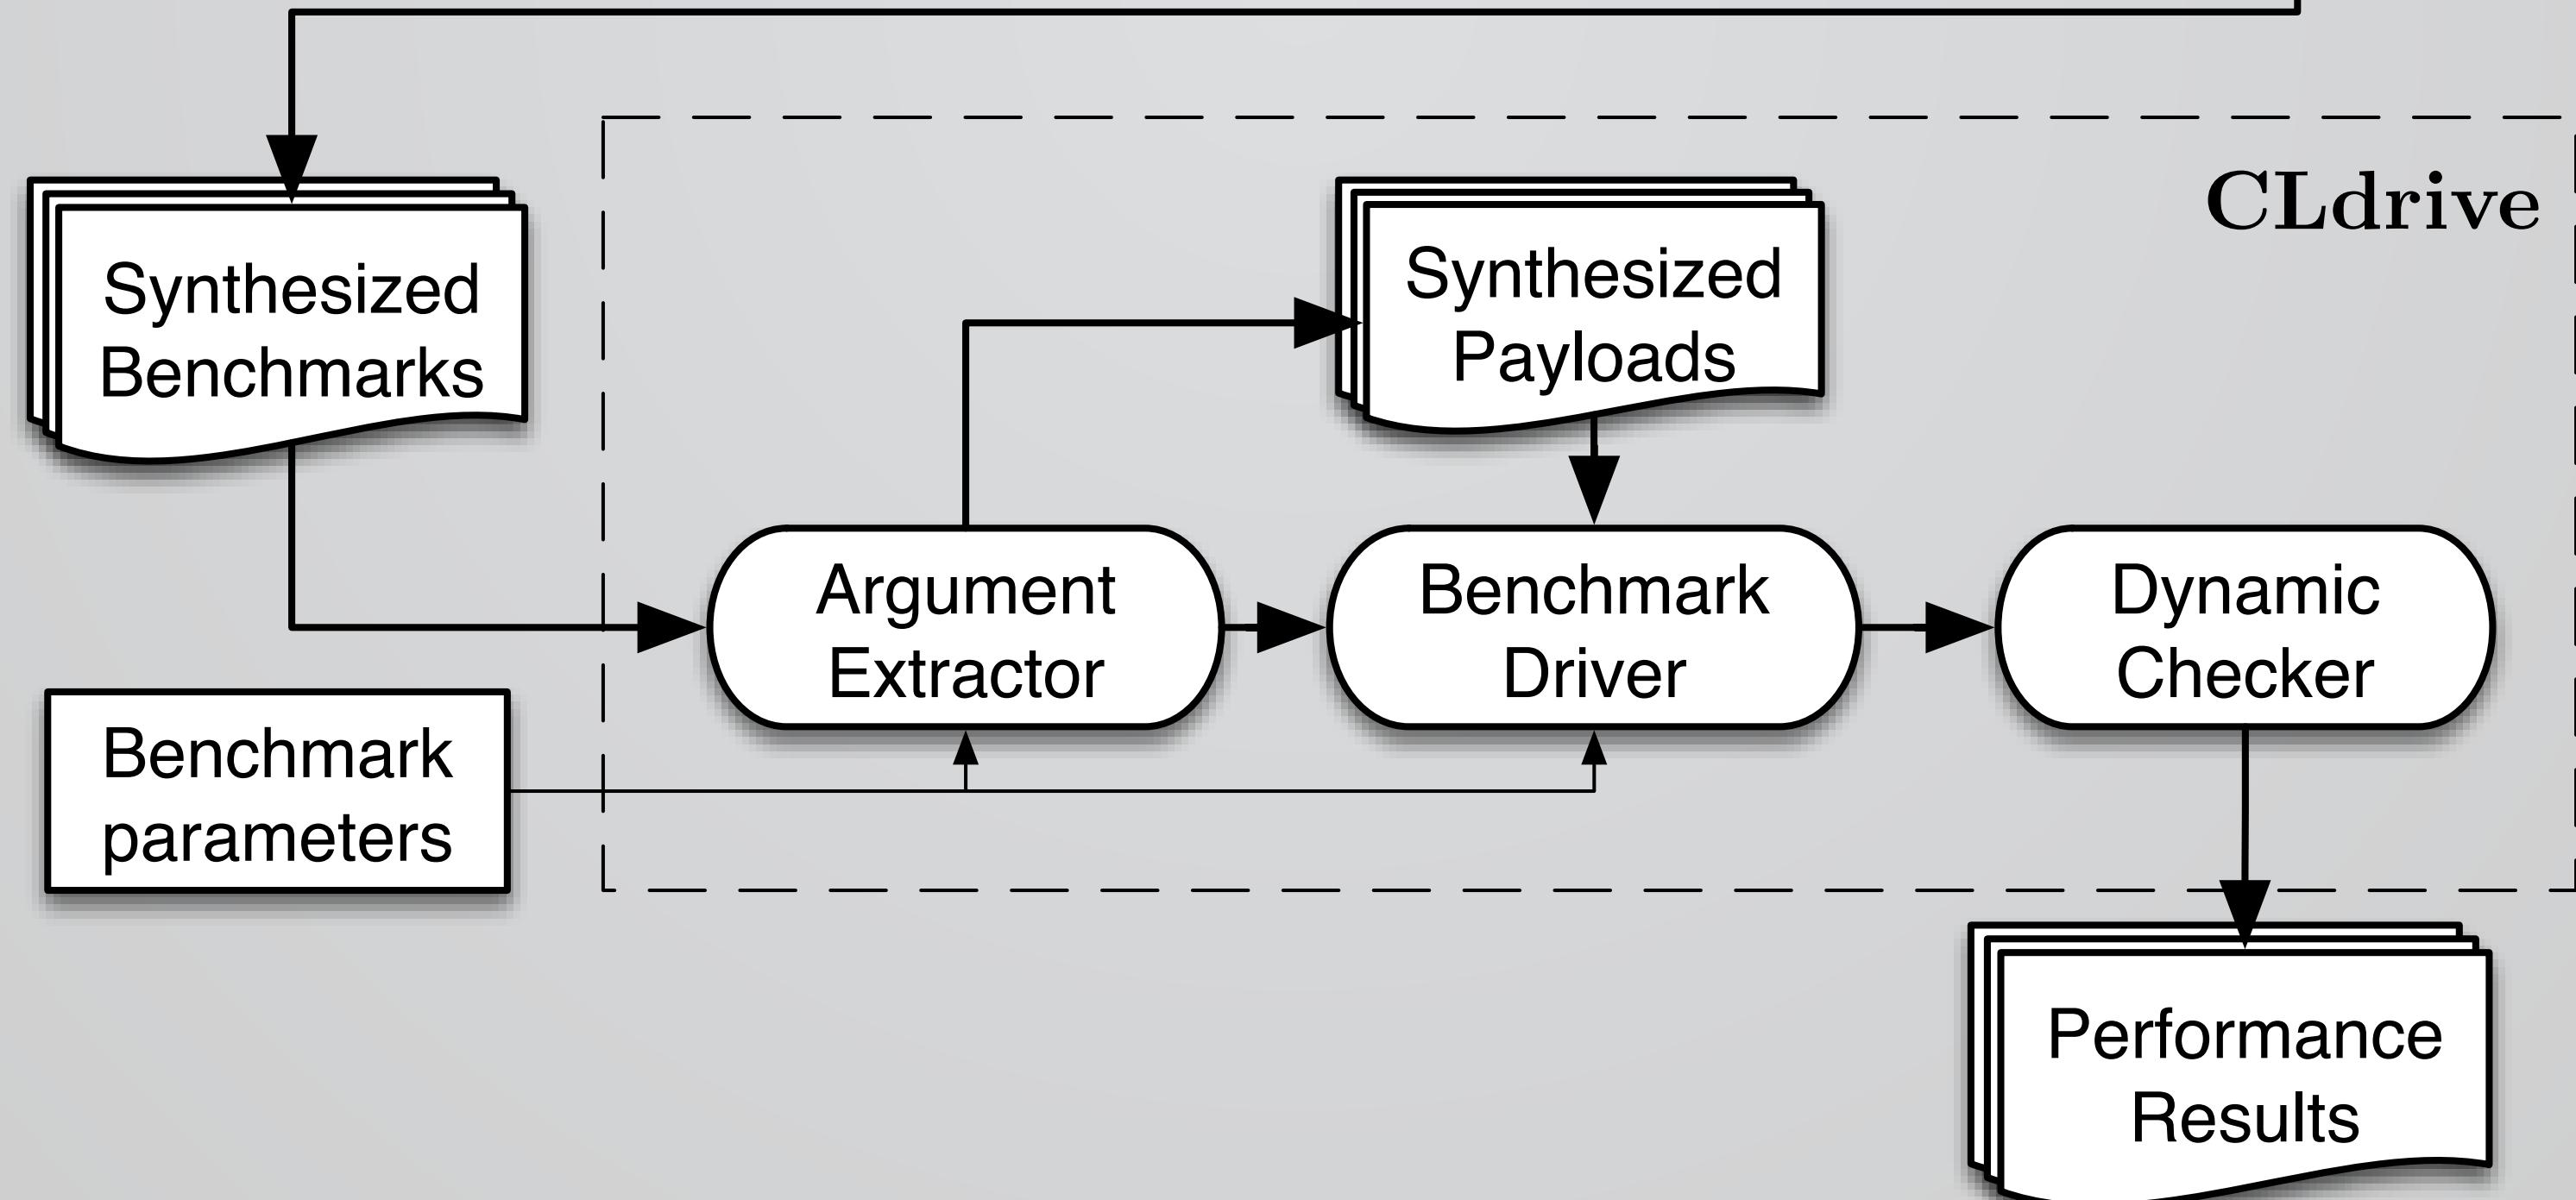
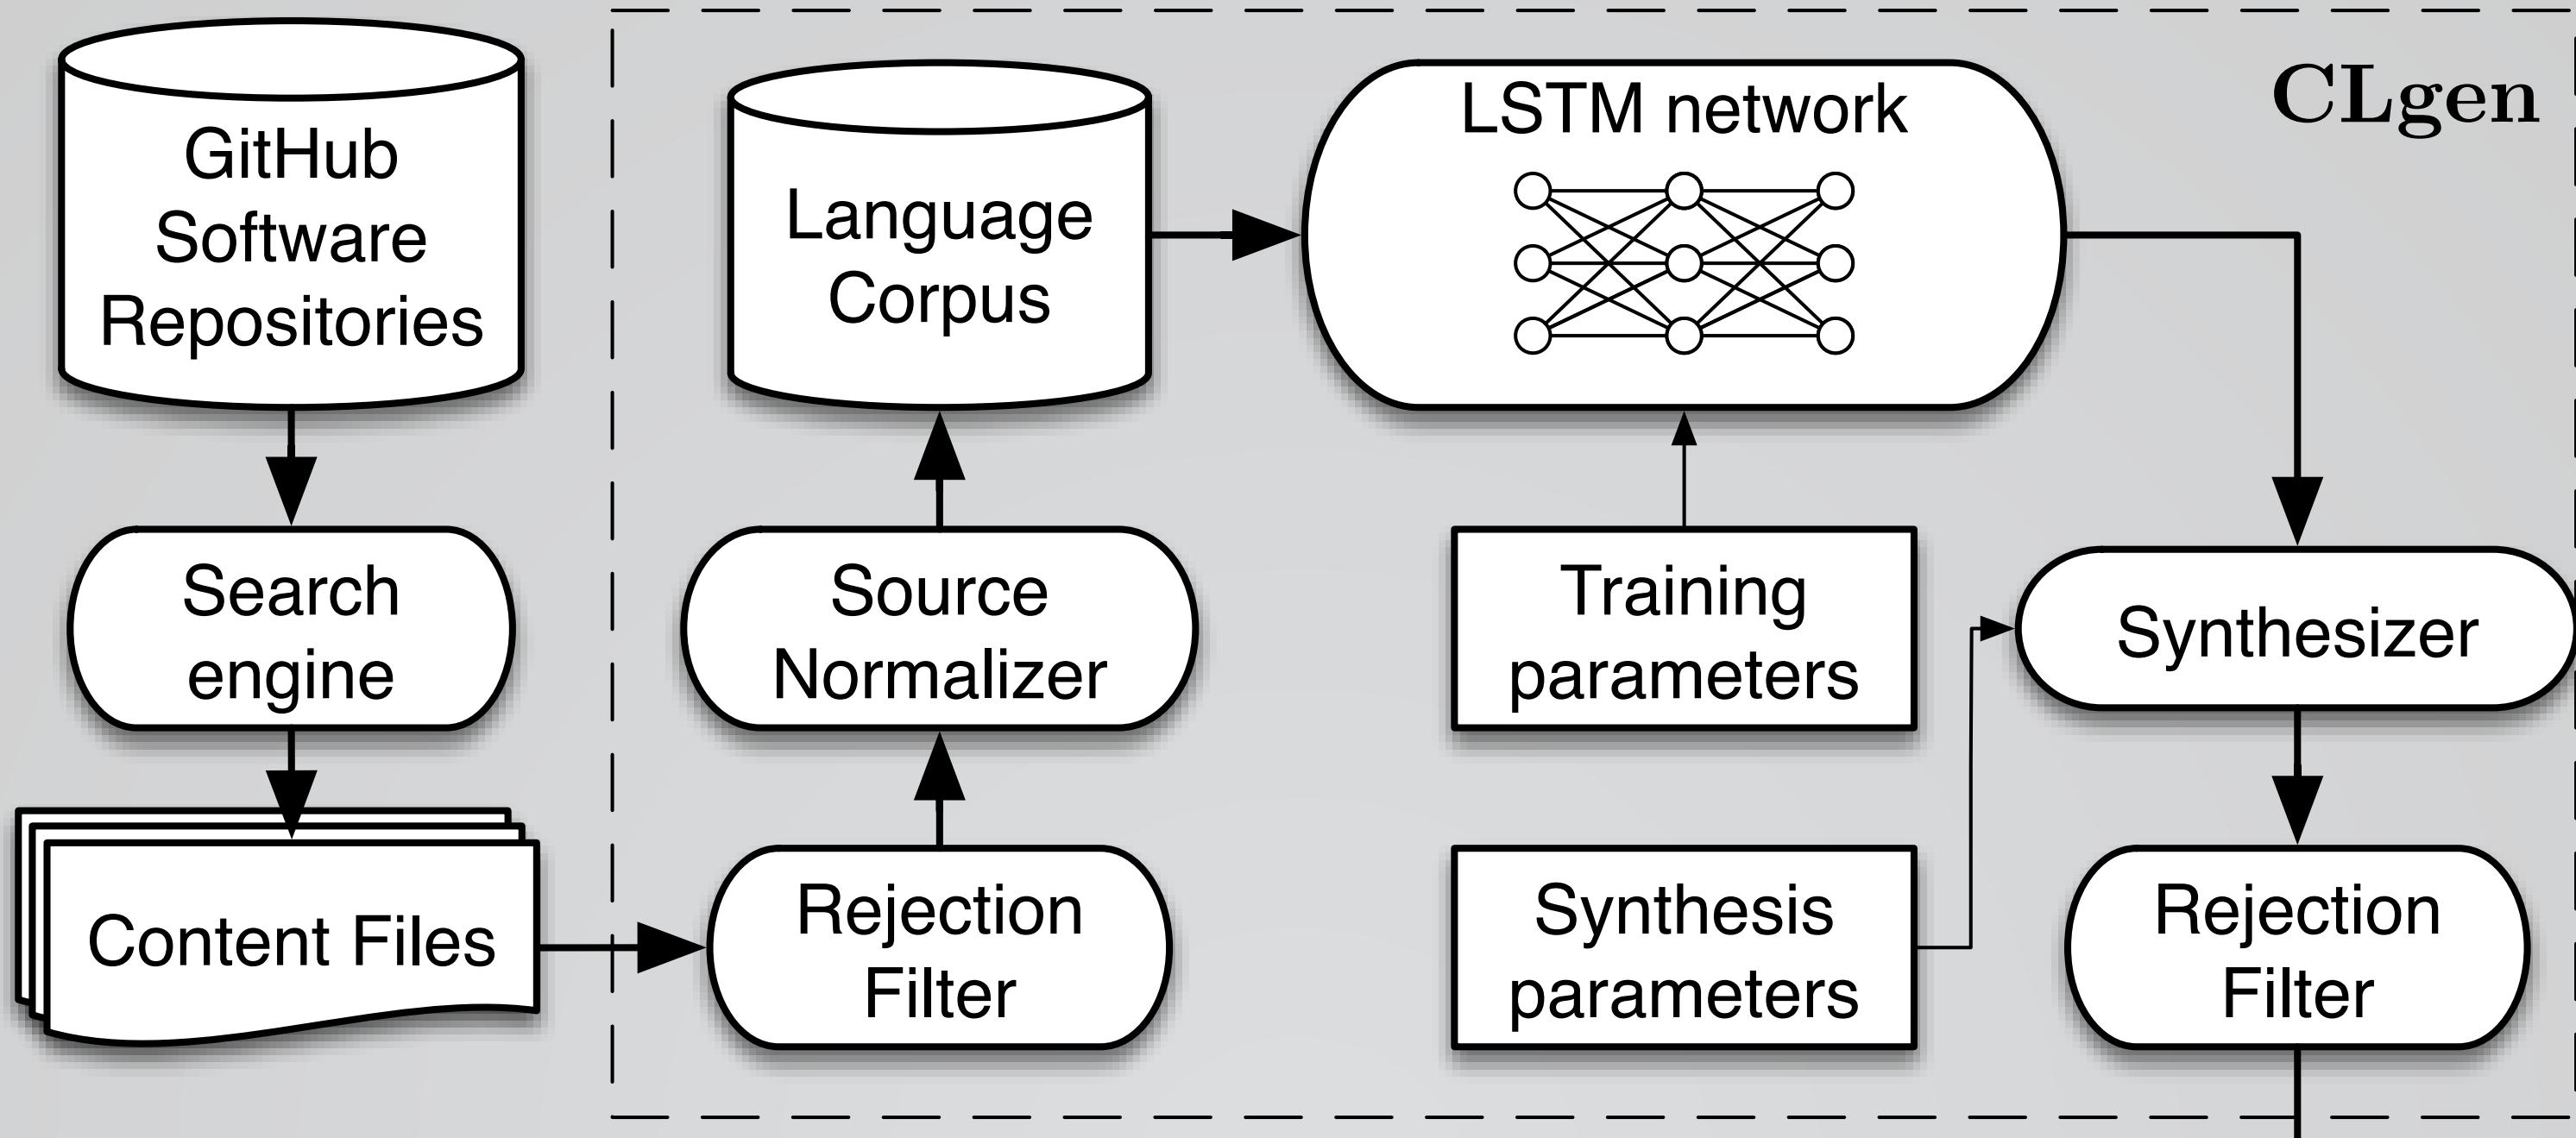
mine code from web

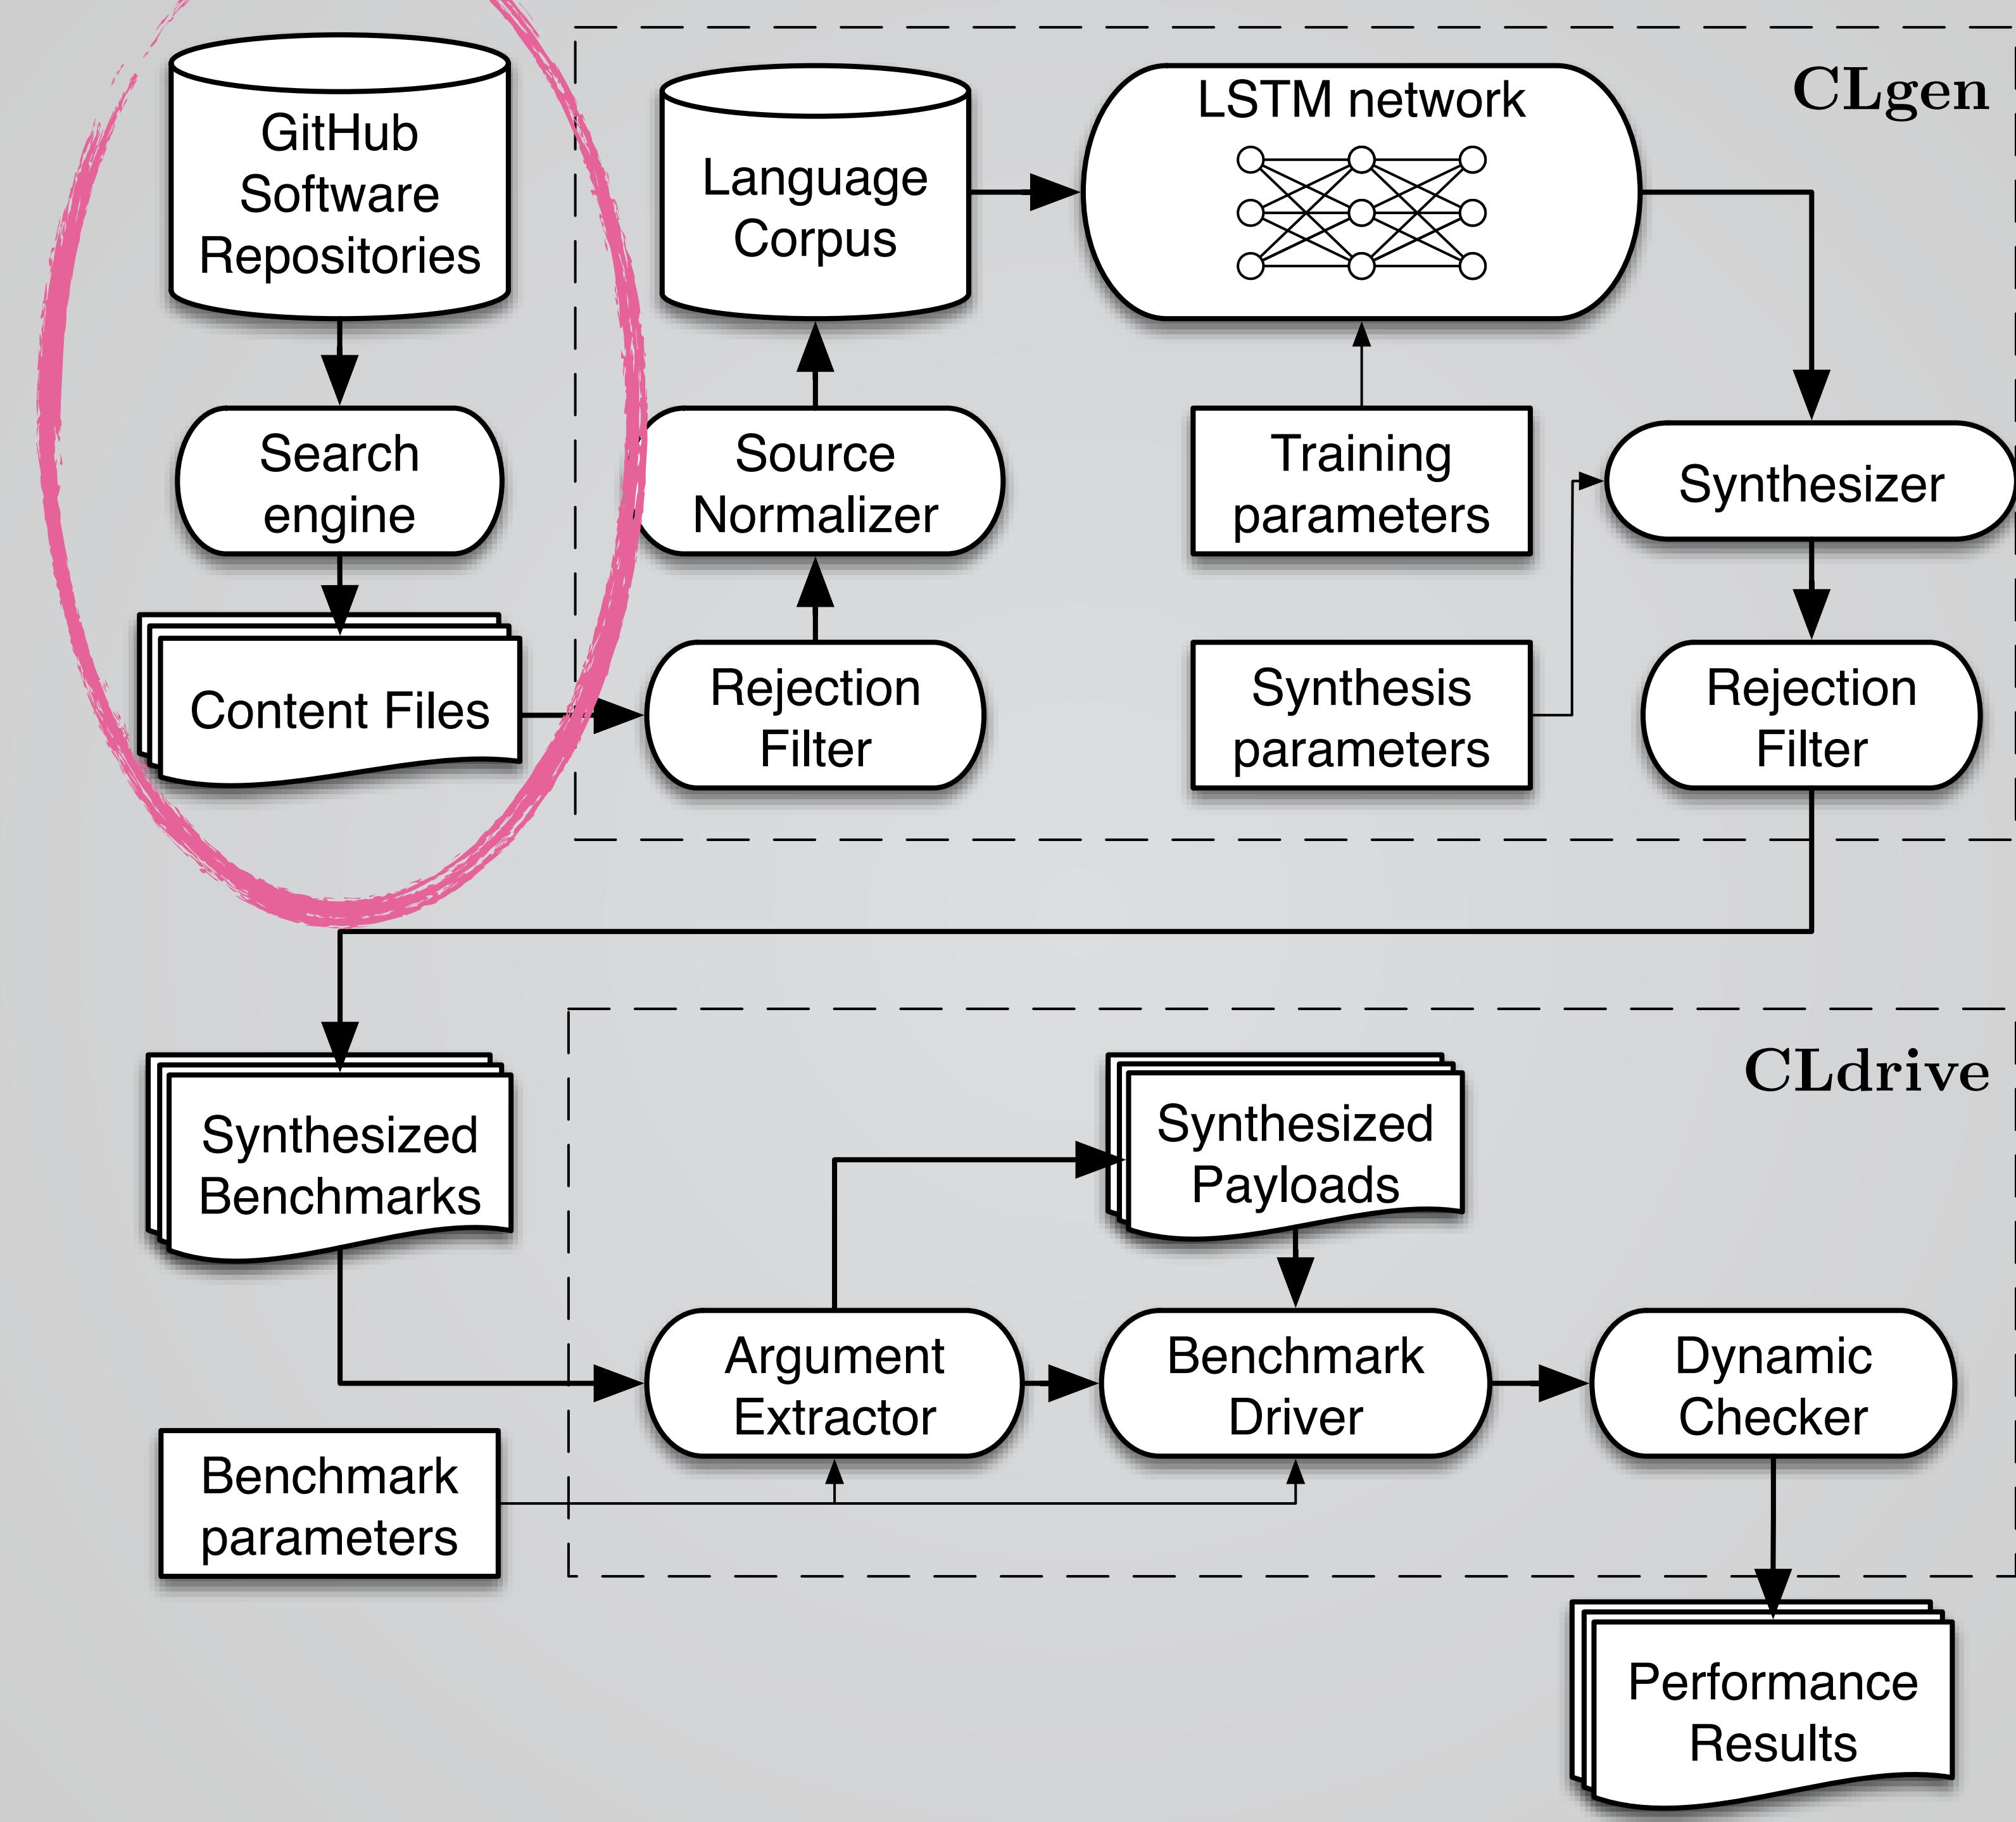


model source distr.



sample and reject





Infer the common usage of a PL from samples.

Huge repository of public knowledge:



And they have an API :-)

```
$ curl https://api.github.com/search/repositories\?
q\=opengl\&sort\=stars\&order\=desc
{
  "total_count": 3155,
  "incomplete_results": false,
  "items": [
    {
      "id": 7296244,
      "name": "lwjgl3",
      "full_name": "LWJGL/lwjgl3",
```

OpenCL is not a first-class language.

Search repositories using loose keyword terms.

e.g. `opencl`, `nvidia`, `gpu`, `cl`, `amd`.

Recursively iterate over git trees to get `.cl` files.

`Foo:MyOpenCLRepo`

 ↳ `/src/gaussian.cl`

 ↳ `#include <common.h>`

(0.6% miss rate) ↳ `/include/common.h`

 ↳ `#include "detail/math.cl"`

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */
//
//          DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION
//
// 0. You just DO WHAT THE FUCK YOU WANT TO.

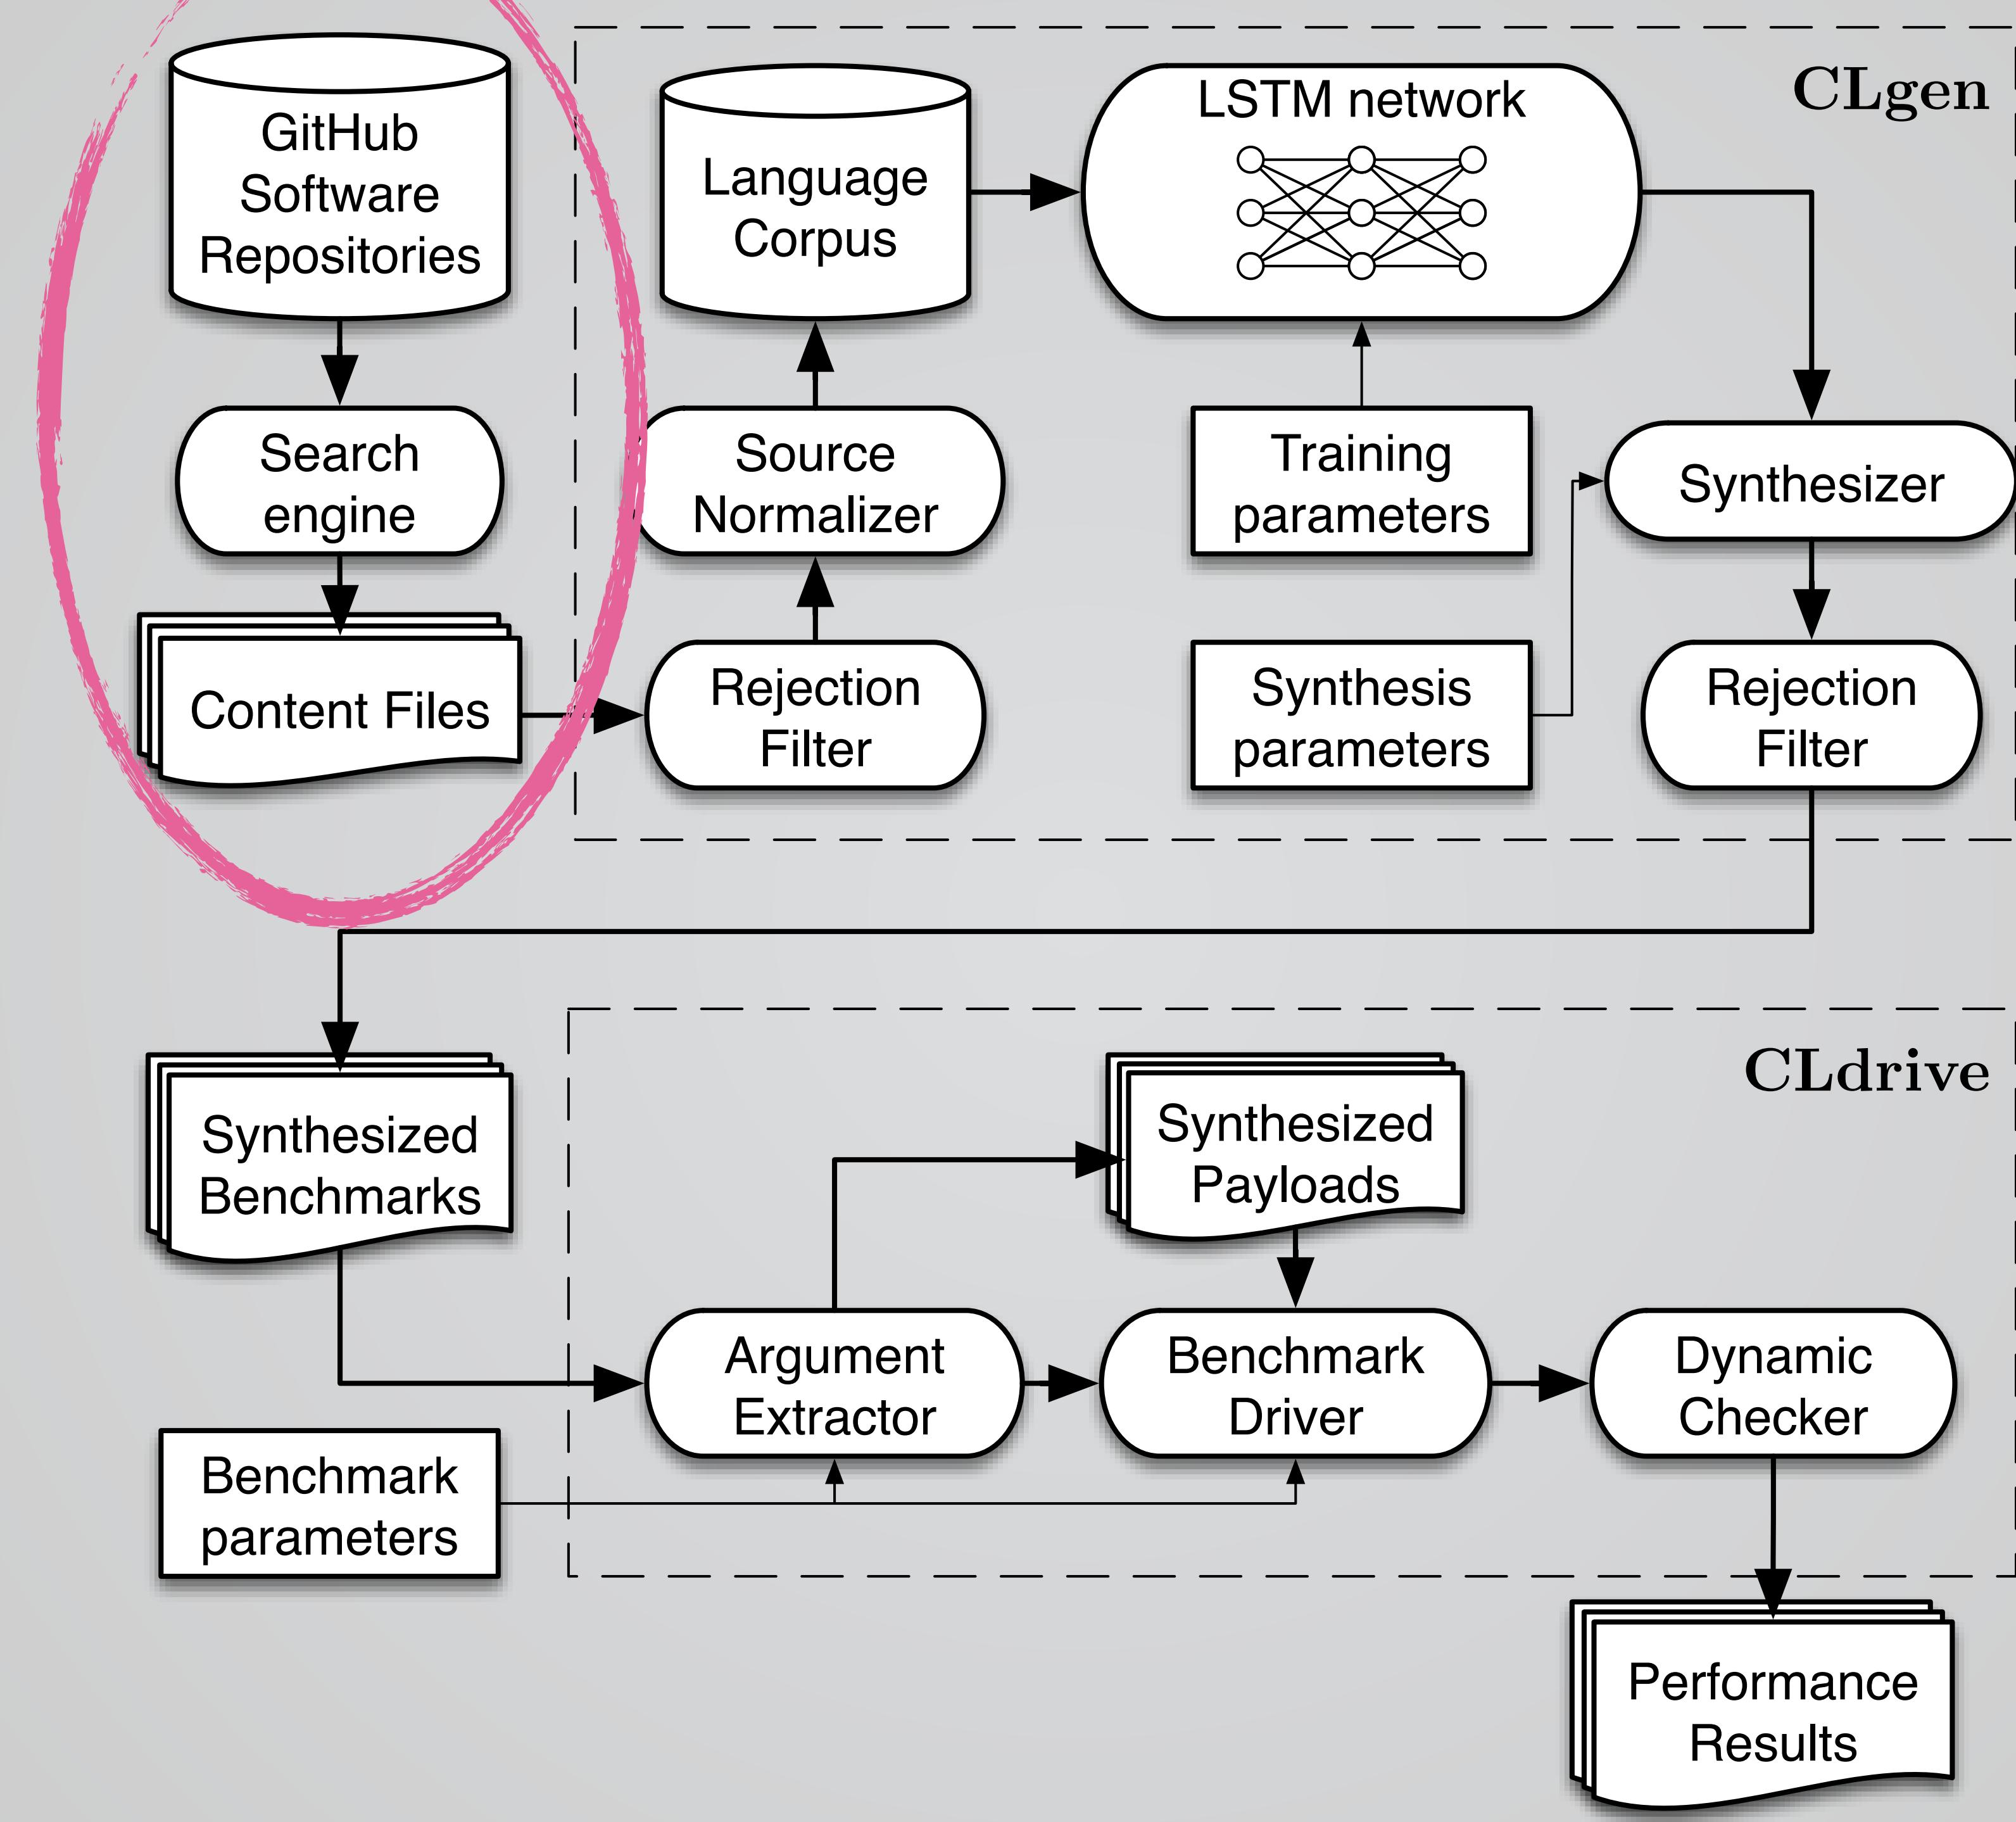
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

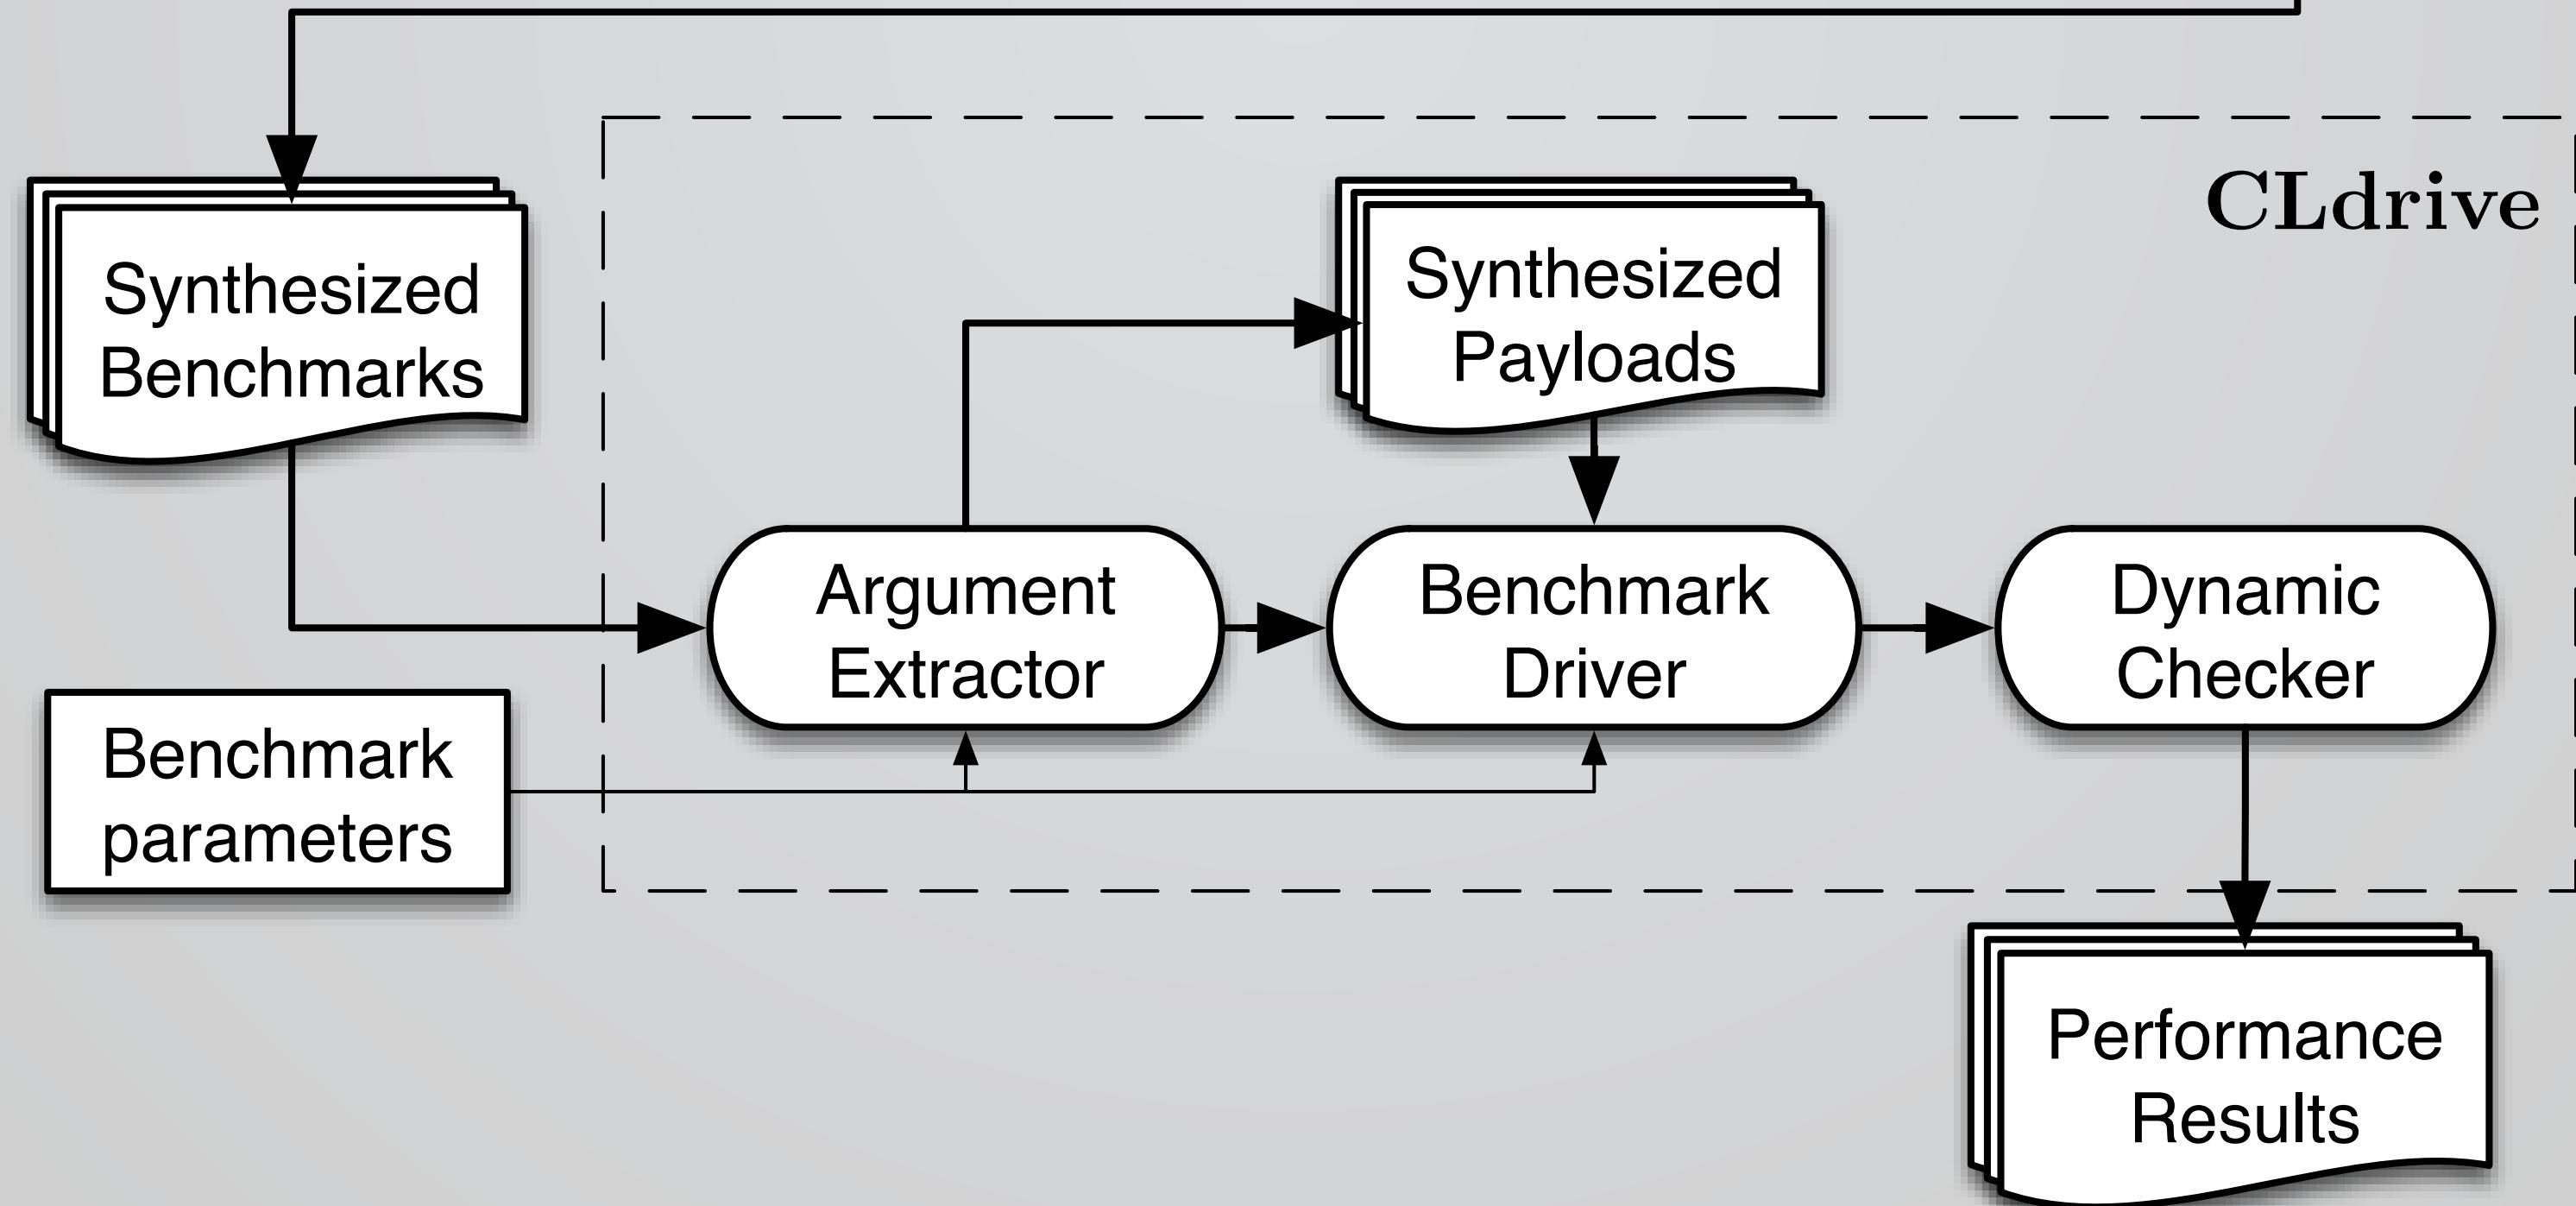
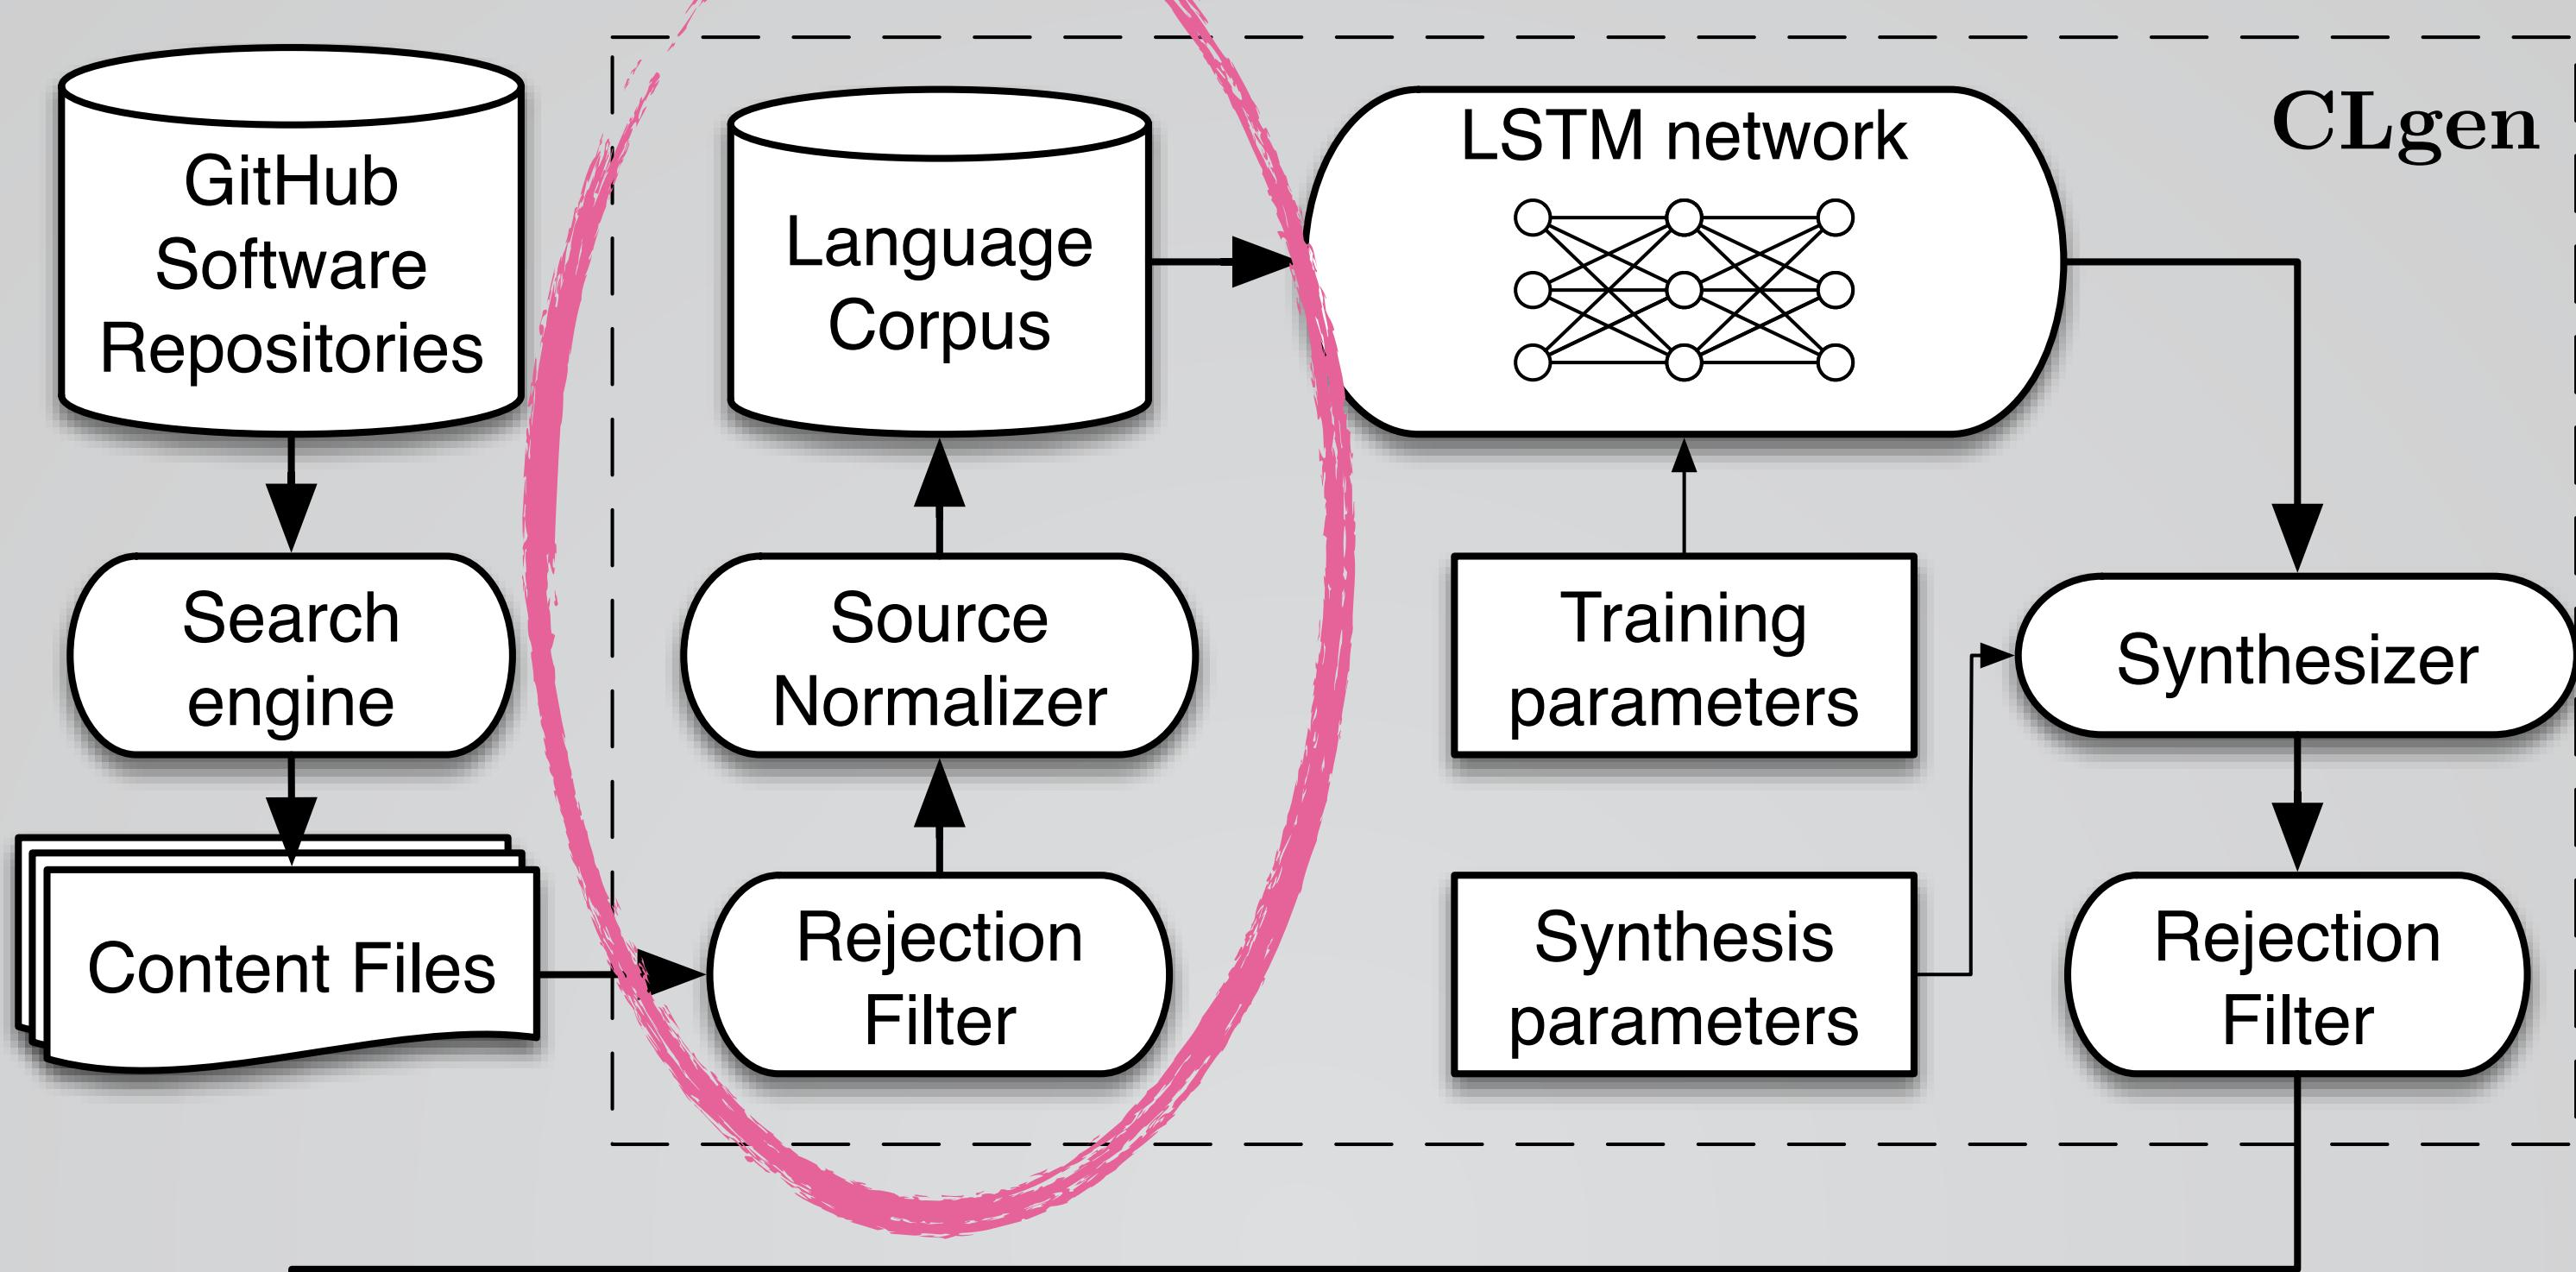
float myclamp(float in) {
#define CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#undef CLAMPING
}

// Do something really flipping cool
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    //
    // 
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

x 8078 files
2.8 million lines





```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */
//
//          DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION
//
// 0. You just DO WHAT THE FUCK YOU WANT TO.

#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#define CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#undef CLAMPING
#define else
    return in;
#undef endif // CLAMPING
}

// Do something really flipping cool
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    //
    //
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */
//
//          DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION
//
// 0. You just DO WHAT THE FUCK YOU WANT TO.

#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#define CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#undef CLAMPING
}

// Do something really flipping cool
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    //
    // 
    int id = get_global_id(0);
    if (id < num_elems)
    {
        out[id] = myclamp(in[id]);
    }
}
```

Is this real, valid OpenCL?
Can we minimise non-functional variance?

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */
//
//          DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE
//  TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION
//
// 0. You just DO WHAT THE FUCK YOU WANT TO.

#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#define CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#undef CLAMPING
}

// Do something really flipping cool
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    //
    // 
    int id = get_global_id(0);
    if (id < num_elems)
    {
        out[id] = myclamp(in[id]);
    }
}
```

Strip comments

Is this real, valid OpenCL?
Can we minimise non-functional variance?

```
/* Copyright (C) 2004 Joe Bloggs <joe@bloggs.io> */  
//  
// DO WHAT THE FUCK YOU WANT TO PUBLIC LICENSE  
// TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION  
//  
// o. You just DO WHAT THE FUCK YOU WANT TO.  
  
#define CLAMPING  
#define THRESHOLD_MIN 1.0f  
#define THRESHOLD_MAX 1.0f  
  
float myclamp(float in) {  
#ifdef CLAMPING  
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;  
#else  
    return in;  
#endif // CLAMPING  
}  
  
// Do something really flipping cool  
__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)  
{  
    //  
    int id = get_global_id(0);  
    if (id < num_elems)  
    {  
        out[id] = myclamp(in[id]);  
    }  
}
```

Strip comments

Is this real, valid OpenCL?
Can we minimise non-functional variance?

~~Strip comments~~

```
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#ifndef CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#else
    return in;
#endif
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~
Preprocess

```
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
#ifndef CLAMPING
    return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
#else
    return in;
#endif
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~
Preprocess

```
#define CLAMPING
#define THRESHOLD_MIN 1.0f
#define THRESHOLD_MAX 1.0f

float myclamp(float in) {
    #ifndef CLAMPING
        return in > THRESHOLD_MAX ? THRESHOLD_MAX : in < THRESHOLD_MIN ? THRESHOLD_MIN : in;
    #else
        return in;
    #endif
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments
Preprocess~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments
Preprocess~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

Does it compile?
Does it contain instructions?

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Does it compile?~~

~~Does it contain instructions?~~

~~Strip comments~~
~~Preprocess~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
Rewrite function names

~~Does it compile?~~
~~Does it contain instructions?~~

```
float myclamp(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void findAllNodesMergedAabb(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = myclamp(in[id]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
Rewrite function names

~~Does it compile?~~
~~Does it contain instructions?~~

```
float A(float in) {  
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;  
}
```

```
__kernel void B(__global float* in, __global float* out, int num_elems)  
{  
    int id = get_global_id(0);  
    if (id < num_elems)  
    {  
  
        out[id] = A(in[id]);  
    }  
}
```

~~Strip comments~~
~~Preprocess~~
~~Rewrite function names~~

~~Does it compile?~~
~~Does it contain instructions?~~

```
float A(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void B(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {

        out[id] = A(in[id]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
~~Rewrite function names~~
~~Rewrite variable names~~

~~Does it compile?~~
~~Does it contain instructions?~~

```
float A(float in) {
    return in > 1.0f ? 1.0f : in < 0.0f ? 0.0f : in;
}

__kernel void B(__global float* in, __global float* out, int num_elems)
{
    int id = get_global_id(0);
    if (id < num_elems)
    {
        out[id] = A(in[id]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
~~Rewrite function names~~
~~Rewrite variable names~~

~~Does it compile?~~
~~Does it contain instructions?~~

```
float A(float a) {
    return a > 1.0f ? 1.0f : a < 0.0f ? 0.0f : a;
}

__kernel void B(__global float* a, __global float* b, int c)
{
    int d = get_global_id(0);
    if (d < c)
    {

        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
~~Rewrite function names~~
~~Rewrite variable names~~

~~Does it compile?~~
~~Does it contain instructions?~~

```
float A(float a) {
    return a > 1.0f ? 1.0f : a < 0.0f ? 0.0f : a;
}

__kernel void B(__global float* a, __global float* b, int c)
{
    int d = get_global_id(0);
    if (d < c)
    {

        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
~~Rewrite function names~~
~~Rewrite variable names~~
Enforce code style

~~Does it compile?~~

~~Does it contain instructions?~~

```
float A(float a) {
    return a > 1.0f ? 1.0f : a < 0.0f ? 0.0f : a;
}

__kernel void B(__global float* a, __global float* b, int c)
{
    int d = get_global_id(0);
    if (d < c)
    {
        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
~~Rewrite function names~~
~~Rewrite variable names~~
Enforce code style

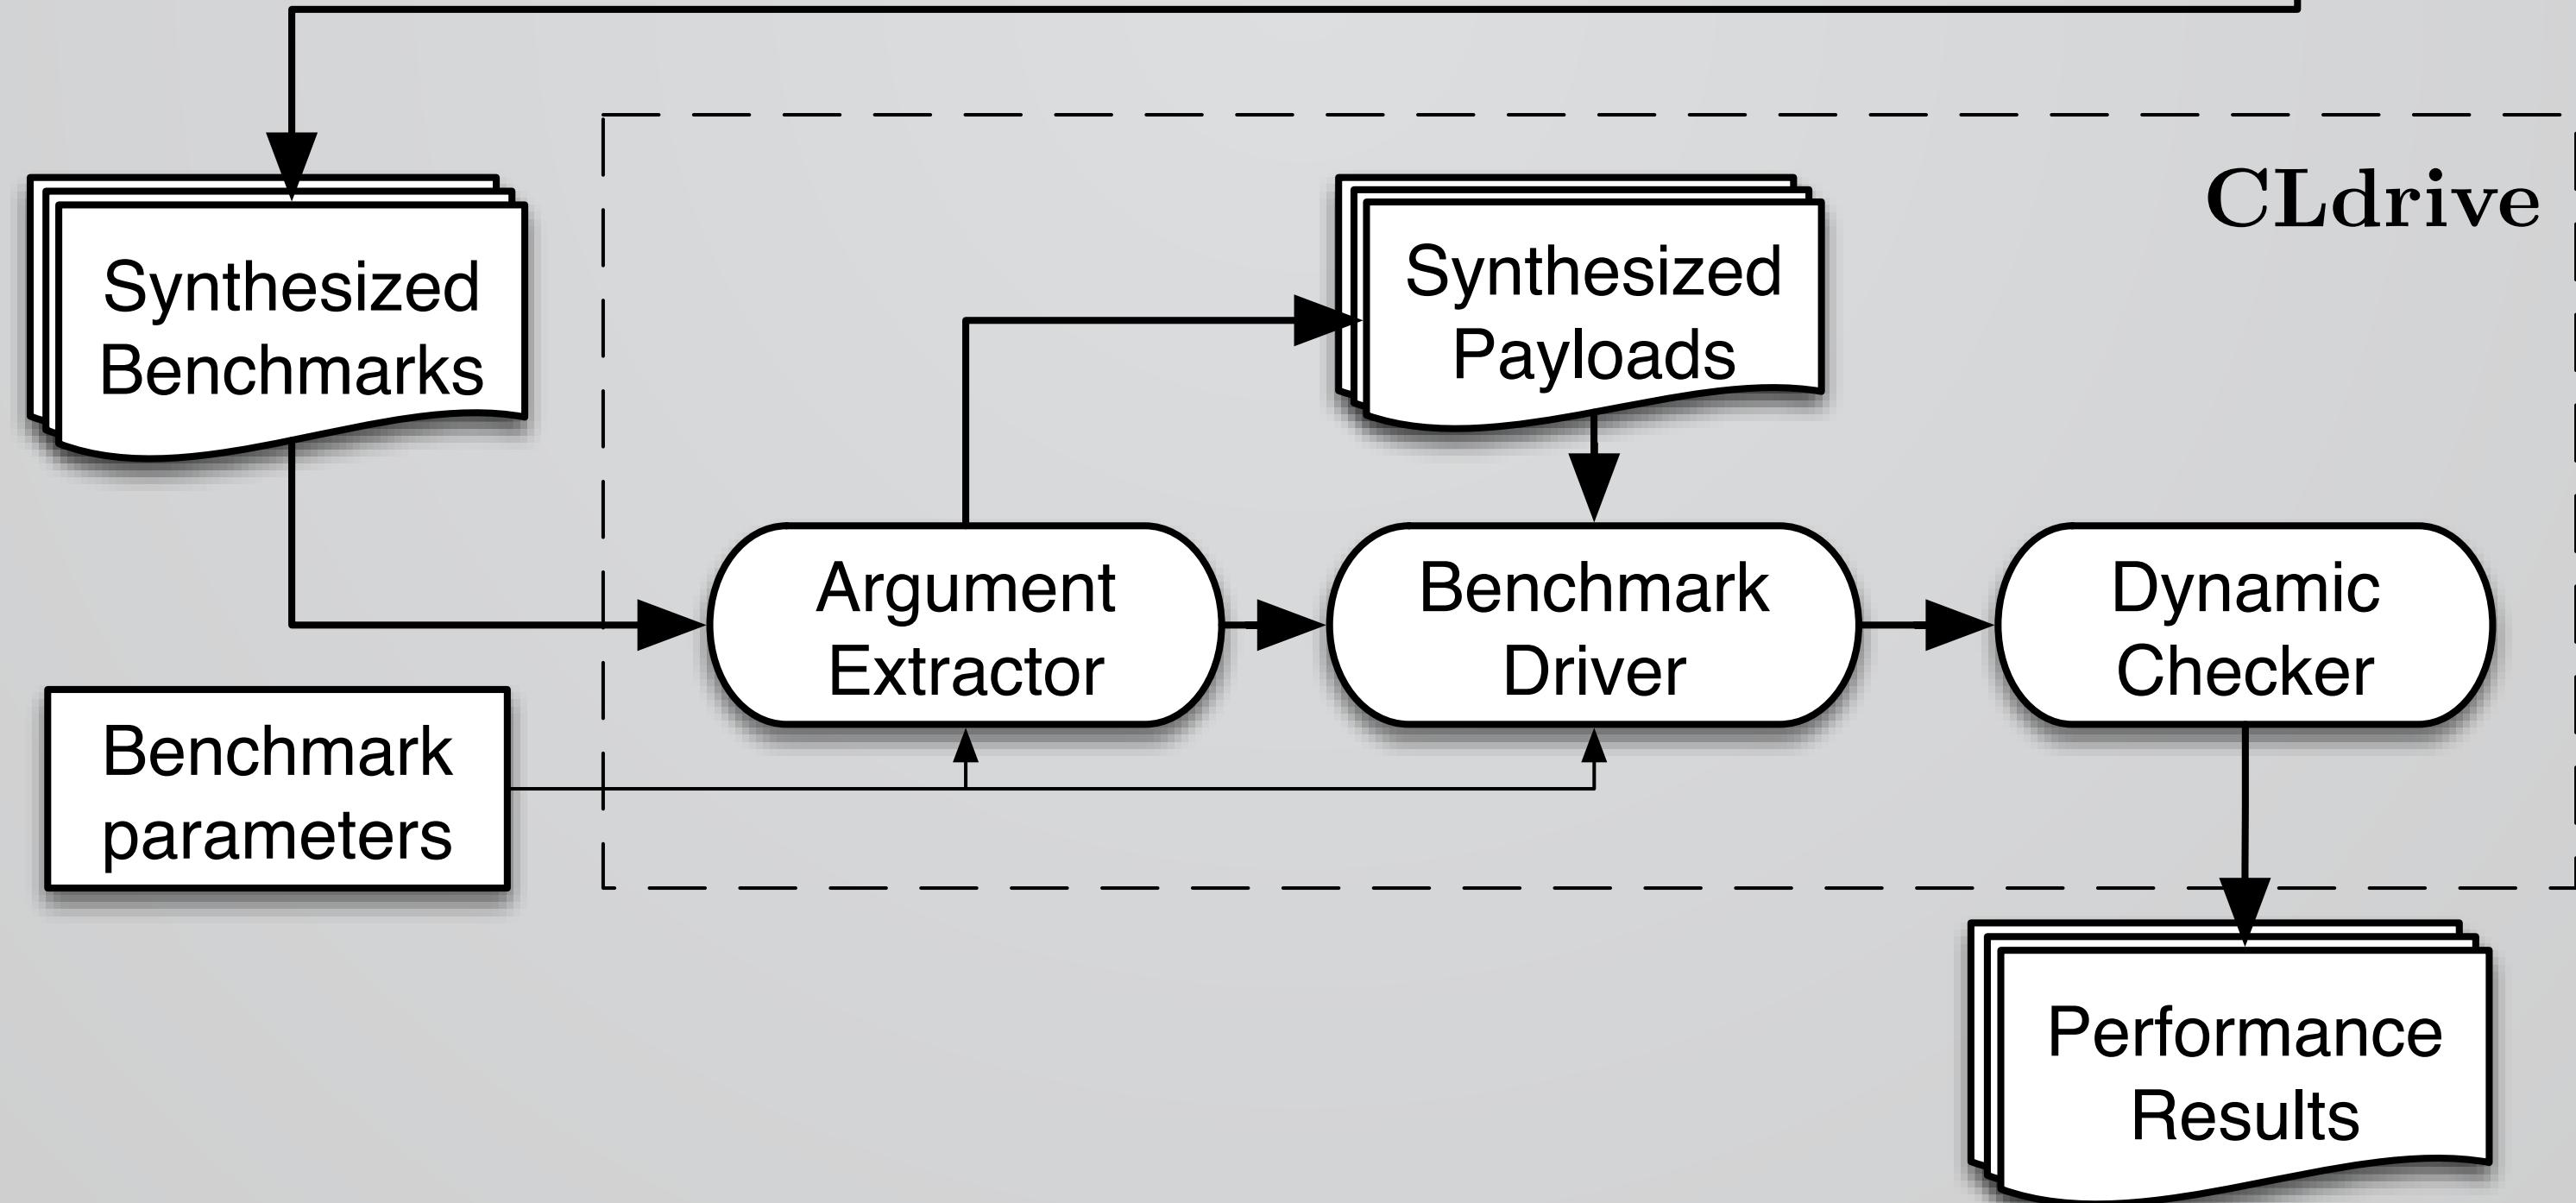
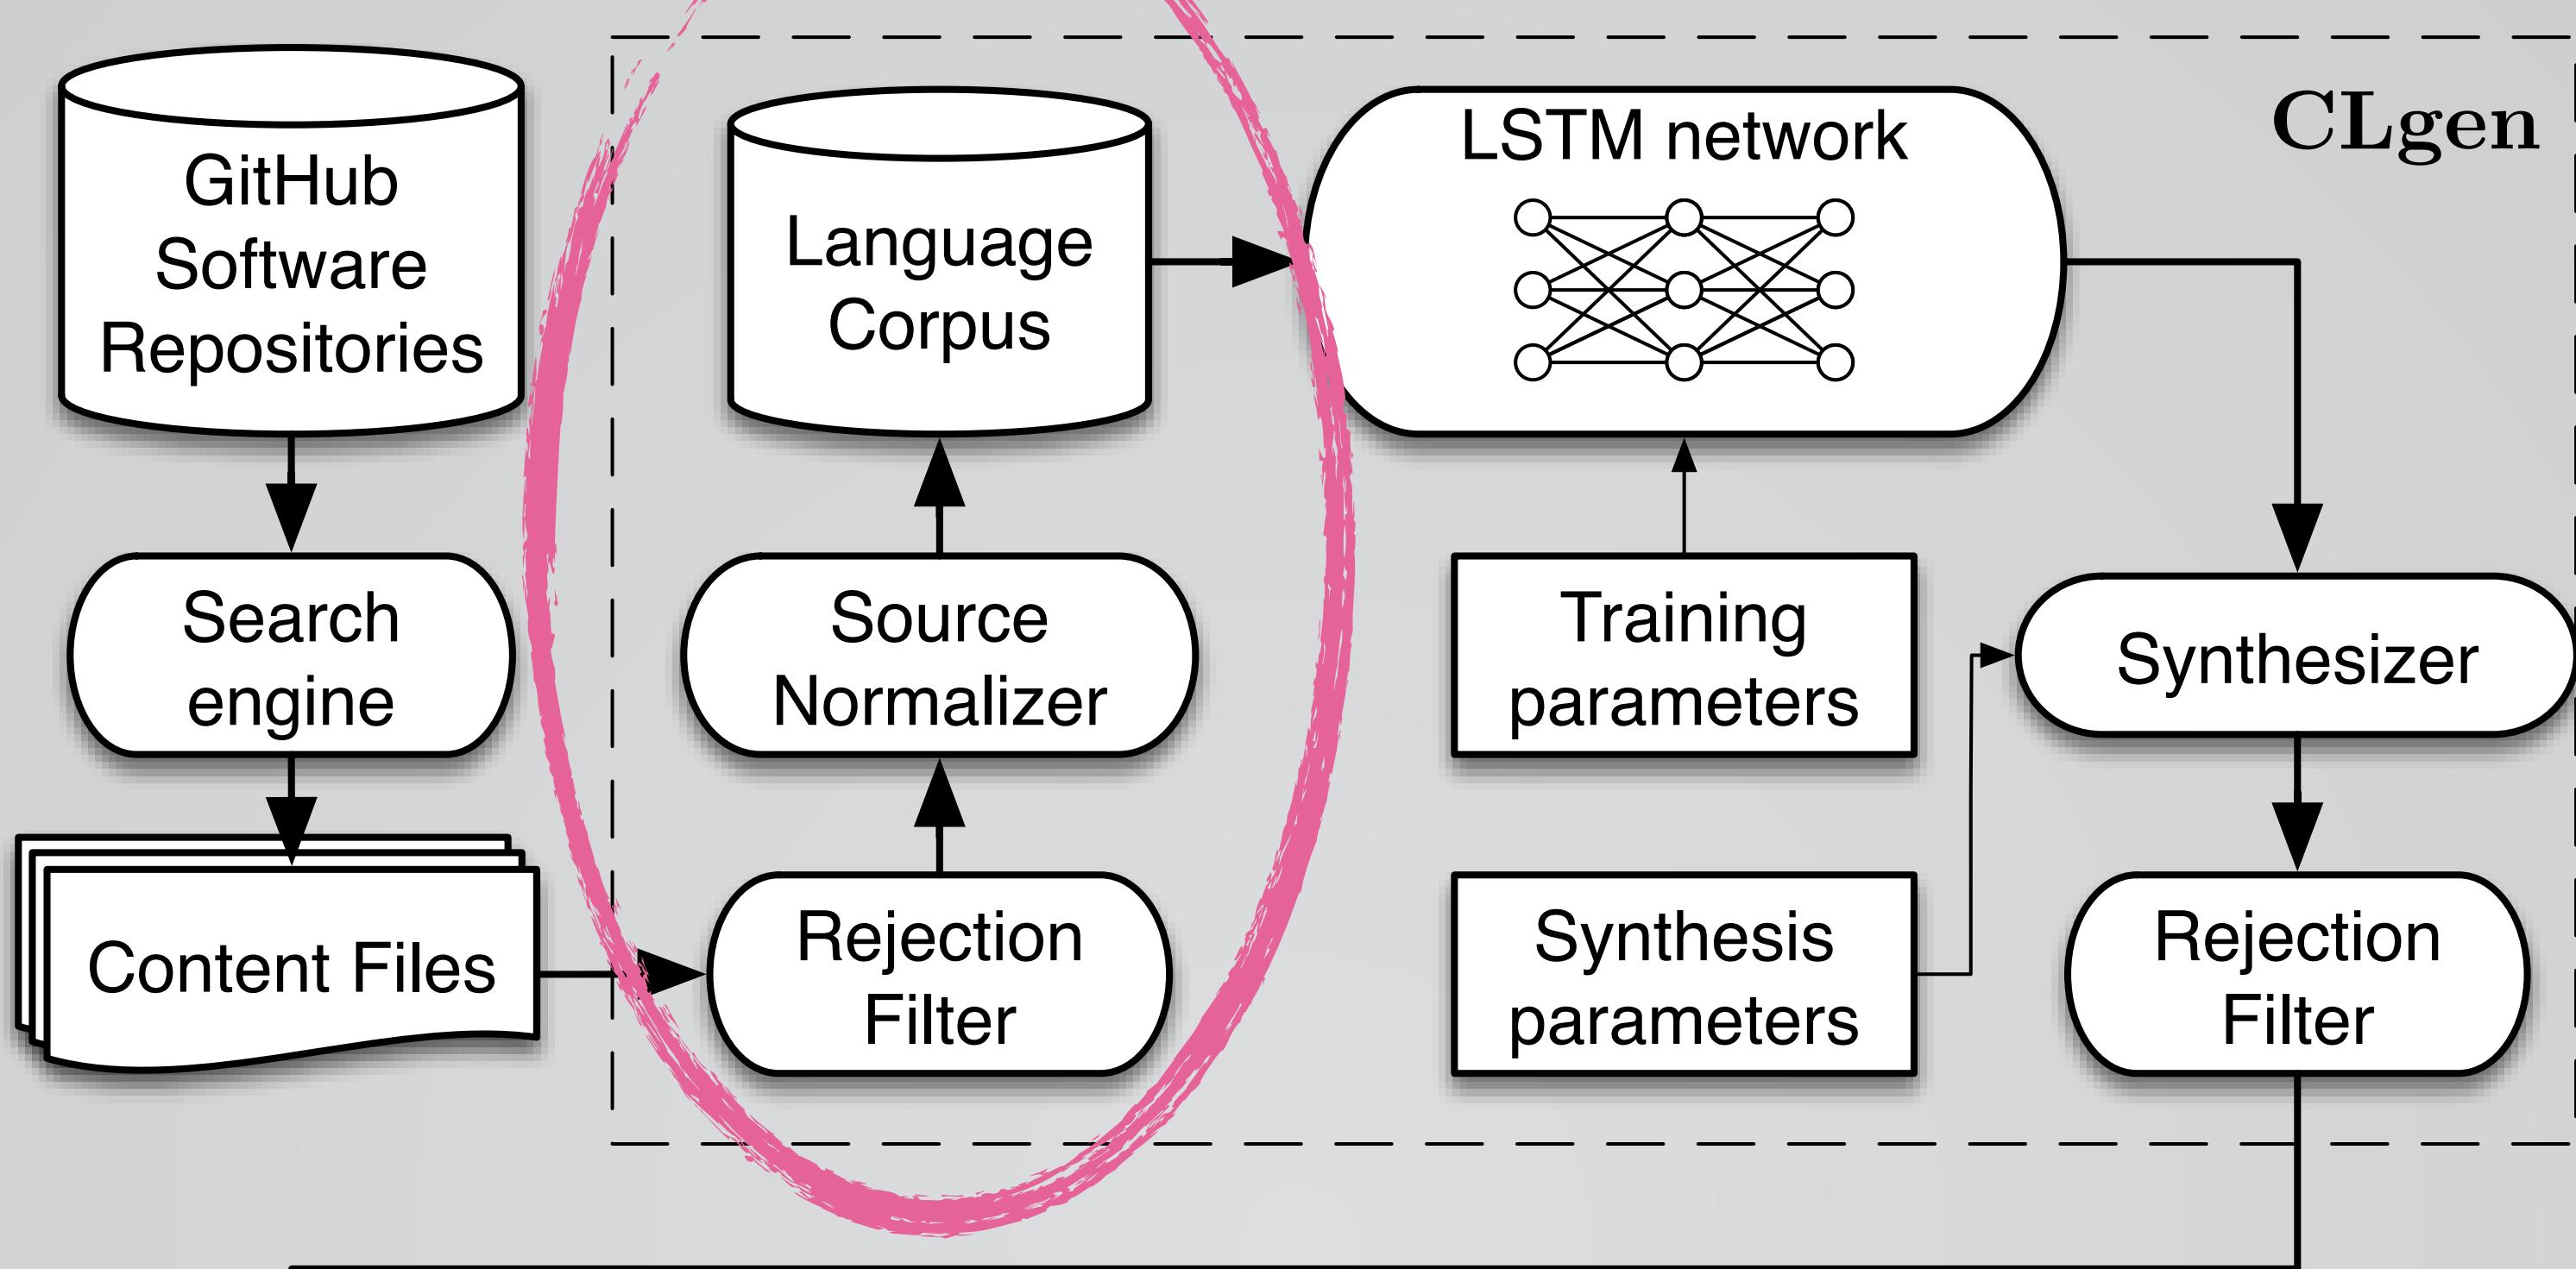
~~Does it compile?~~
~~Does it contain instructions?~~

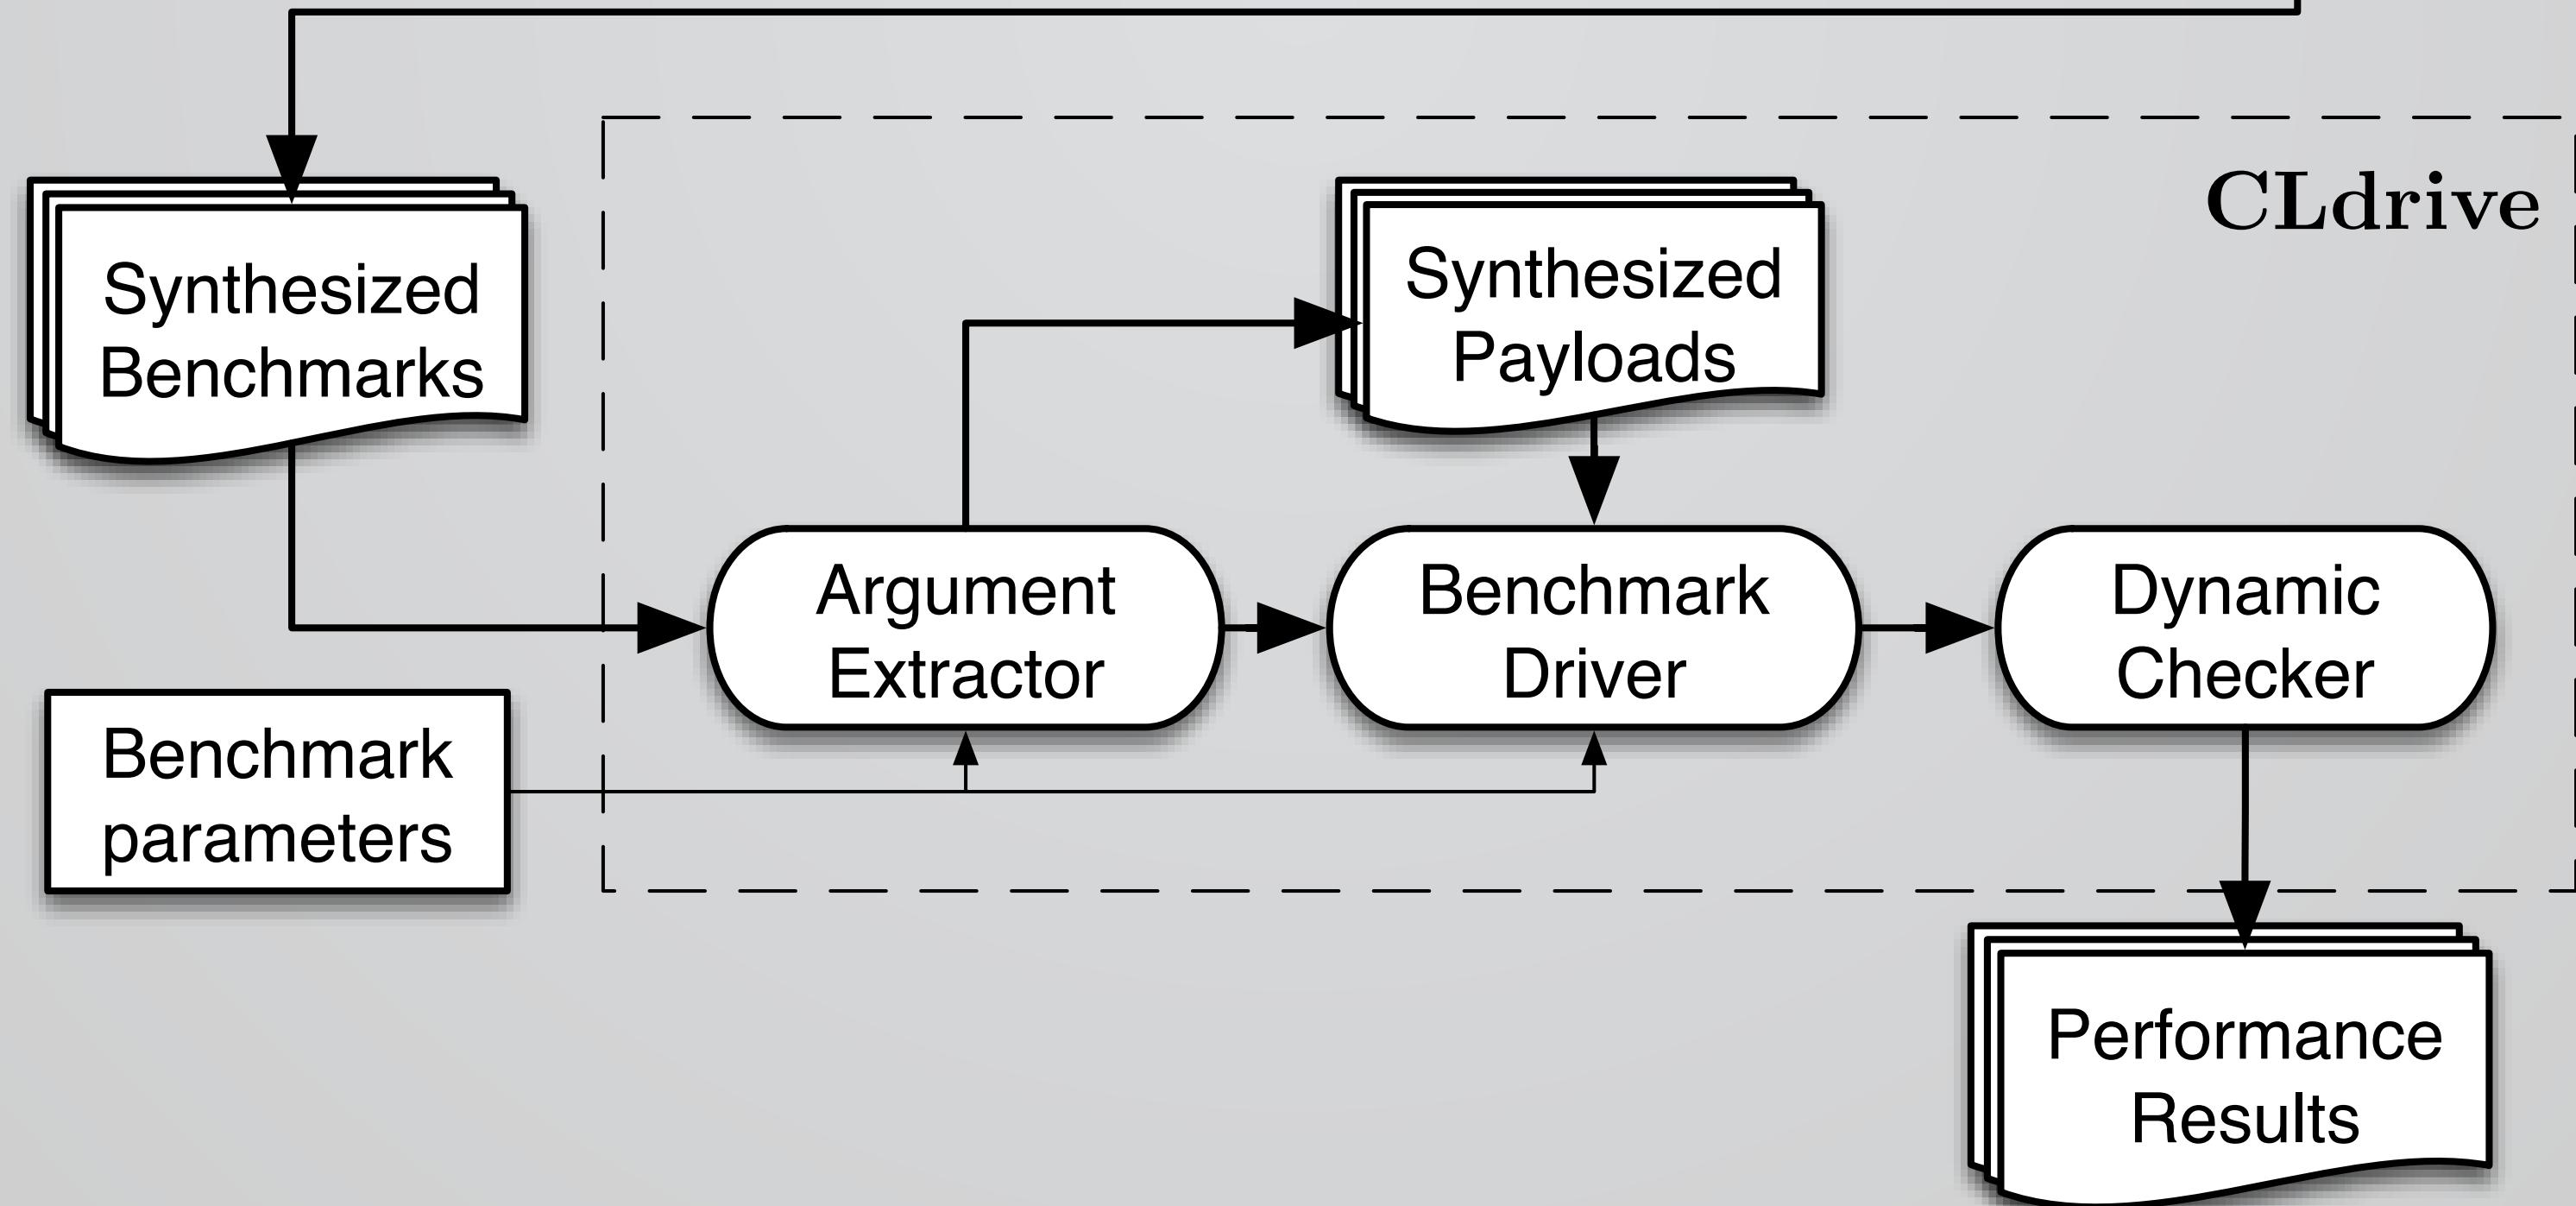
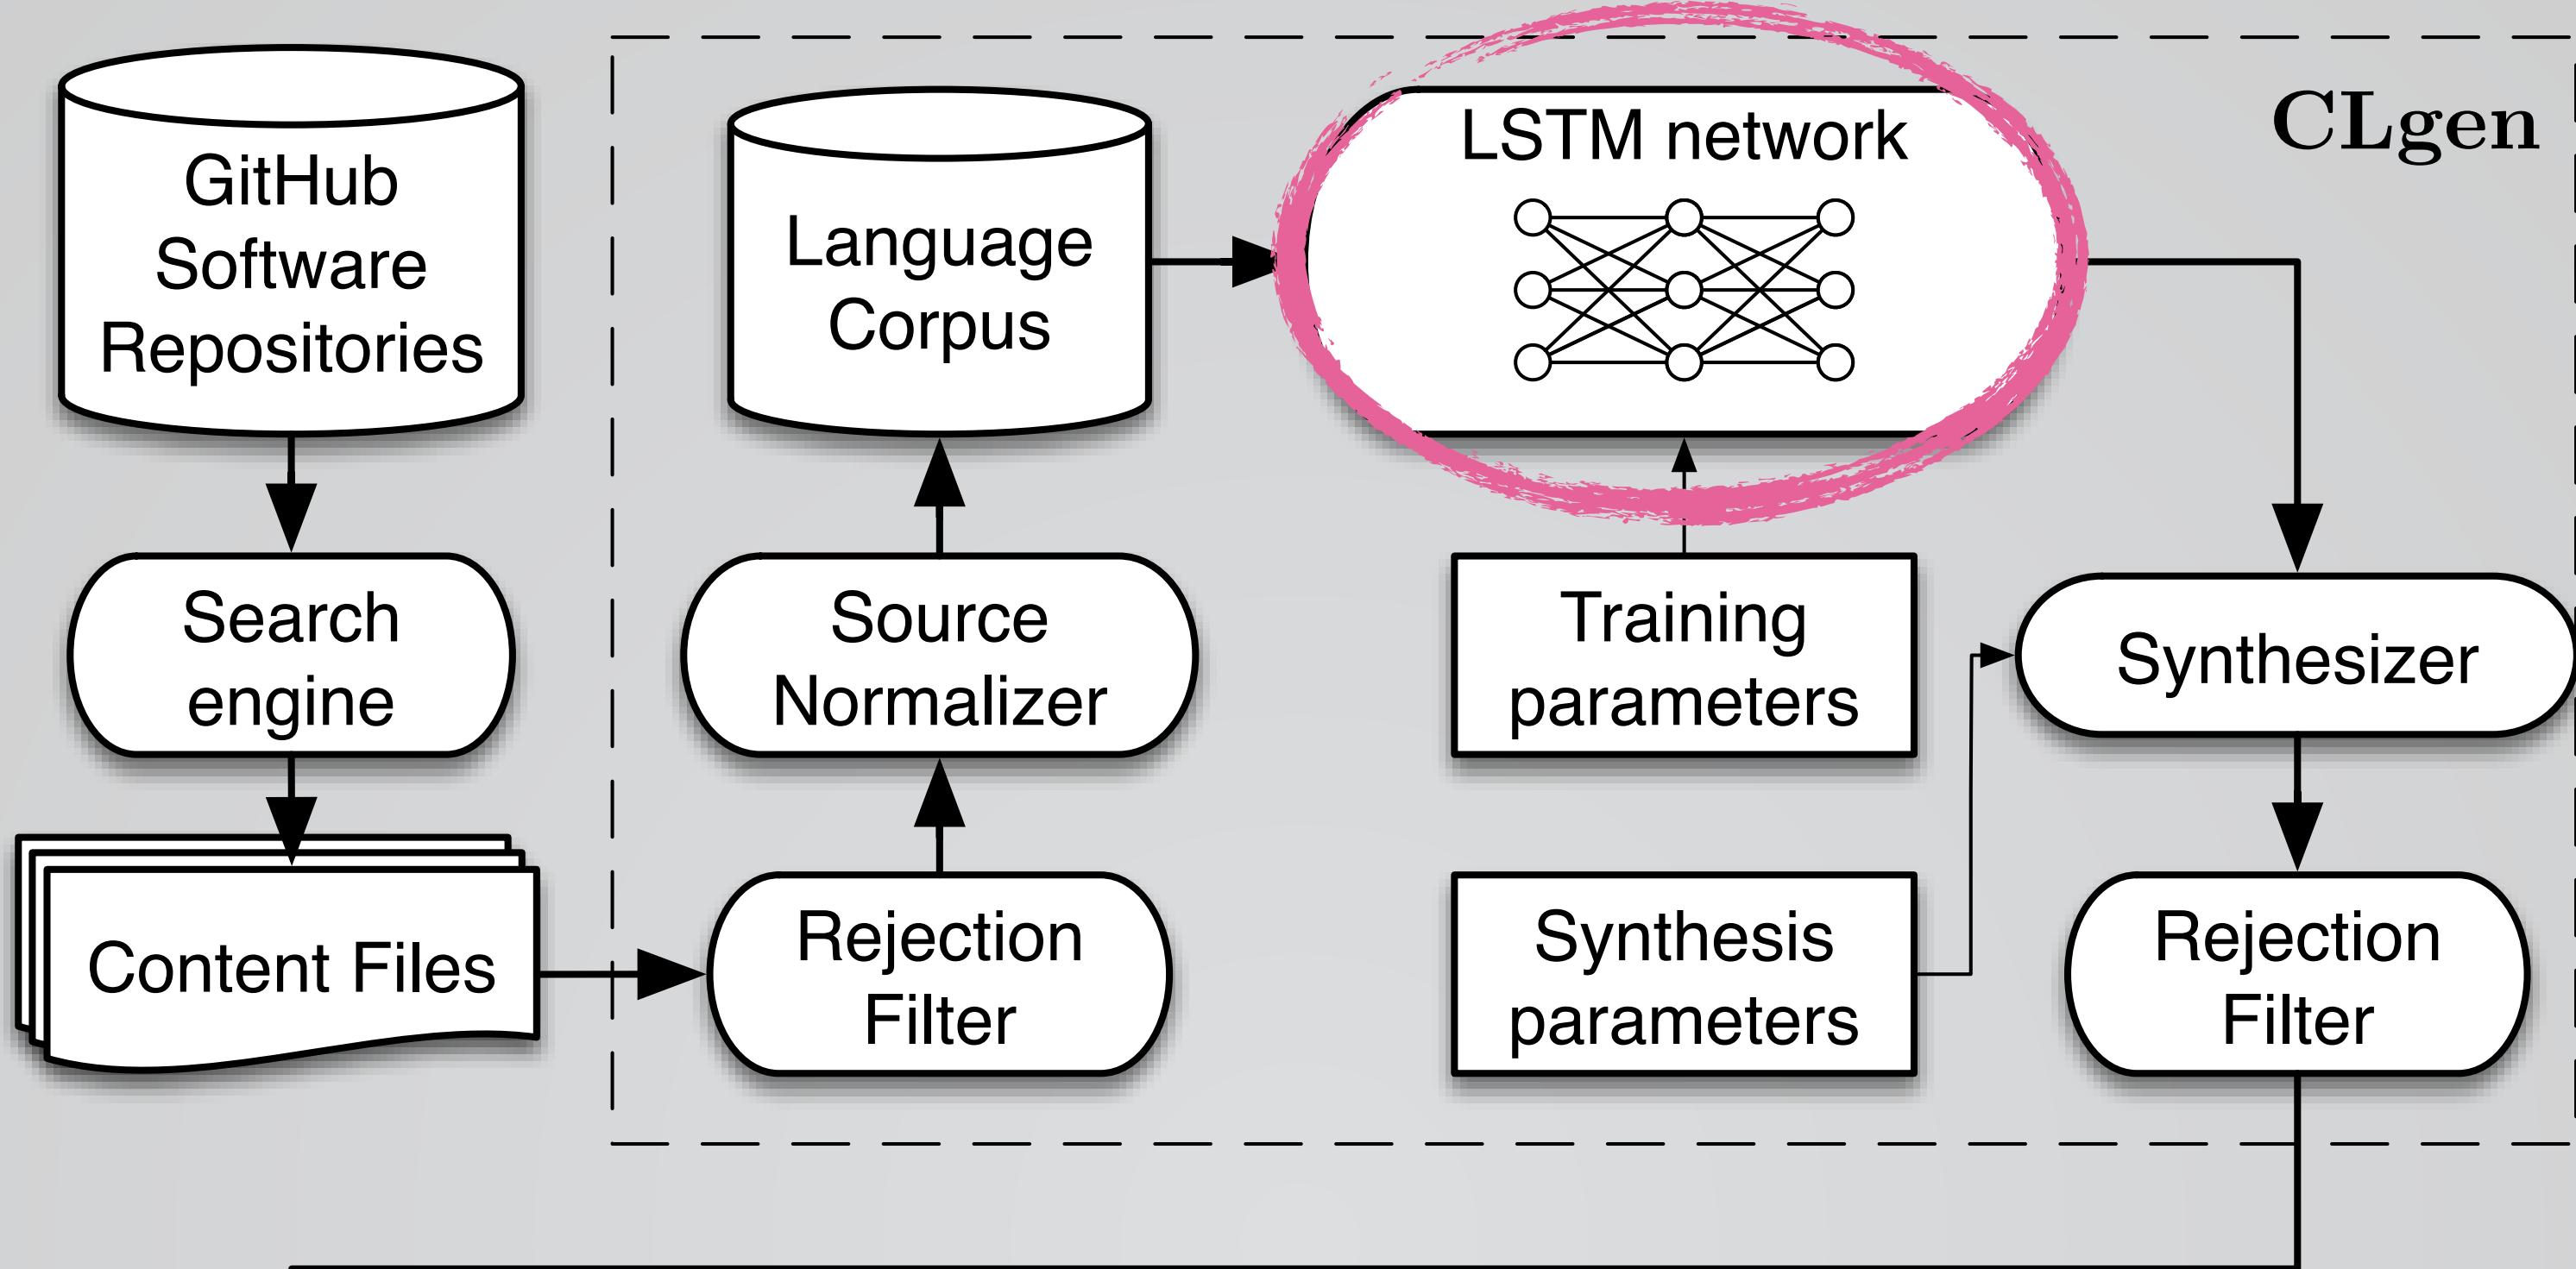
```
float A(float a) {
    return a > 1.0f ? 1.0f : in < 0.0f ? 0.0f : a;
}

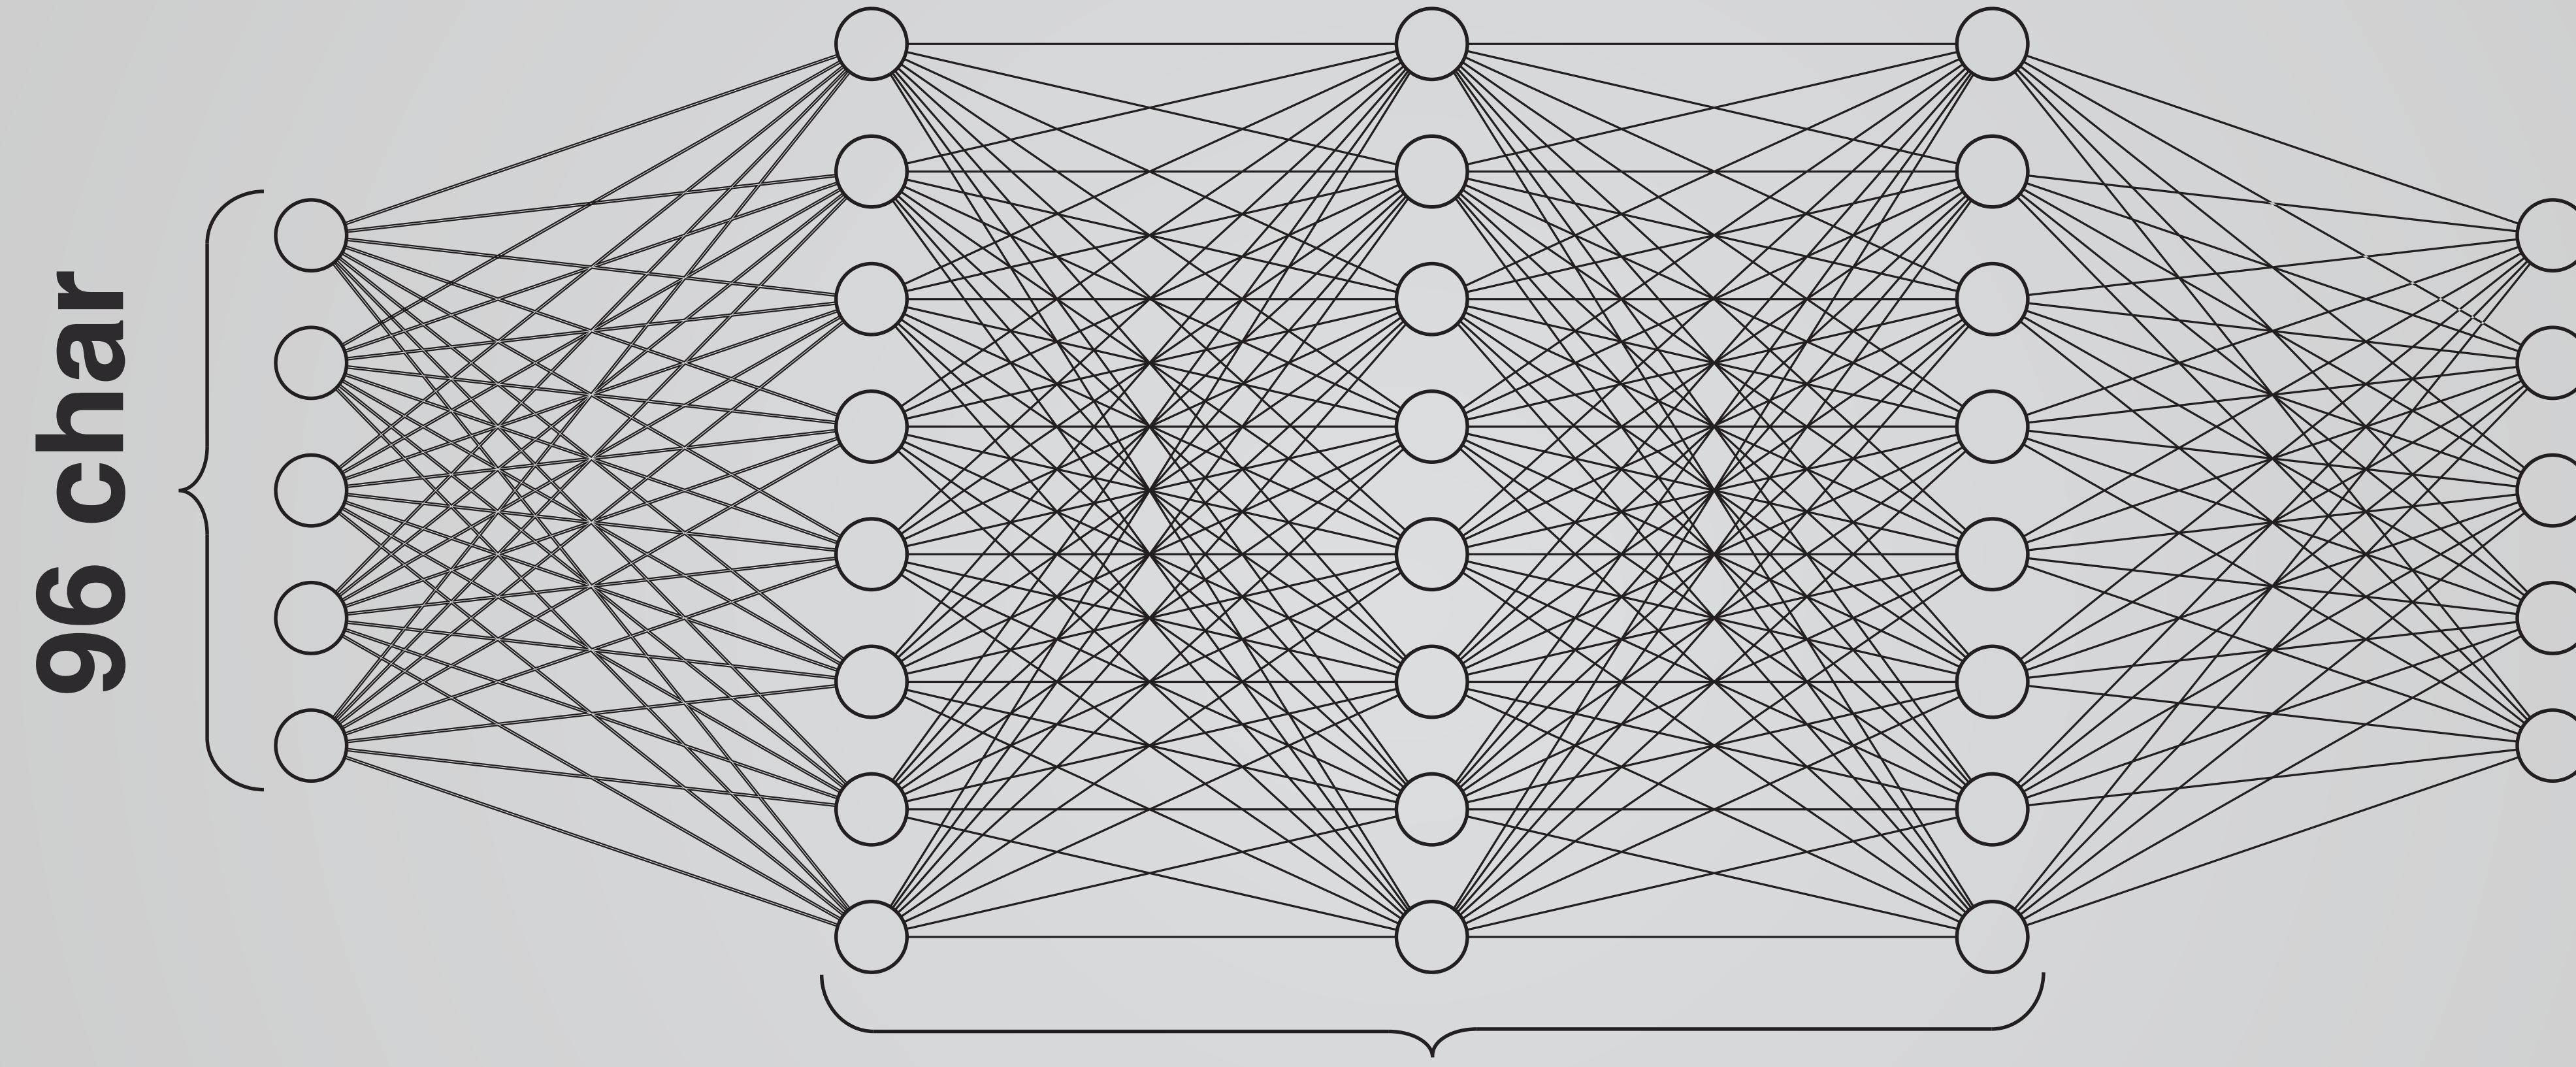
__kernel void B(__global float* a, __global float* b, int c) {
    int d = get_global_id(0);
    if (d < c) {
        b[d] = A(a[d]);
    }
}
```

~~Strip comments~~
~~Preprocess~~
~~Rewrite function names~~
~~Rewrite variable names~~
~~Enforce code style~~

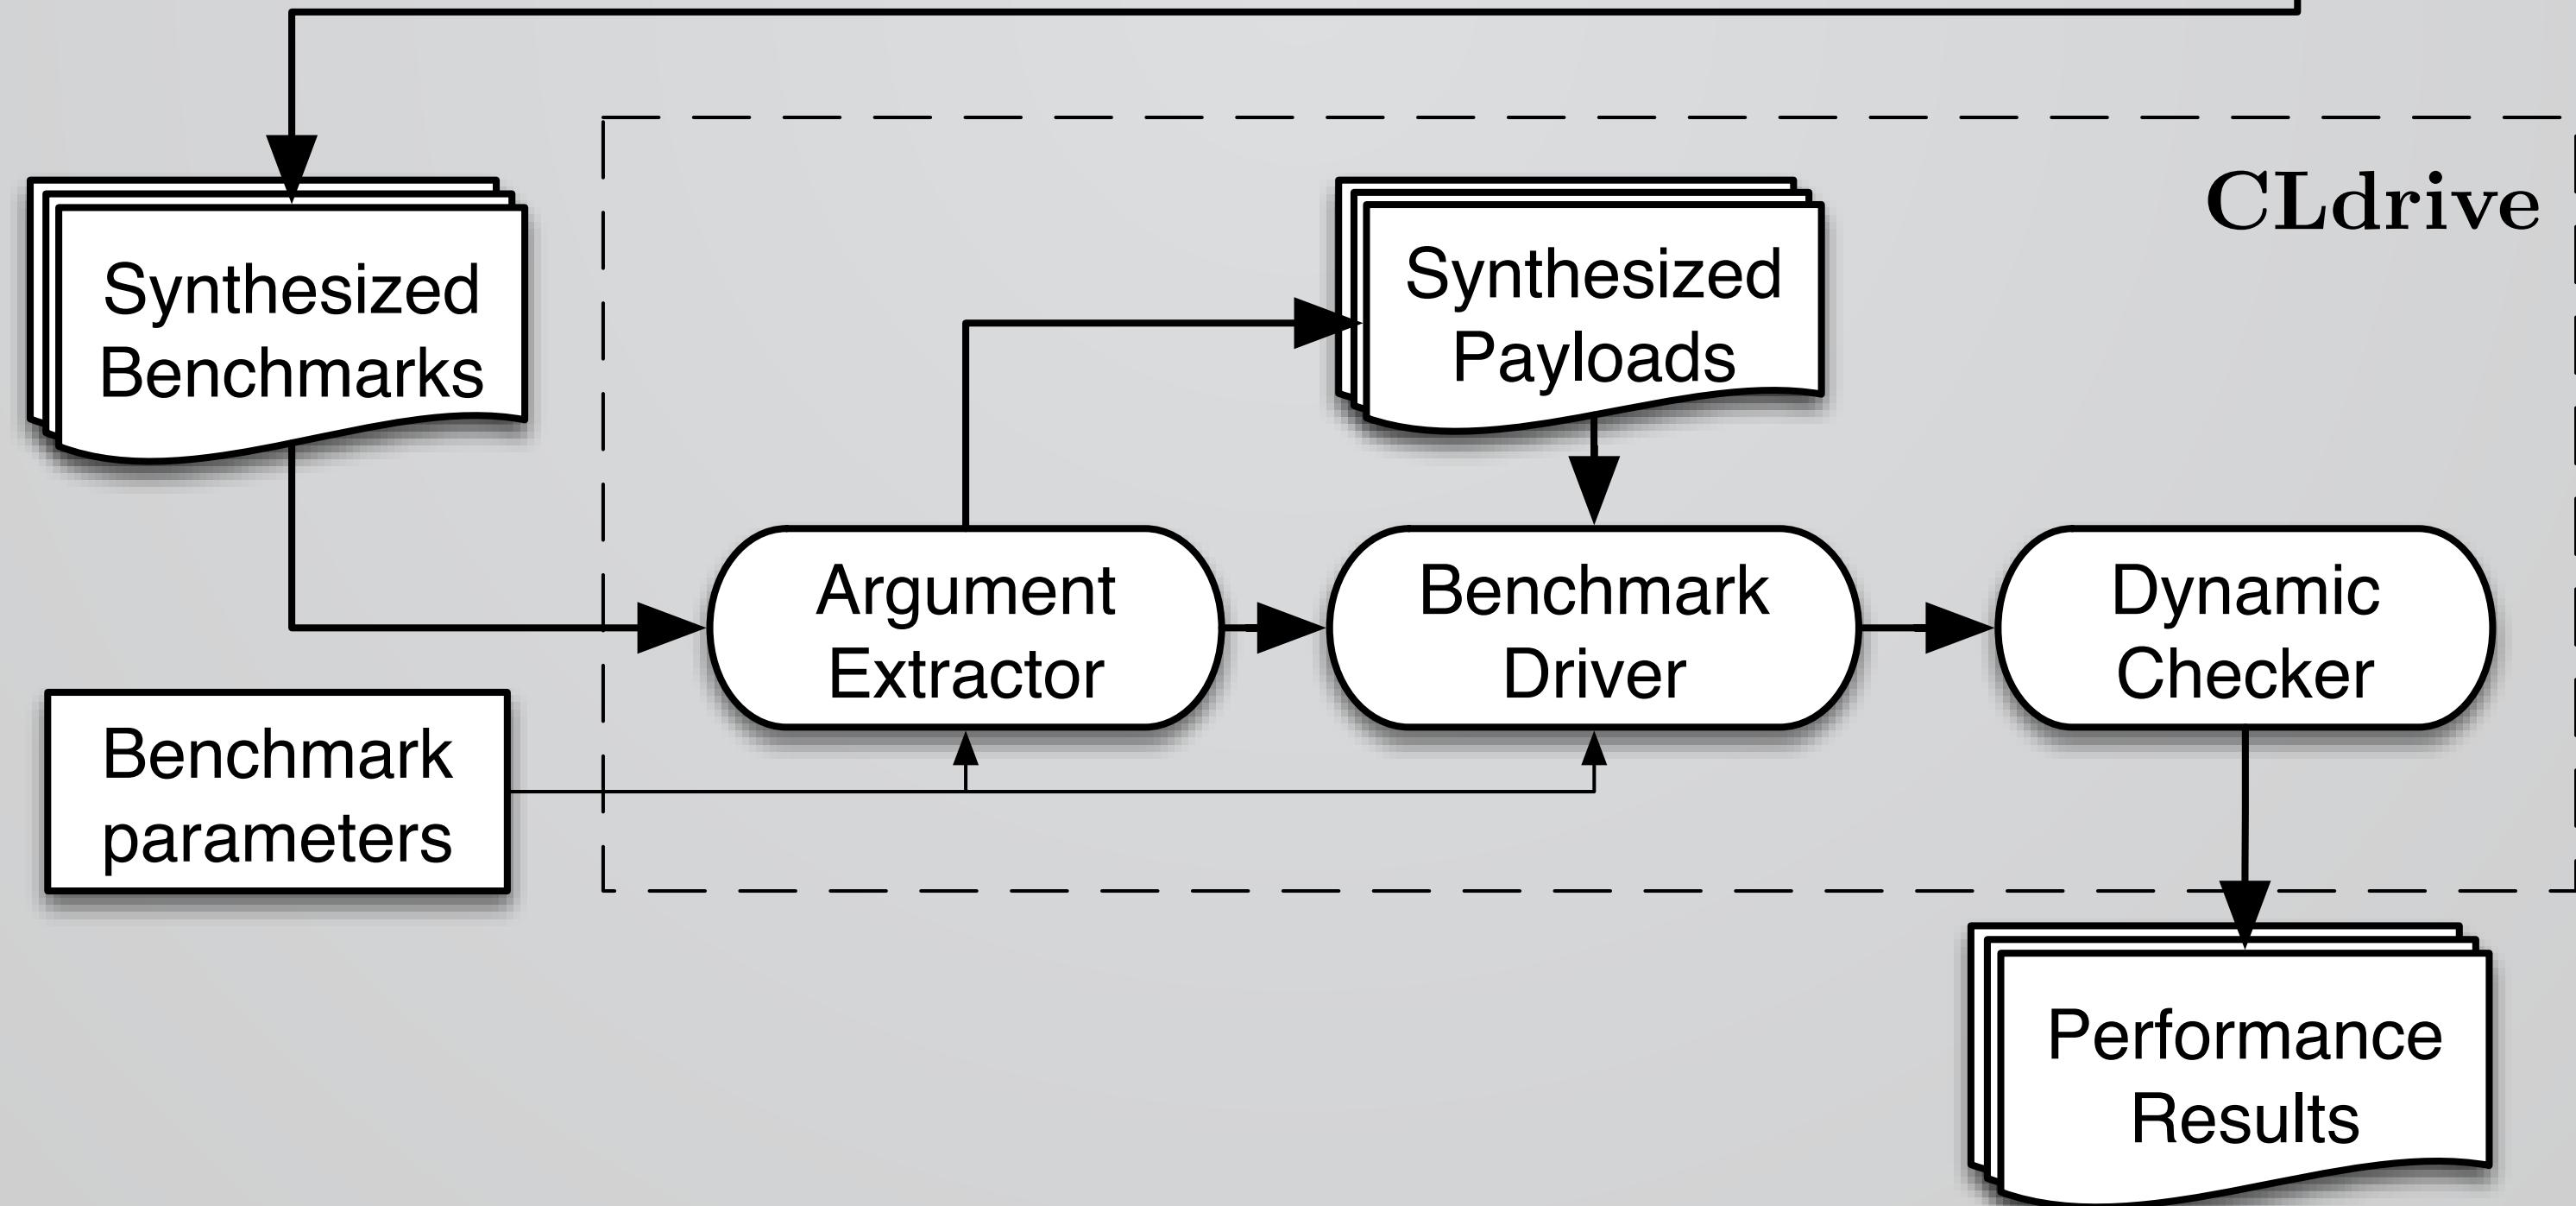
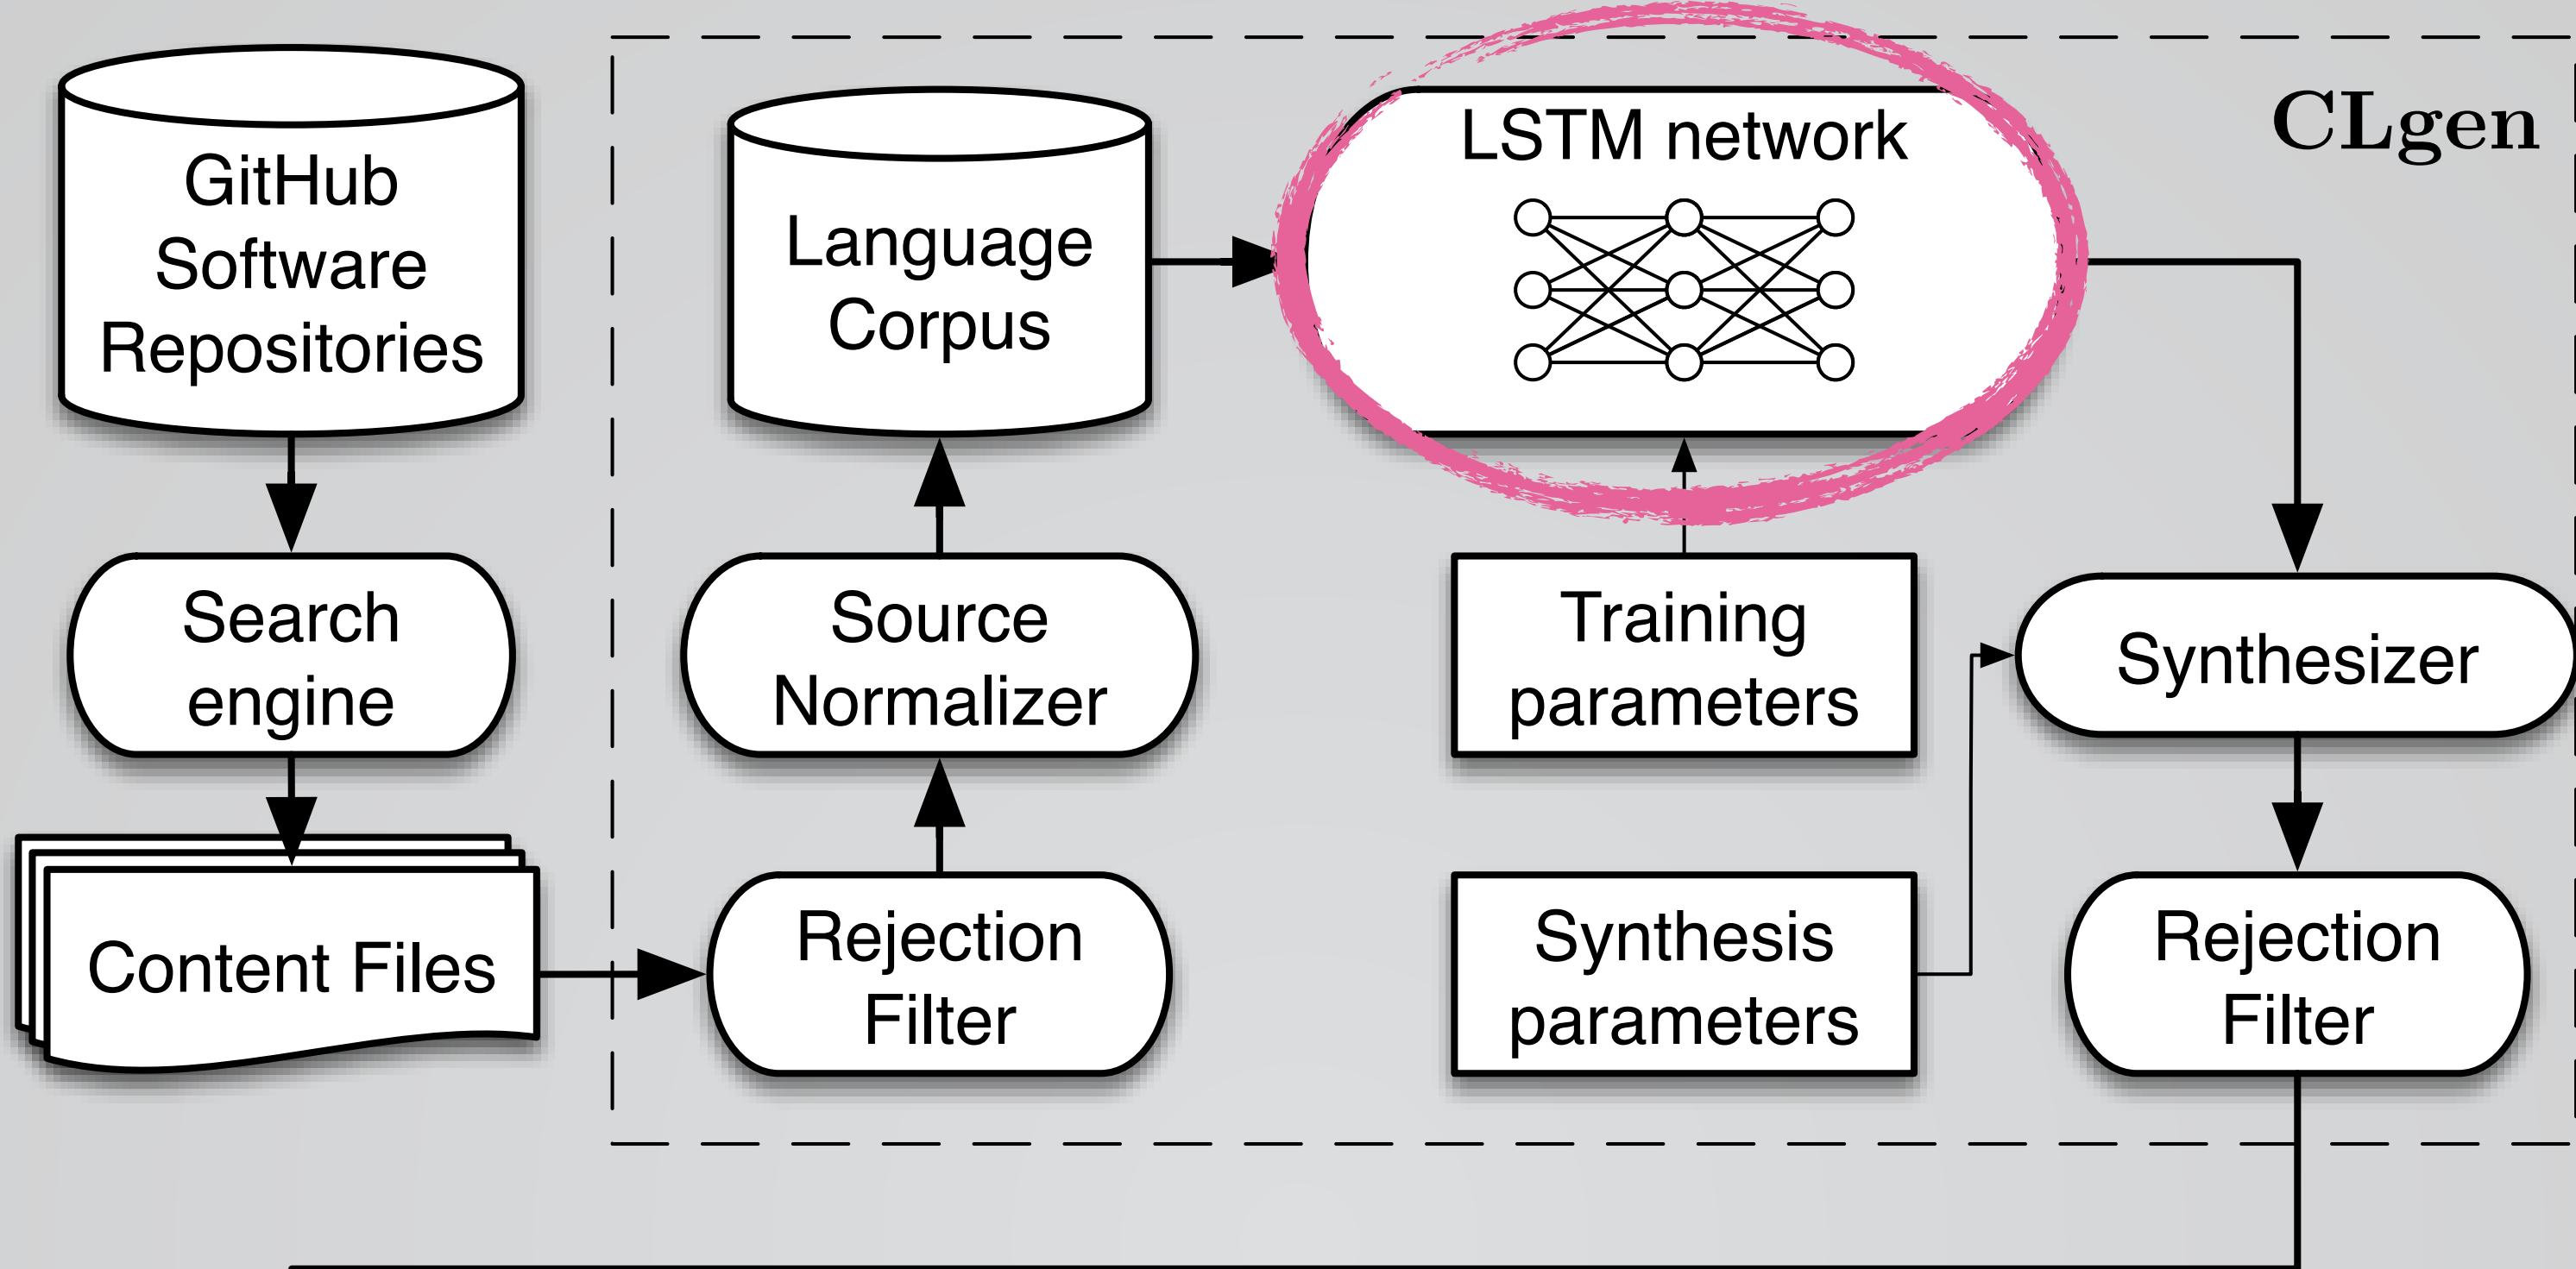
~~Does it compile?~~
~~Does it contain instructions?~~

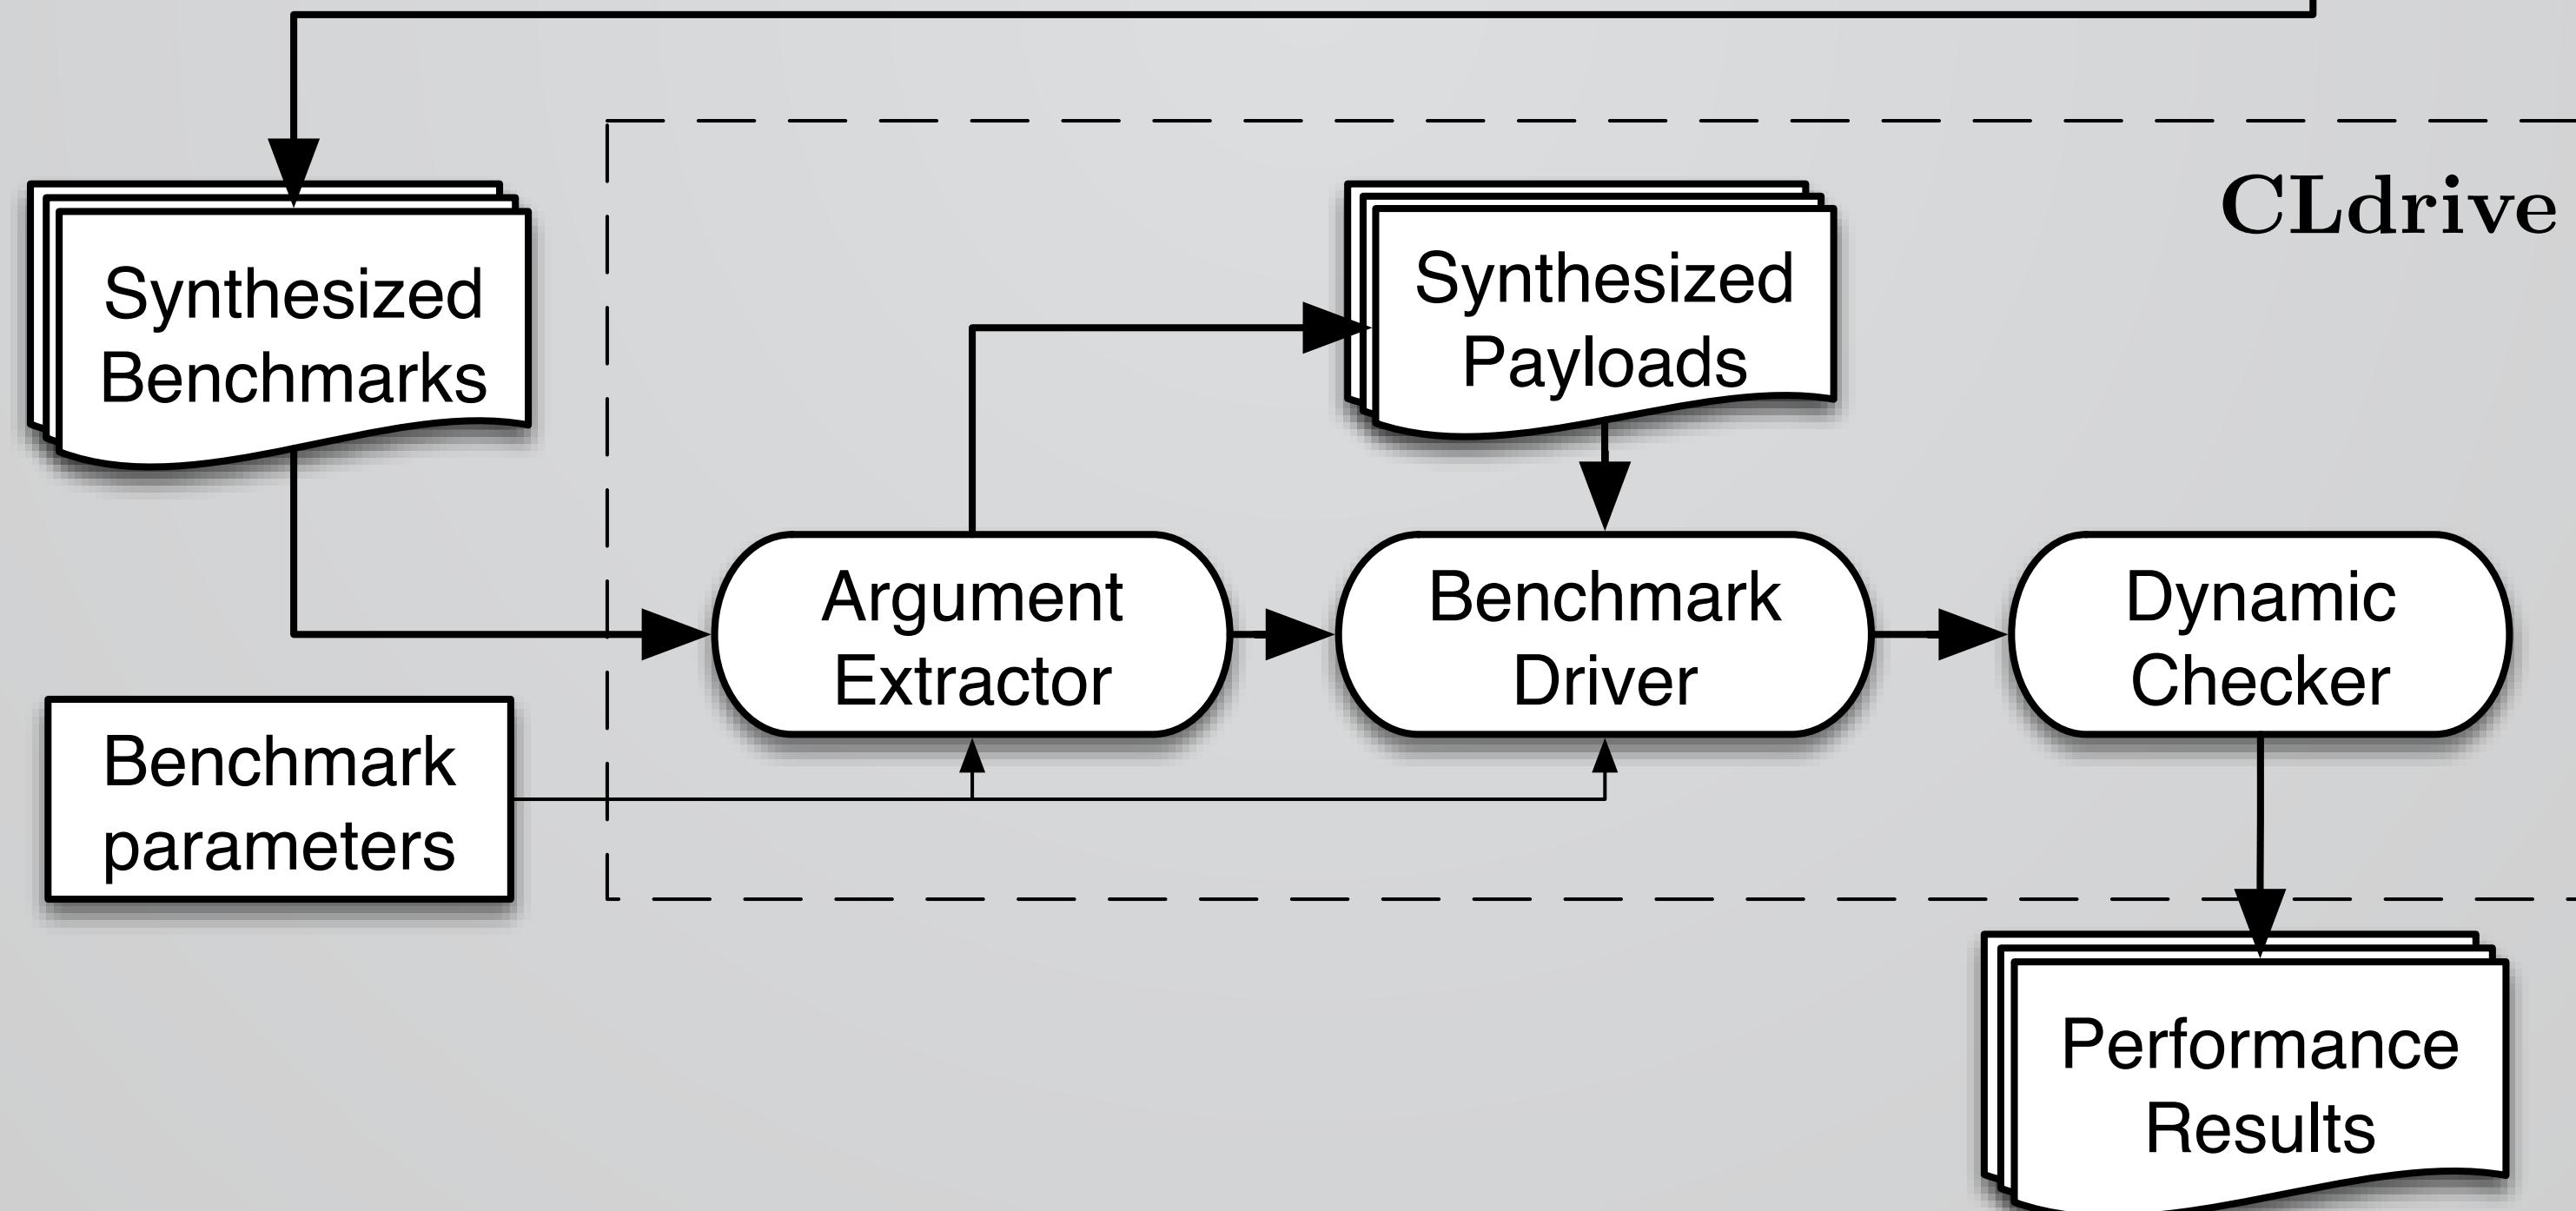
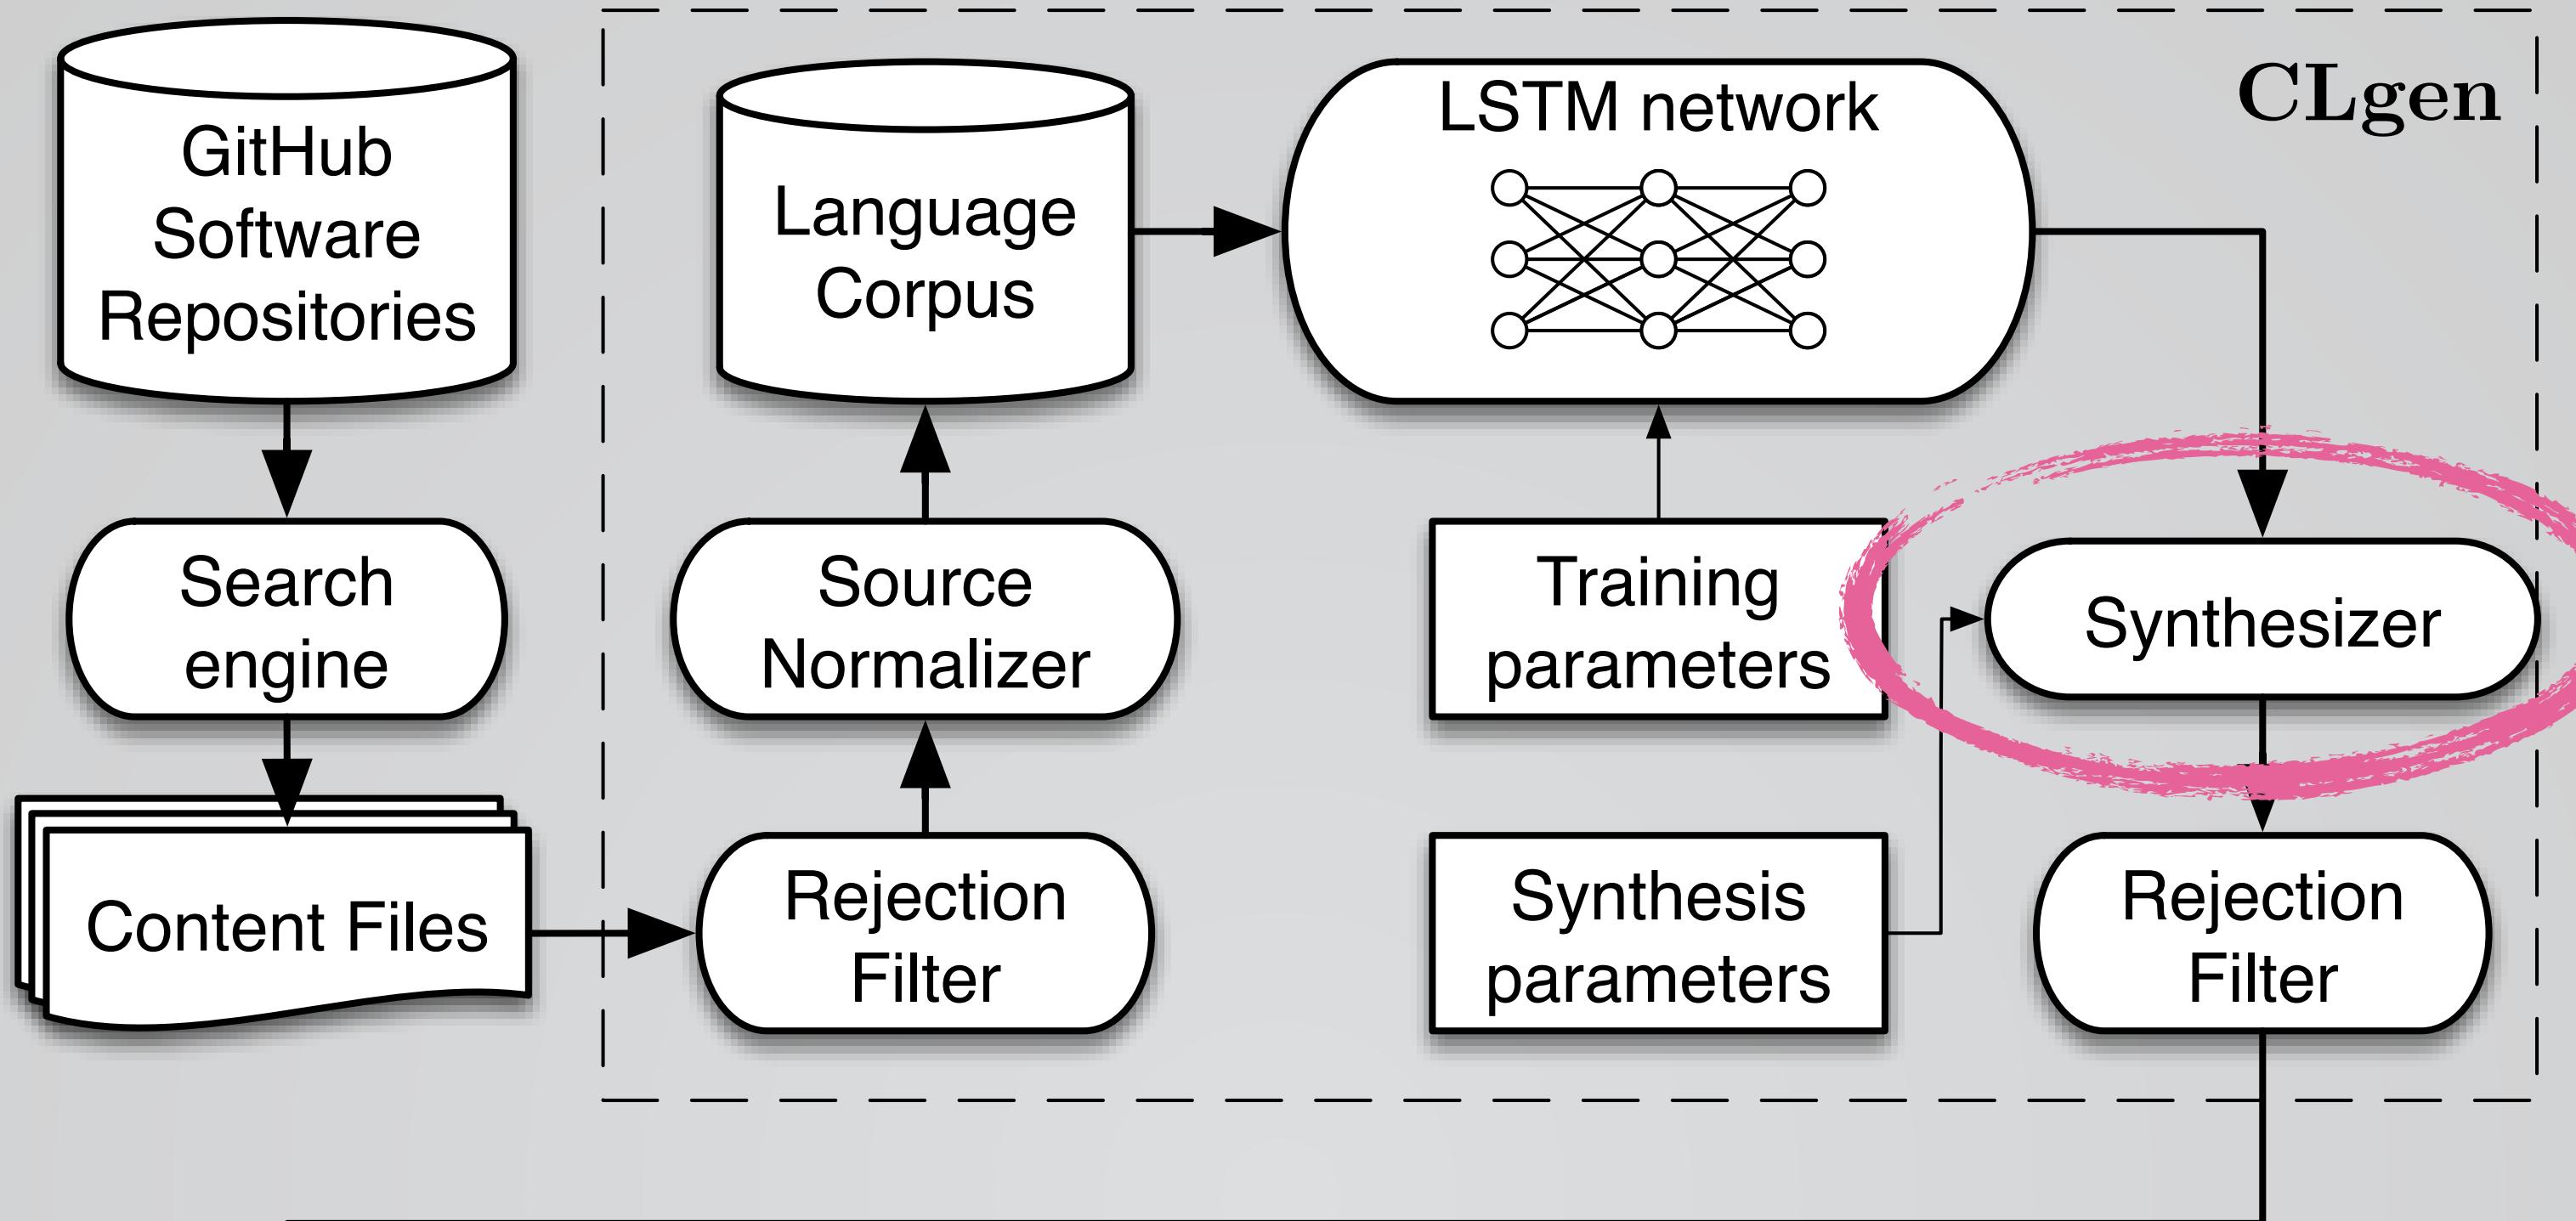






2048 node, 3 layer





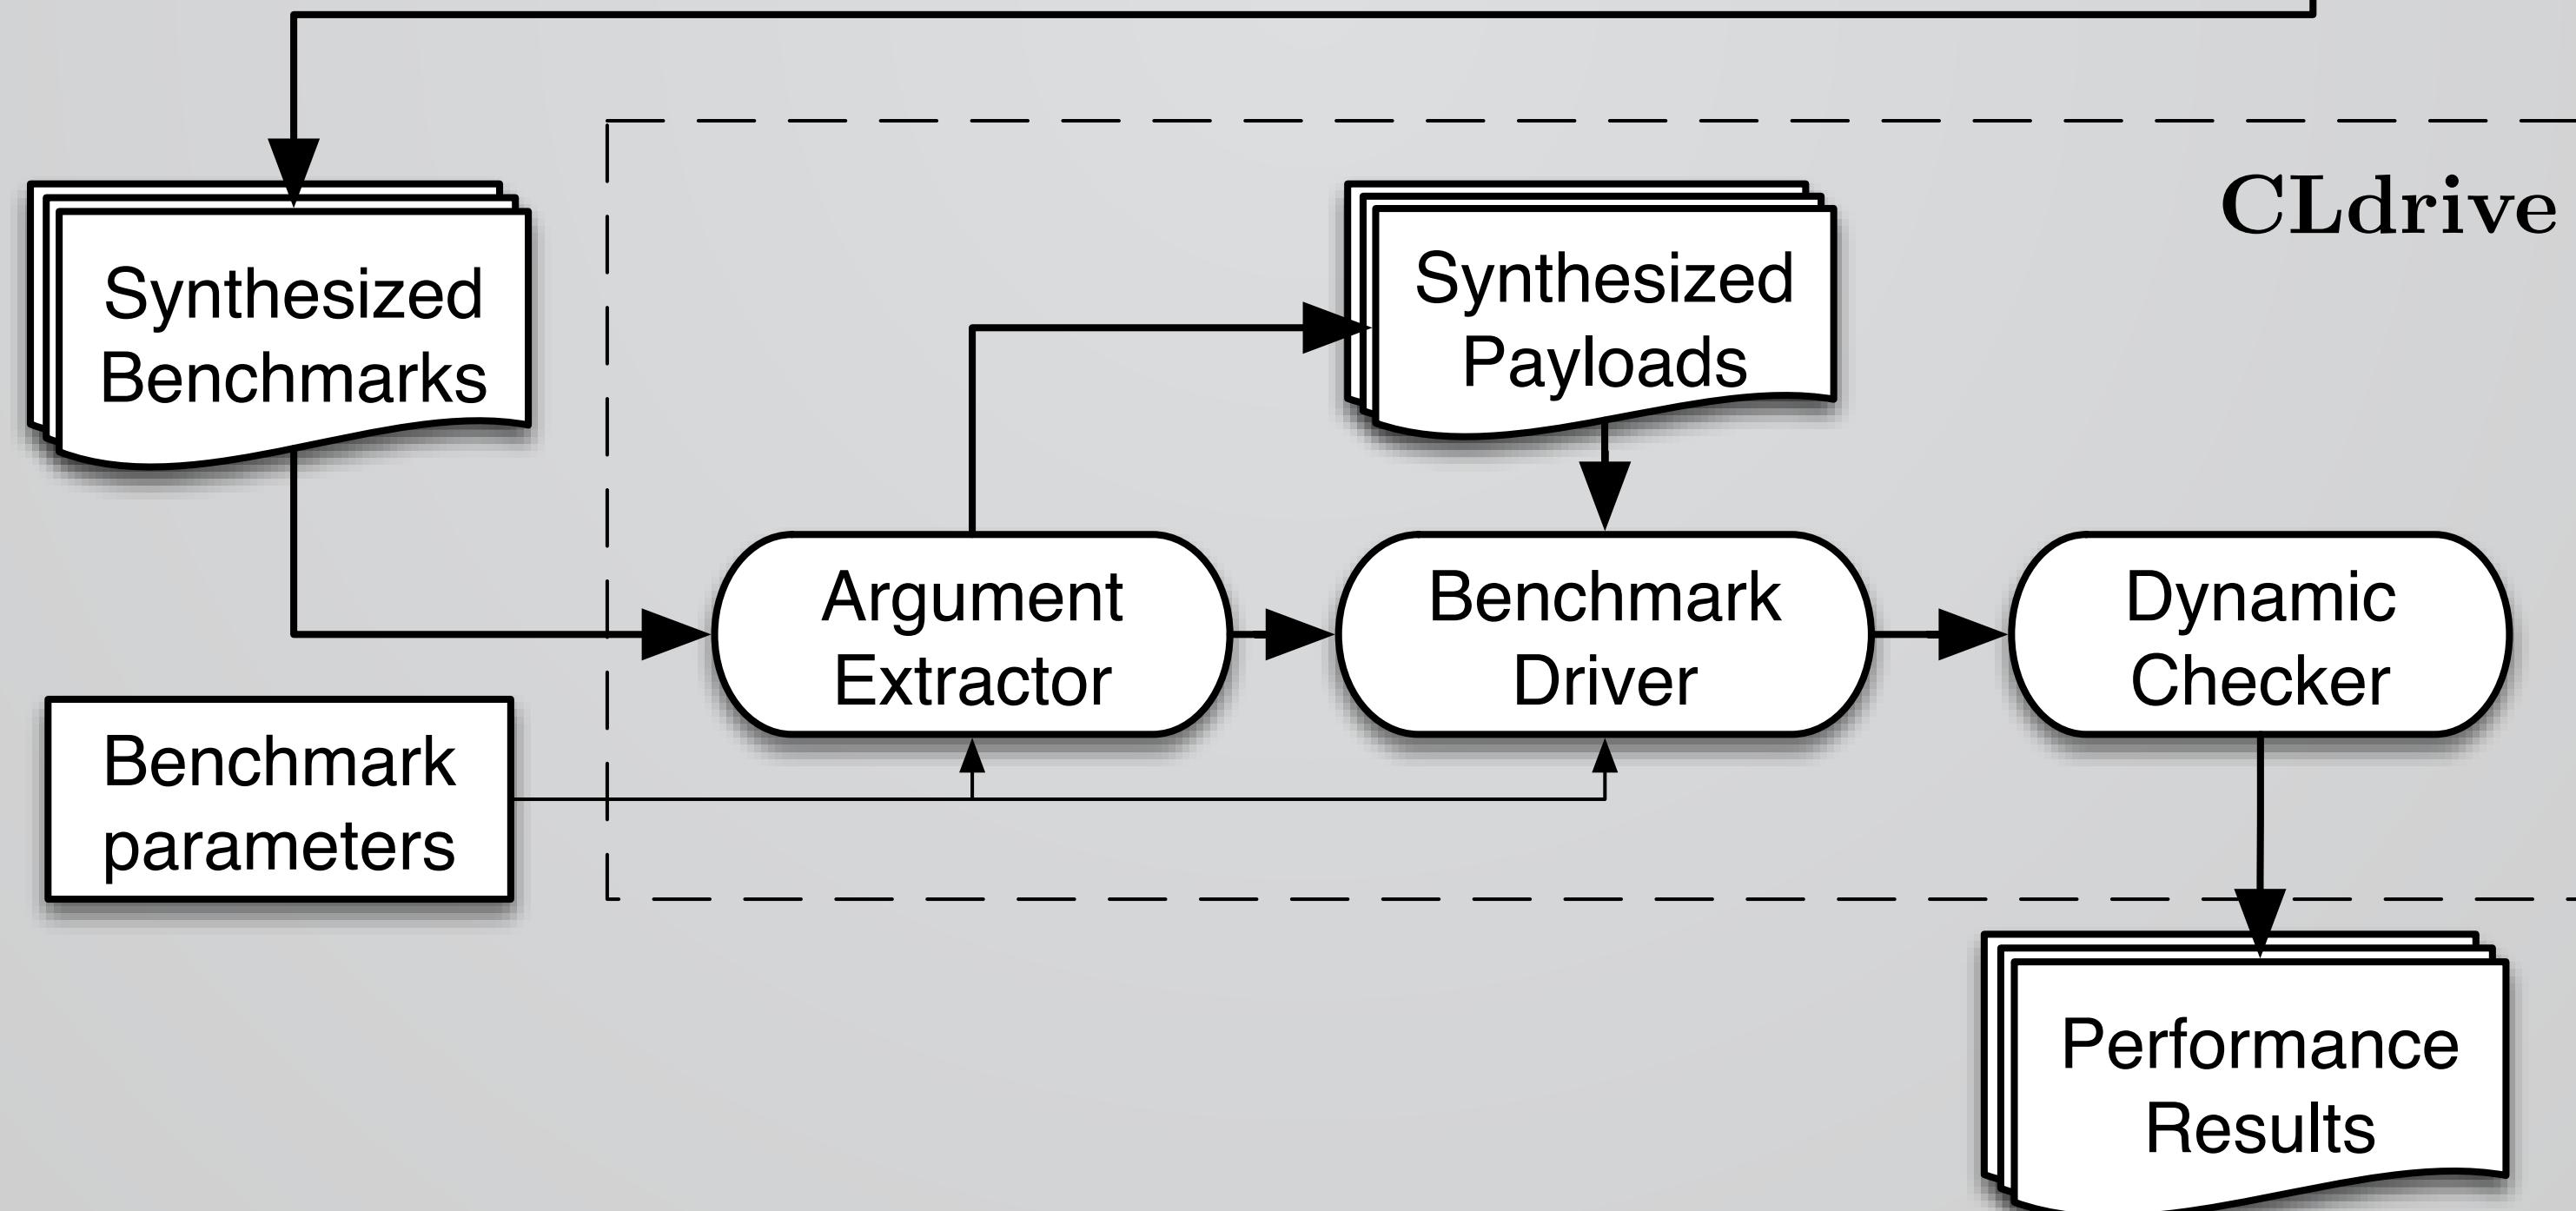
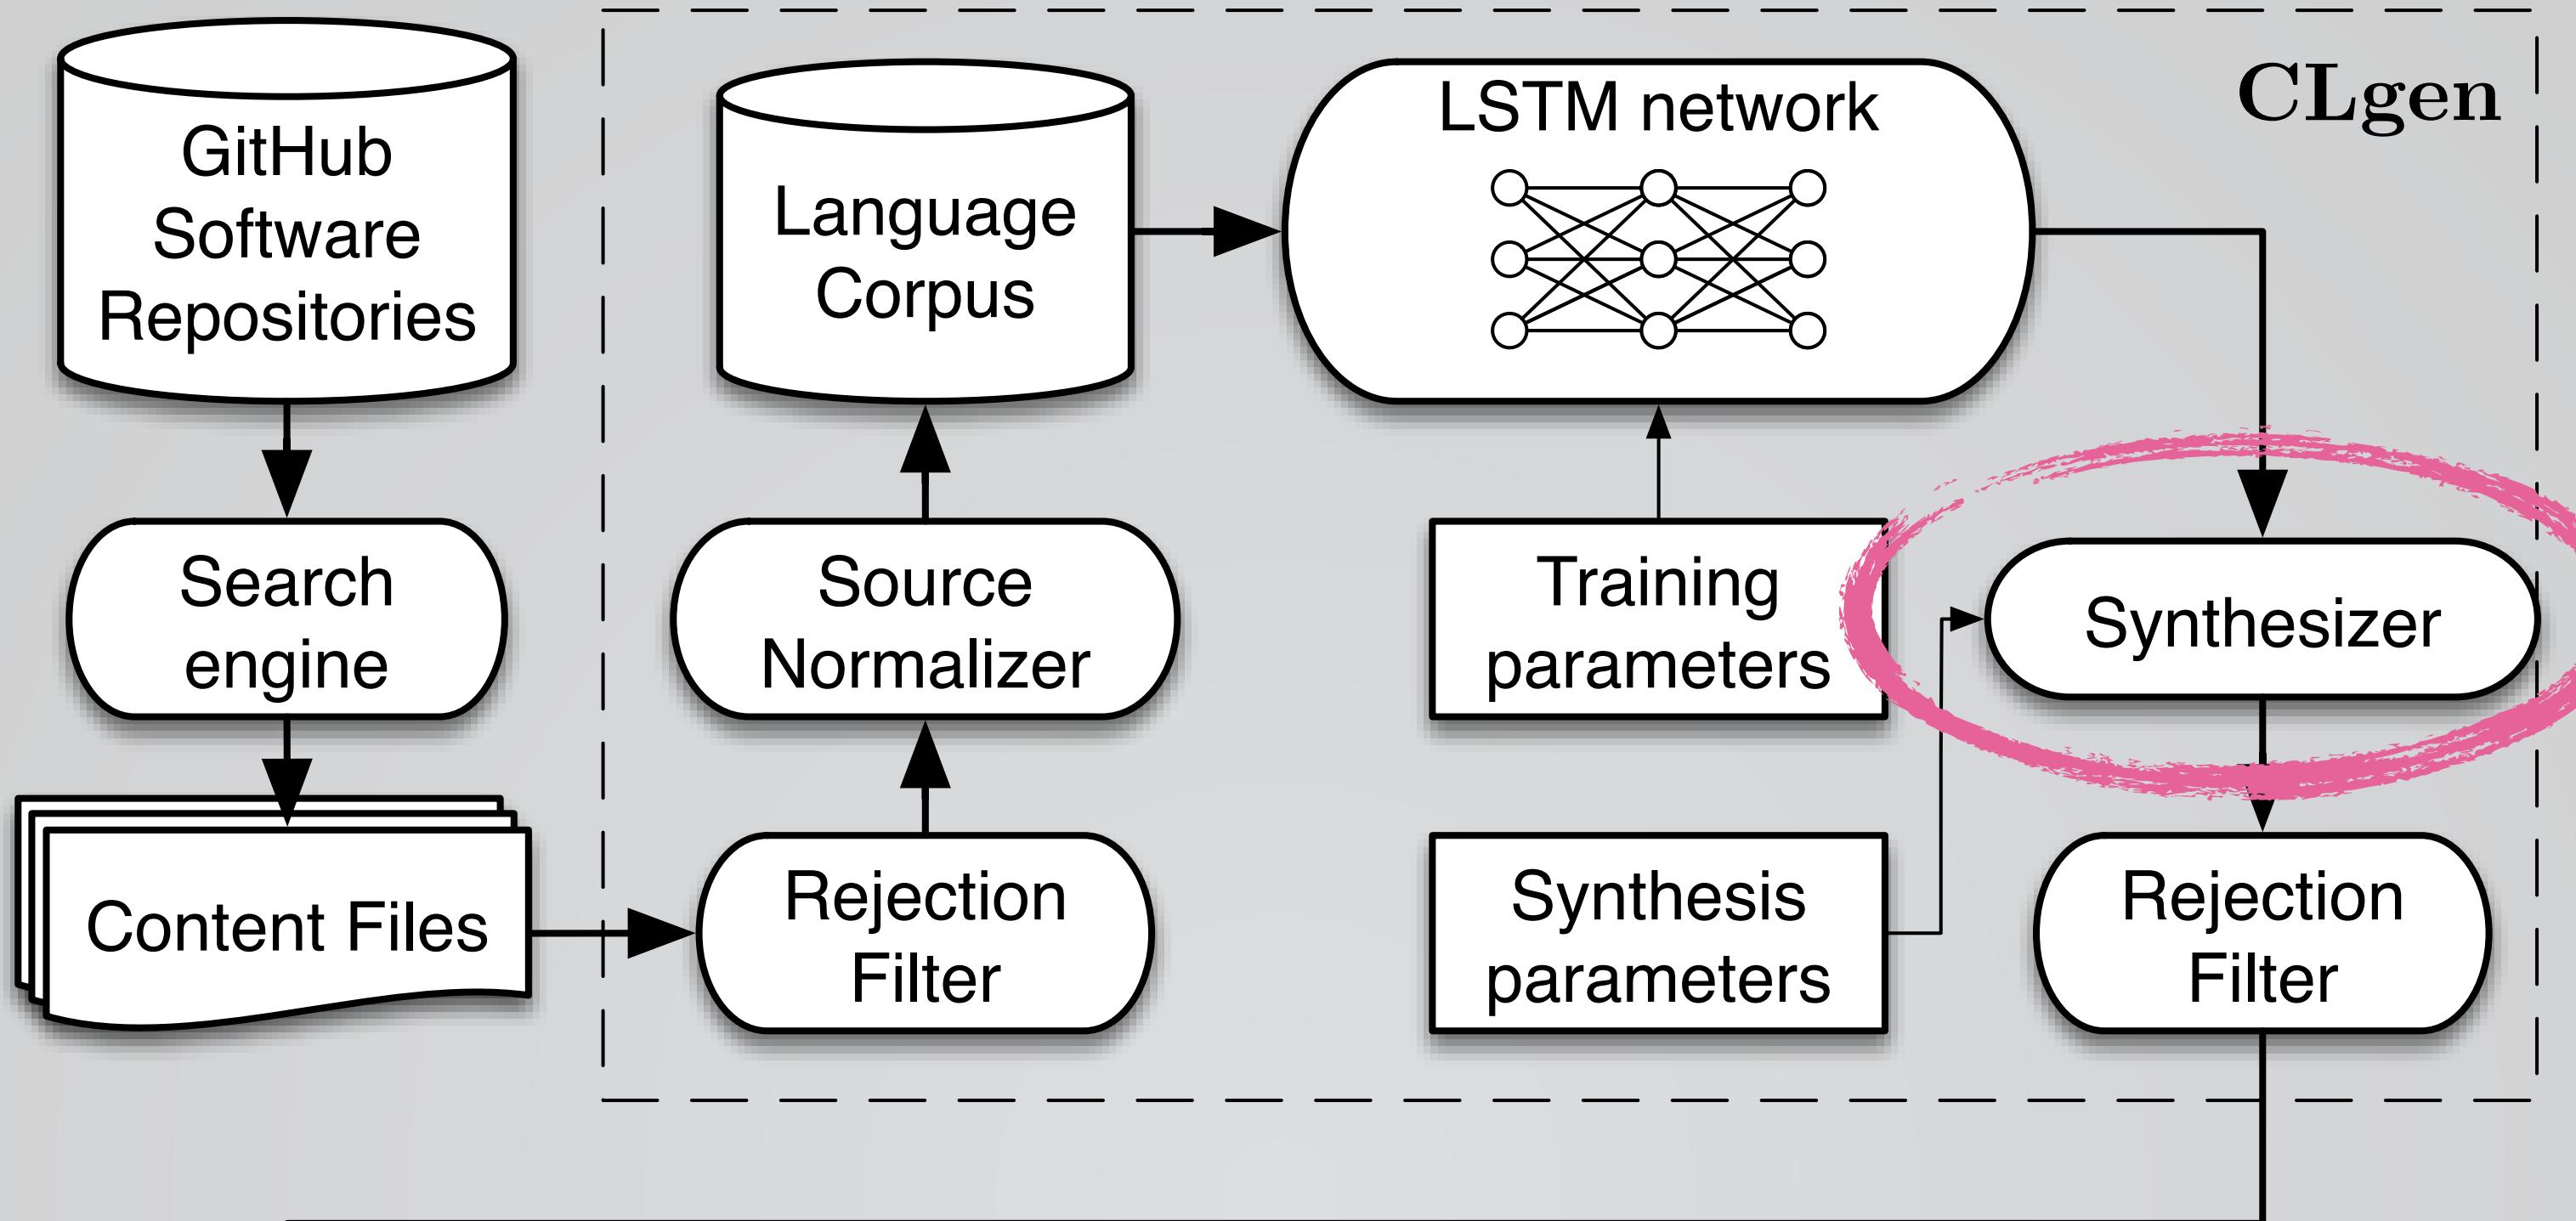
kernel synthesis

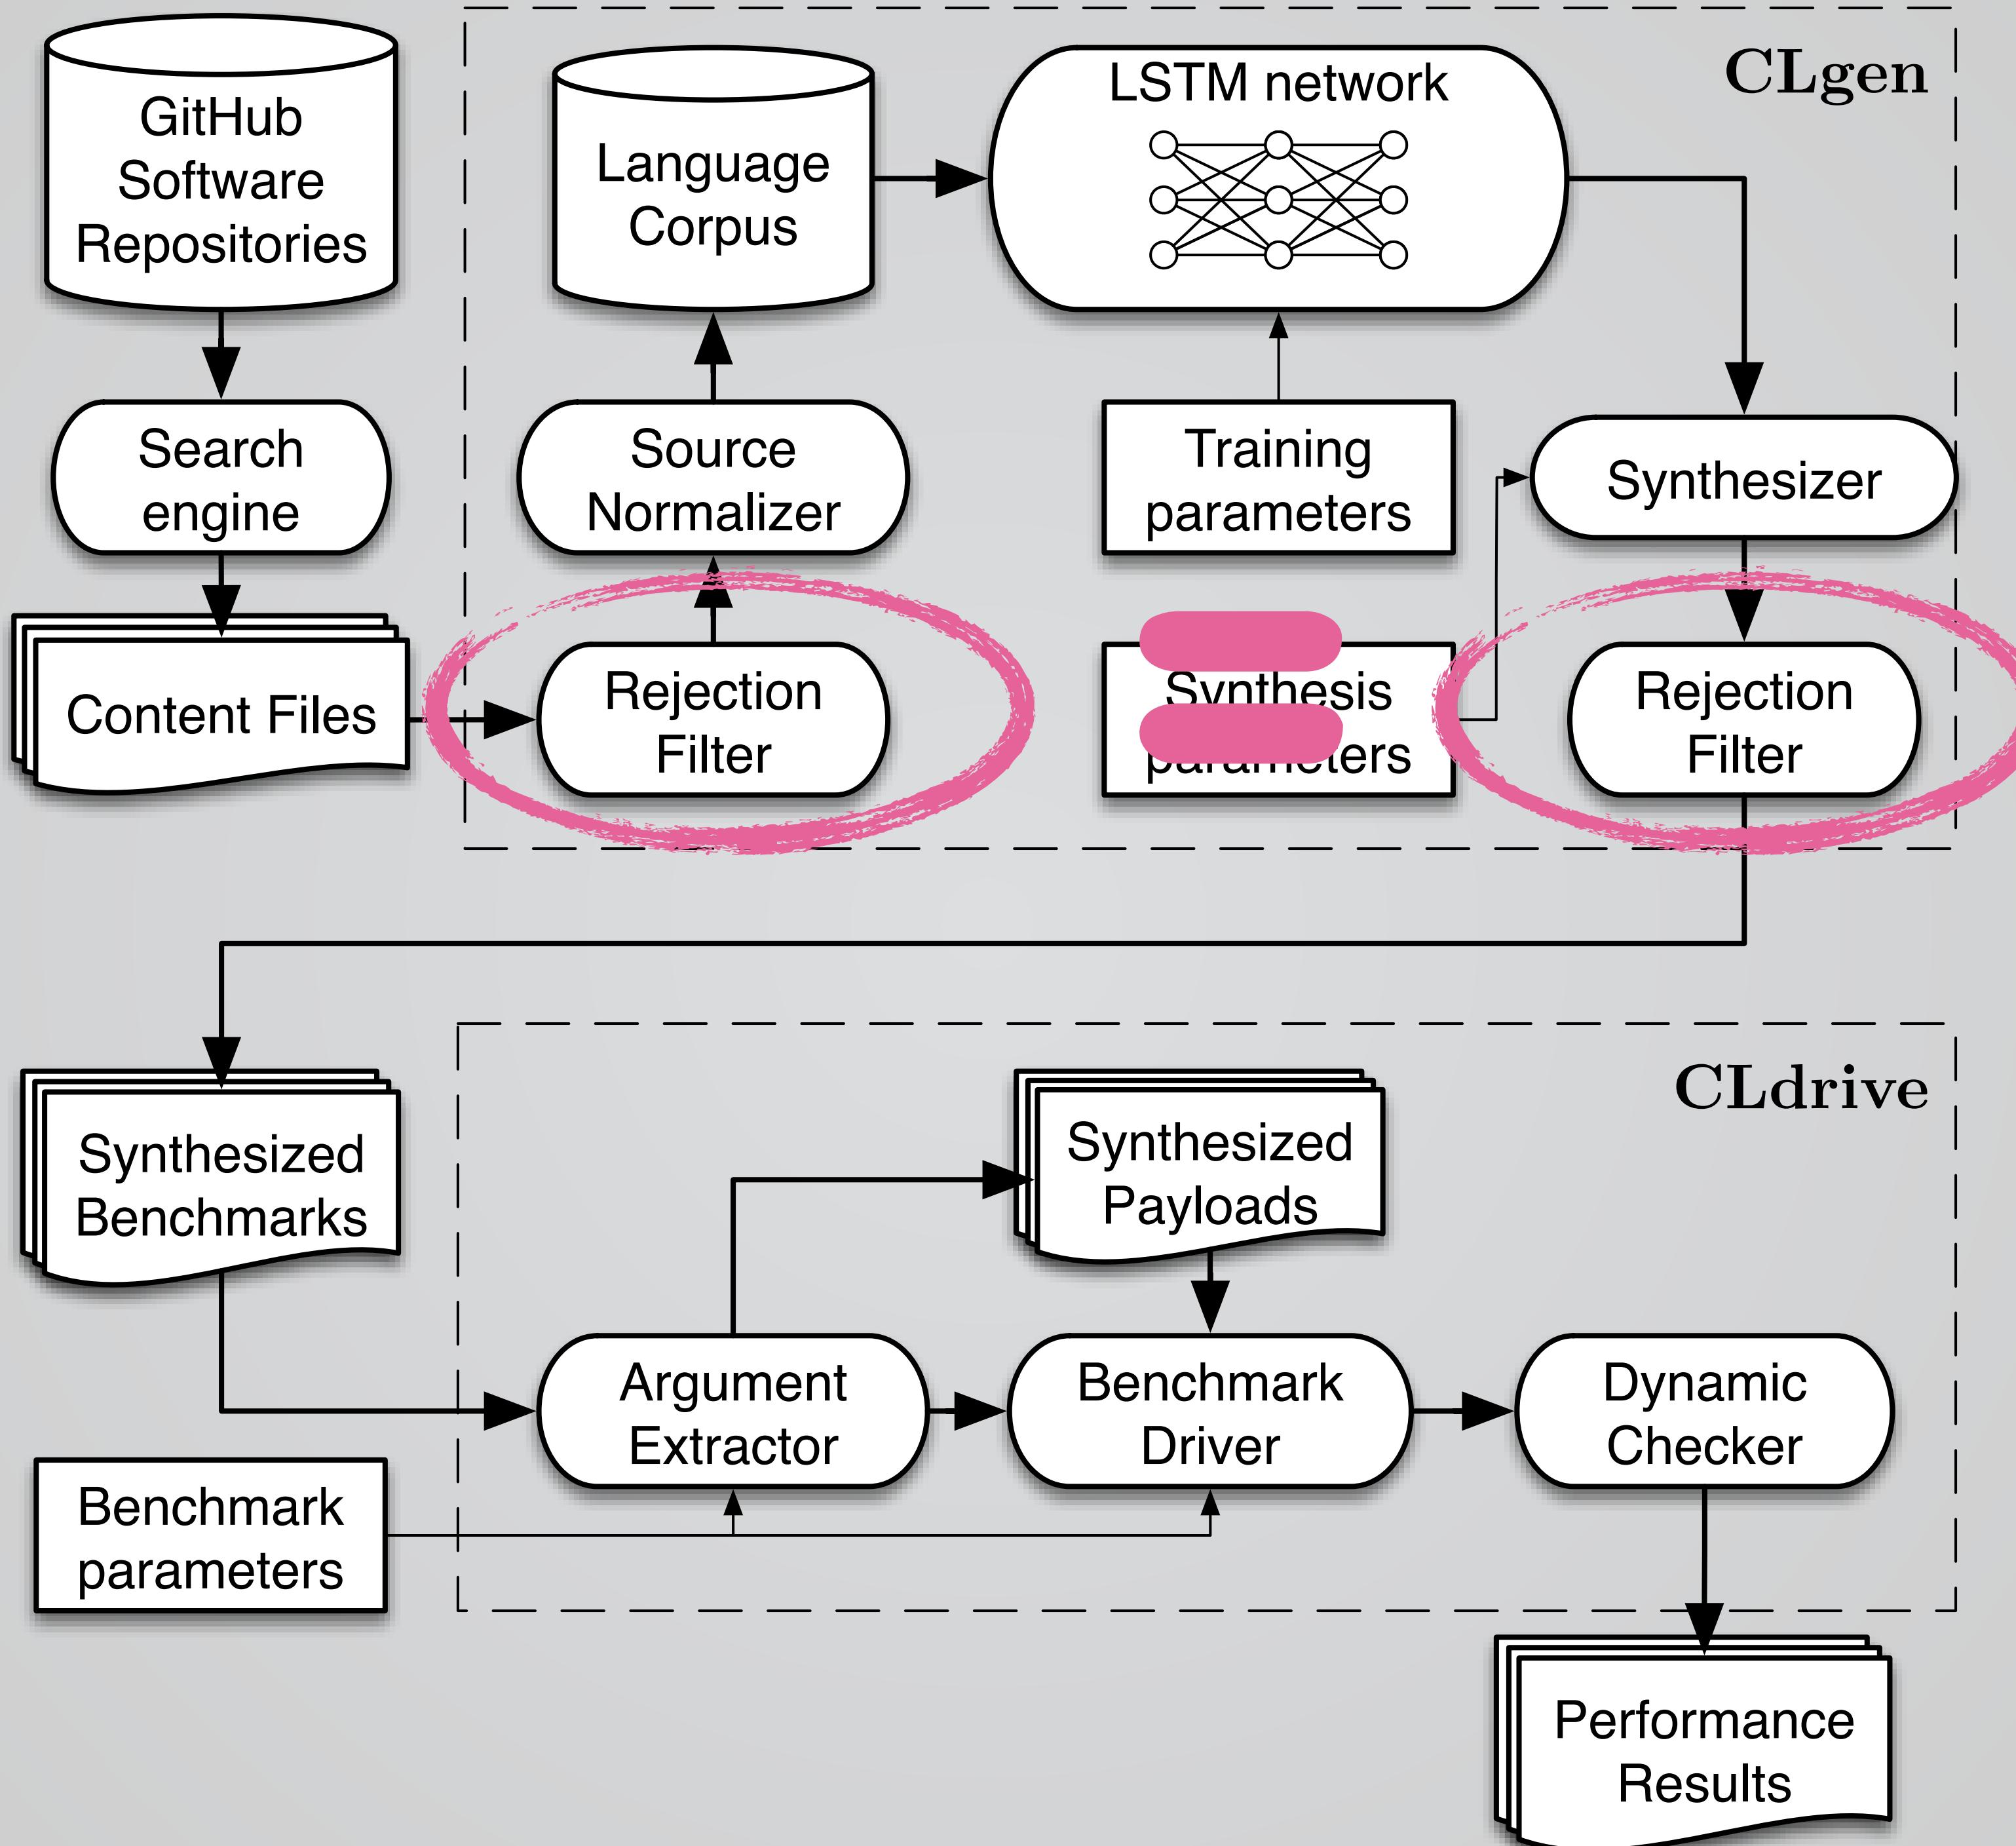
```
S = '__kernel void A(__global float* a) {'  
depth = 1  
while depth > 0:  
    c = predict_next_character(S)  
    if c == '{':  
        depth += 1  
    if c == '}':  
        depth -= 1  
    S += c  
  
return S
```

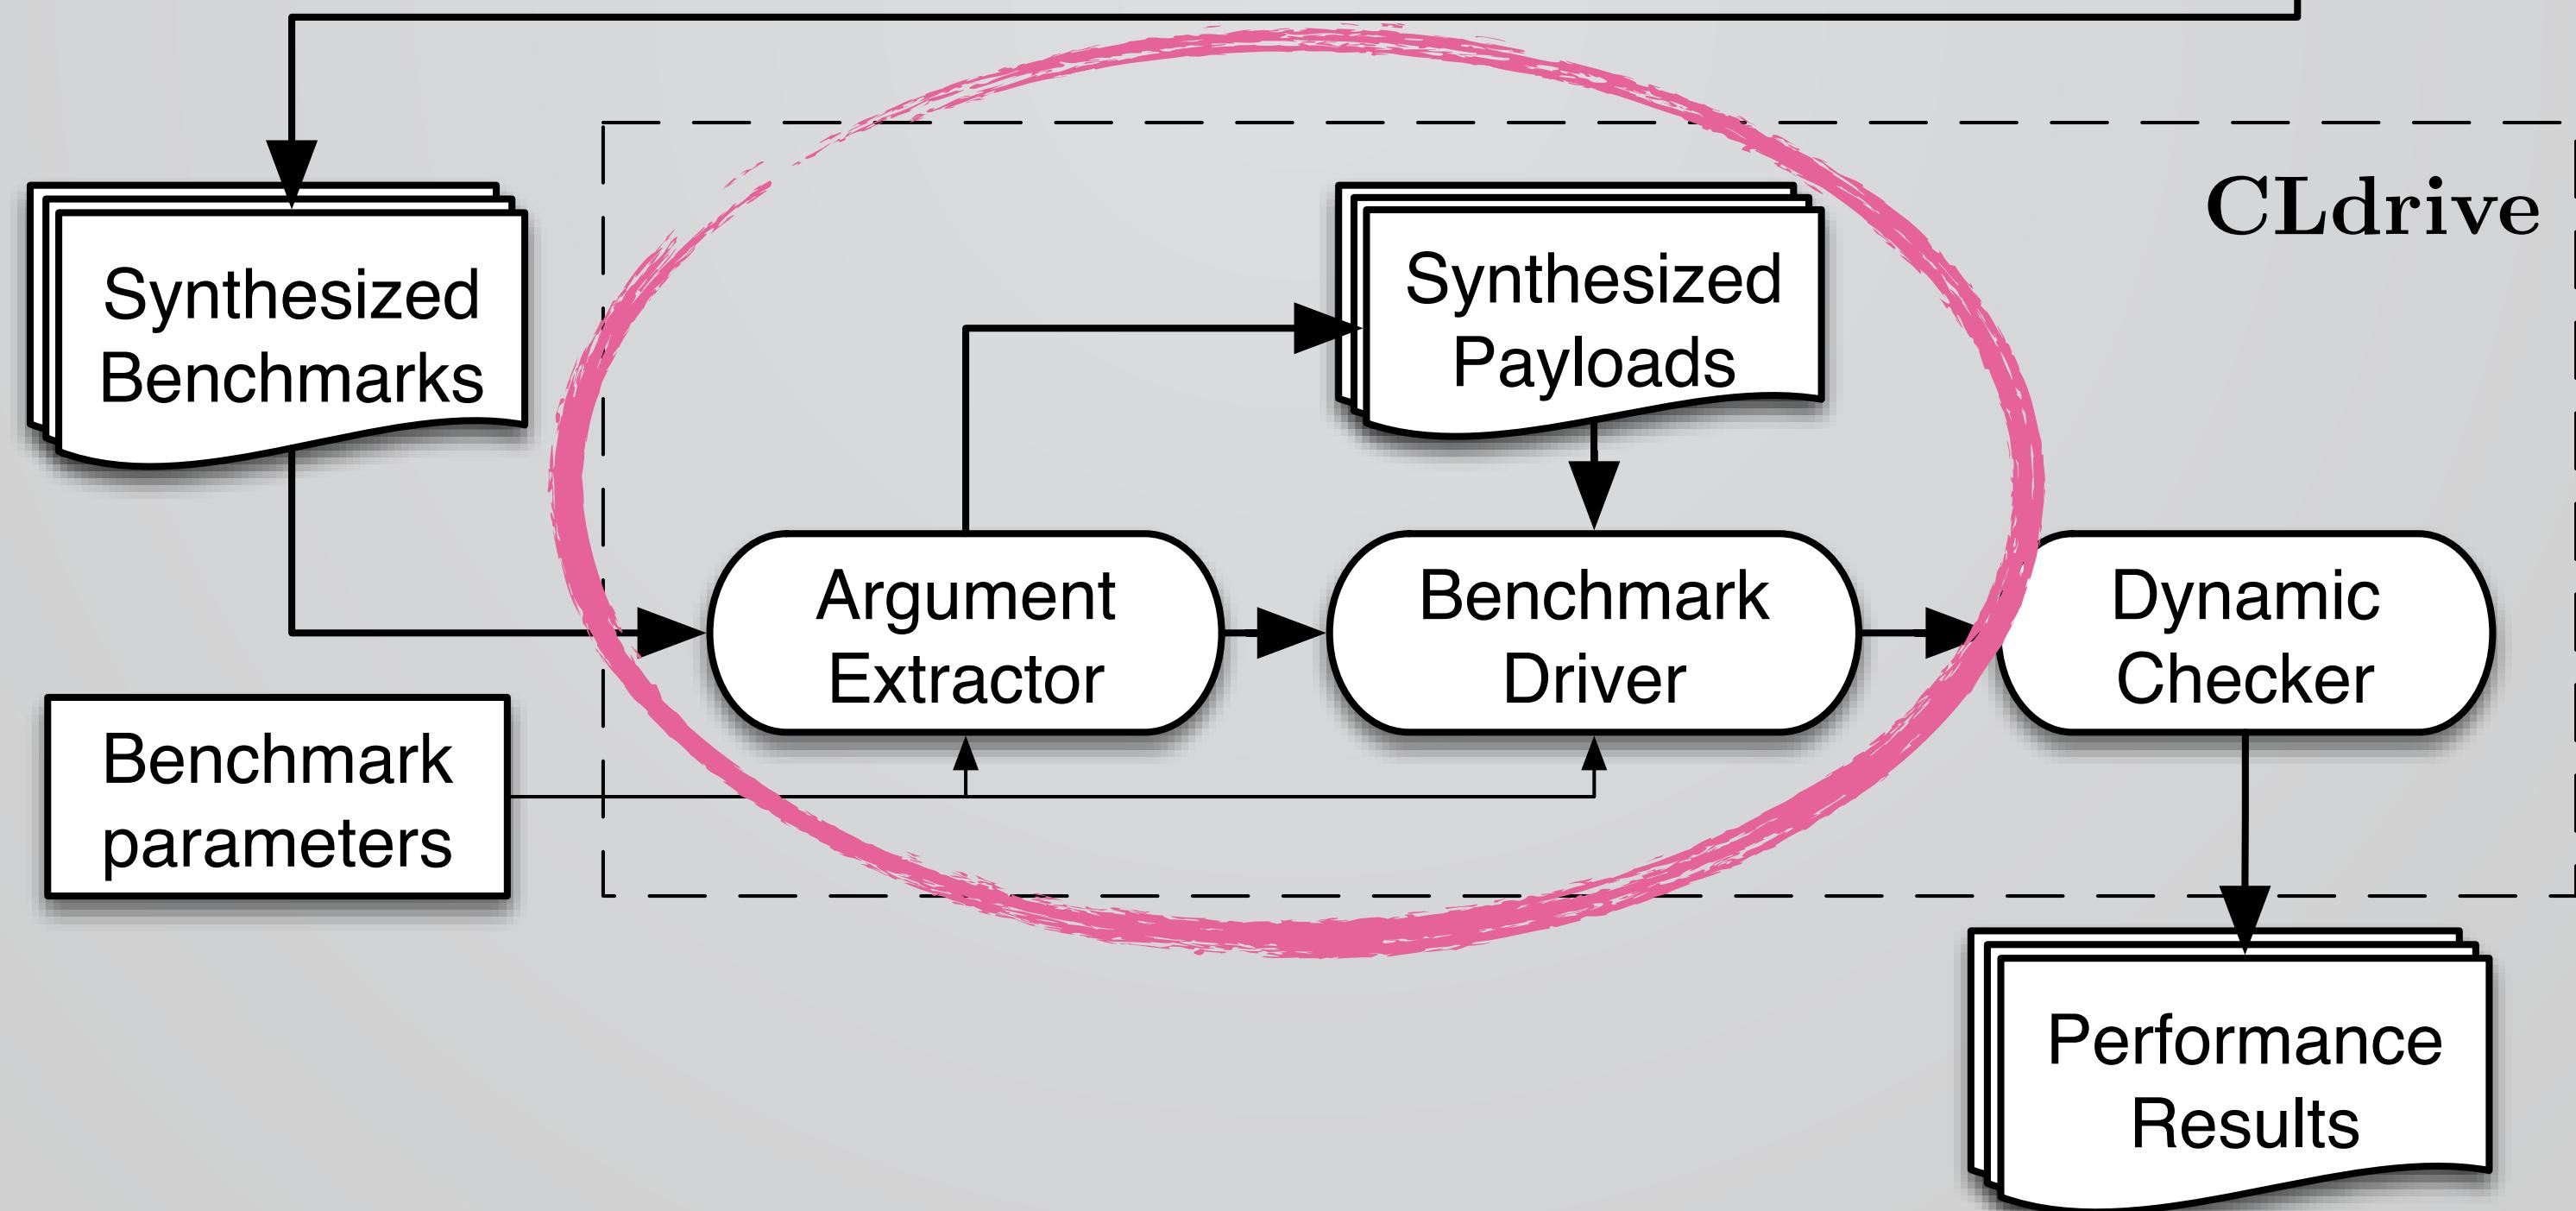
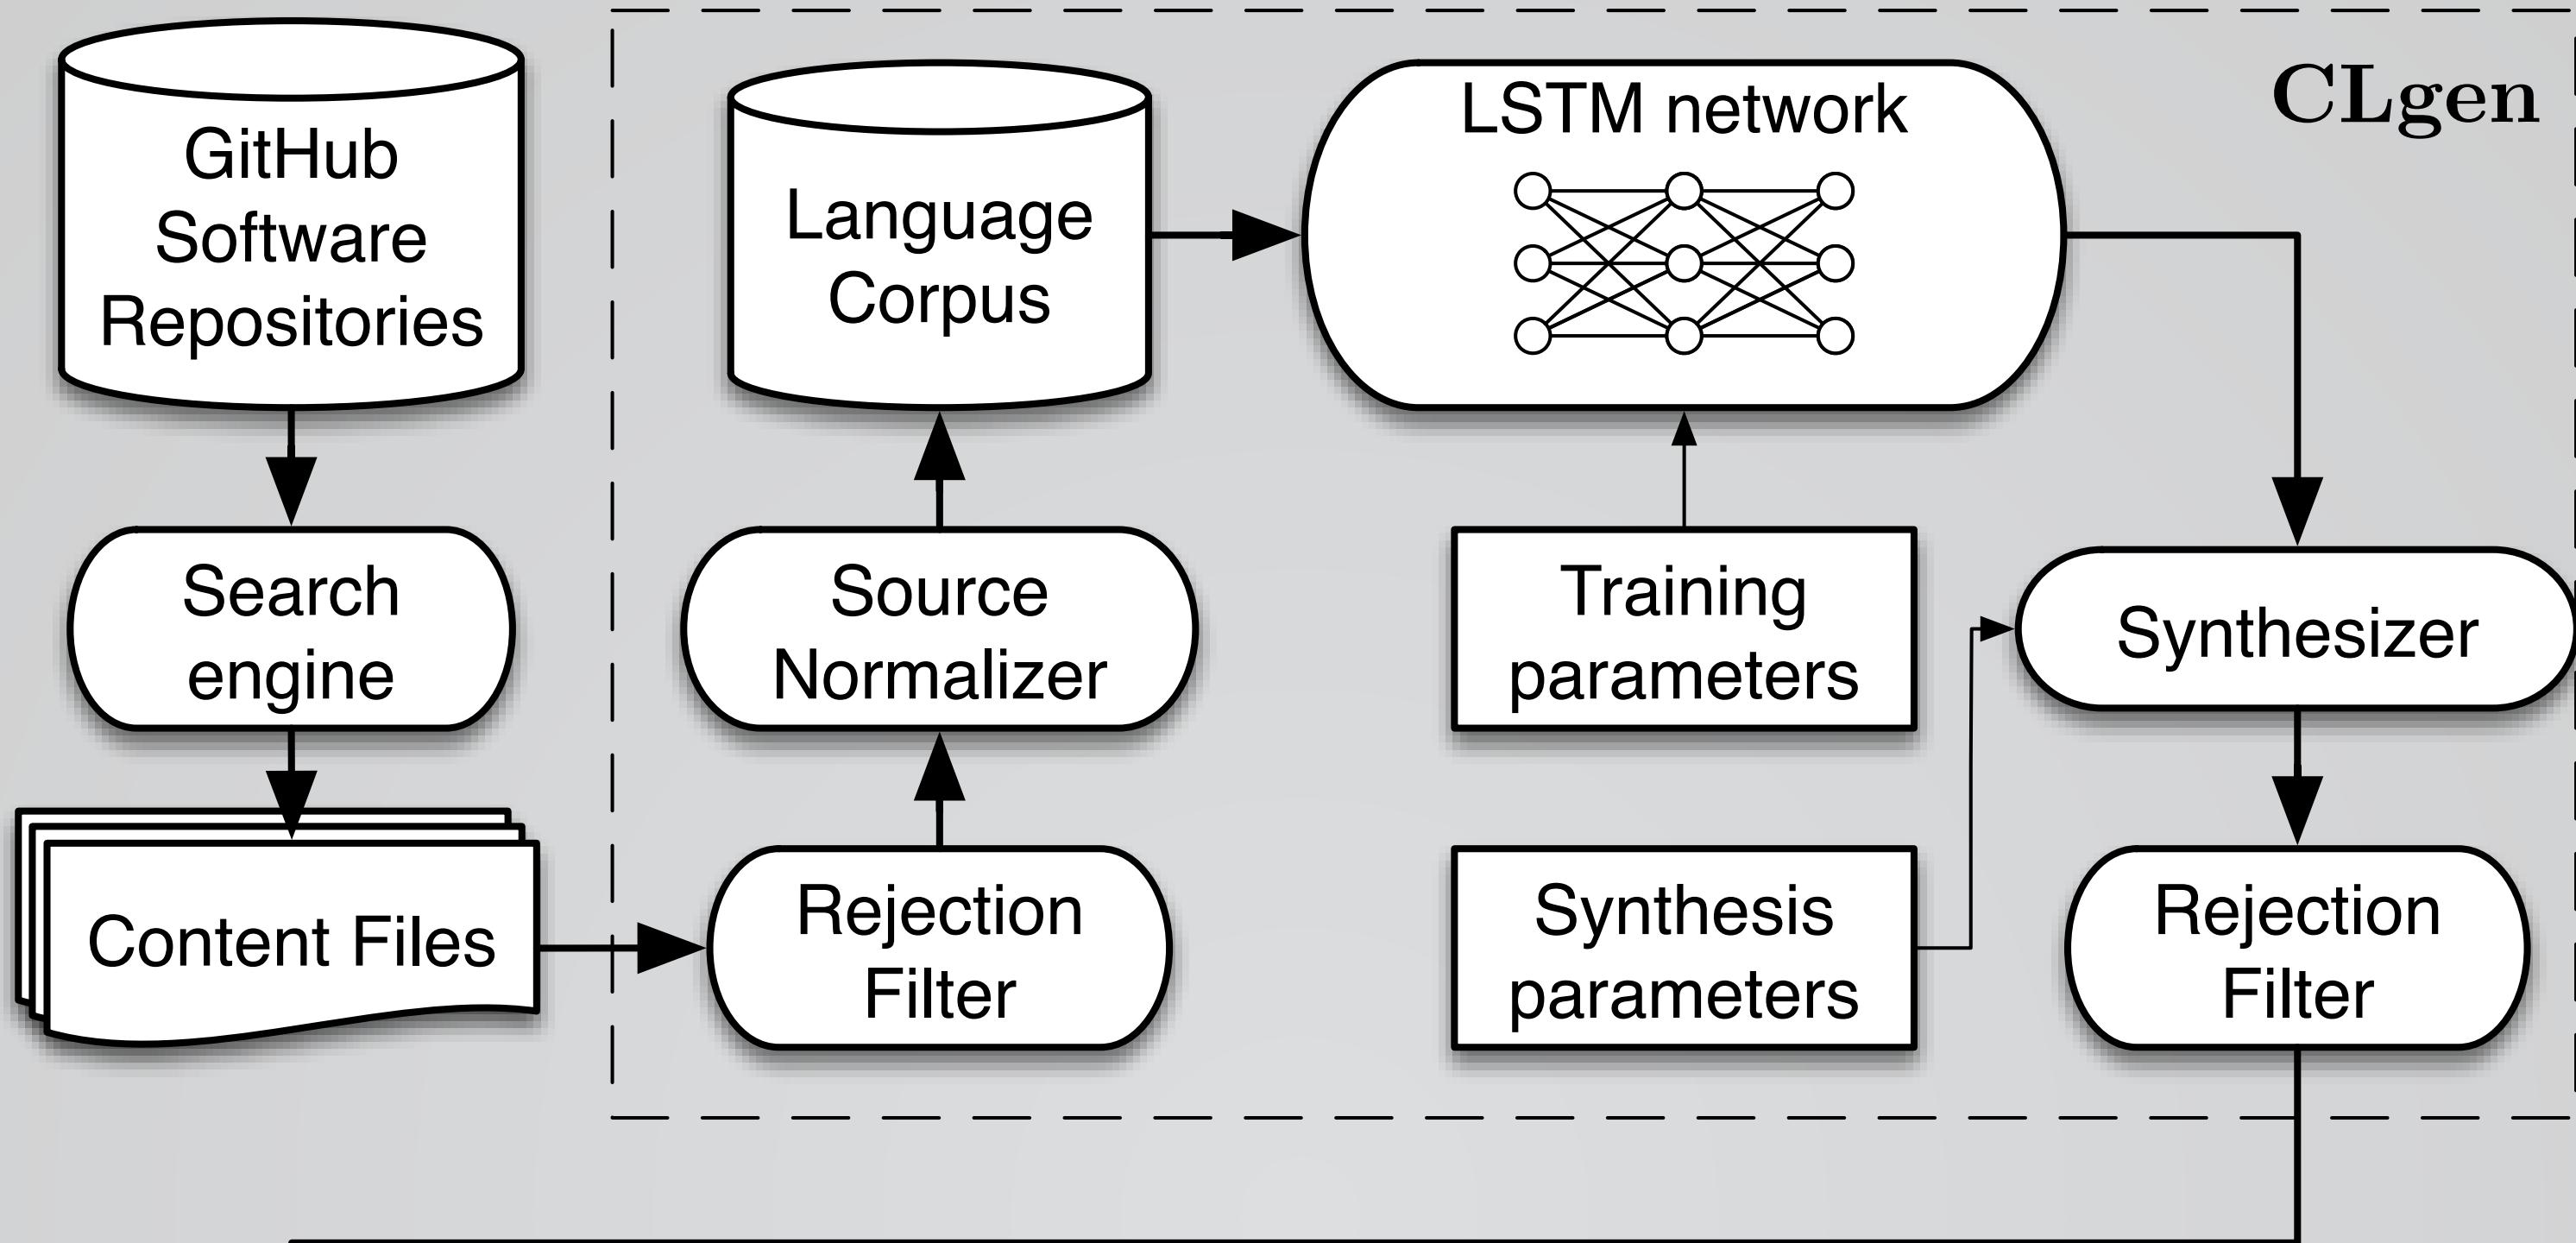


Deep Learning Program Generator.

\$ clgen

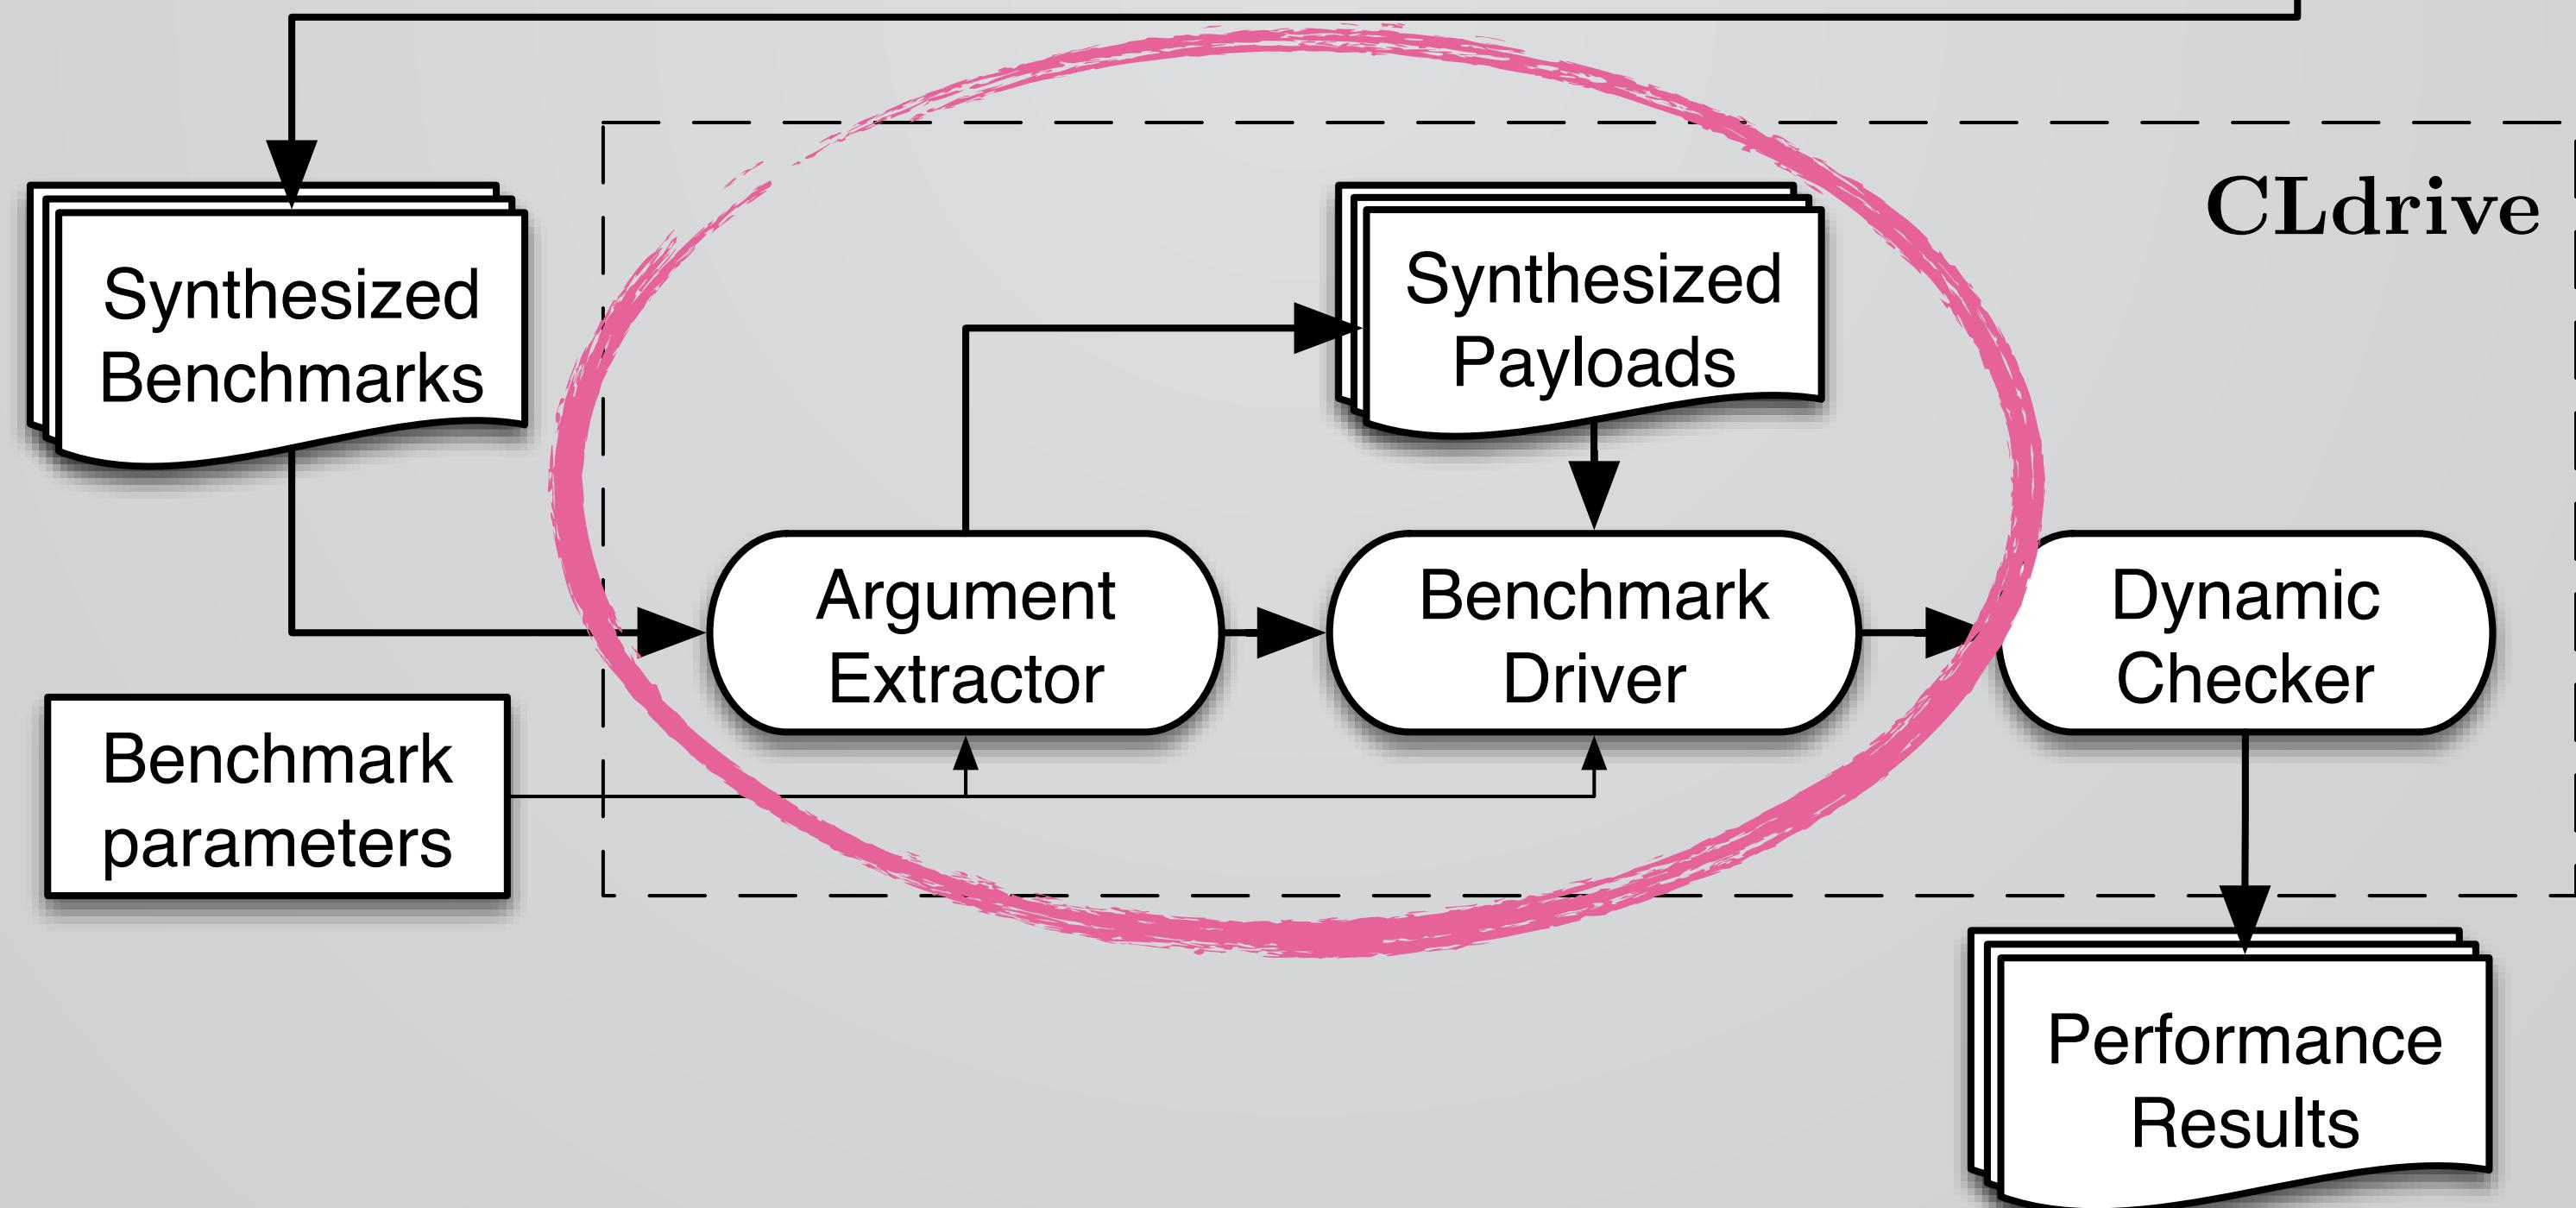
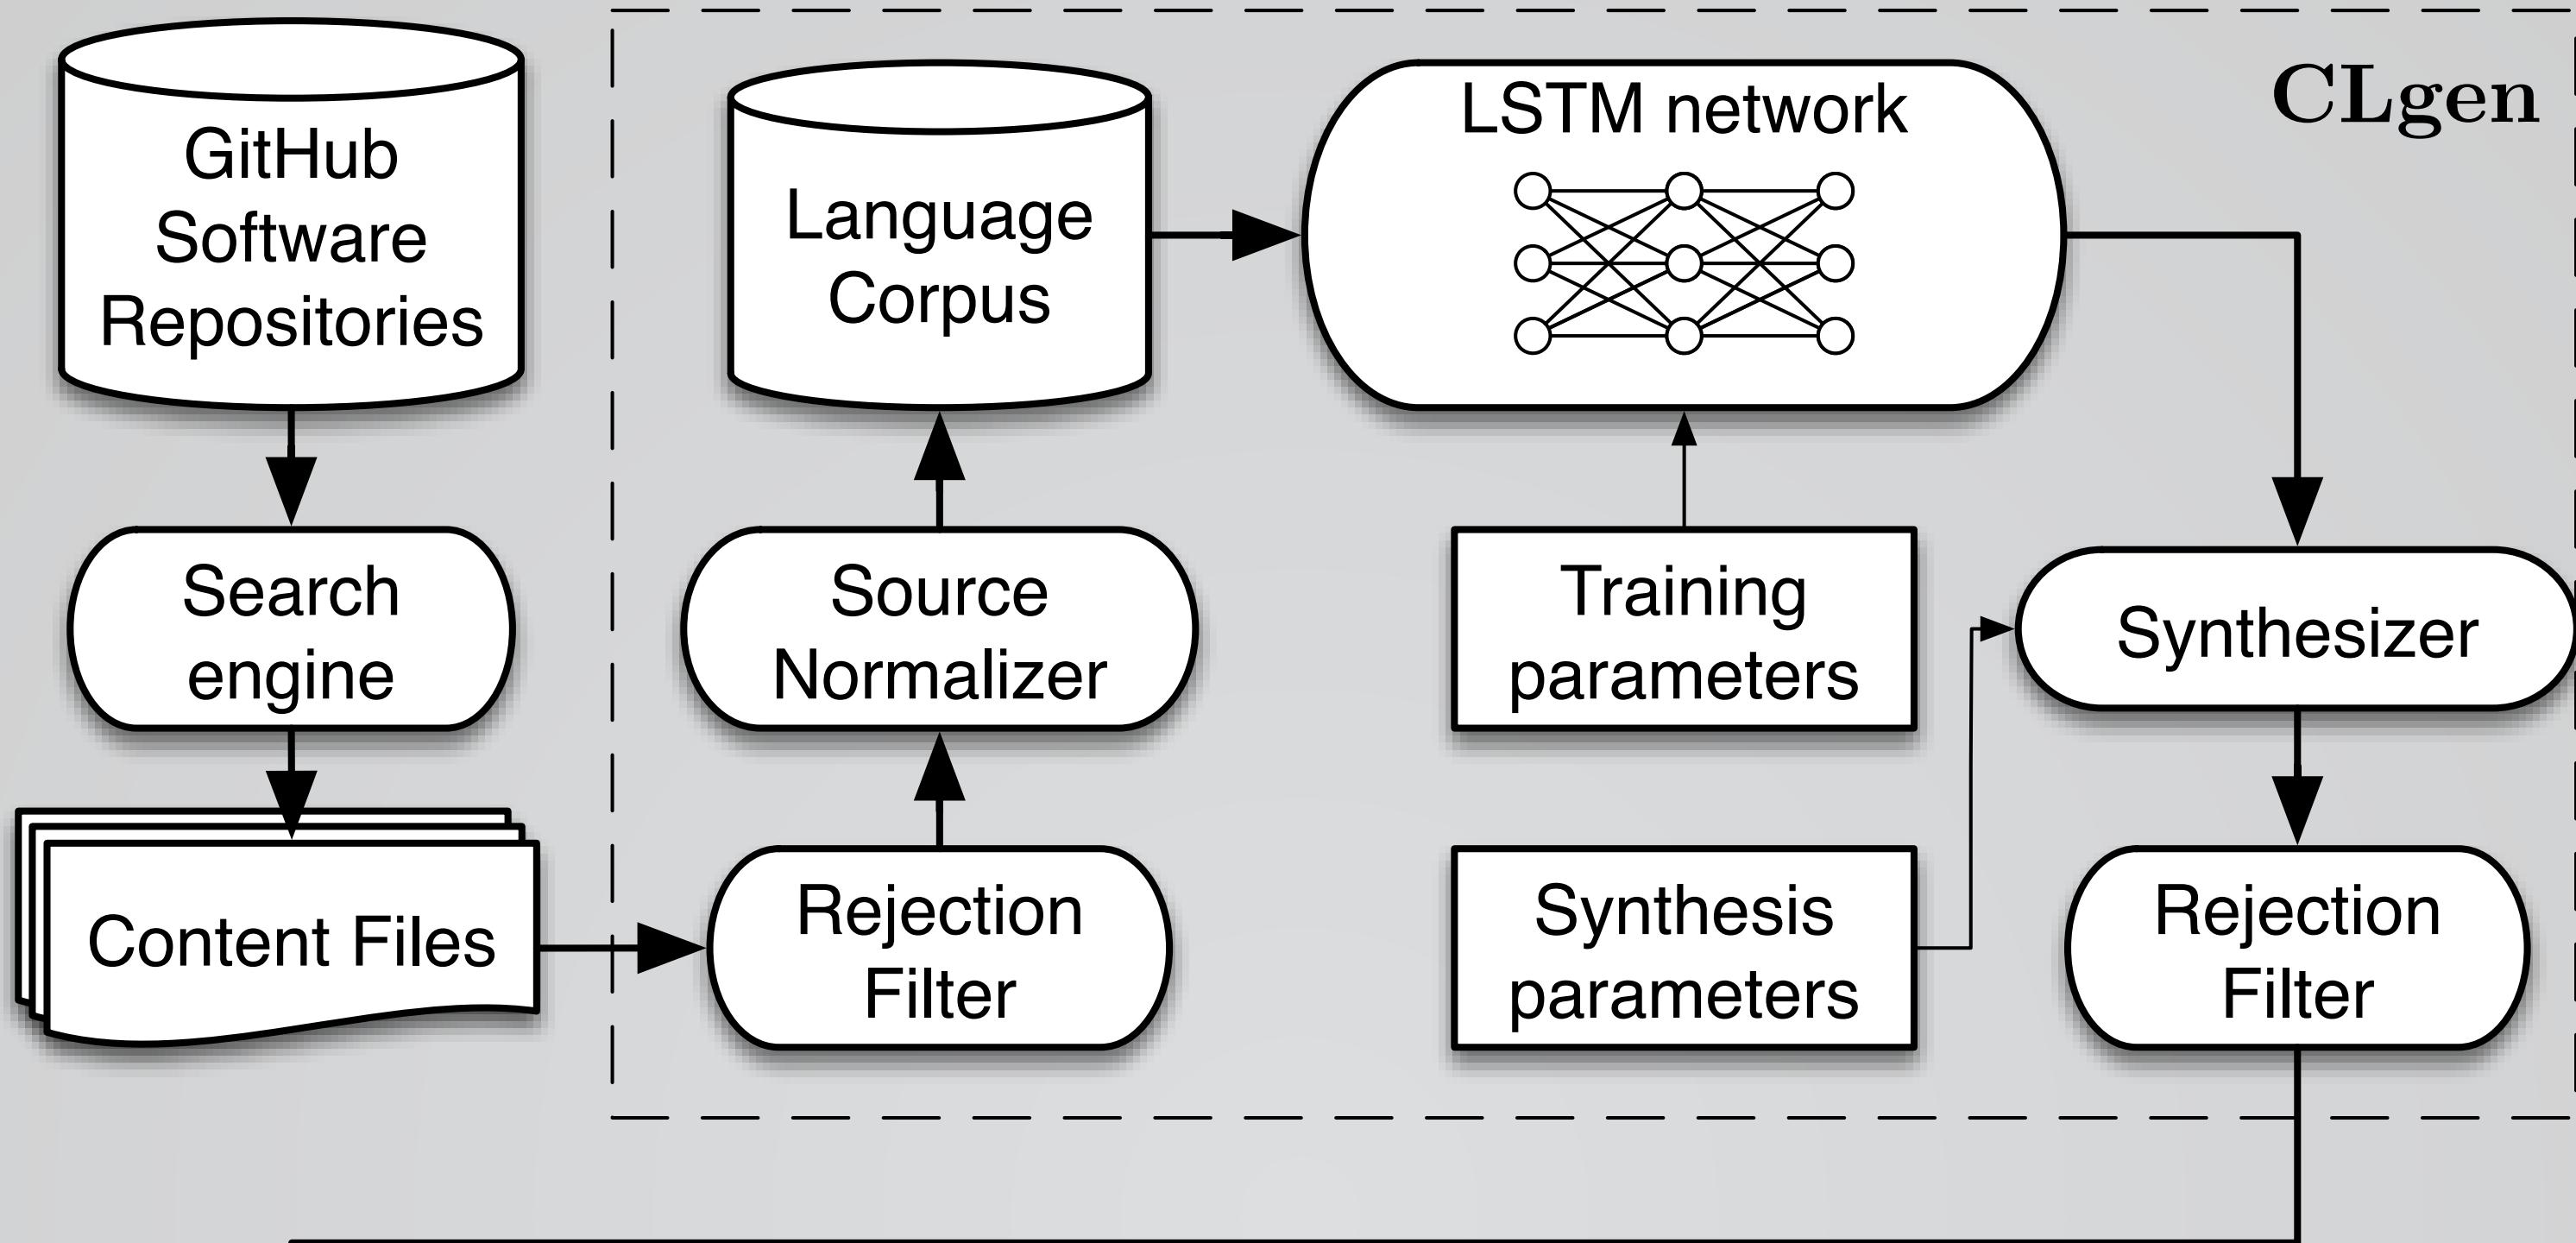


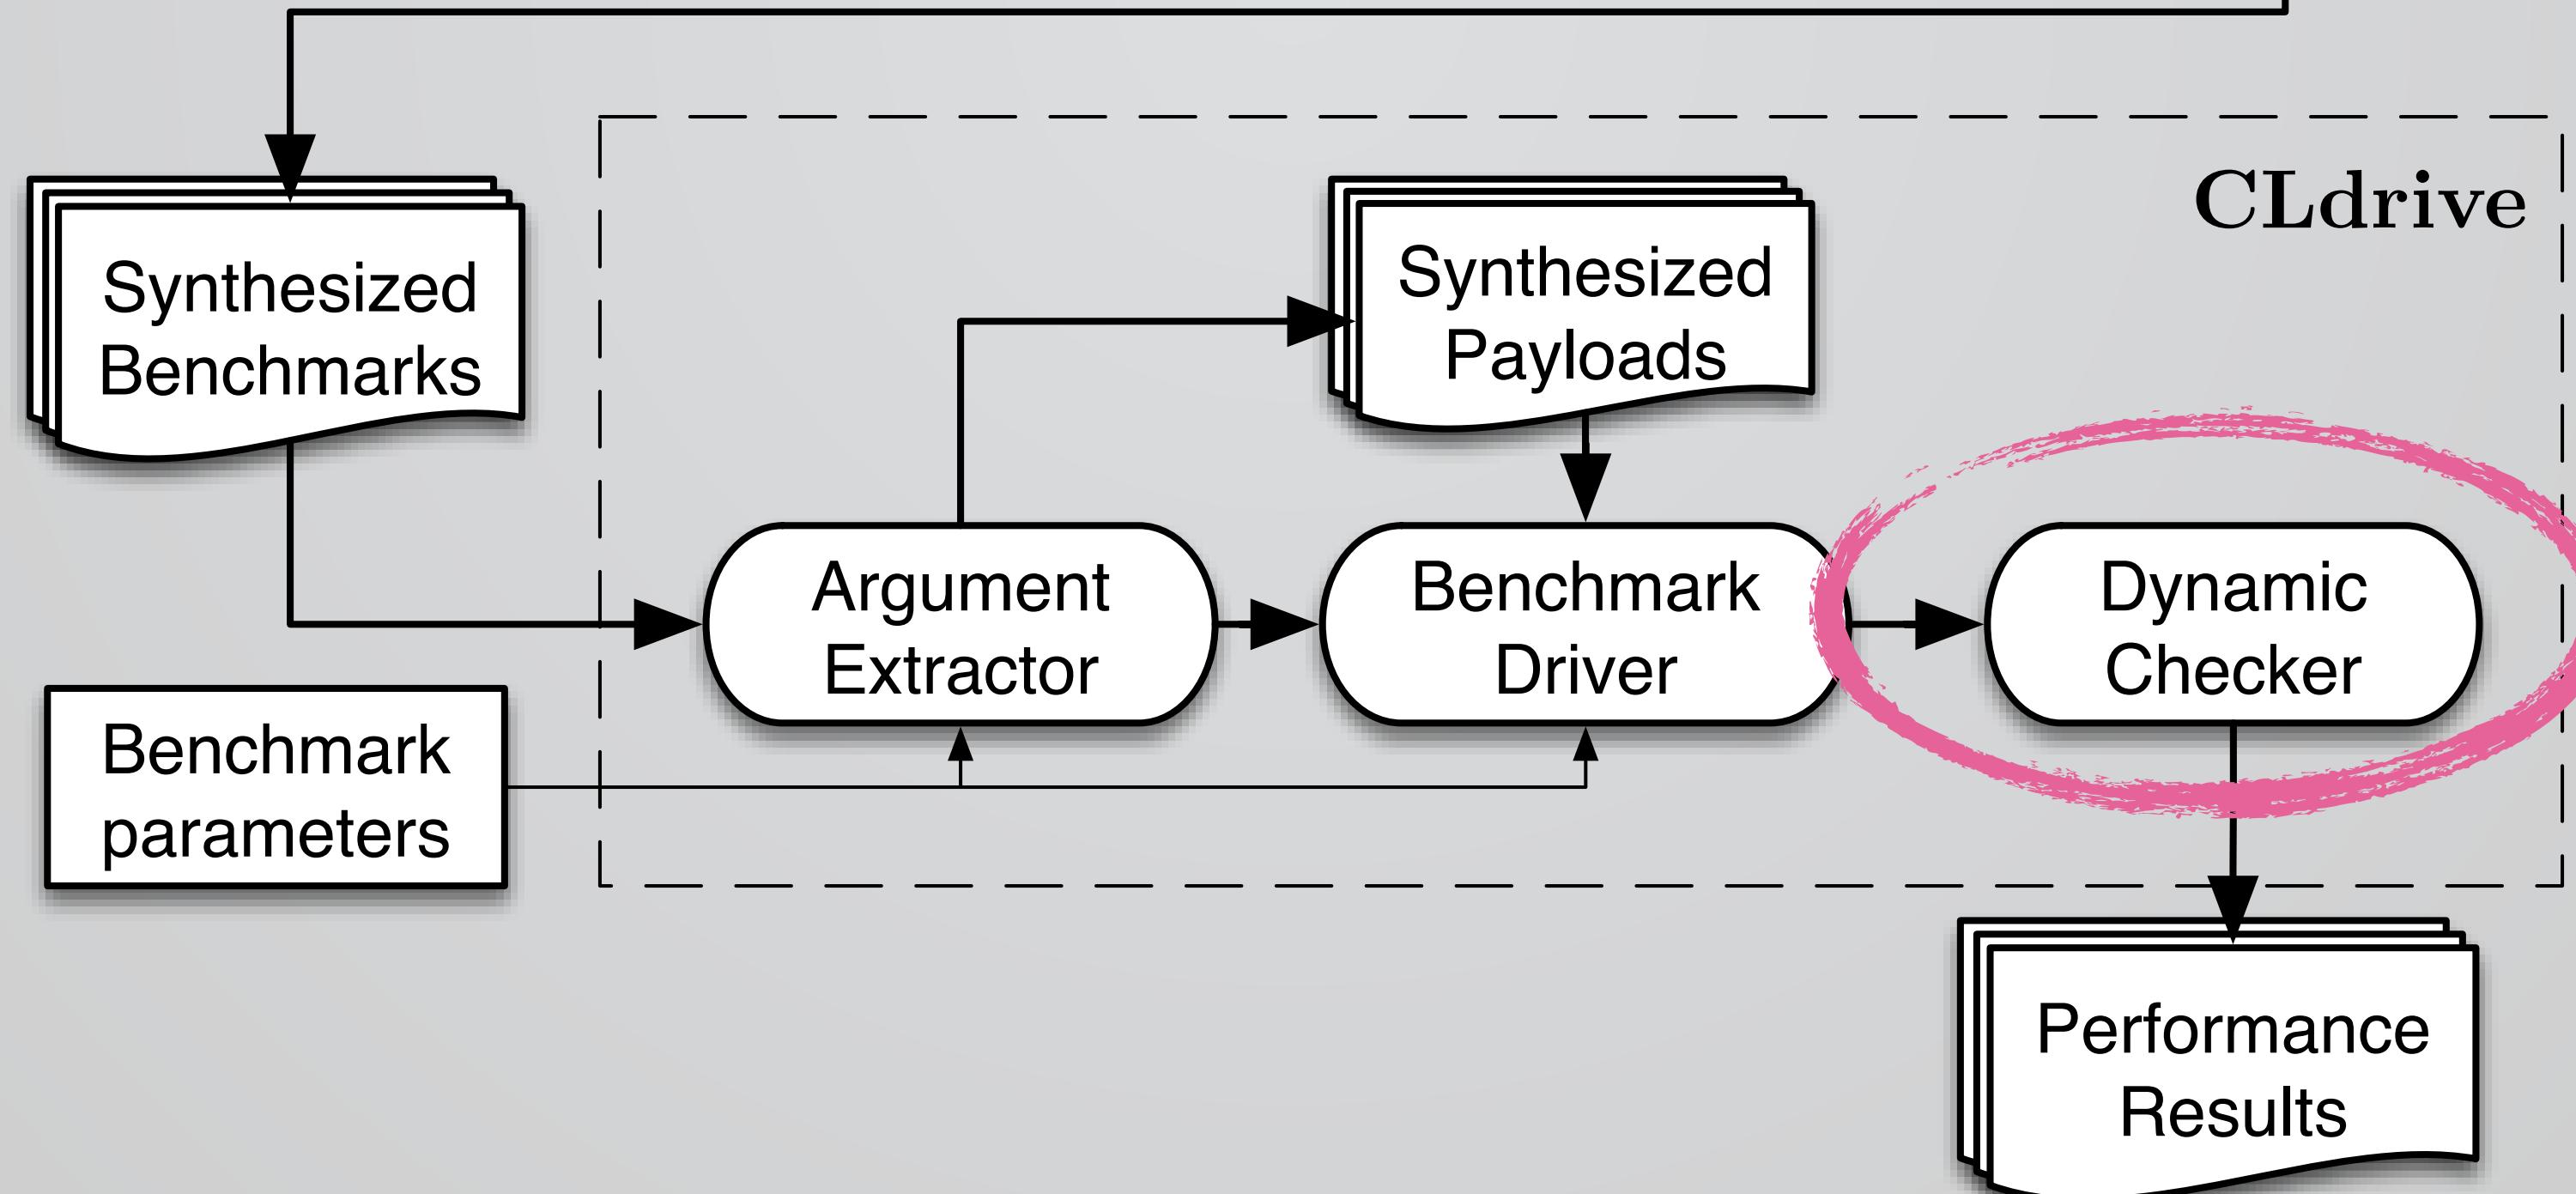
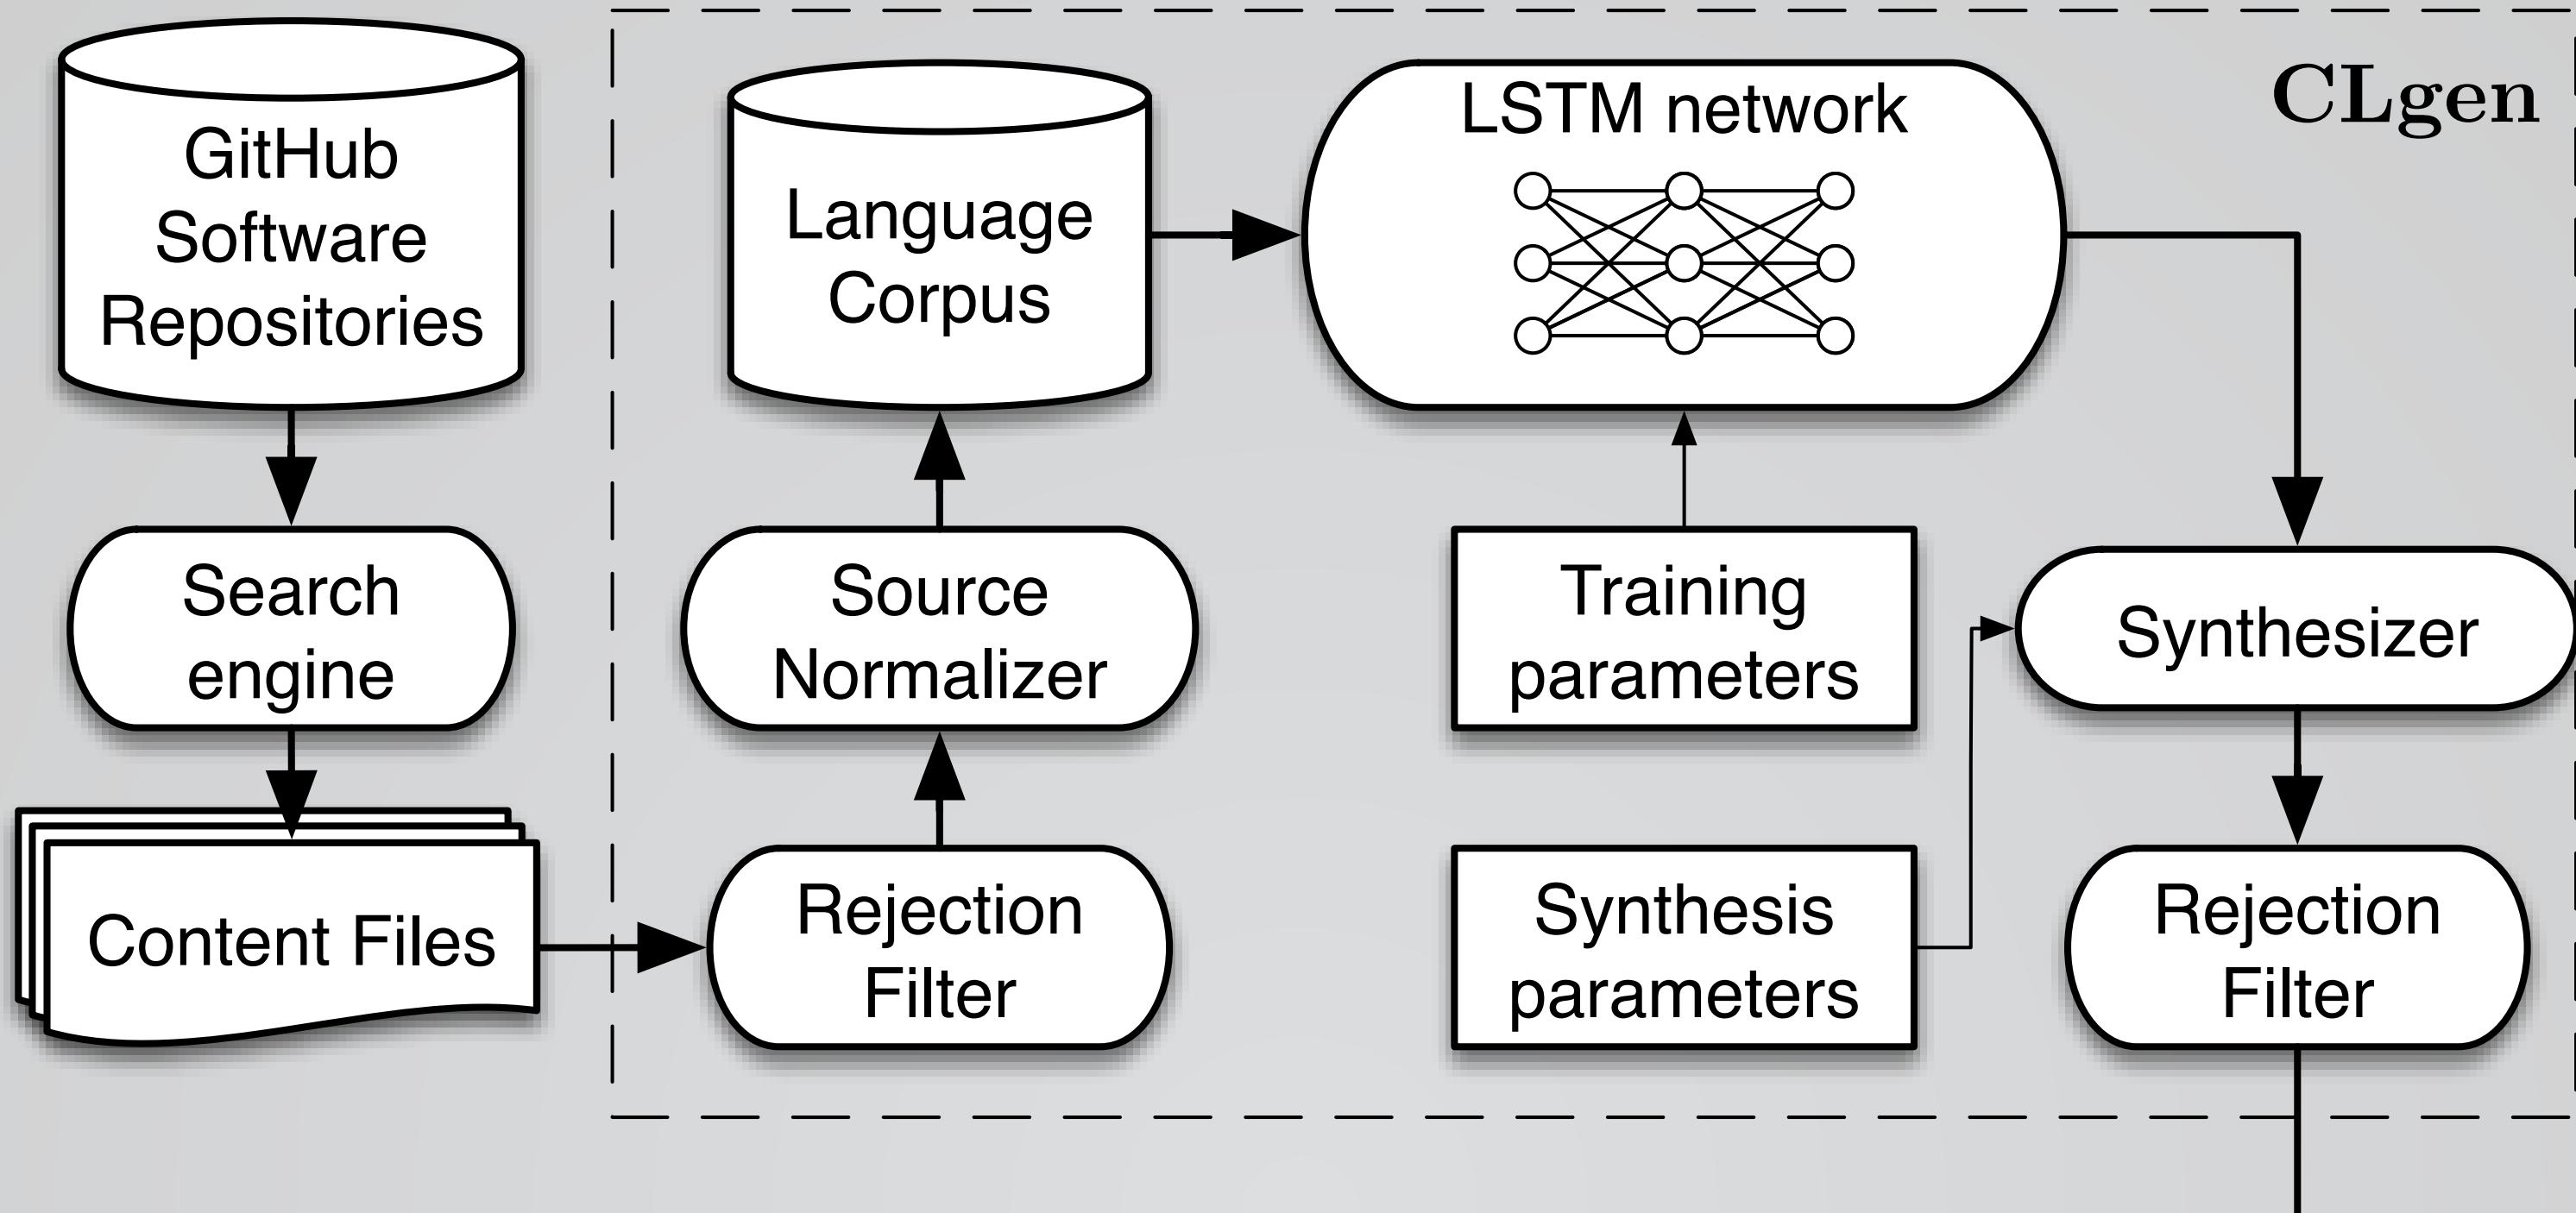




```
_kernel void A(__global float* a,  
          __global float* b,  
          __global float* c,  
          const int d) {  
  
    int e = get_global_id(0);  
    if (e >= d) {  
        return;  
    }  
    c[e] = a[e] + b[e] + 2 * a[e] + b[e] + 4;  
}
```

```
_kernel void A(__global float* a,
          __global float* b,
          __global float* c,
          const int d) {
    int e = get_global_id(0);
    if (e >= d) {
        return;
    }
    c[e] = a[e] + b[e] + 2 * a[e] + b[e] + 4;
}
```

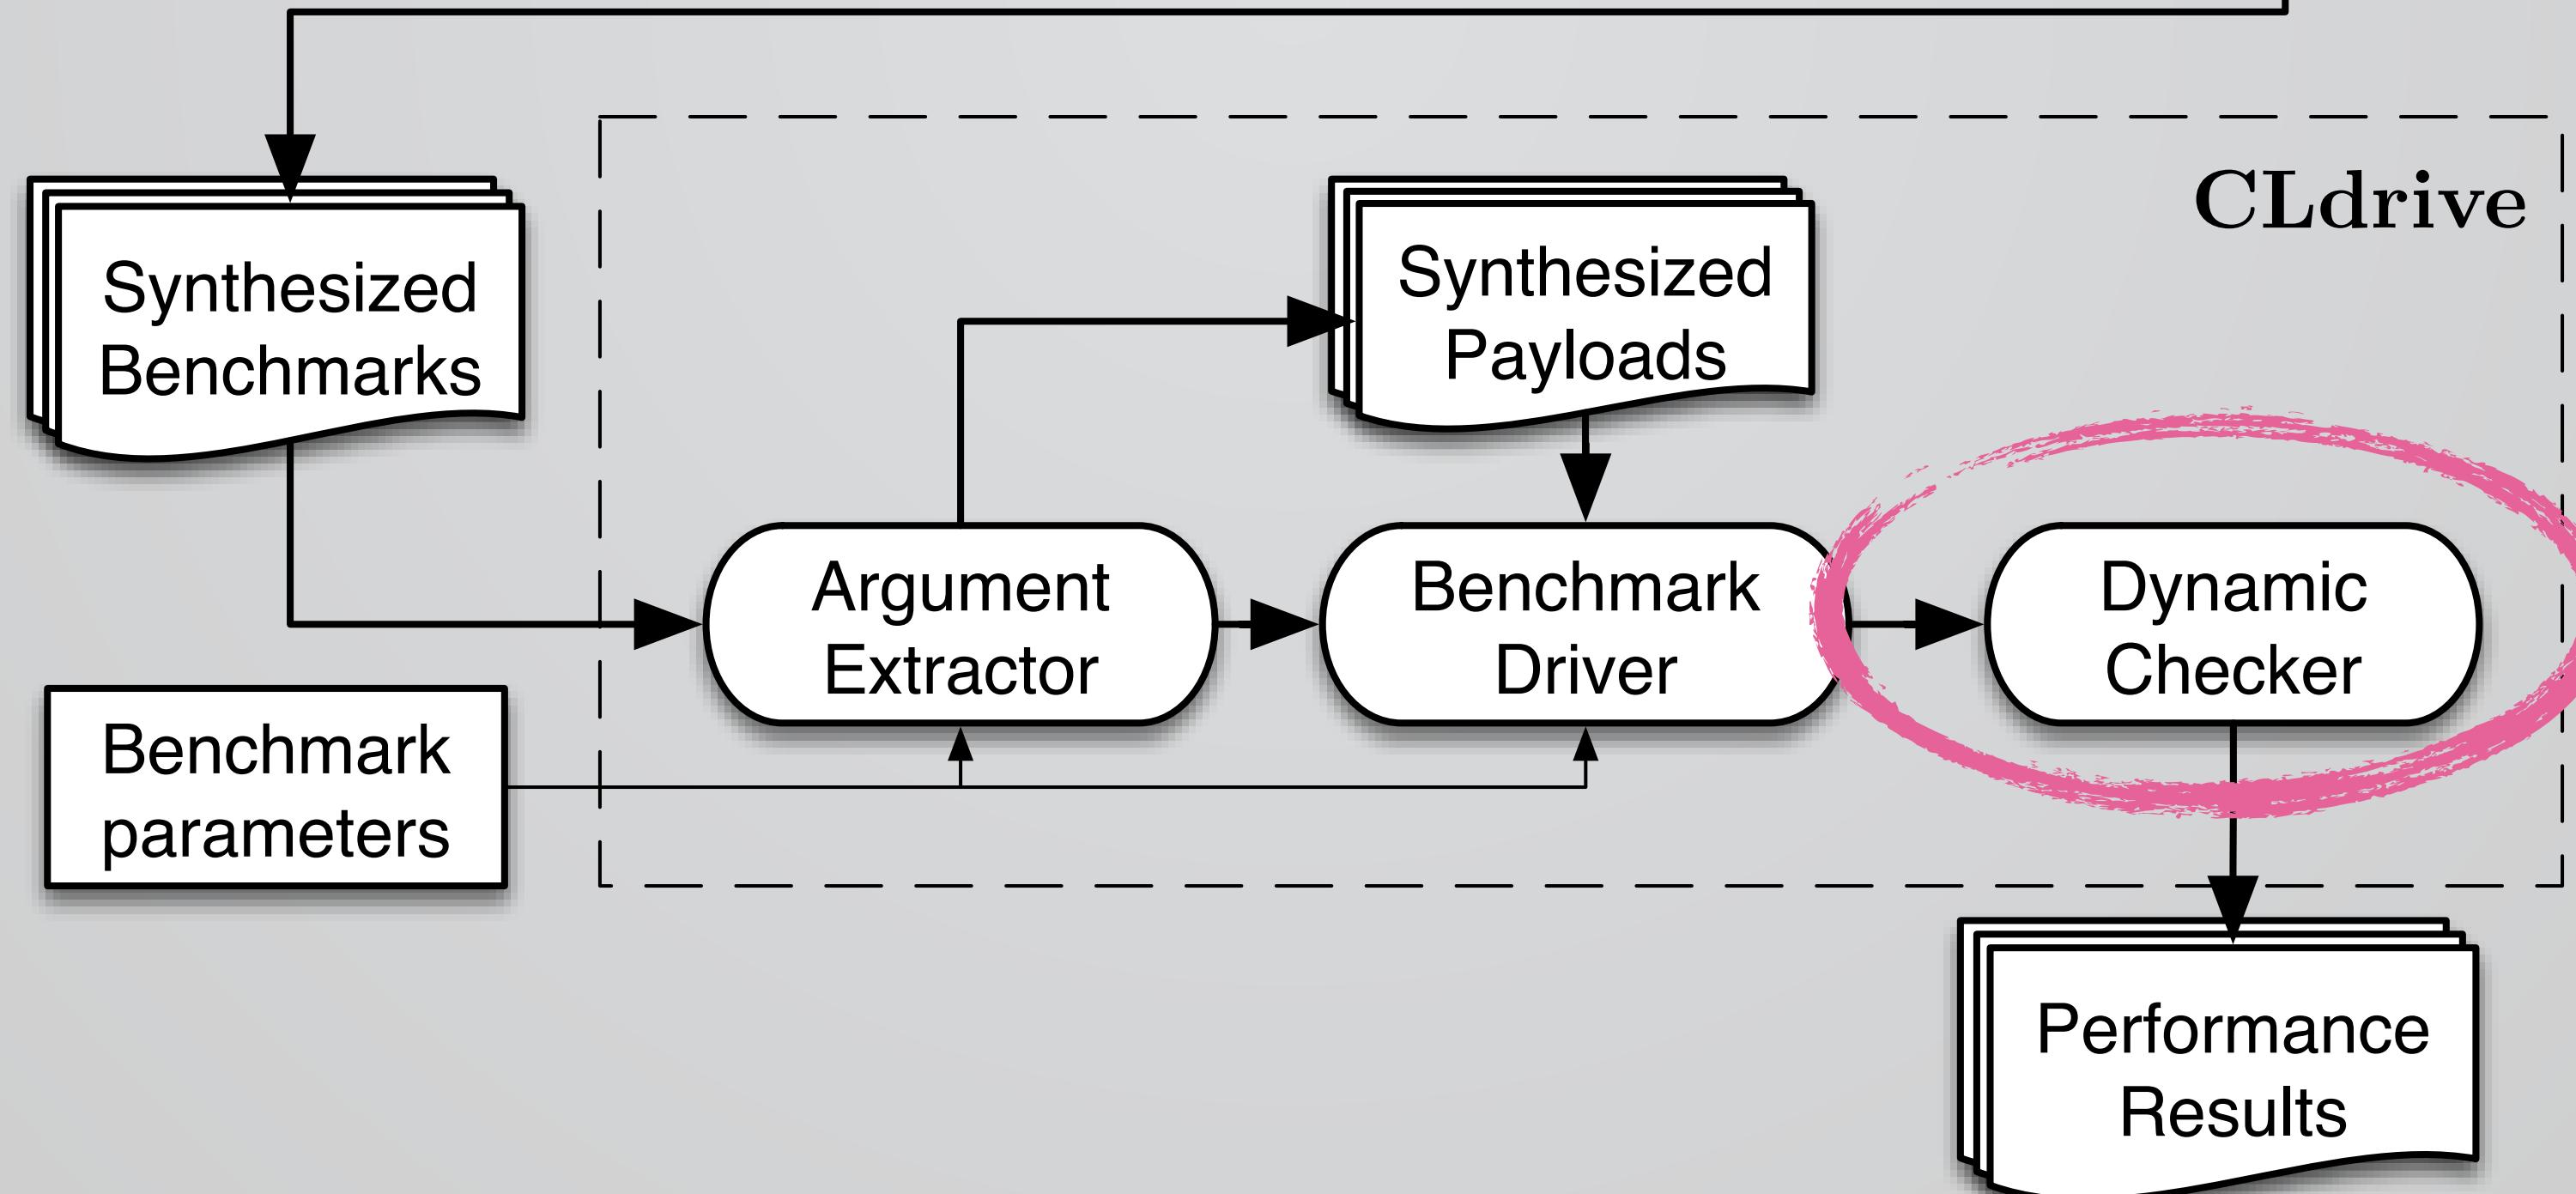
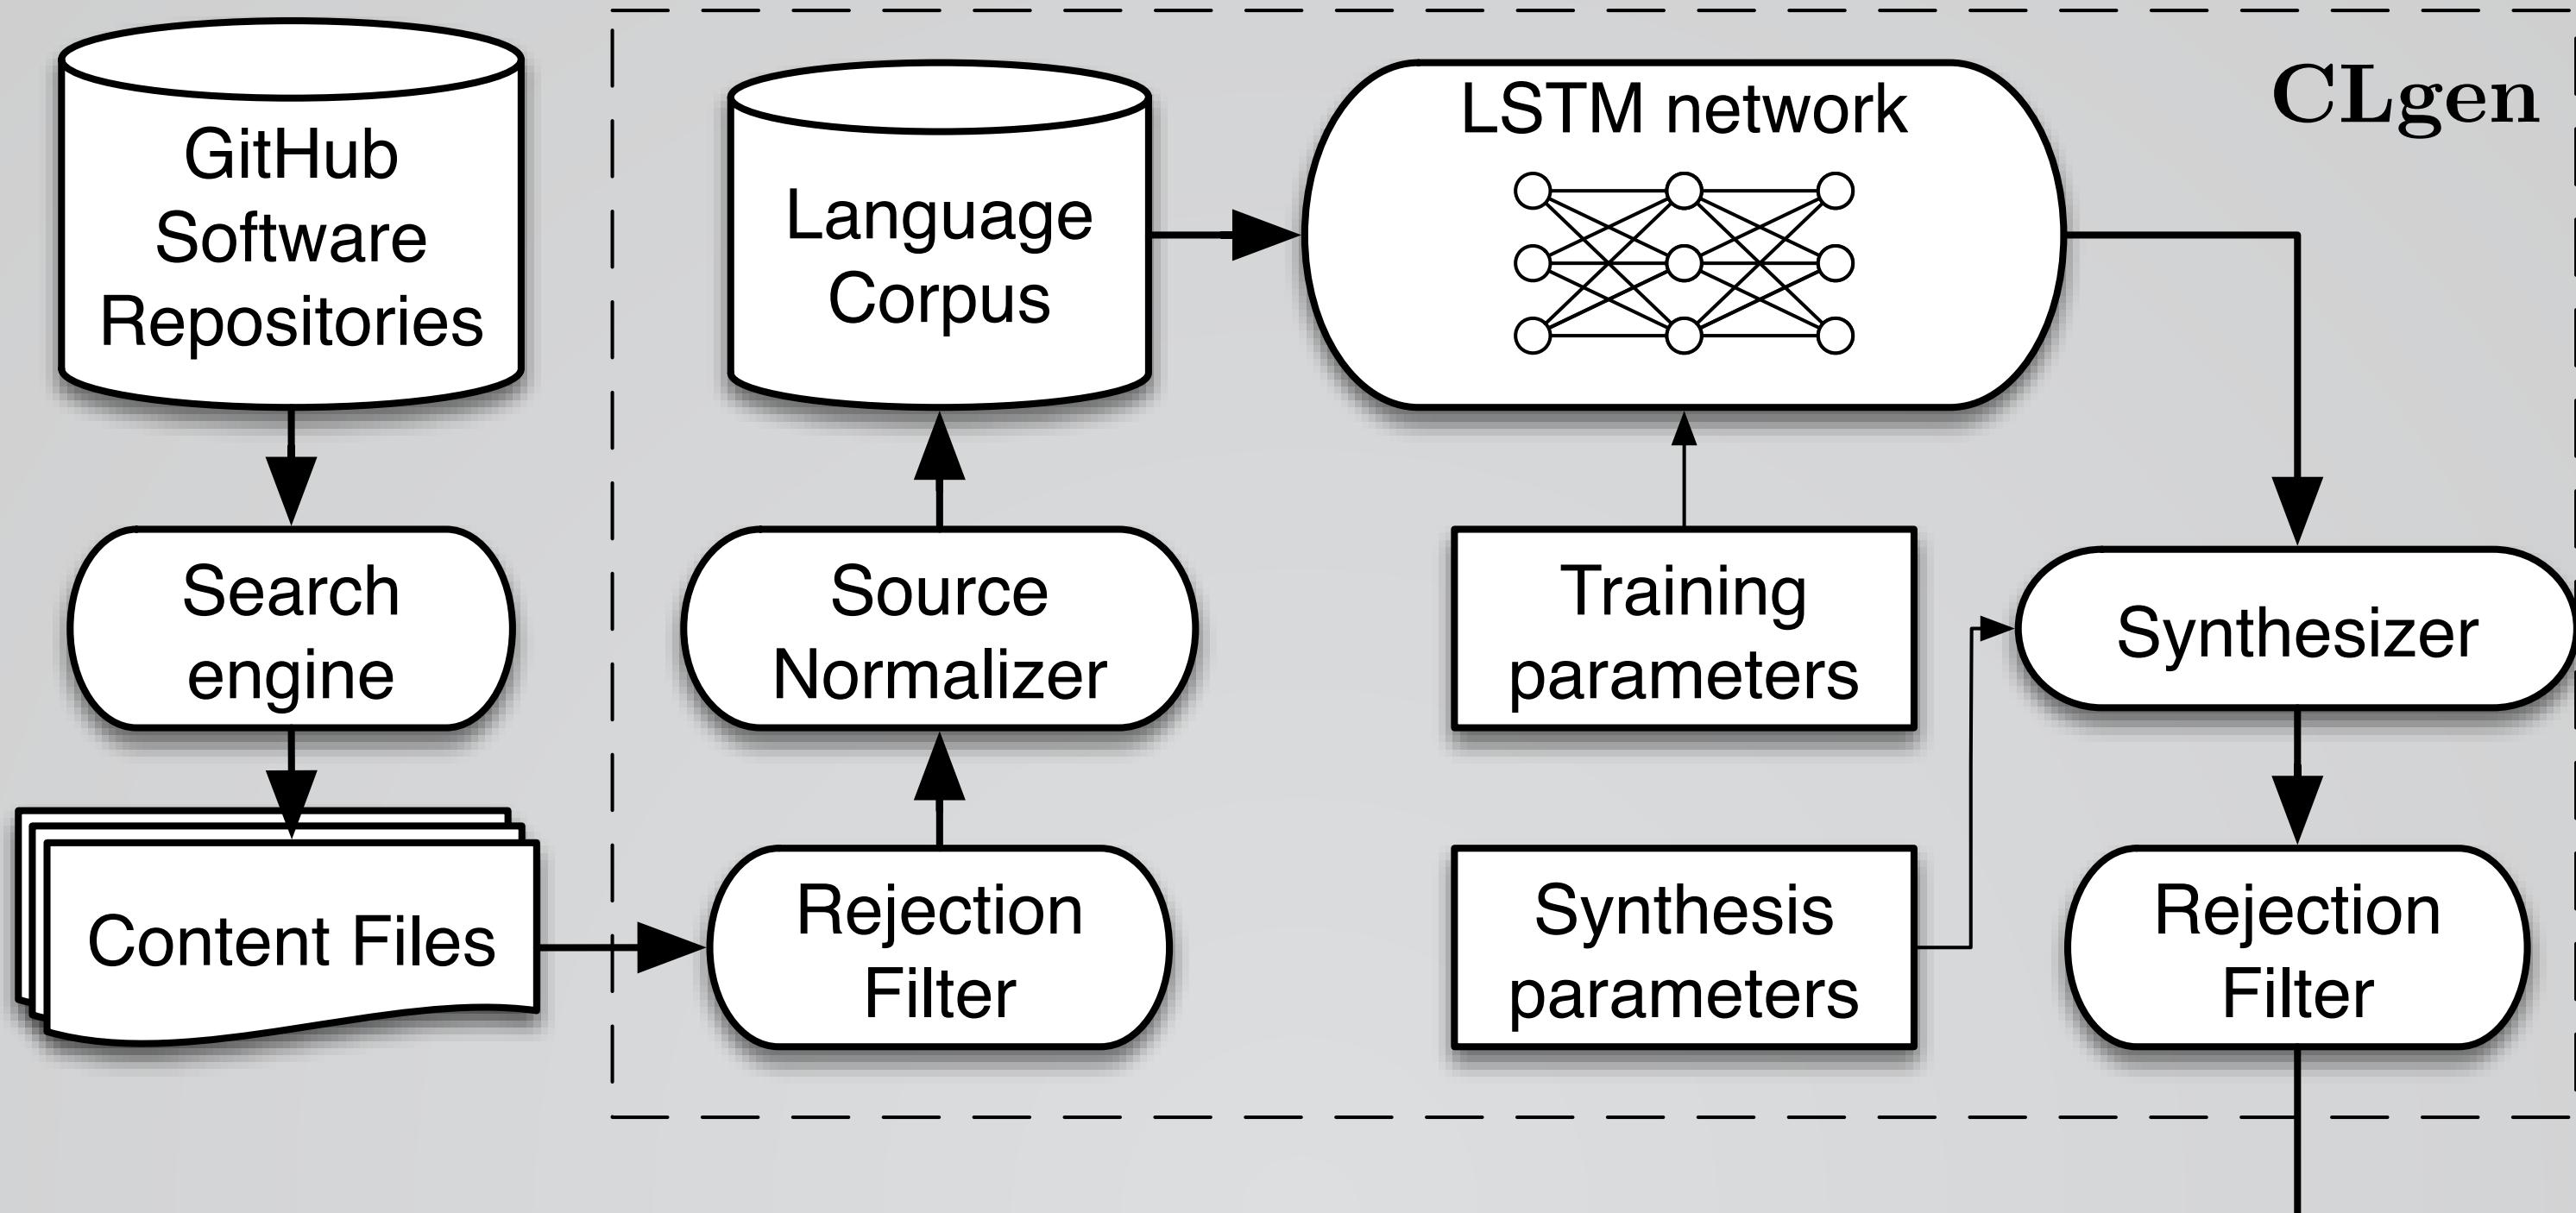




```
_A = random_payload(_A)    # generate inputs
_B = random_payload(_B)
_C = copy(_C)
_D = copy(_D)

A_ = k(_A)    # compute outputs
B_ = k(_B)
C_ = k(_C)
D_ = k(_D)

# differential test
assert (A_ != _A || B_ != _B) else NO_OUTPUTS
assert (A_ != B_ || C_ != D_) else INPUT_INSENSITIVE
assert (A_ == C_ && B_ == D_) else NON_DETERMINISTIC
```



does it
work?

Round 1

Player: 1002, Robot: 998

★★★★★ What a piece of junk!

By George S. Mitchell "gsmitchell"

I checked around Amazon as well as some other sites, and decided these were what I needed. I read the reviews, looked at the pictures and bought two of these. I received them today. I unpacked one of them, read the instructions and immediately filled out the return request for Amazon. These mounts are very, very heavy and require 6, yes SIX, screws to mount to the wall. Are you kidding me? The TVs are only 22 lbs! Two screws in the studs and then more damage to the wall than is necessary. What a joke! And Lag bolts? Ever heard of screws? The other issue is that these are very hard to move. I bought an articulating wall mount for a reason, and I expect them to be able to move without a wrench in my hand.

★★★★★ Excellent for the price

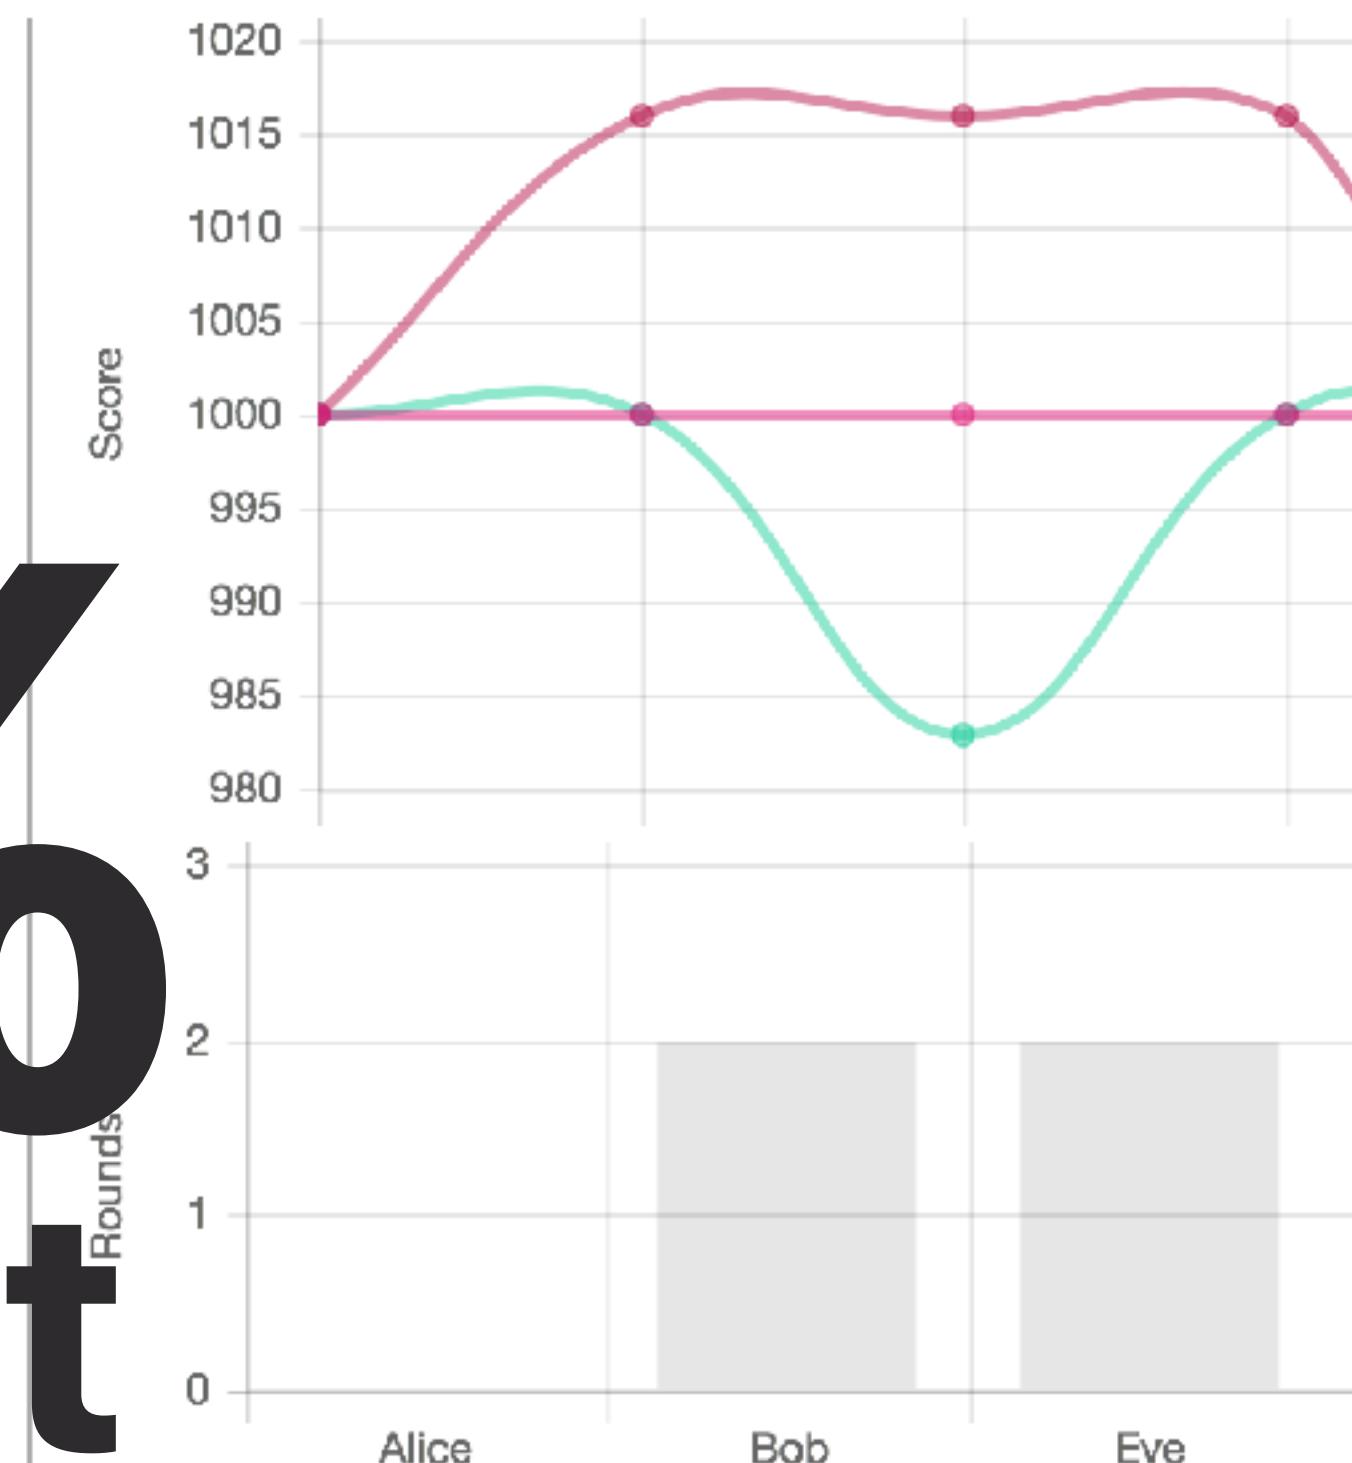
By A. Lopez "David Roberts"

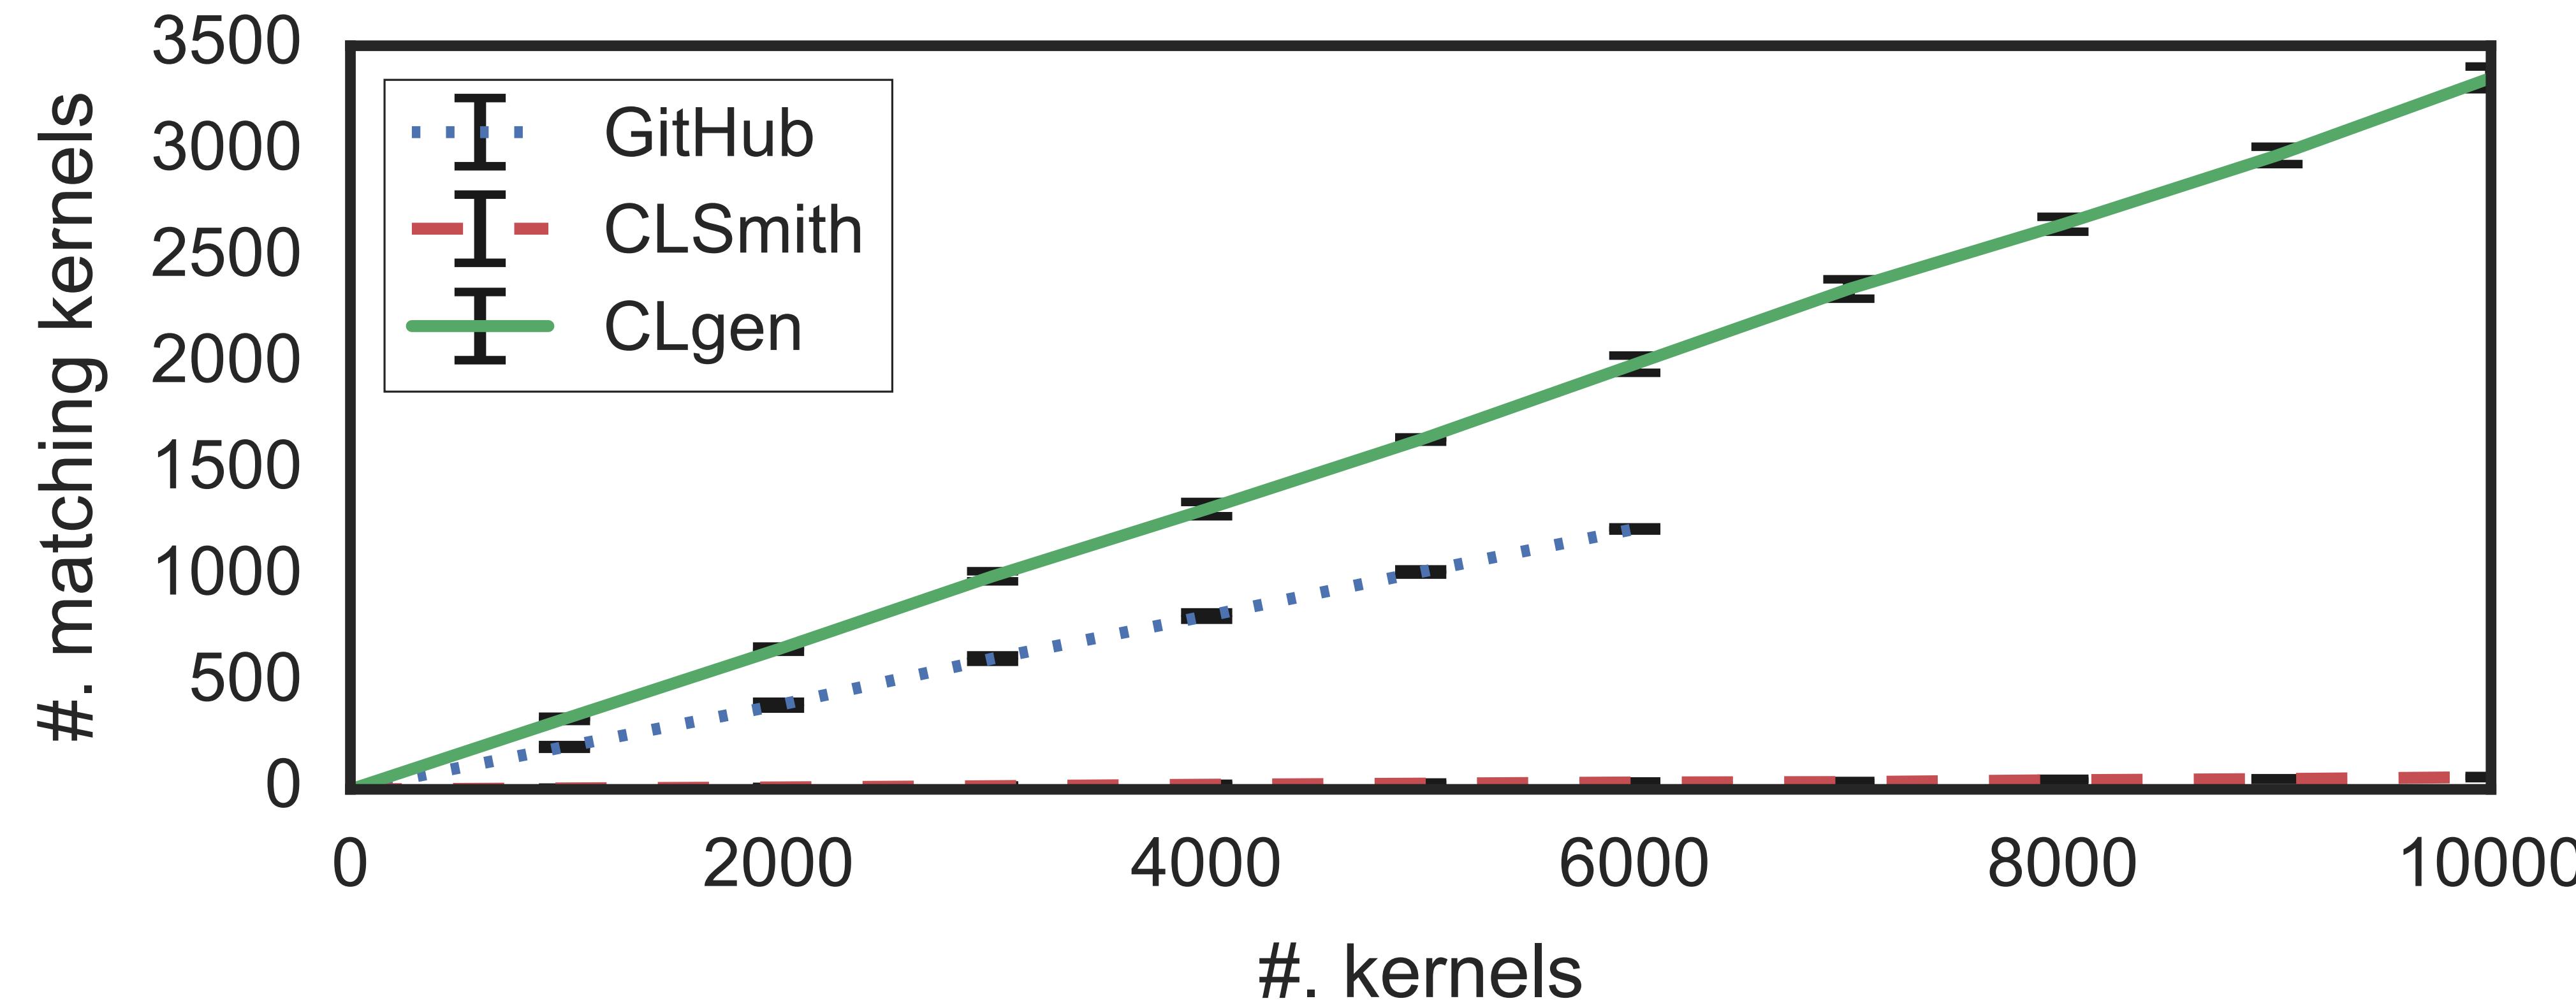
This thing is fantastic. It has the same size and speed and sharp sound. I like the power turntable. I use it every time you connect your computer and trying to figure out you want to get a problem with a ti
a problems with a converter that come out. Good price for it and a couple of days ago we
software installed so far I am very happy with this product. Works great for what we
wanted to do. The way the laptop company will a lot
more expensive. I can't tell that any other device
this is a protective product that works well
for a reasonable price and works great for the price
and it works.

52% blind test

This side is the Robot

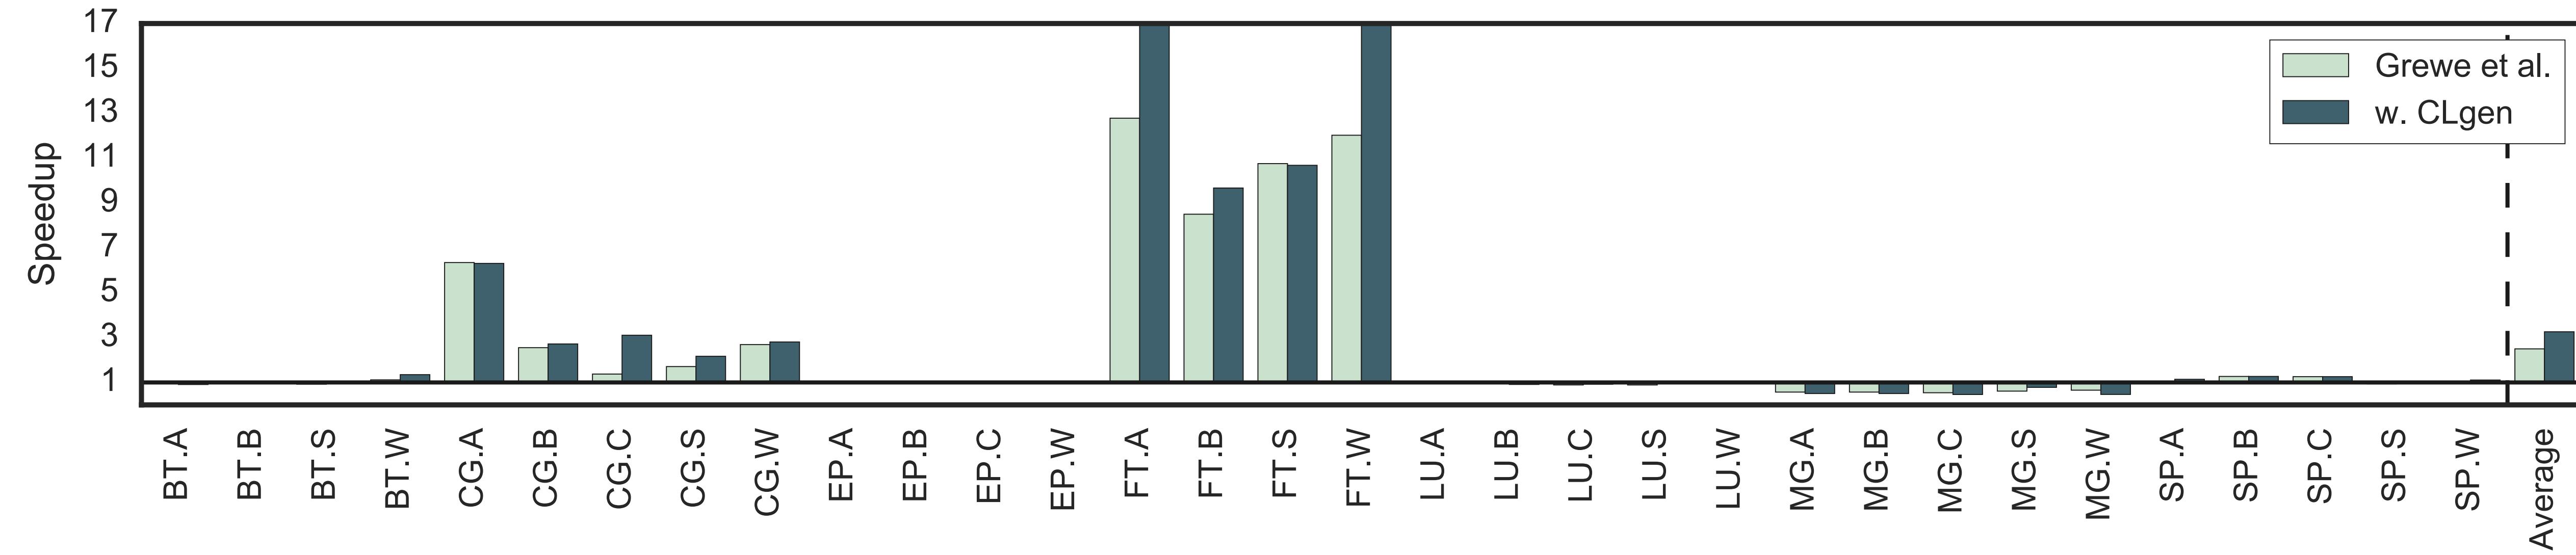
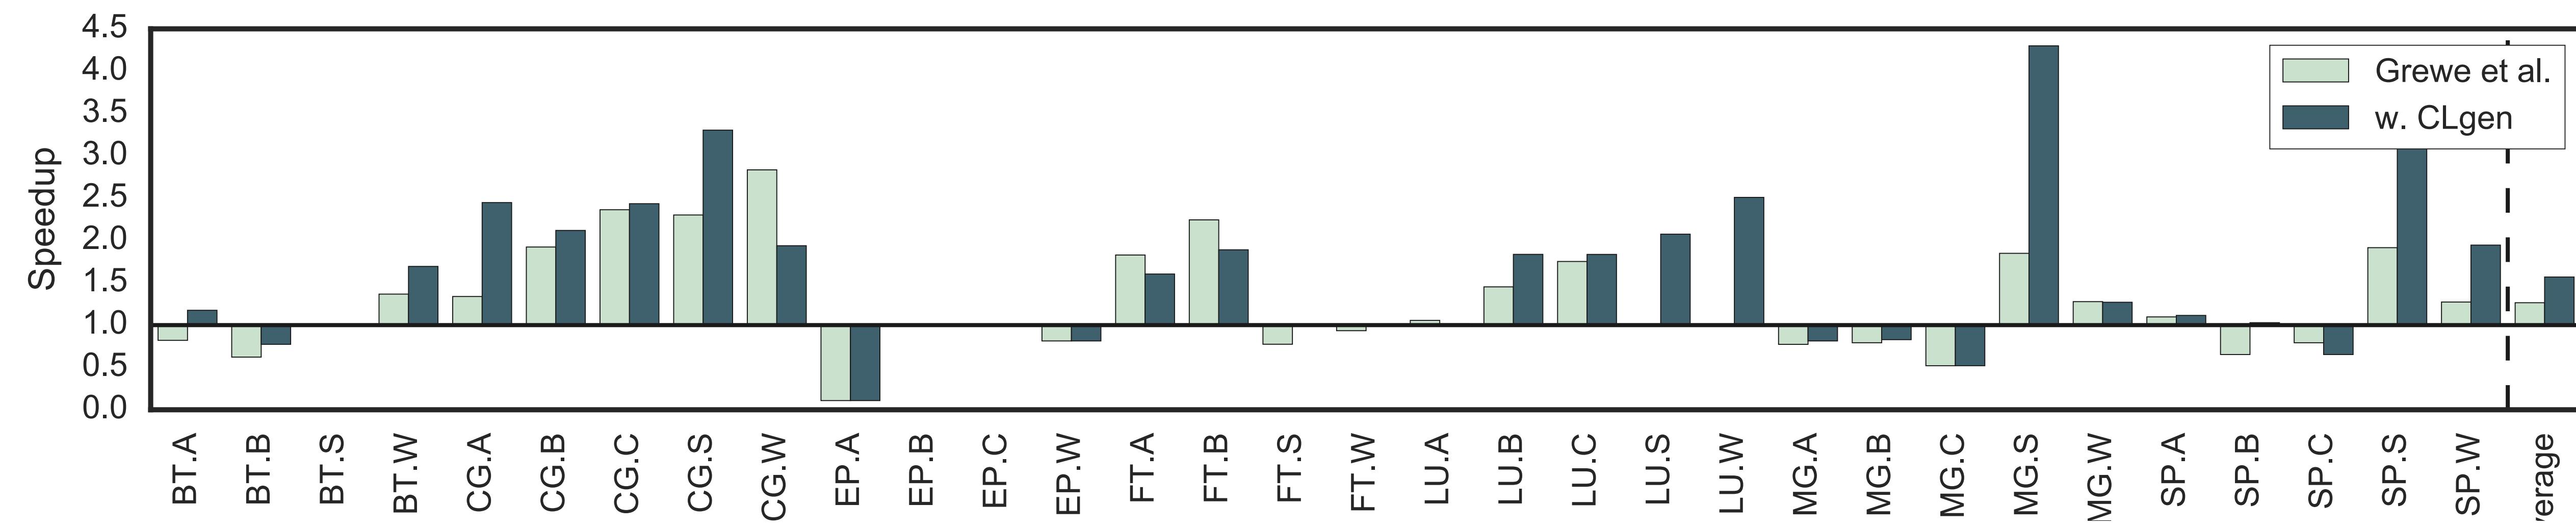
This side is the Robot



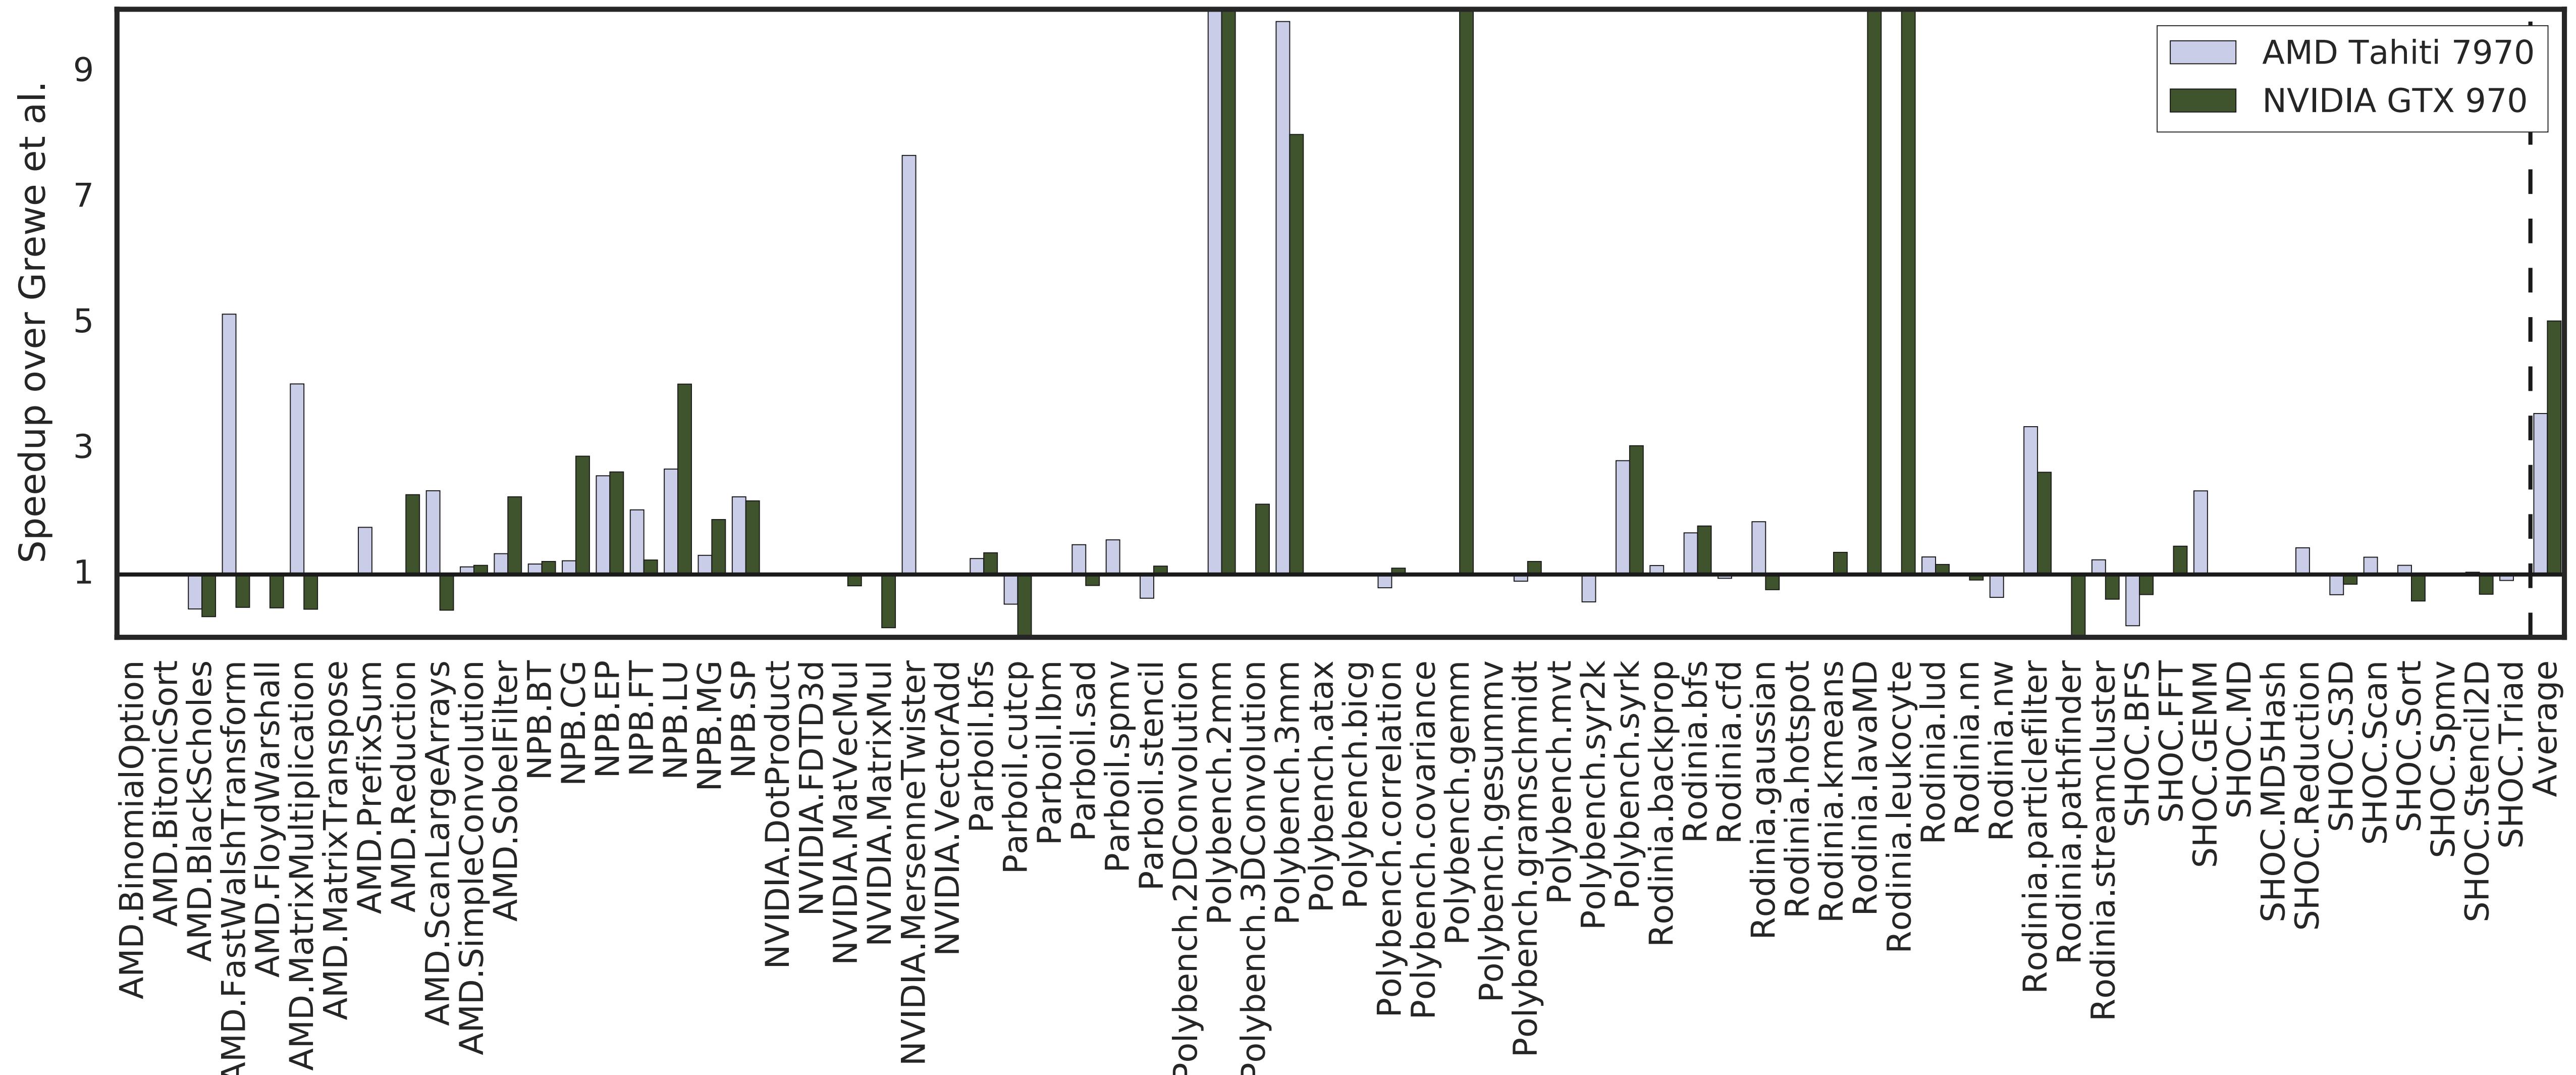


Raw Code Features

comp	static	#. compute operations
mem	static	#. accesses to global memory
localmem	static	#. accesses to local memory
coalesced	static	#. coalesced memory accesses



7 programs, 1,000 synthetic benchmarks. 1.27x faster



71 programs, 1,000 synthetic benchmarks. 4.30x faster

Good Things

Basically* language agnostic.

35 million repos on GitHub. We're using 0.00004%.

Generates 2000 OpenCL benchmarks per machine per day.

Bad Things

No support for things declared outside of kernel scope.

Undirected almost to a fault.

What

next?

very short term

paper + ae
google
eurolab
taco

next six months

discard rate < 90%

searching feature space
grammar

next six months

discard rate < 90%

searching feature space
grammar

next six months

discard rate < 90%

1. alternate vocabularies
2. corpus encodings

next six months

discard rate < 90%

searching feature space
grammar

next six months

discard rate < 90%

searching feature space
grammar

next six months

1. hill climb mutated seeds and RNG
2. train models with source + features

searching feature space

Experiments:

convergence to 256 benchmarks
can it beat GitHub for all kernels?

next six months

discard rate < 90%

searching feature space
grammar

next six months

discard rate < 90%

searching feature space

grammar

next six months

Oracle:

- + generate_identifier(S, regexp)
- + chose_next(S, production_rules[])

grammar

theses

- 
1. introduction
 2. literature review
 3. language modelling for program code
 4. synthesising programs
 5. searching the feature space
 6. learning the compiler optimisation pipeline
 7. conclusions

reccap

paper

adapt
hipgpu
cgo

posters

google
hipecac
pldi
acaces

talks

amazon
ocado
codeplay

idea Directed
random program
generation

idea Gamifying
hyperparameter
tuning through
non-interactive
turing tests

idea Optimisation
pass selection and
ordering using
reinforcement
learning